Introduction

http://www.cs.cmu.edu/~16385/

16-385 Computer Vision
Spring 2018, Lecture 1

Overview of today's lecture

- Teaching staff introductions
- What is computer vision?
- Course fast-forward and logistics

Teaching staff introductions

Instructor: Ioannis (Yannis) Gkioulekas

I won’t hold it against you if you mispronounce my last name

Originally from Greece

National Technical University of Athens (2004-09)

Harvard University (2009-17)

Carnegie Mellon University (2017-now)

me at Harvard in 2011
(obviously need new photo)

My website: http://www.cs.cmu.edu/~igkioule

Building a scatterometer

Camera for measuring parameters of scattering materials

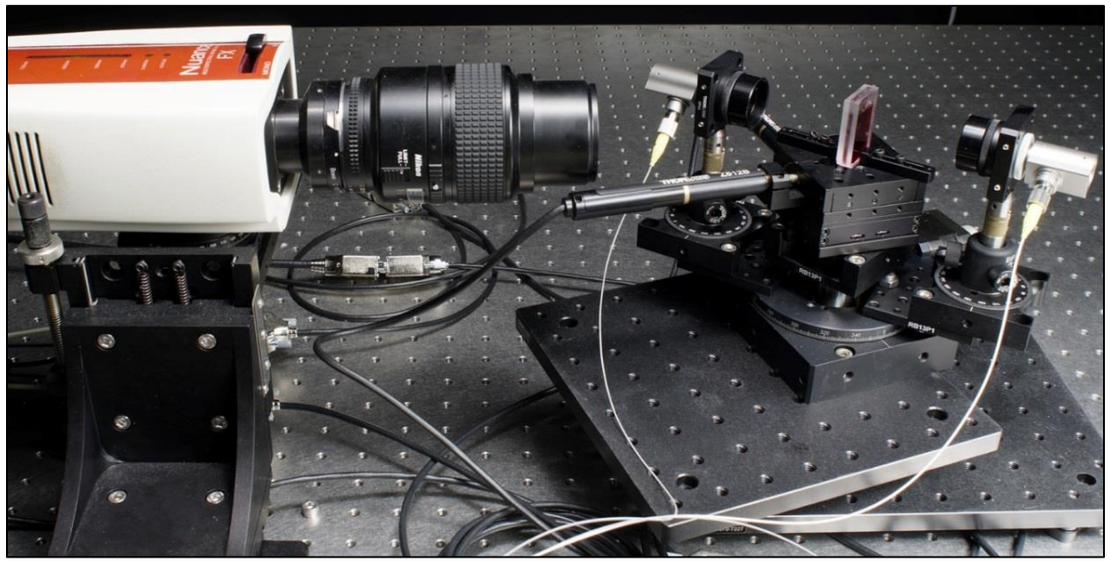

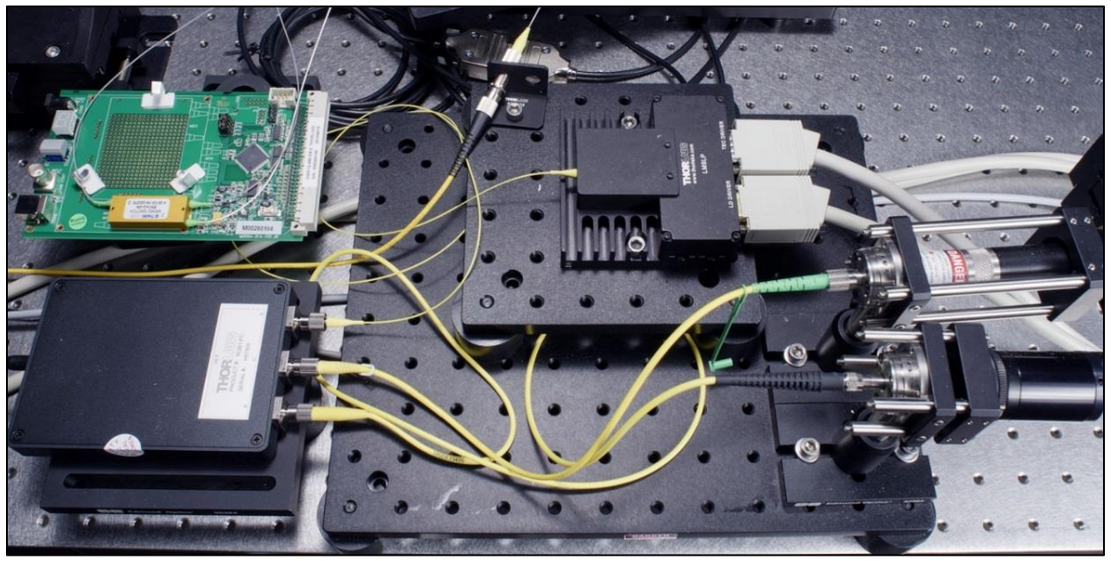

image synthesized from measurements

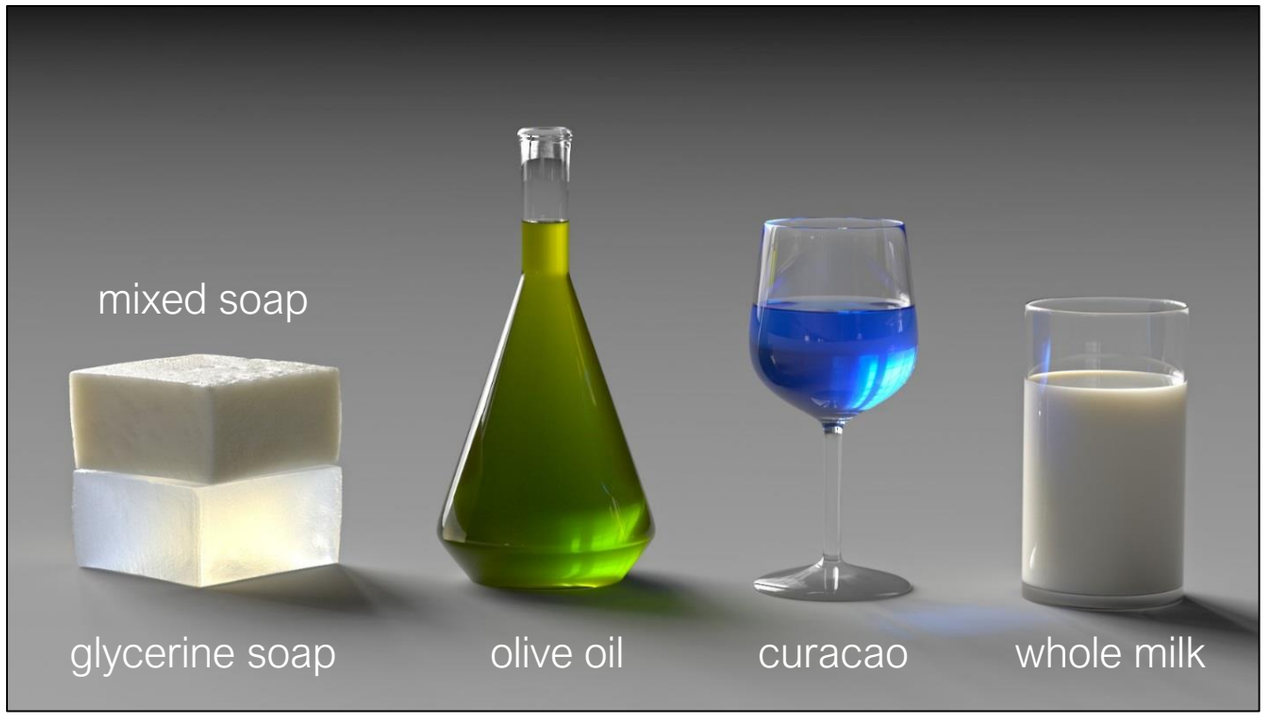

Quadrillion FPS video

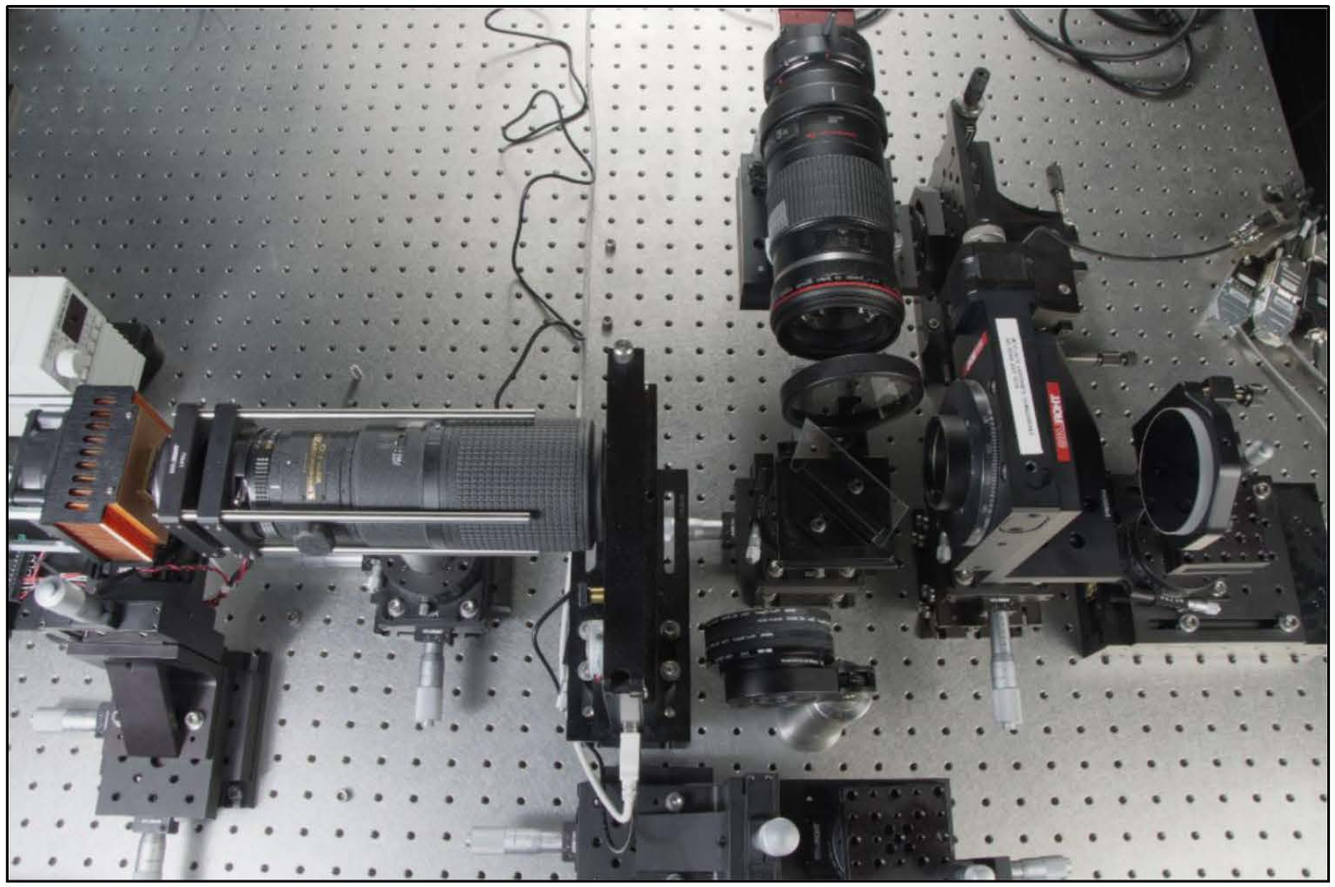

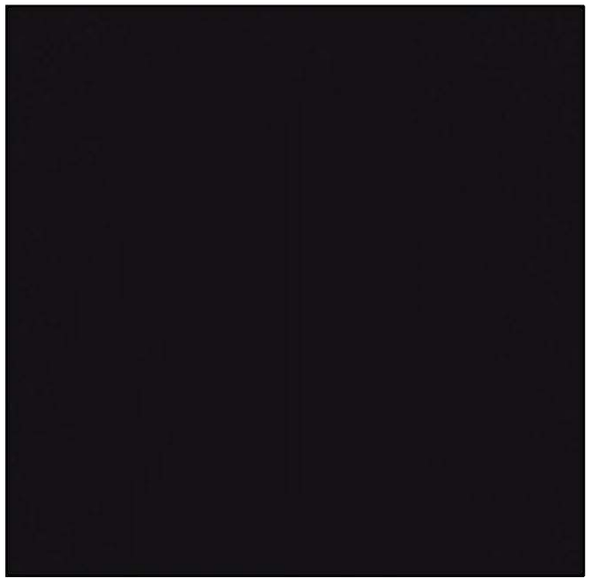

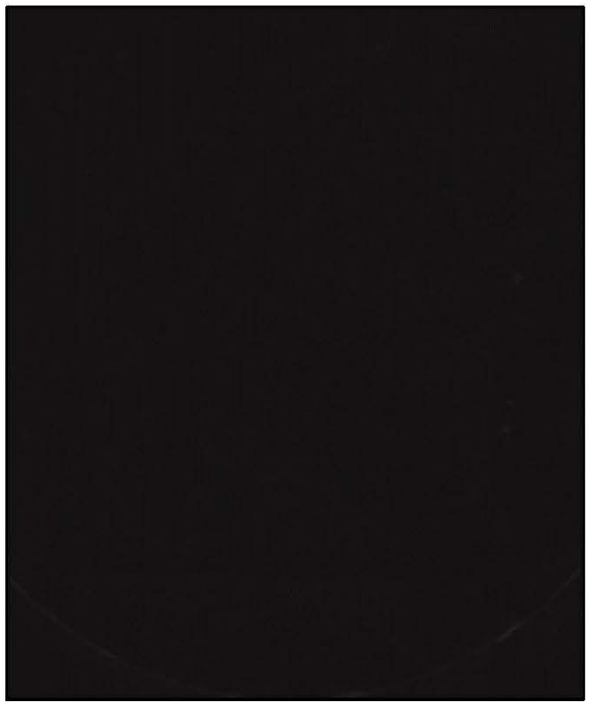

Seeing inside objects

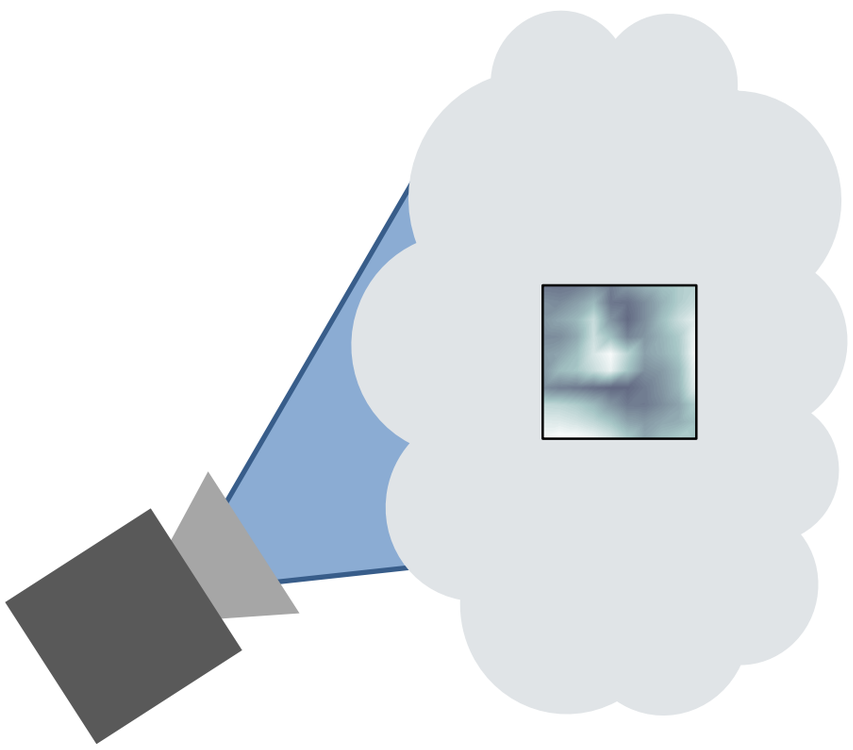

camera

thick smoke cloud

what a regular camera sees

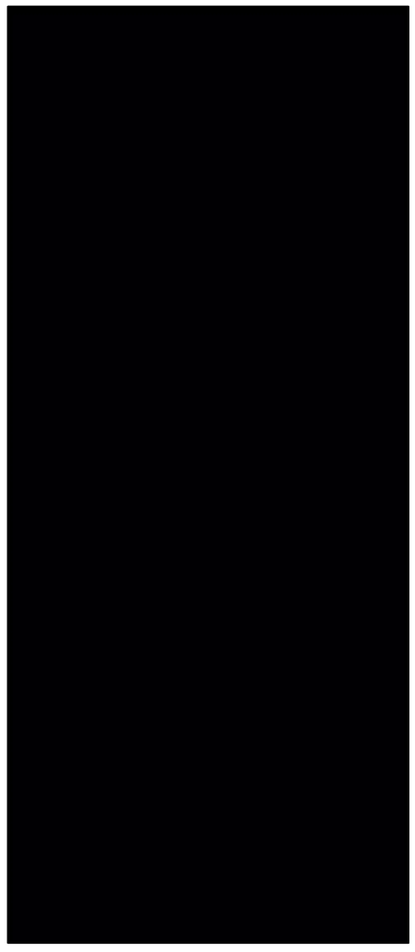

what our camera sees

a slice through the cloud

Seeing around walls

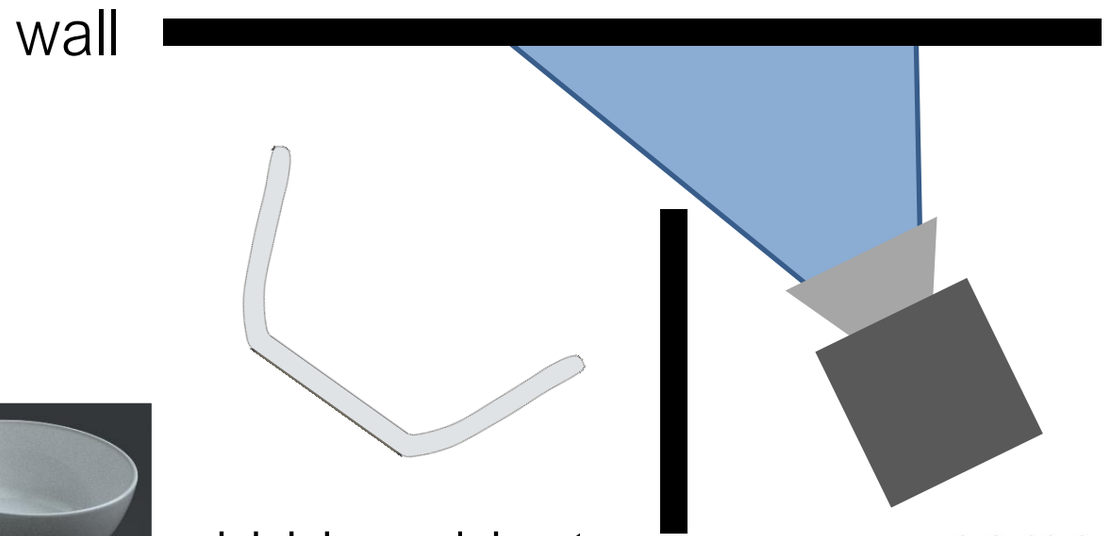

hidden object

camera

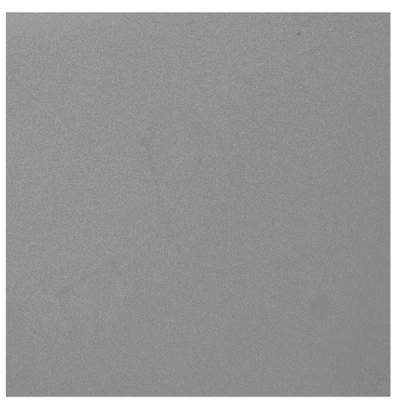

what a regular camera sees

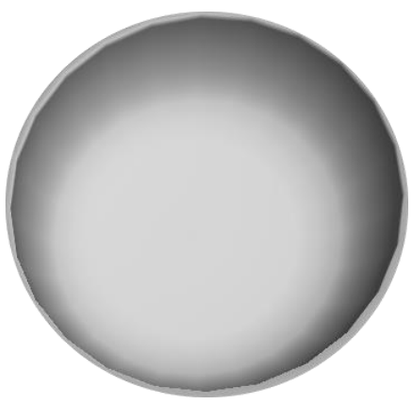

what shape our camera sees

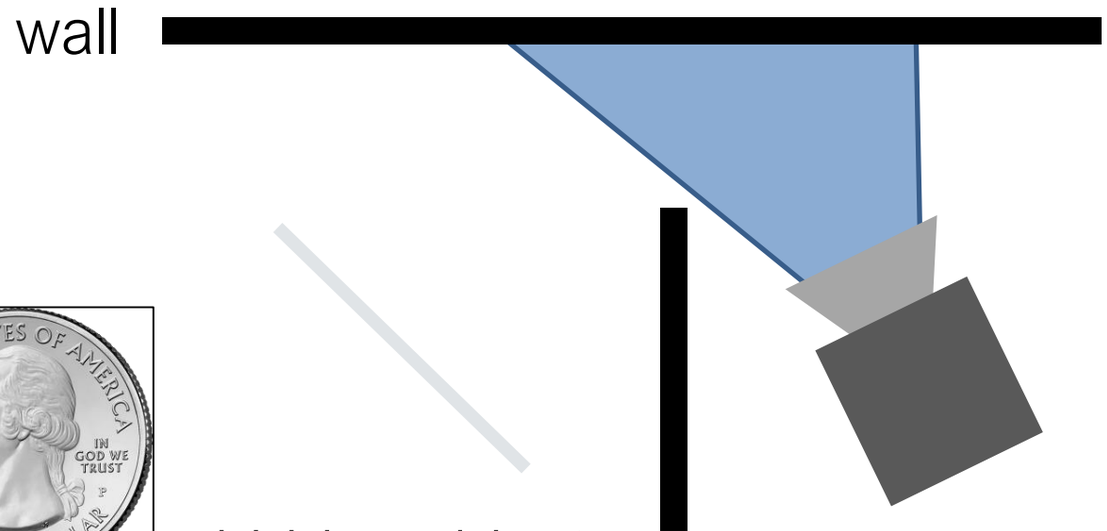

hidden object

camera

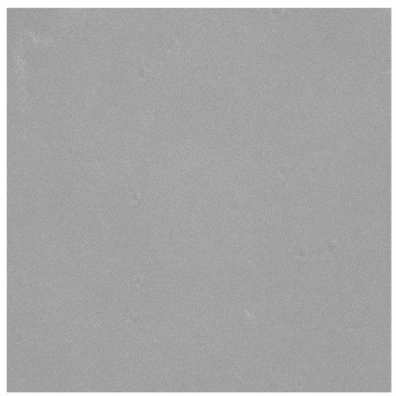

what a regular camera sees

what depth our camera sees

TA: Gaurav Mittal

Master of Science in Computer Vision (MSCV)

Research Interests

Image scene understanding/generation, deep learning, language and visual intelligence

Current Area of Research

Neural network compression using Reinforcement Learning (under Prof. Kris Kitani)

Past Research

- Video generation using captions with Variational Autoencoders
- Detection and coarse segmentation of garbage in images (SpotGarbage)

TA: Shashank Tripathi

Master of Science in Computer Vision (MSCV)

Research Interests

Object detection/tracking, Deep Learning, Medical Image Processing

Current Area of Research

Neural network compression, Reinforcement Learning for vision

Past Research

- Deep spectral-matching based shape features for efficient Alzheimer's disease classification
- Superpixel-augmented Convolutional Neural Networks
- Vision control system for automated micromanipulation

What is computer vision?

What a person sees

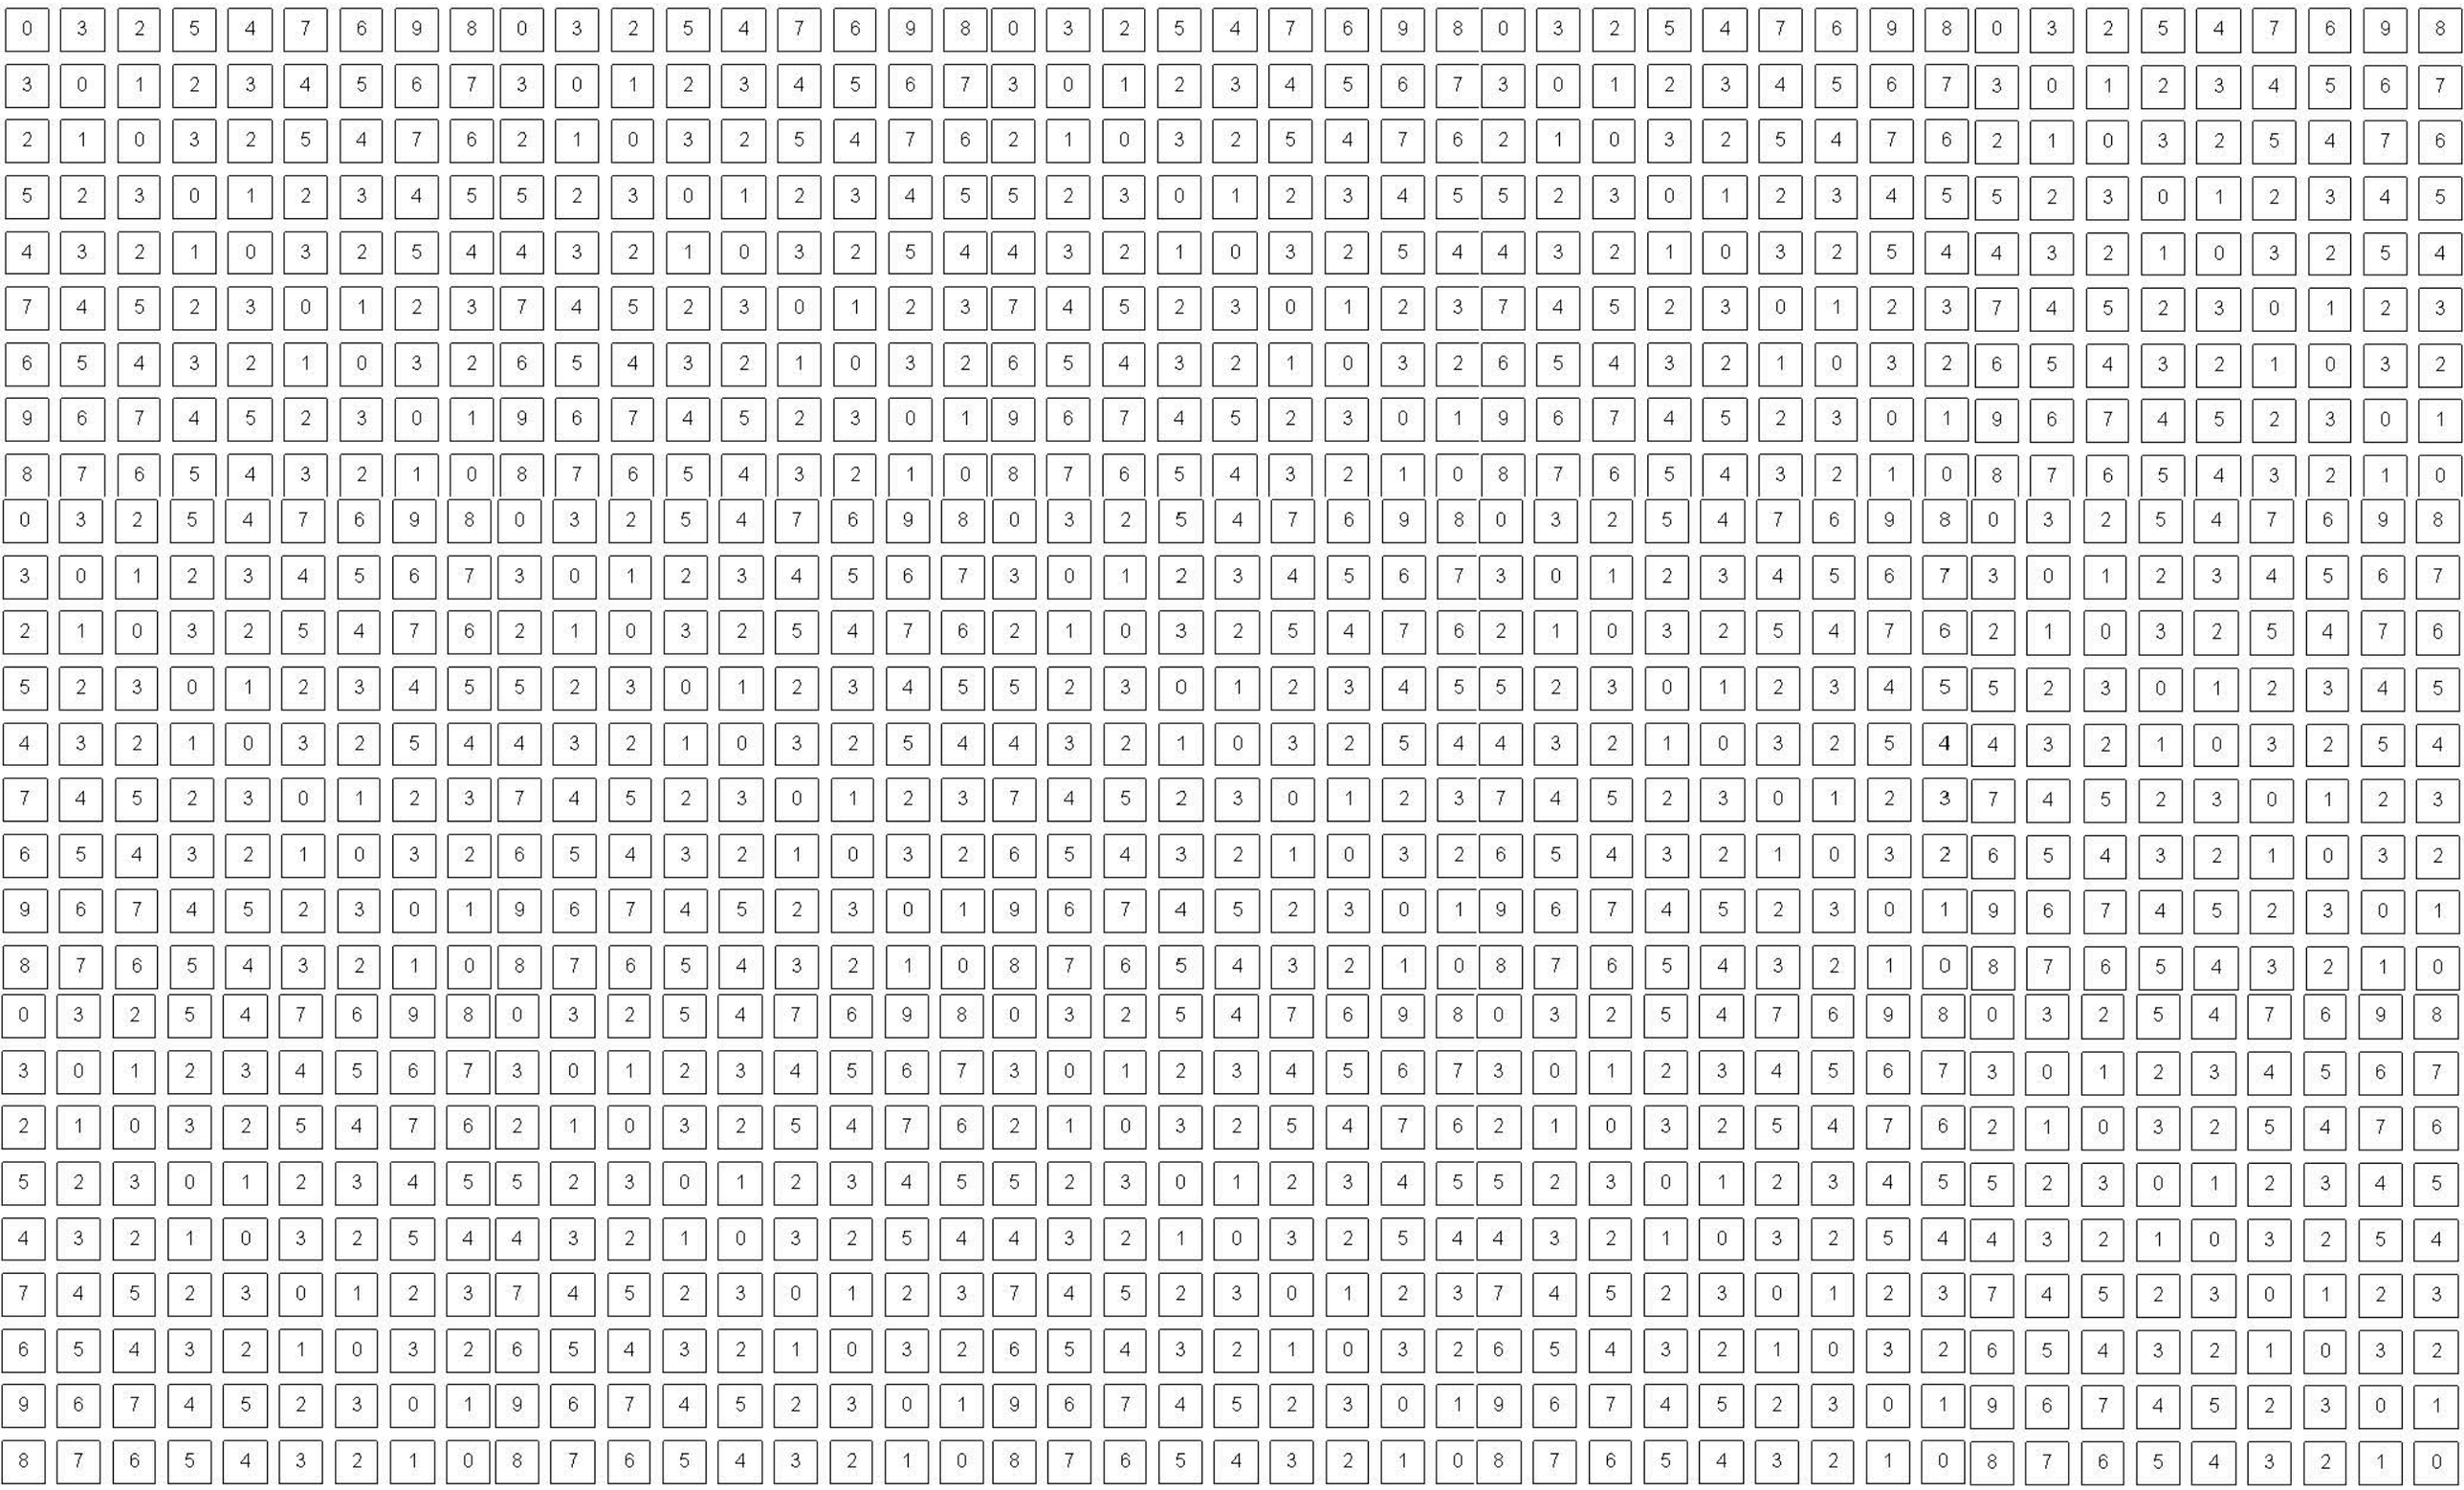

What a computer sees

Why are we able to interpret this image?

The goal of computer vision is

to give computers

(super) human-level perception

typical perception pipeline

representation

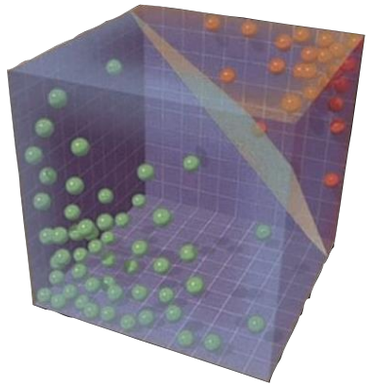

‘fancy math’

output

typical perception pipeline

representation

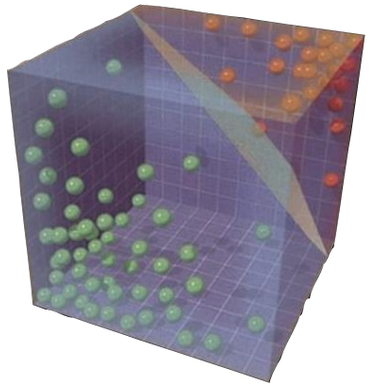

‘fancy math’

what should we look at?
(image features)

output

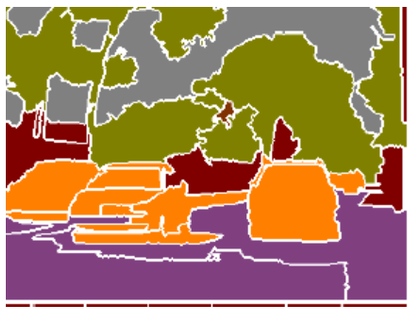

what can we understand?
(semantic segmentation)

typical perception pipeline

representation

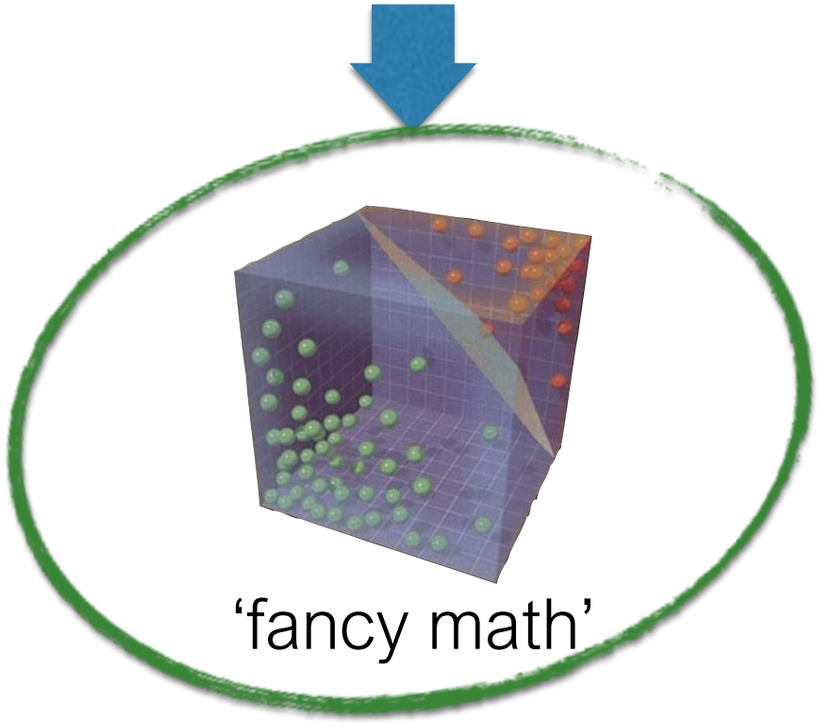

output

what should we look at?
(image features)

easy to get lost in
the techniques

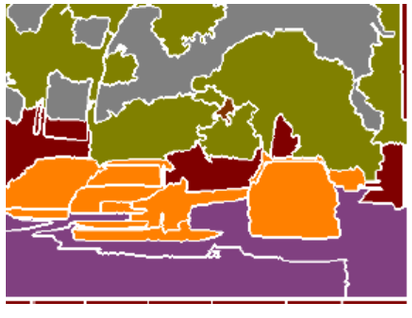

what can we understand?
(semantic segmentation)

typical perception pipeline

representation

what should we look at?
(image features)

The parts that we are most interested in

'fancy math'

output

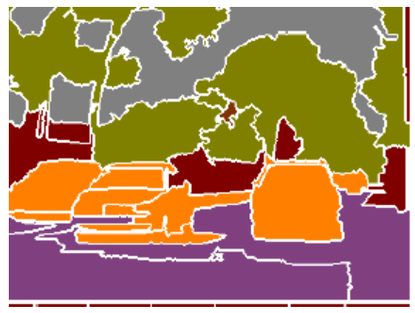

what can we understand?
(semantic segmentation)

Important note:

In general, computer vision does not work

Important note:

In general, computer vision does not work

(except in certain situations/conditions)

Applications of computer vision

Machine vision

Automated visual inspection

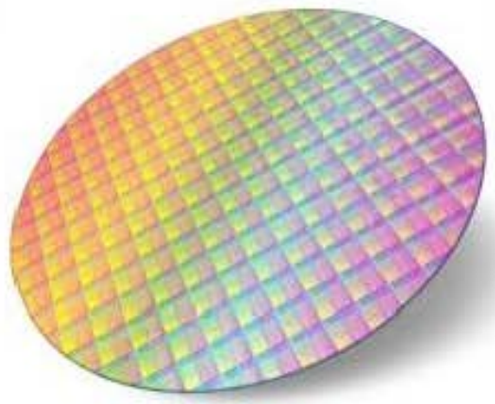

Object Recognition

Toshiba Tech IS-910T 2013

DataLogic LaneHawk LH4000 2012

Face detection

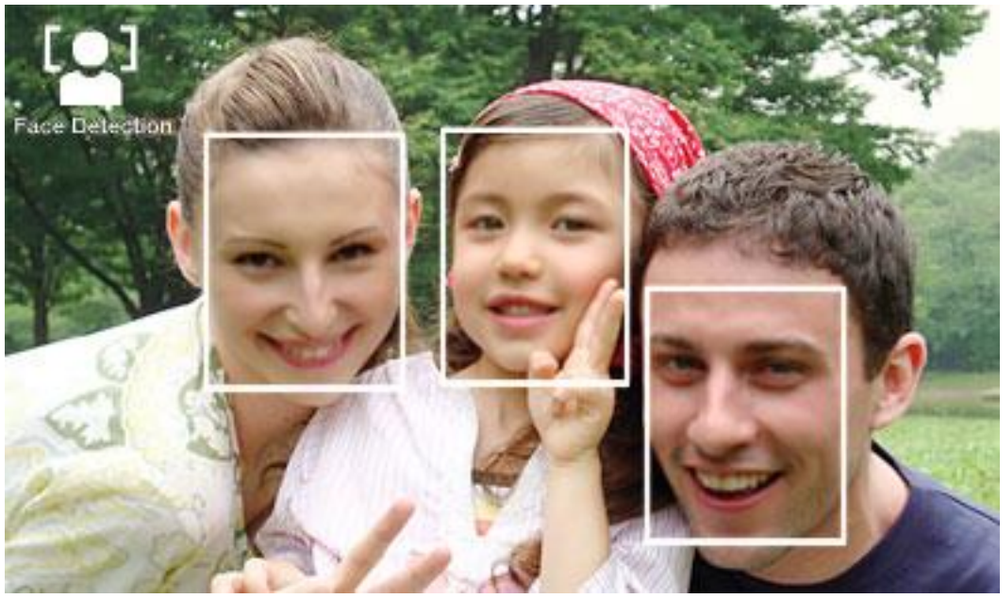

Sony Cyber-shot

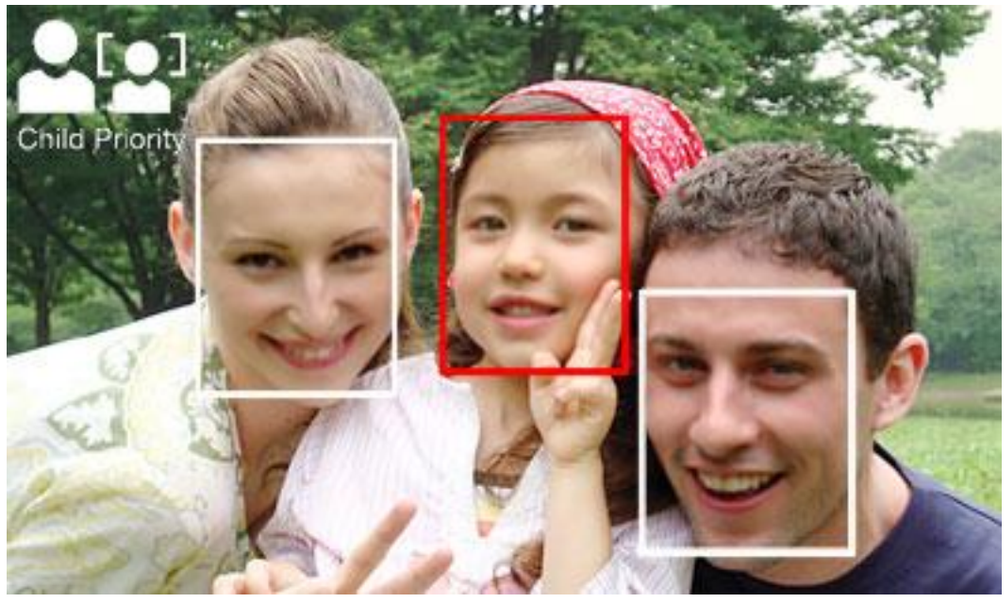

Age recognition

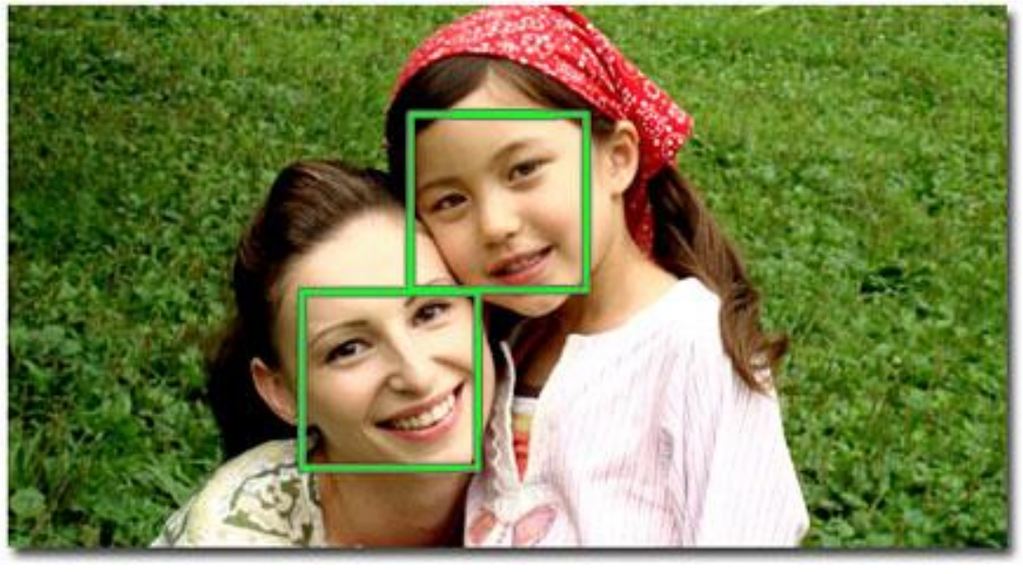

Smile recognition

Face makeovers

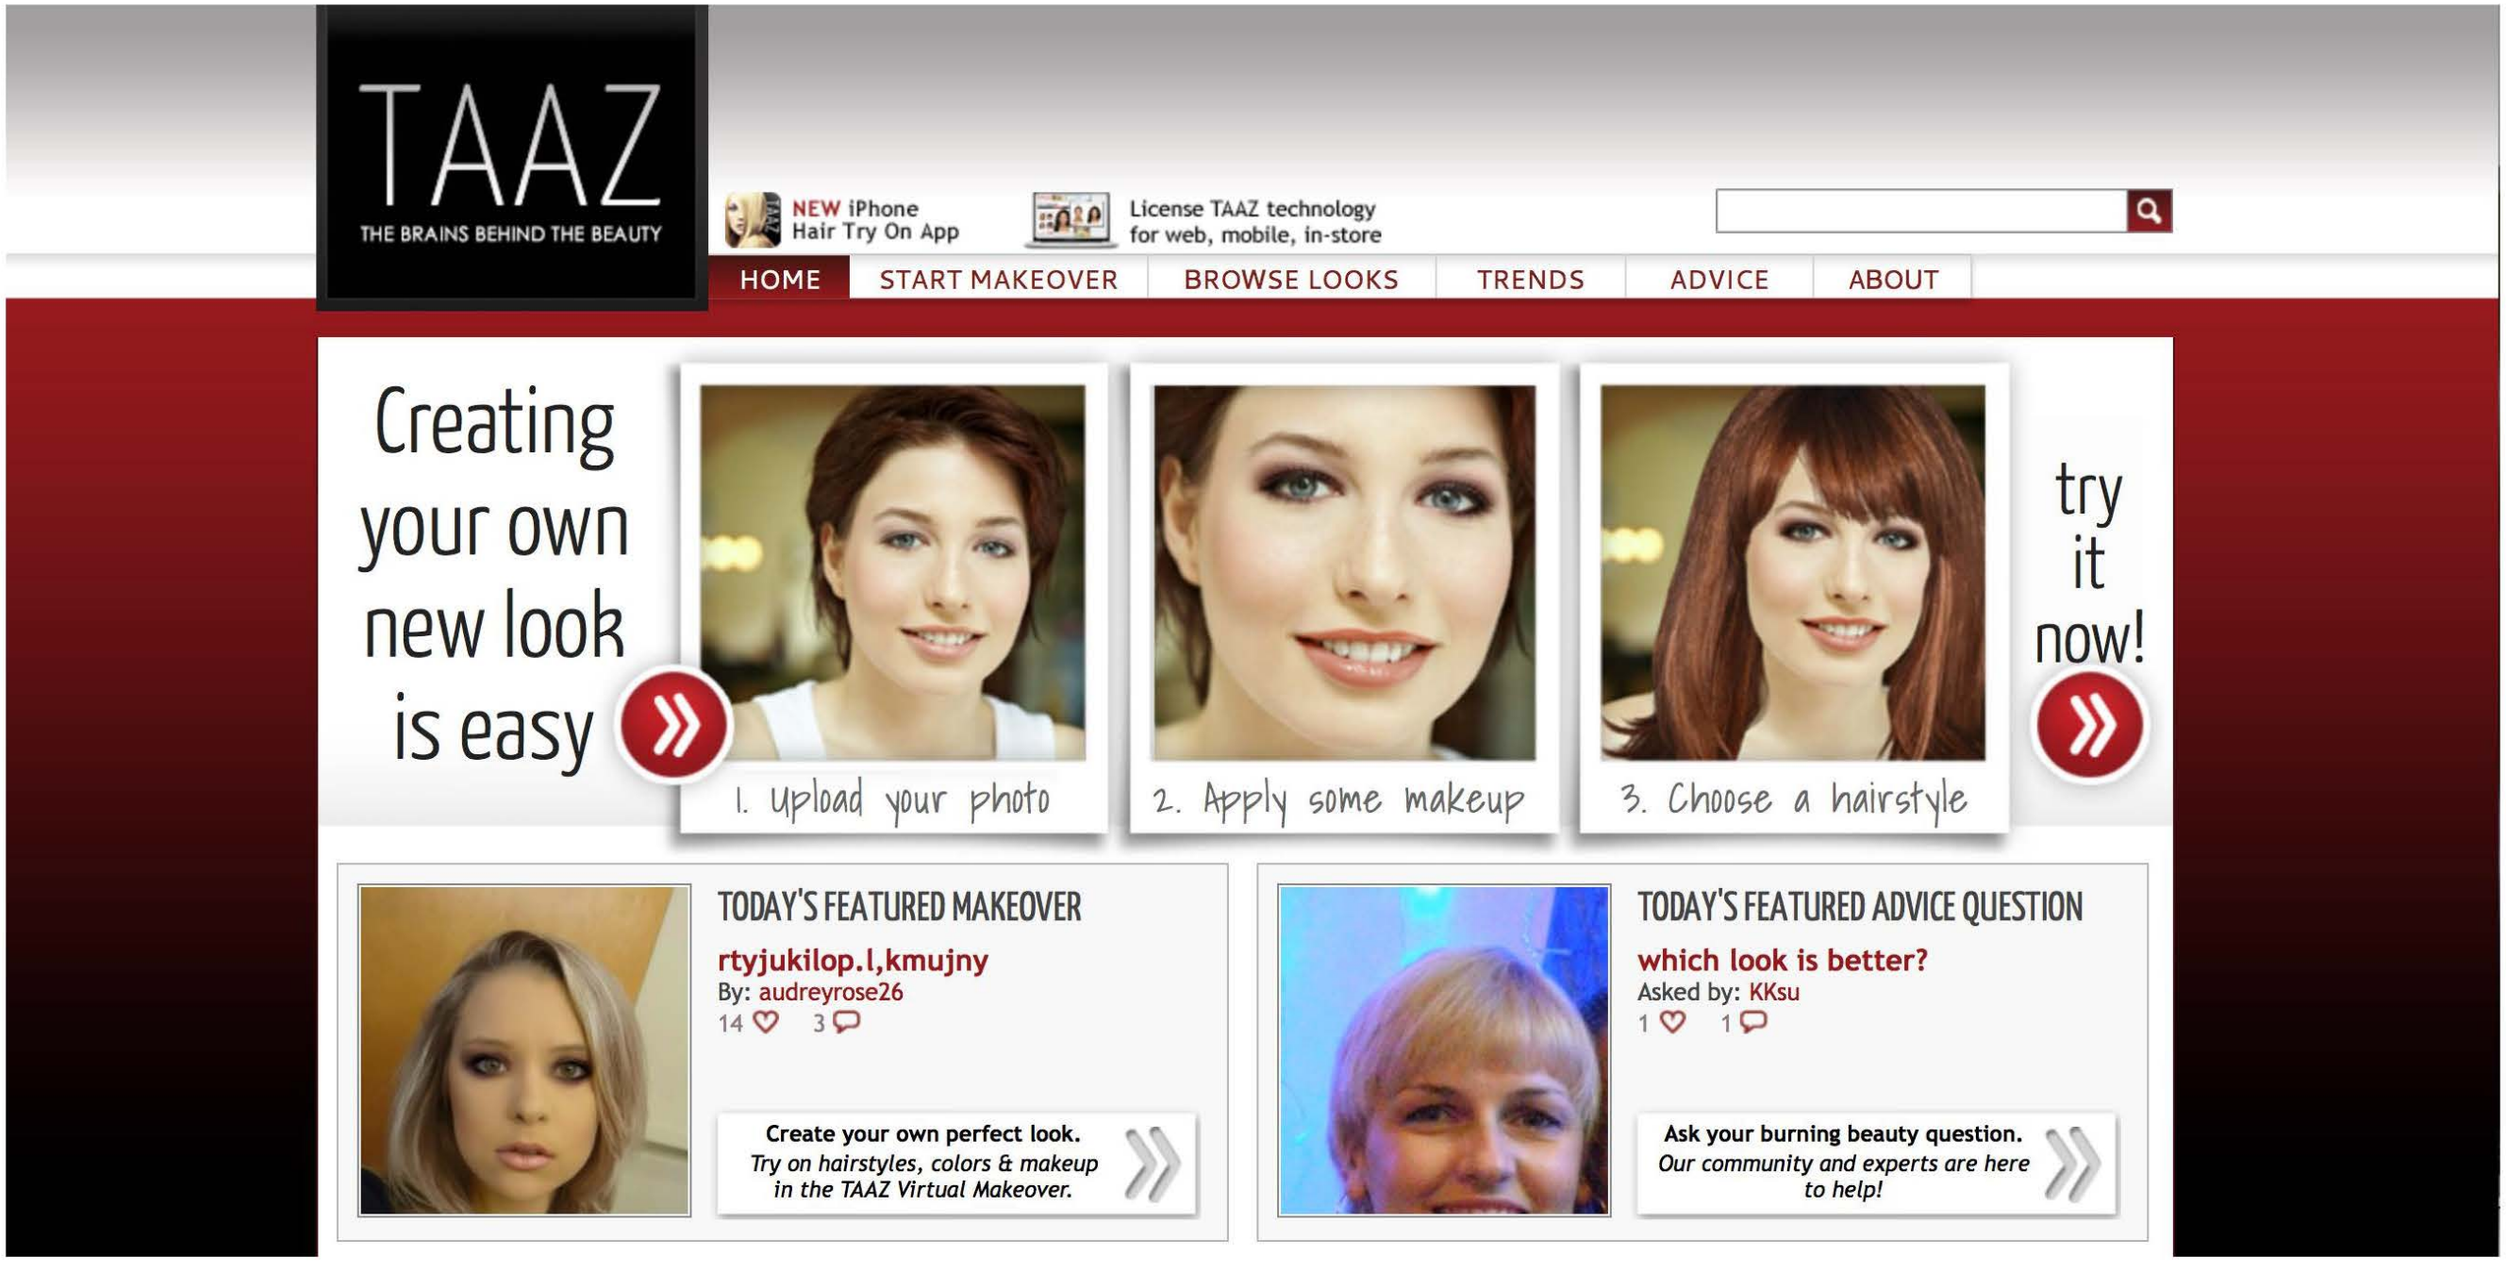

leafsnap


Verizon
12:38 PM
Back
Results
Map
Snap It! Results
1
Red Maple
Acer rubrum
2
Striped Maple
Acer pensylvanicum
3
Sycamore Maple
Acer pseudoplatanus
12:20 PM
100%
First
Last
Scientific
Ilex opaca
American Hornbeam
Carpinus caroliniana
American Linden
Tilia americana
American Sycamore
Platanus occidentalis
Amur Corktree
Phellodendron amurense
Q A B C D E F G H J K K L M N O P Q R S T U W Y
Home
Browse
Collection
Options
Snap It!

Word Lens

Word Lens

www.QuestVisual.com

First-down line

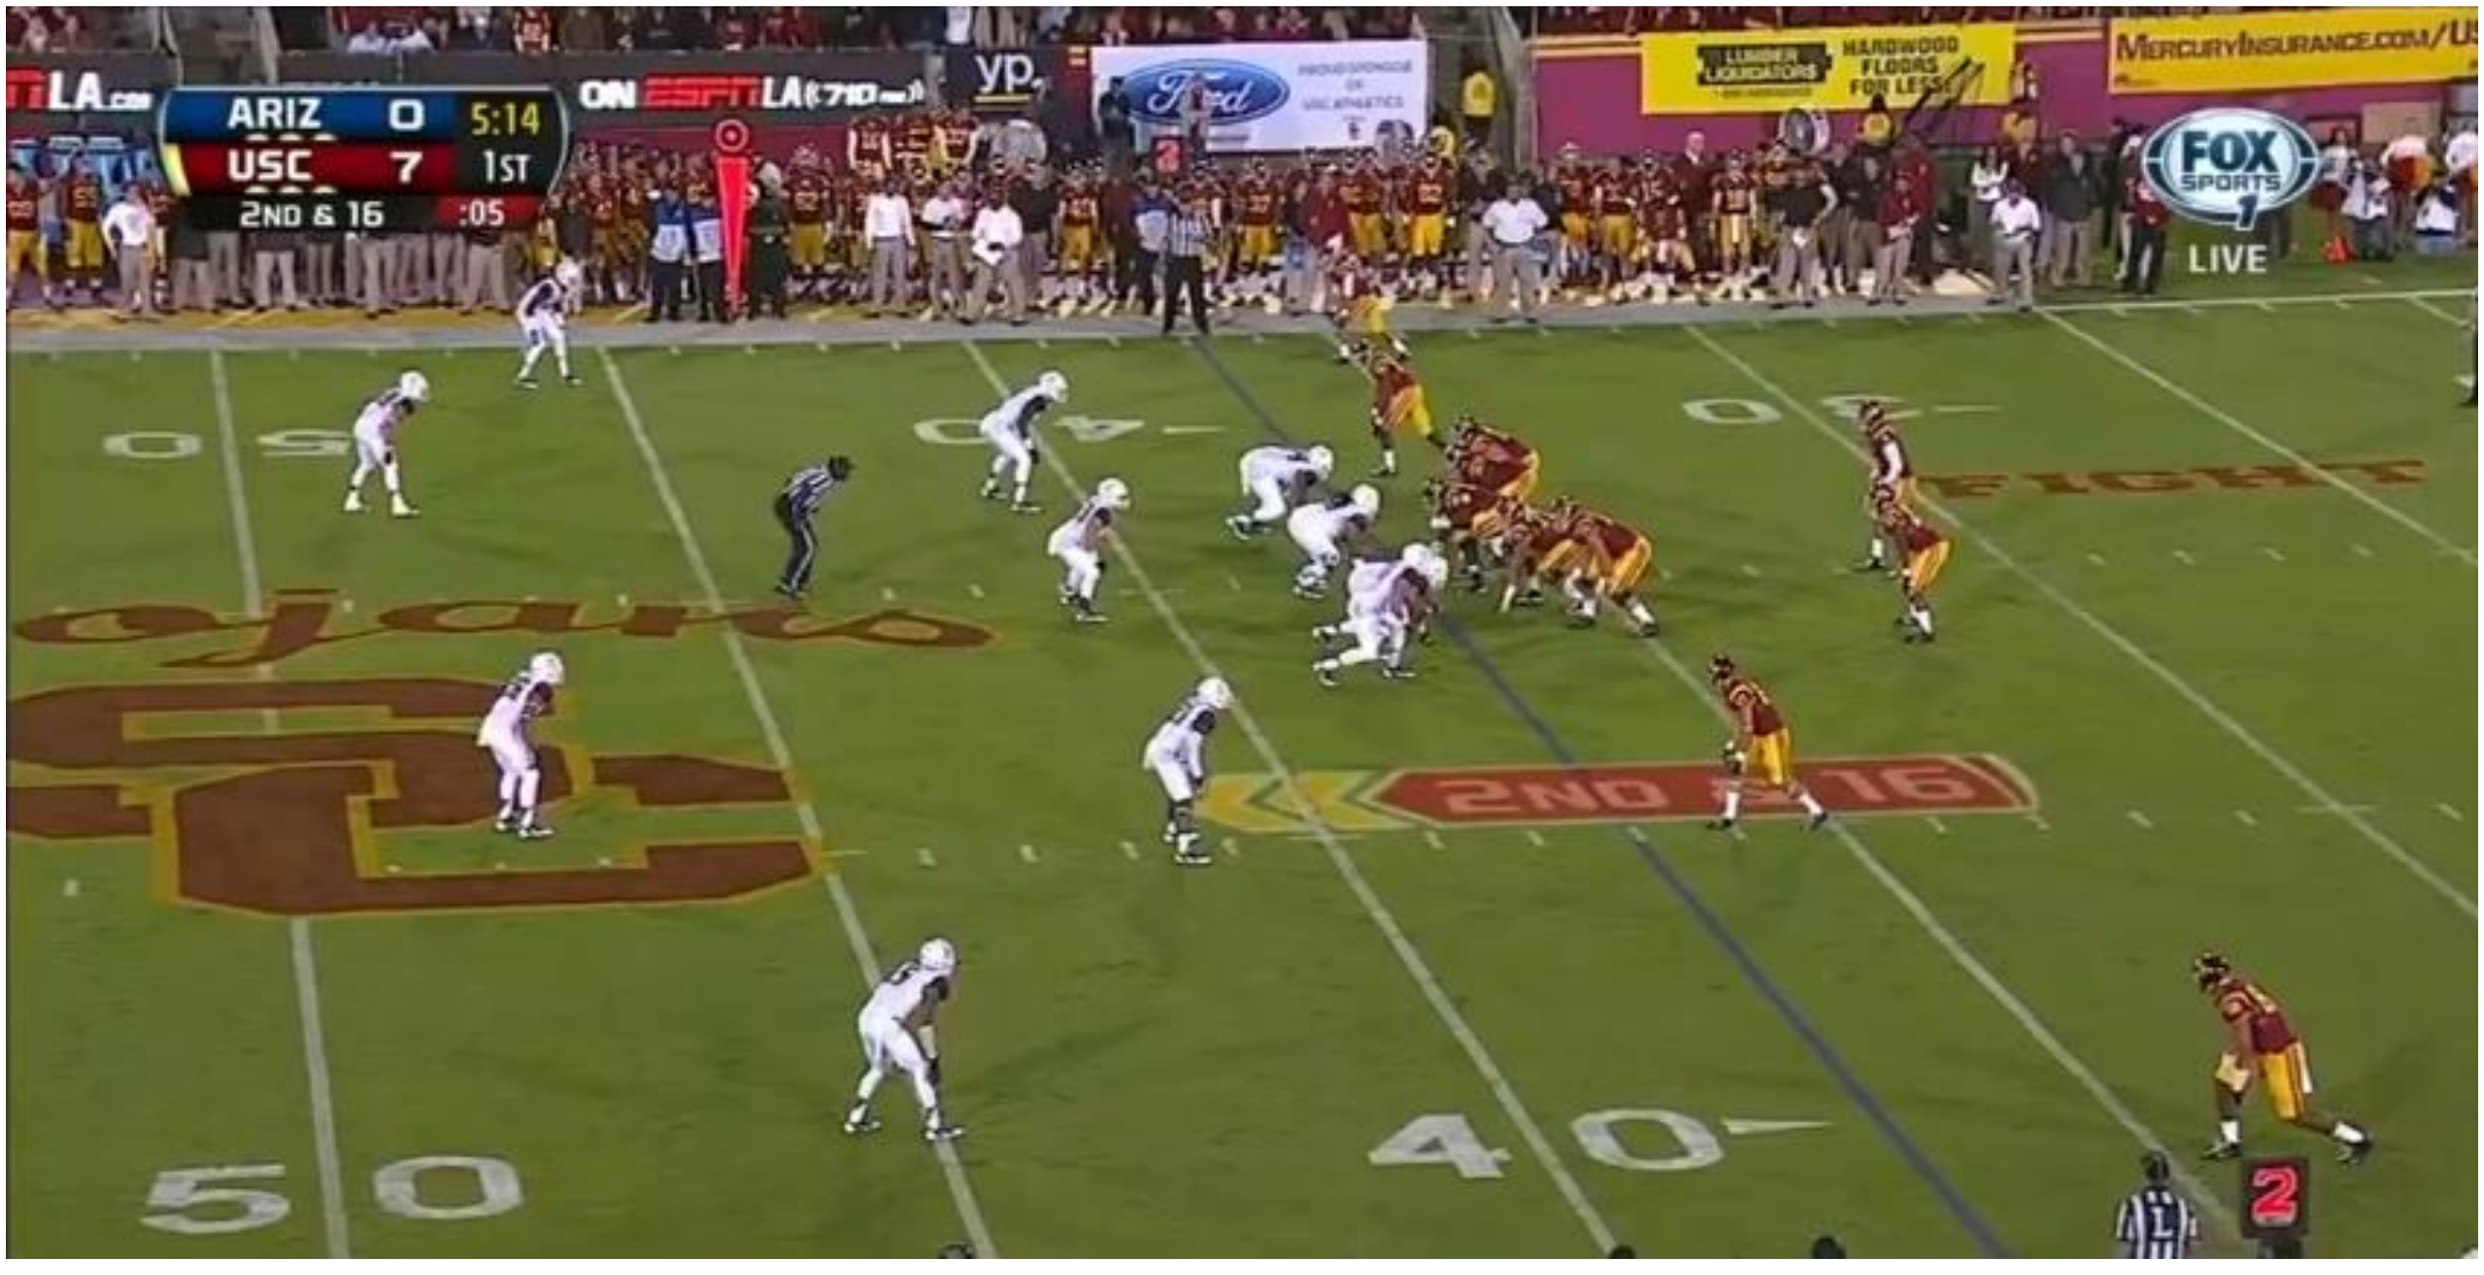

BMW 5 series

BMW night vision

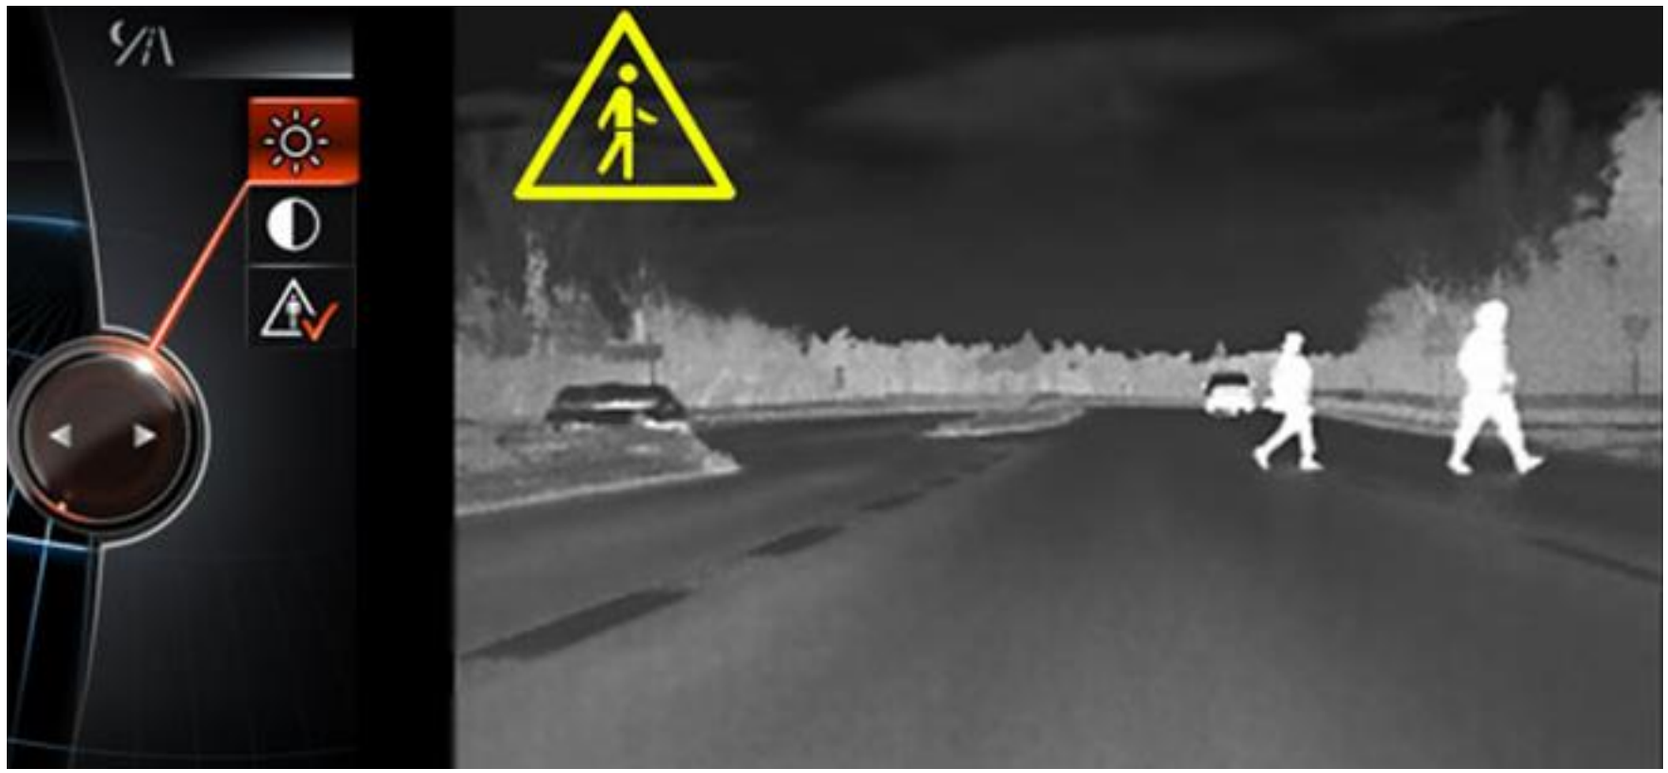

Infinity EX

“Around view” camera

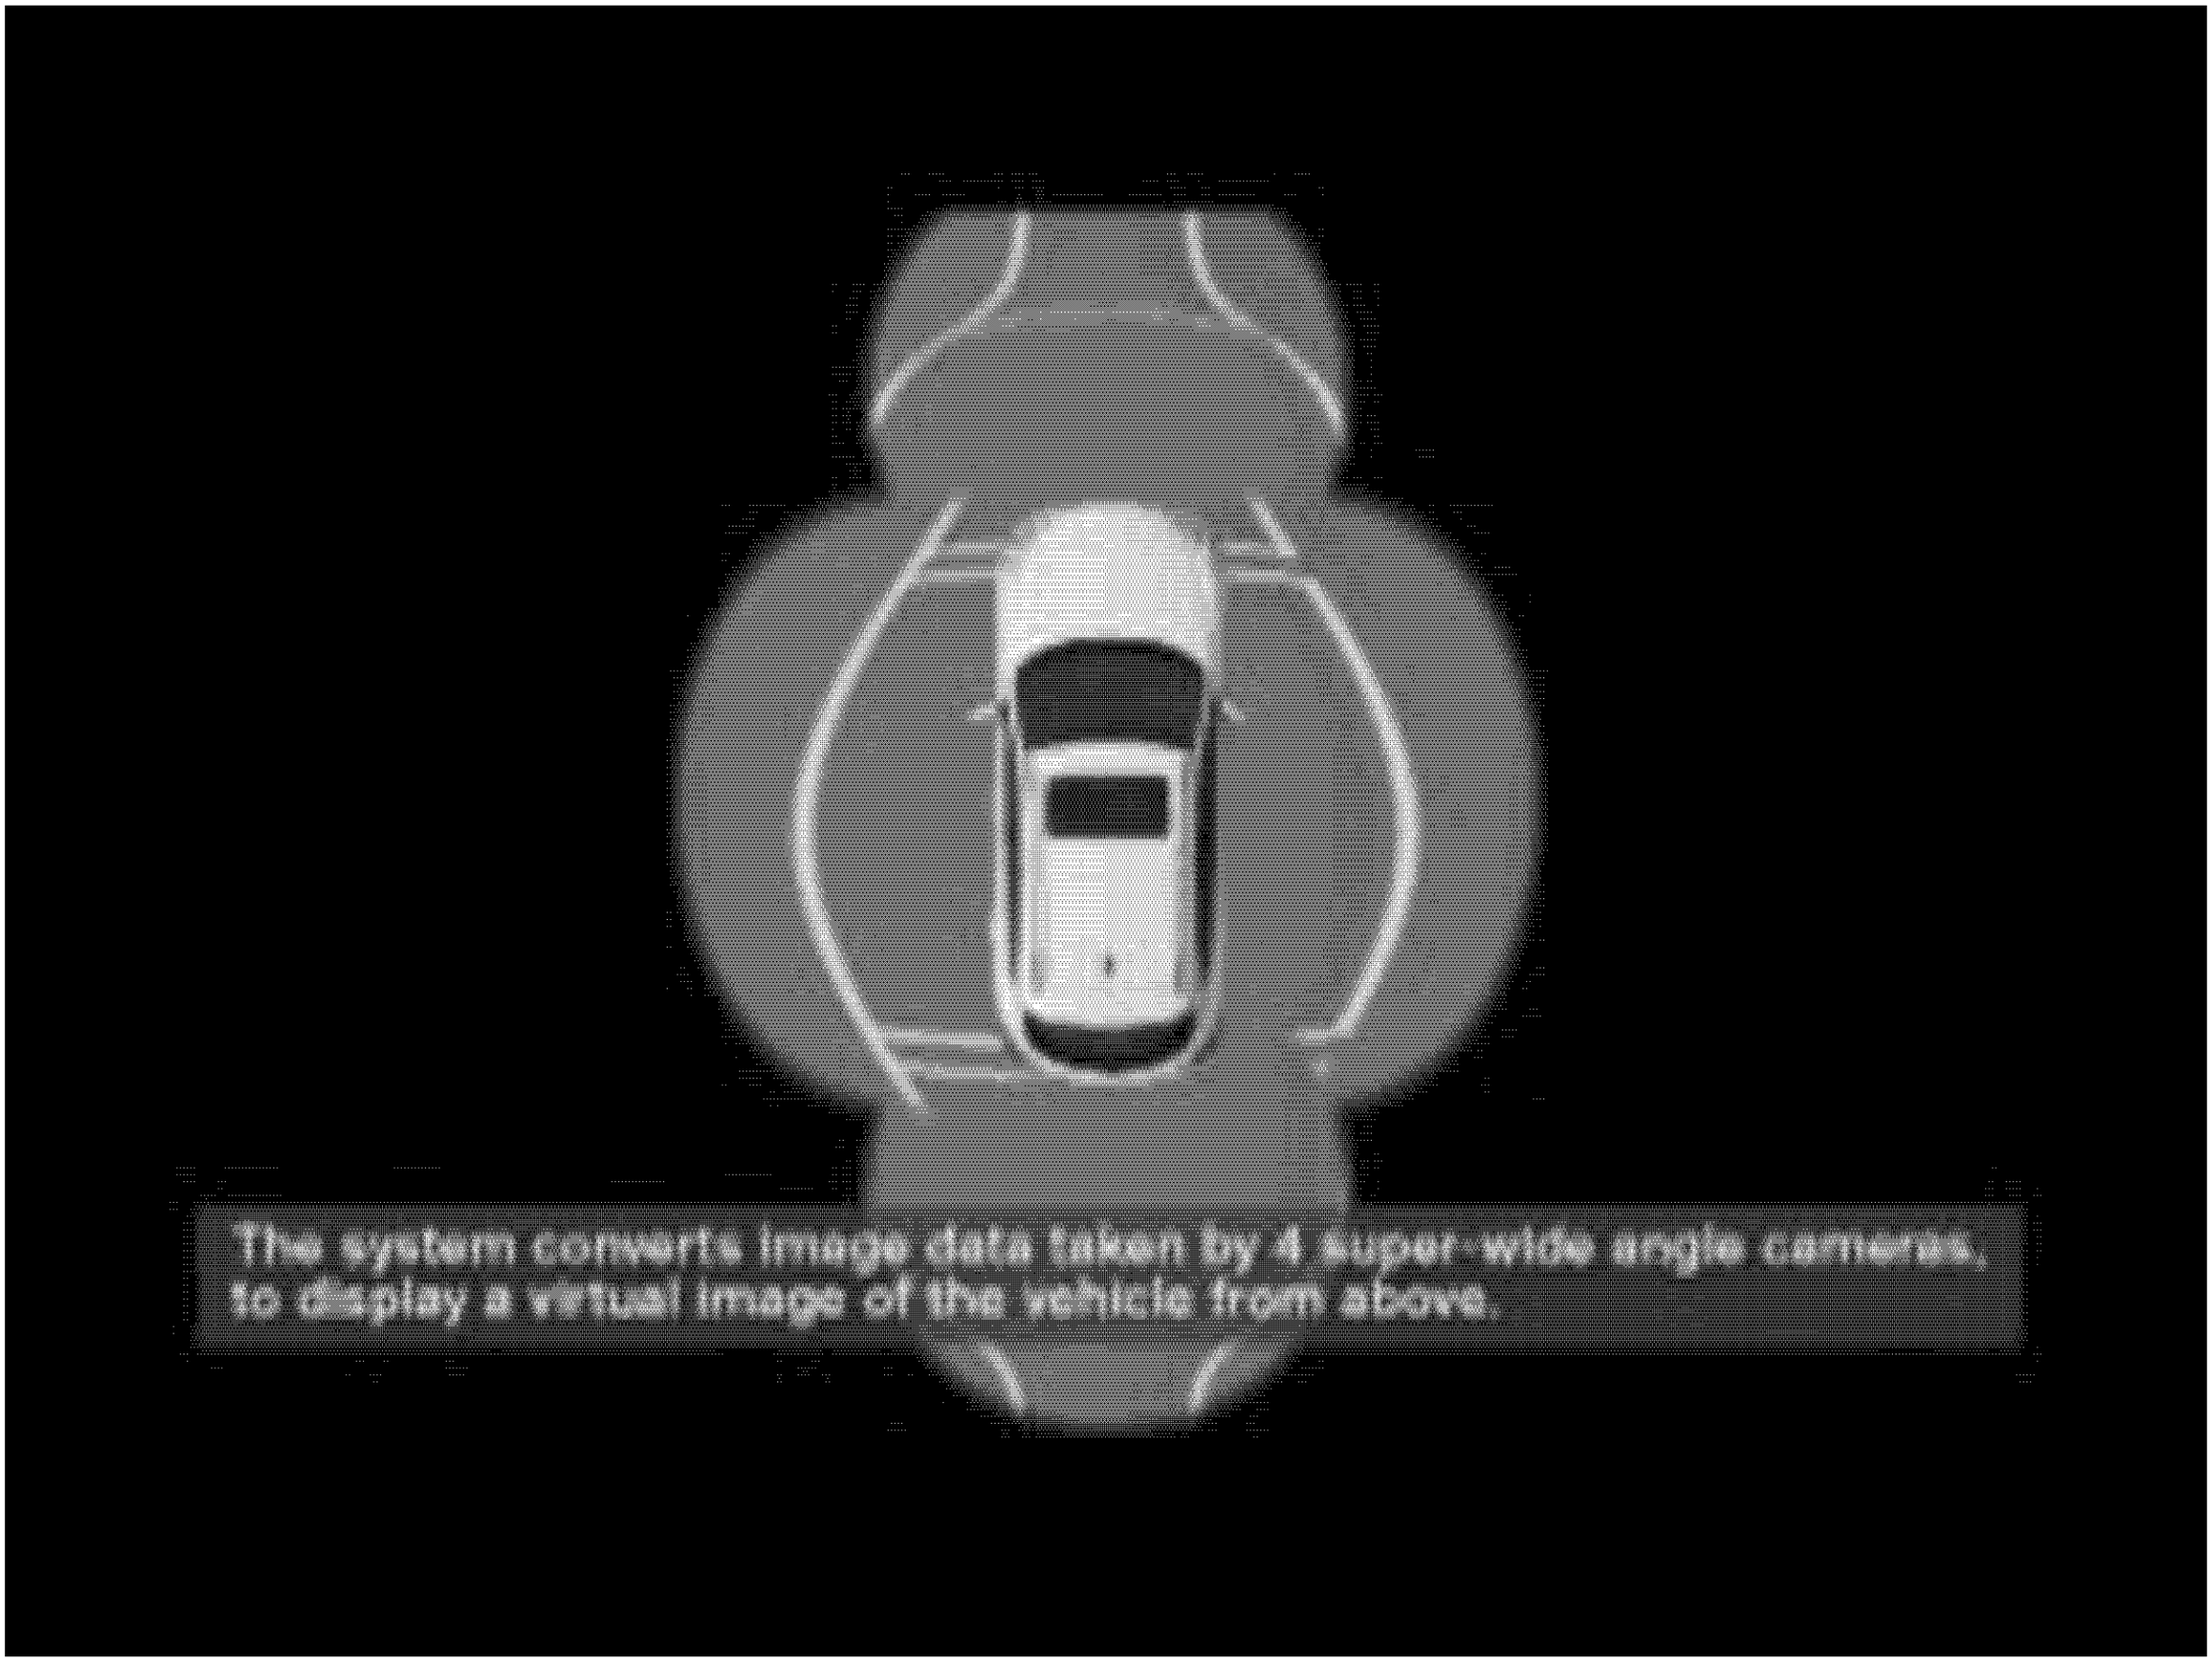
The system converts image data taken by 4 super-wide angle cameras,
to display a virtual image of the vehicle from above.

2015

Vision in Cars

Image stitching

Photosynth

Tango

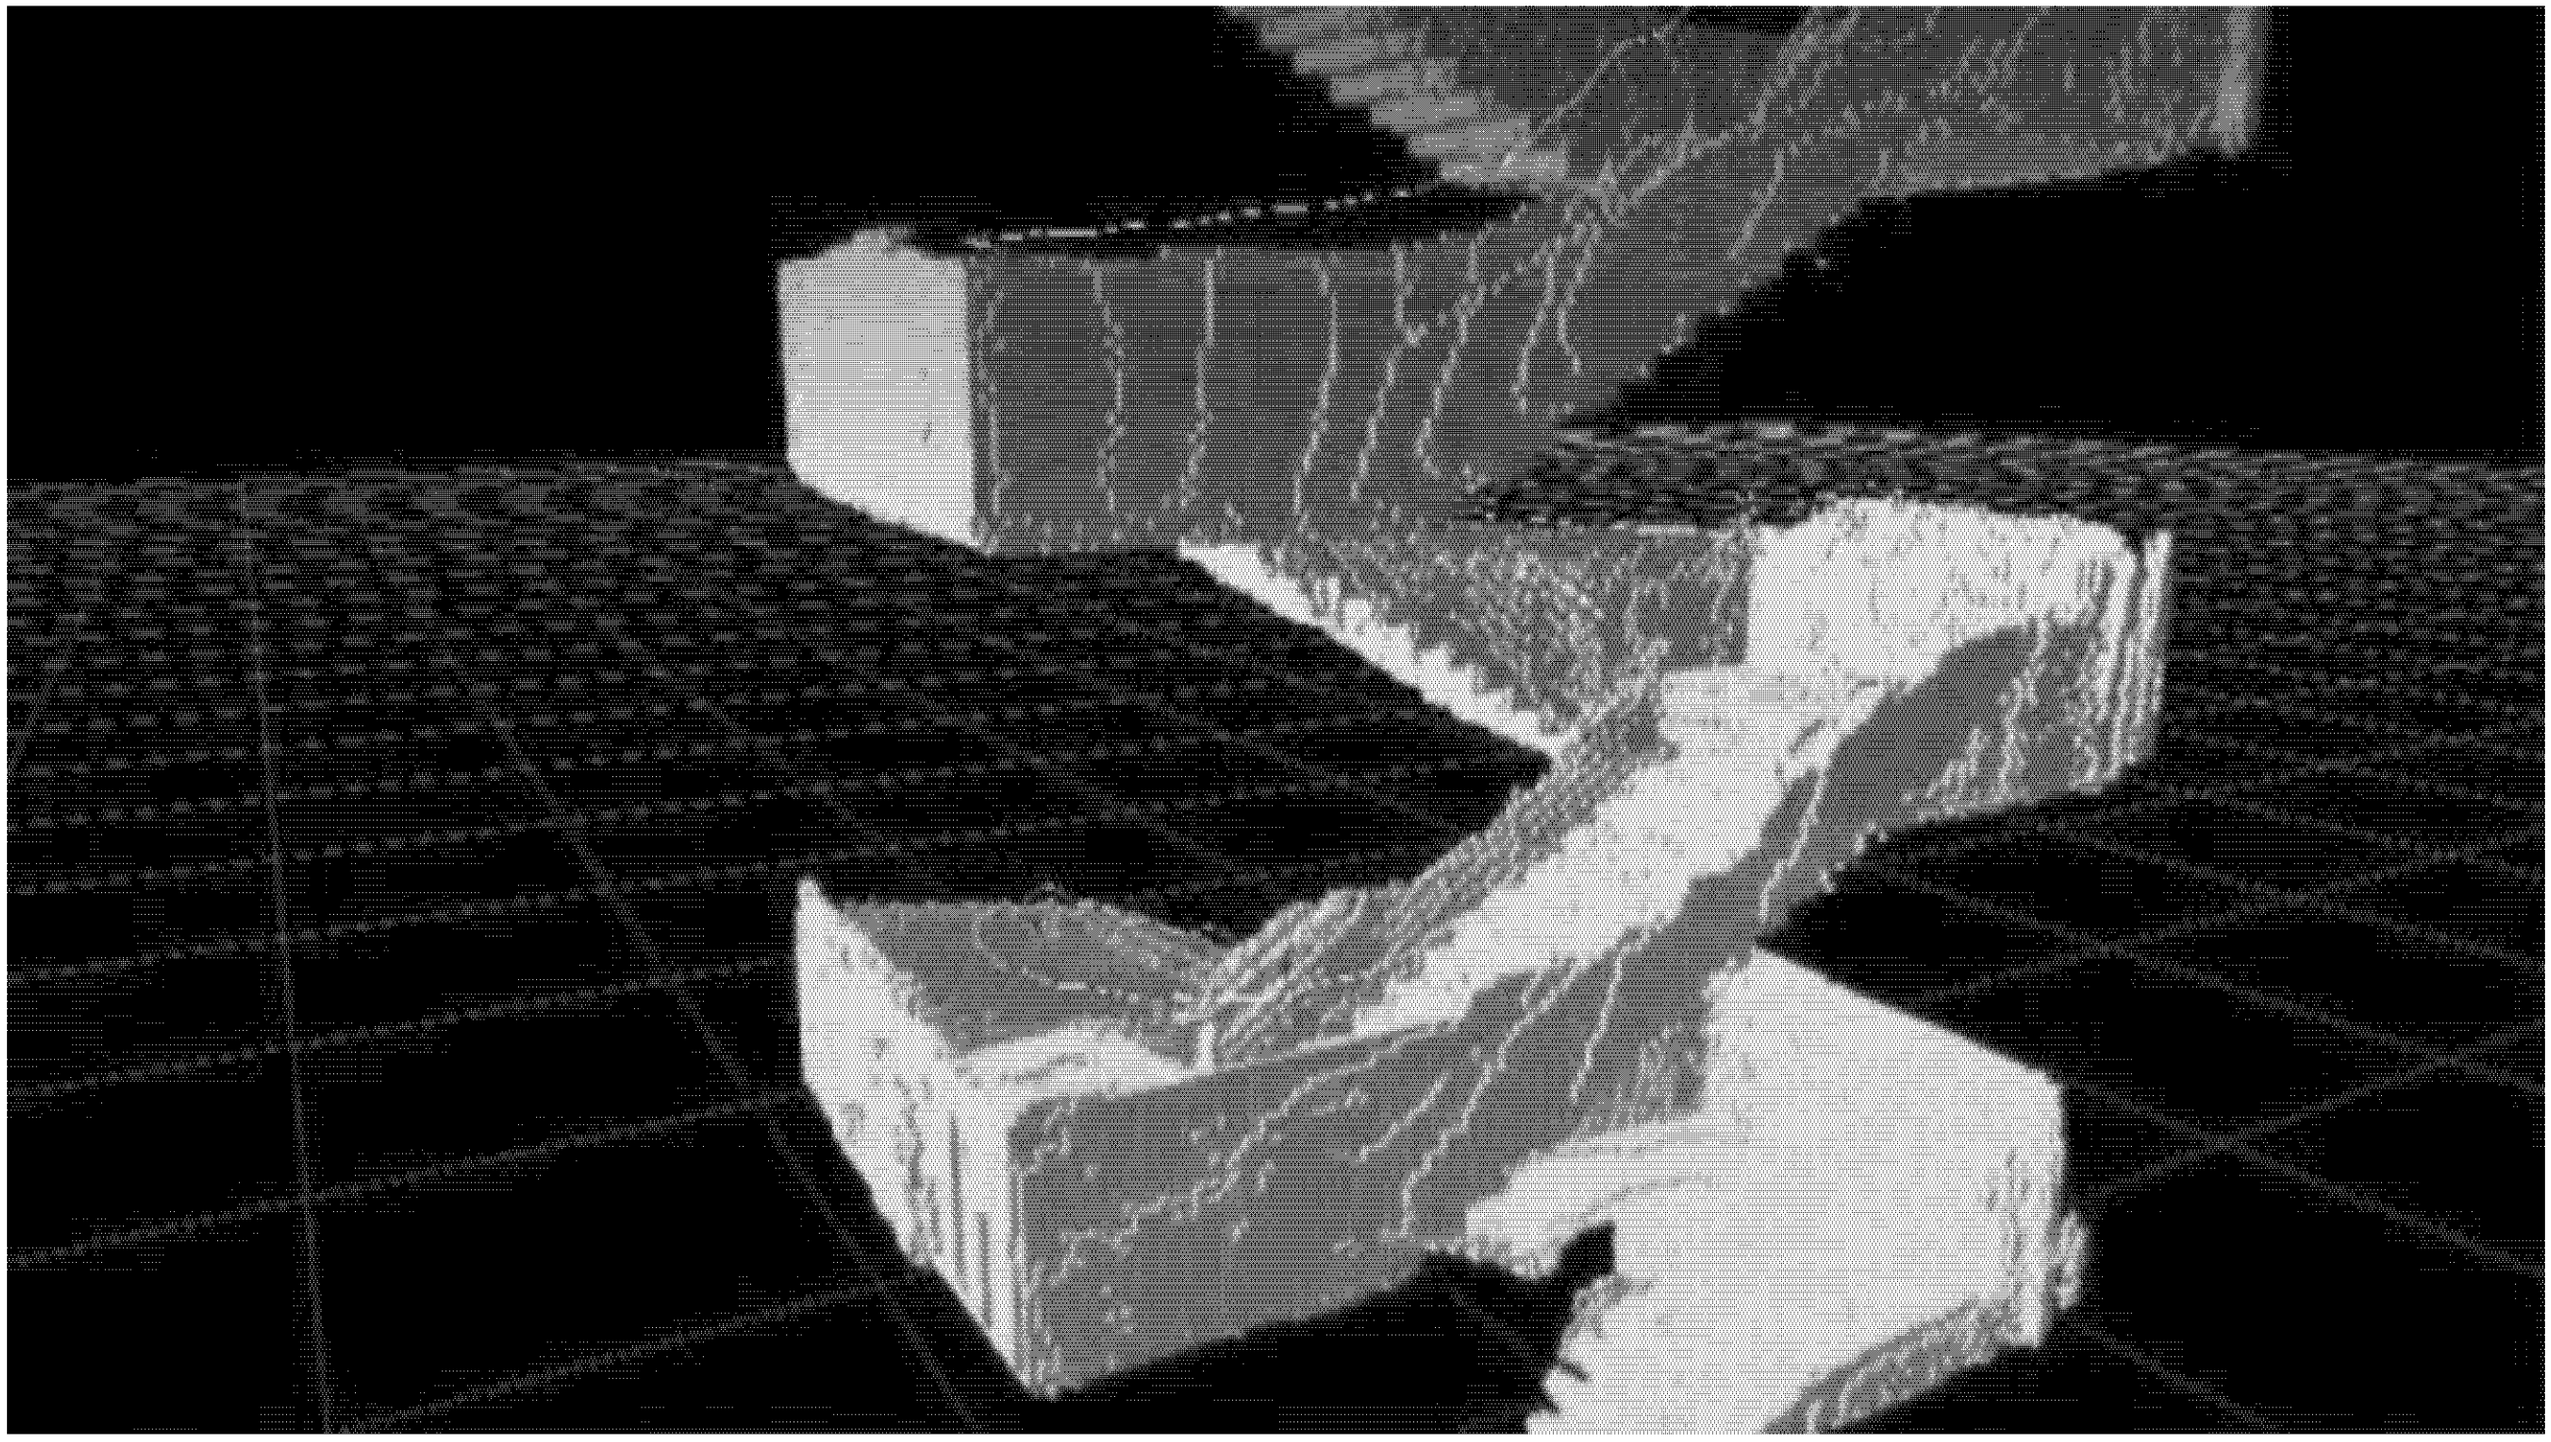

Virtual Fitting

2015

Computer Vision for VR

Deep Face

Deep Dream

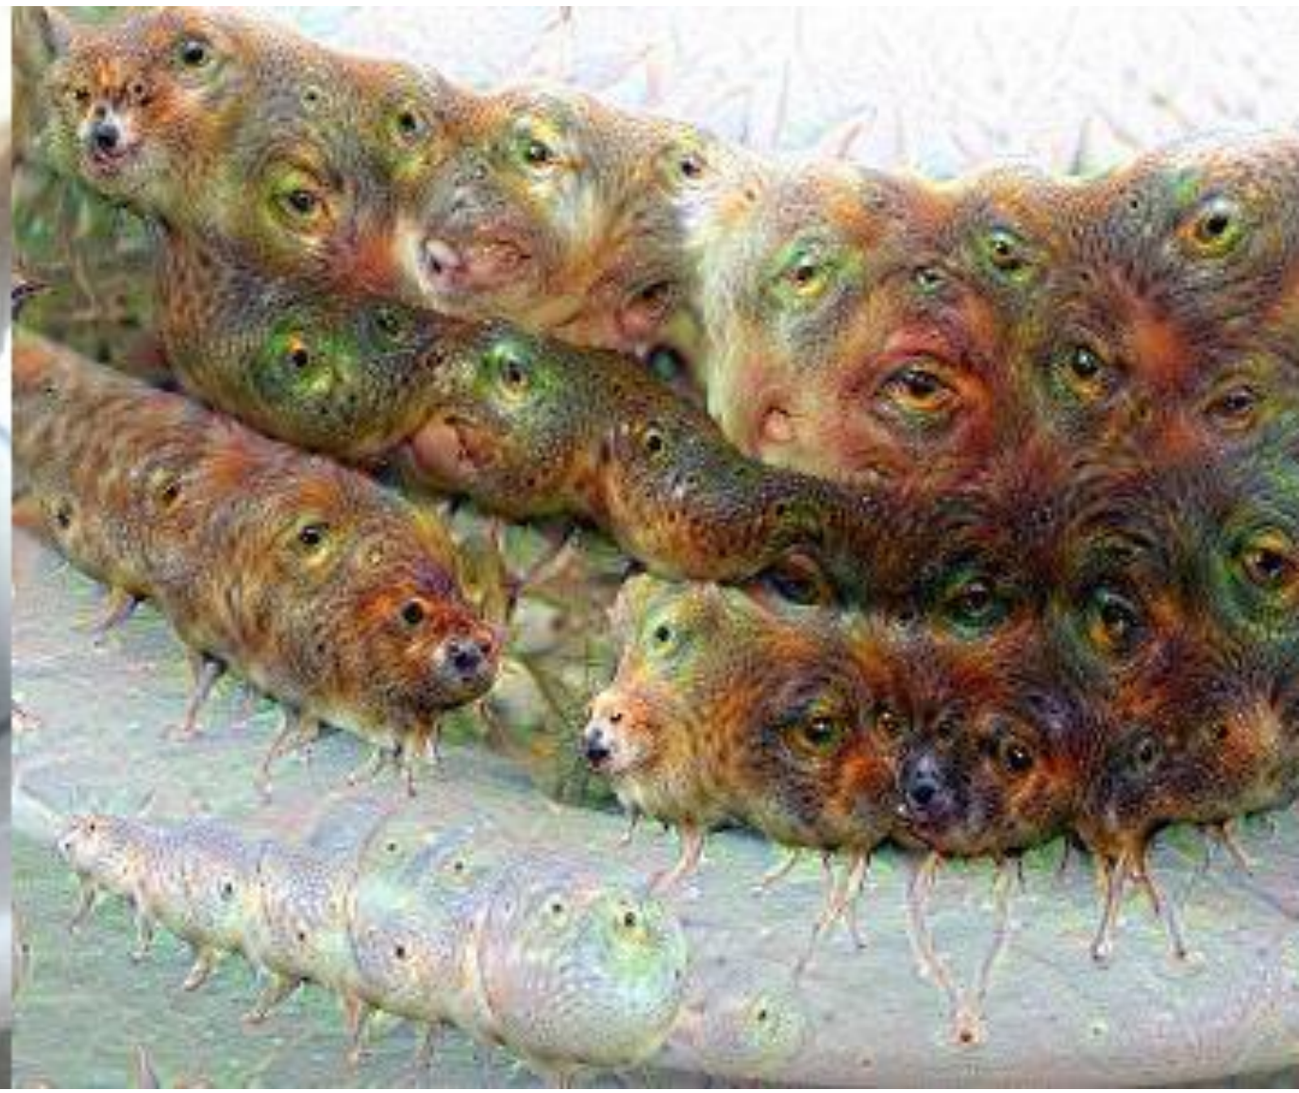

Facebook video style transfer 2016

Face2Face: Real-time Face Capture and Reenactment of RGB Videos

Justus Thies[1], *Michael Zollhöfer*[2], *Marc Stamminger*[1], *Christian Theobalt*[2], *Matthias Nießner*[3]

[1]University of Erlangen-Nuremberg
[2]Max-Planck-Institute for Informatics
[3]Stanford University

CVPR 2016 (Oral)

It’s a good time to do computer vision

Industry aggressively hiring CV faculty from universities

(this slide is already out of date by at least 3 CMU faculty)

NYU

UW

NYU

USC

UW

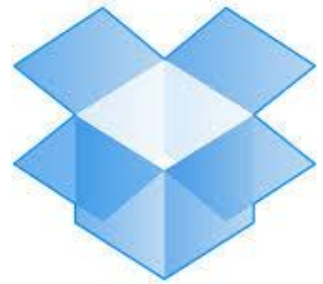

UCSD

Columbia

Industry aggressively hiring CV graduates, or even students!

(~~strong~~ dominant industrial presence at conferences for recruitment)

facebook research | Research Areas | Publications | People | Programs | Downloads | Careers | Blog

JULY 21, 2017

Advancing computer vision technologies at CVPR 2017

By: Facebook Research

Google Research Blog

The latest news from Research at Google

Google at CVPR 2017

Friday, July 21, 2017

CVPR GROWTH

Number of attendees at CVPR

Original slide courtesy of CVPR 2016

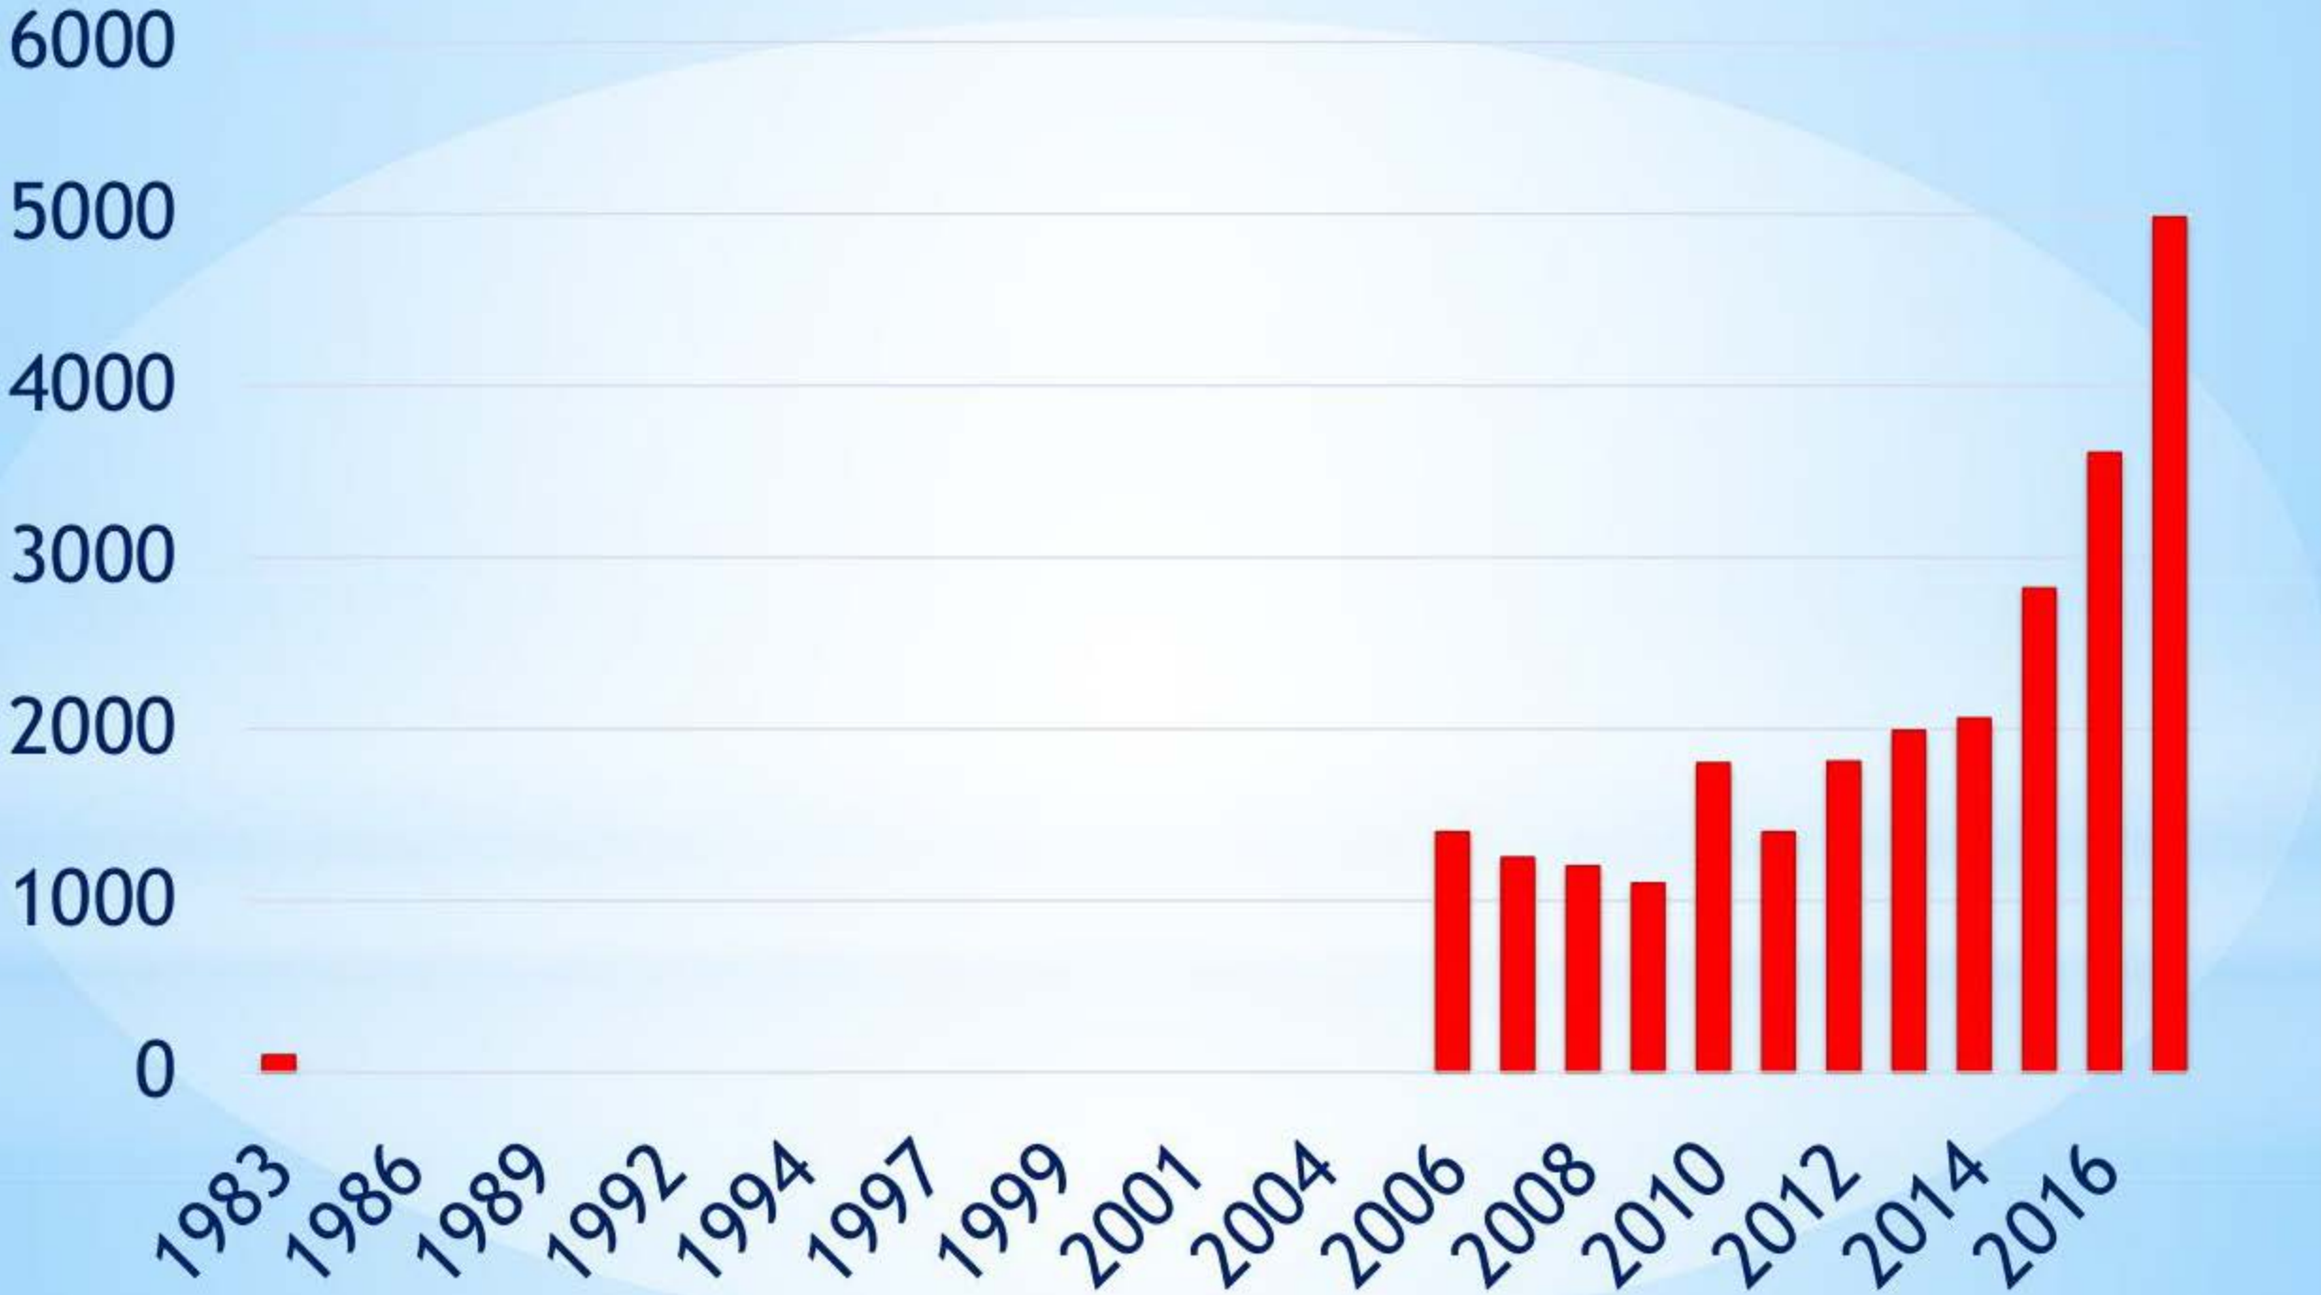

CVPR GROWTH

Original slide courtesy of CVPR 2016

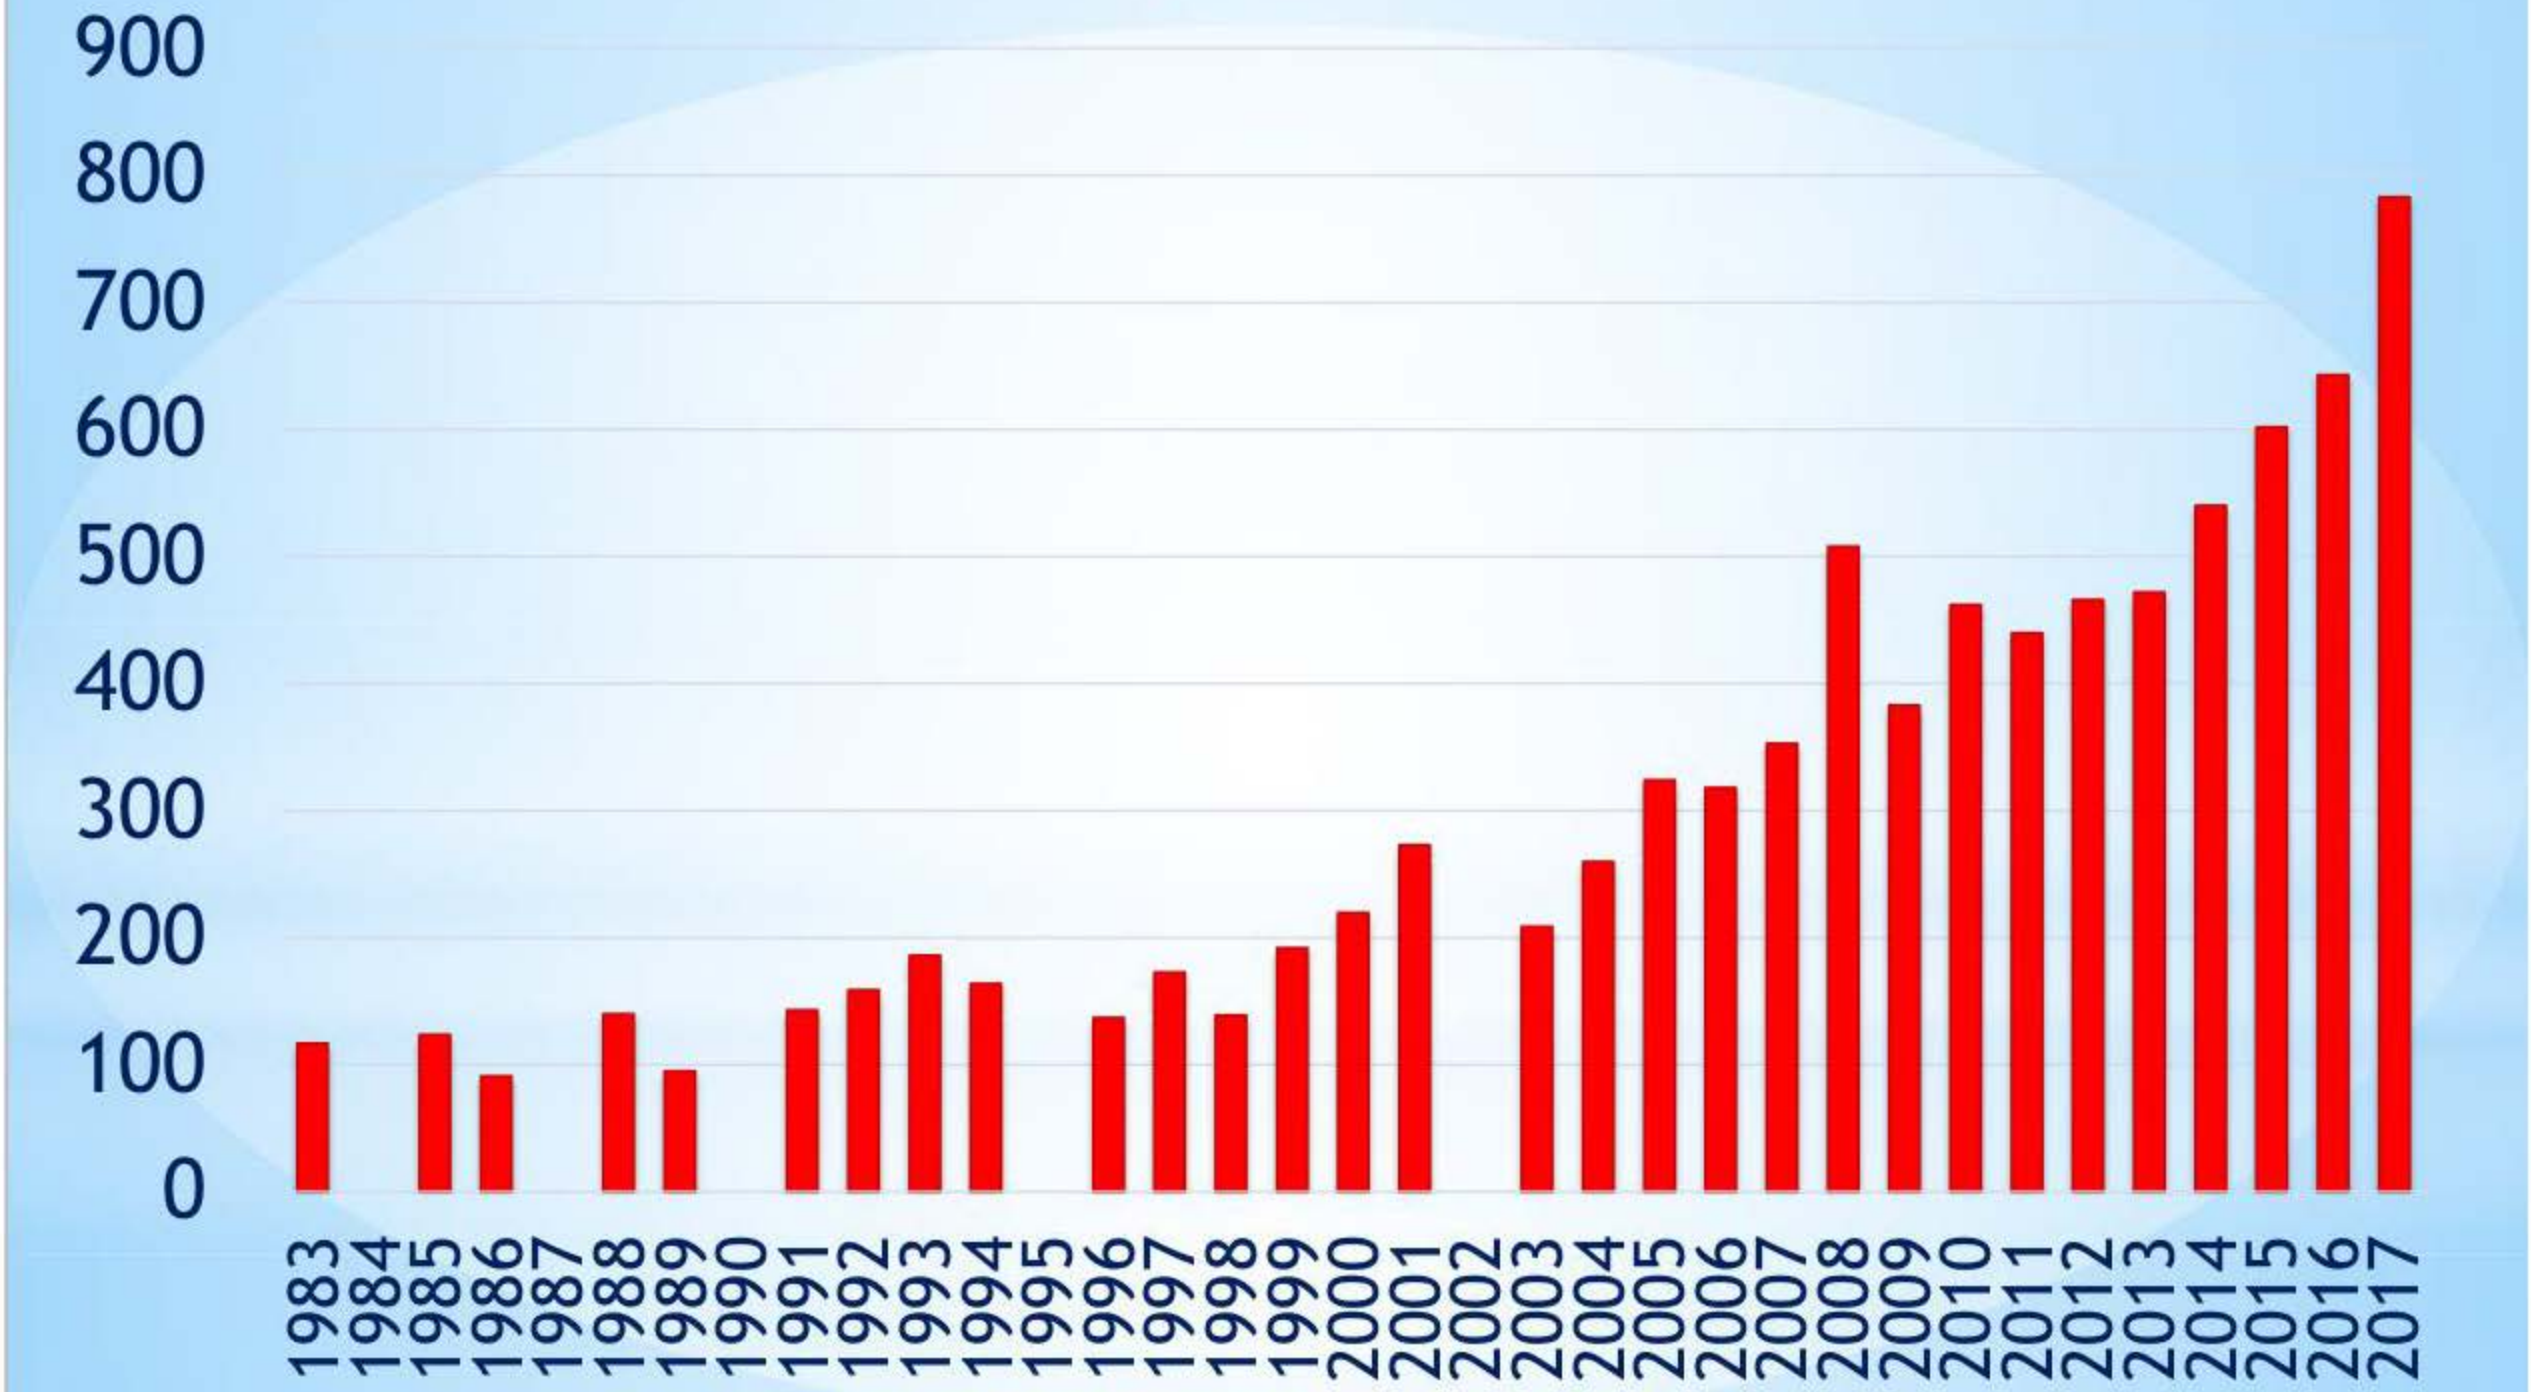

Computer vision at CMU

Dedicated courses for each subject we cover in this class:

- Physics-based Methods in Vision
- Geometry-based Methods in Computer Vision
- Computational Photography
- Visual Learning and Recognition
- Statistical Techniques in Robotics
- Sensors and sensing

… plus an entire department's worth of ML courses.

ICCV 2017: CMU was the second most common academic affiliation among authors

(can you guess the first?)

Master in Computer Vision at CMU

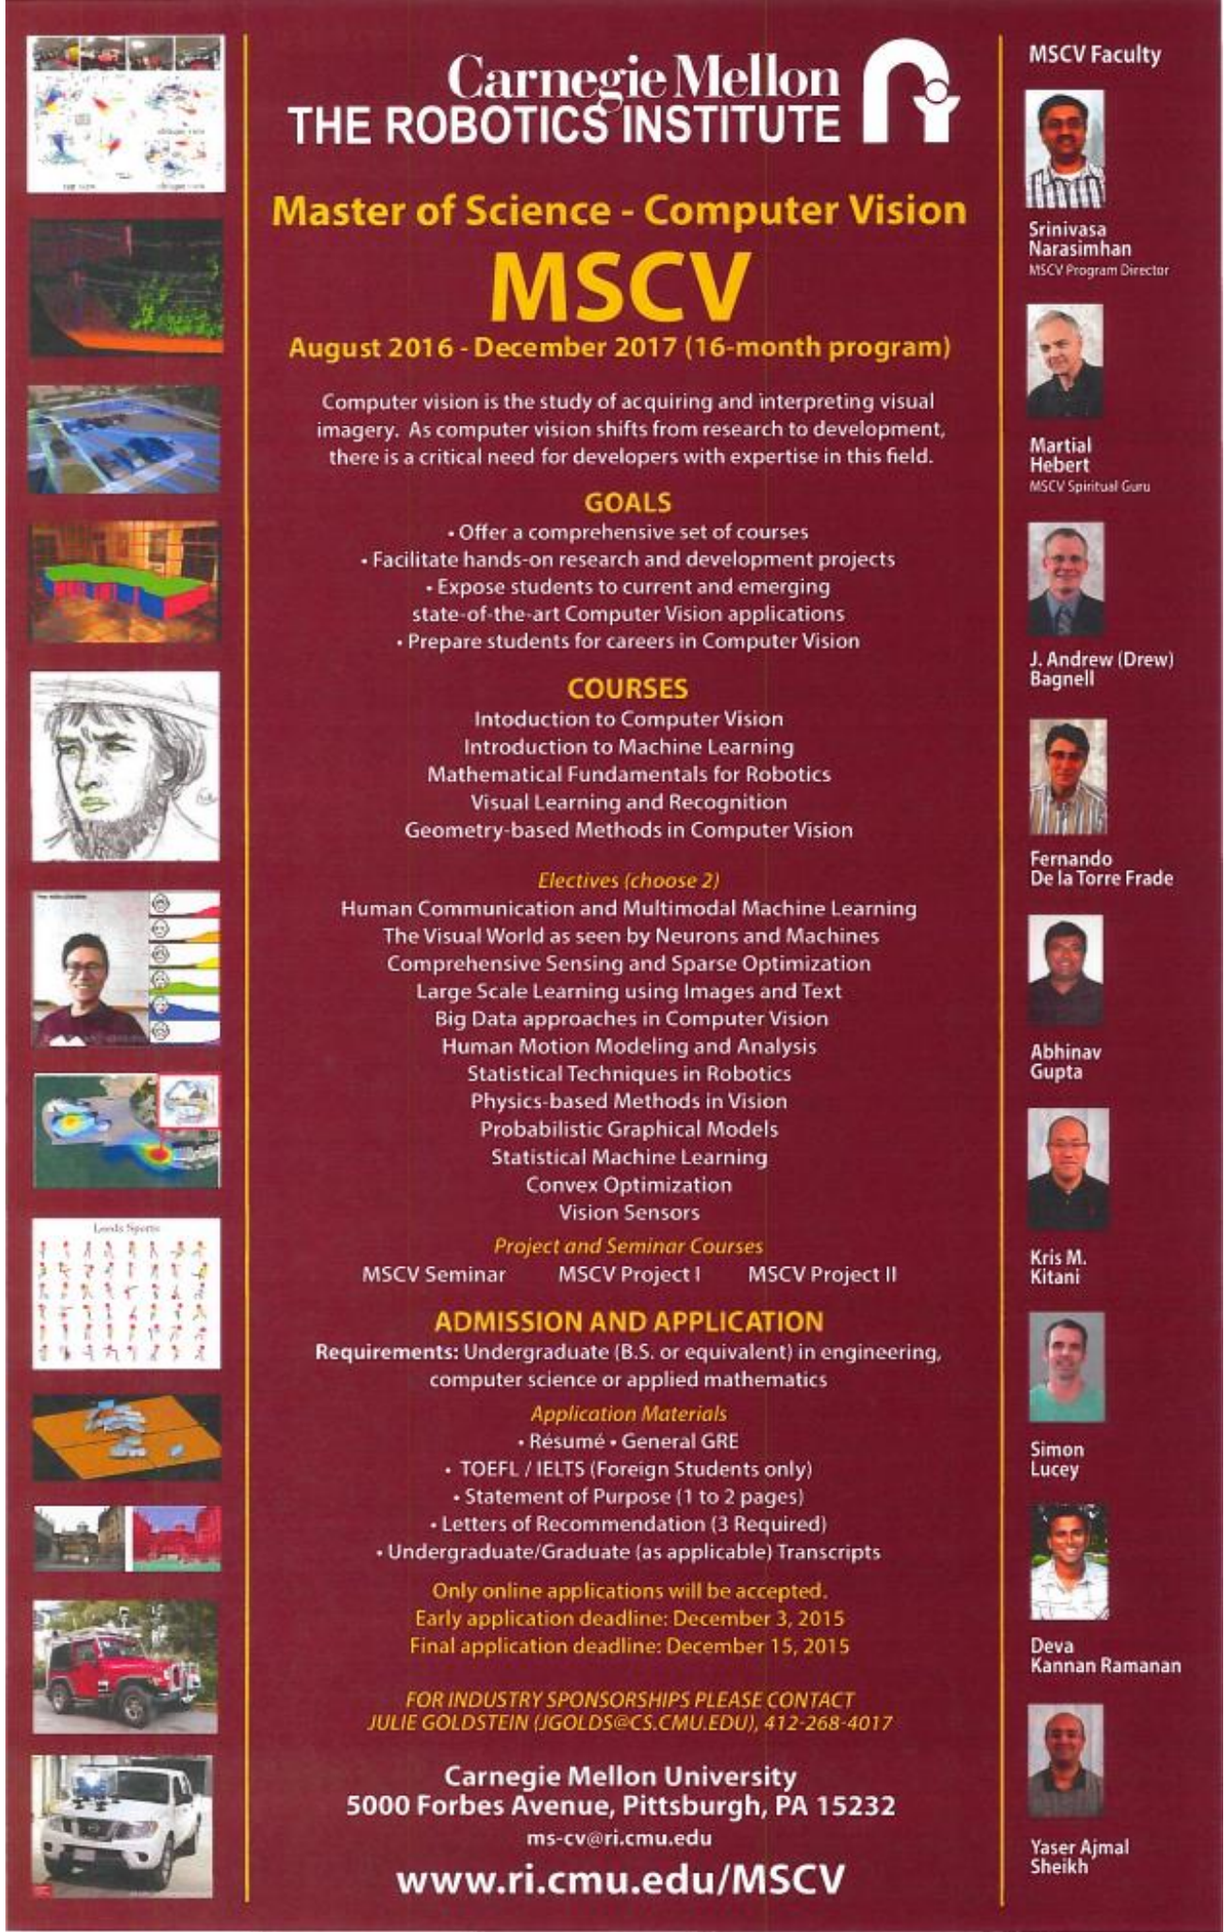

Course logistics

Website

http://www.cs.cmu.edu/~16385/

Assignments Canvas

https://canvas.cmu.edu/courses/3561

Discussion&notes piazza

https://piazza.com/class/jbfezjo9yao2jj

(you should sign up here on your own)

Topics to be covered

Image processing:

- Basics of filtering.
- Image pyramids.
- Gradients and lines.
- Hough transforms.

Topics to be covered

Feature detection and correspondences:

- Corner detection.
- SIFT et al.
- Feature descriptors.
- RANSAC.

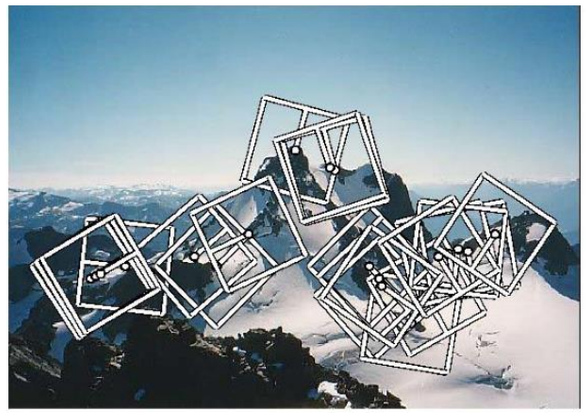
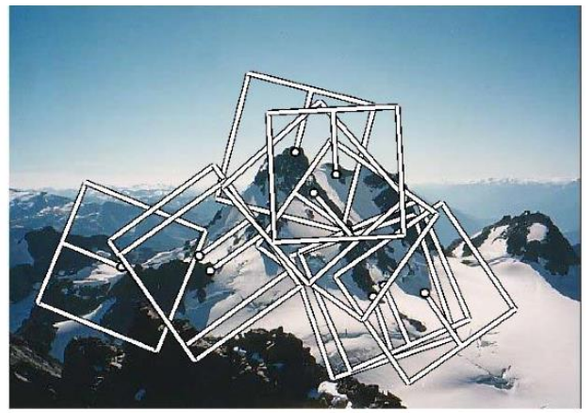
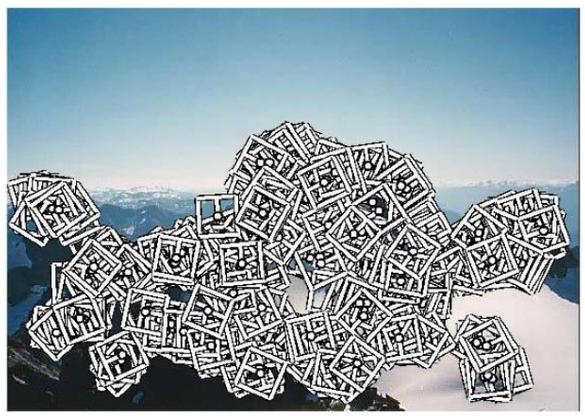
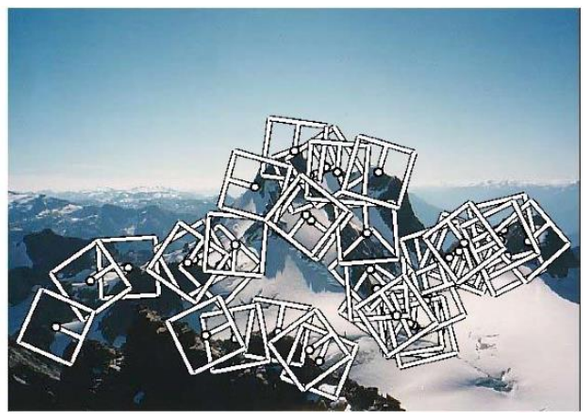

Topics to be covered

Transformations and geometry:

- Homographies and image alignment.
- Camera models.
- Fundamental matrix.
- Epipolar geometry and stereo.
- Structure from motion.

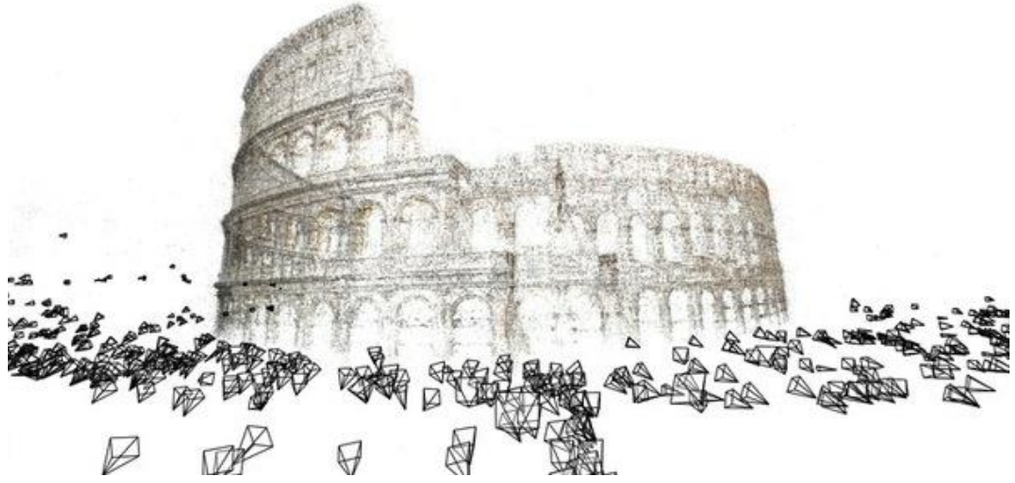

Topics to be covered

Physics-based vision:

- Reflectance and image formation.
- Radiometry.
- Shape from shading.
- Photometric stereo.
- Color.

Topics to be covered

Objects, faces, and learning:

- Basics of probability.
- K-means, KNN, PCA, SVM.
- Bag of words.
- Viola-Jones face detection.
- Perceptron, backpropagation.
- Convolutional neural networks.

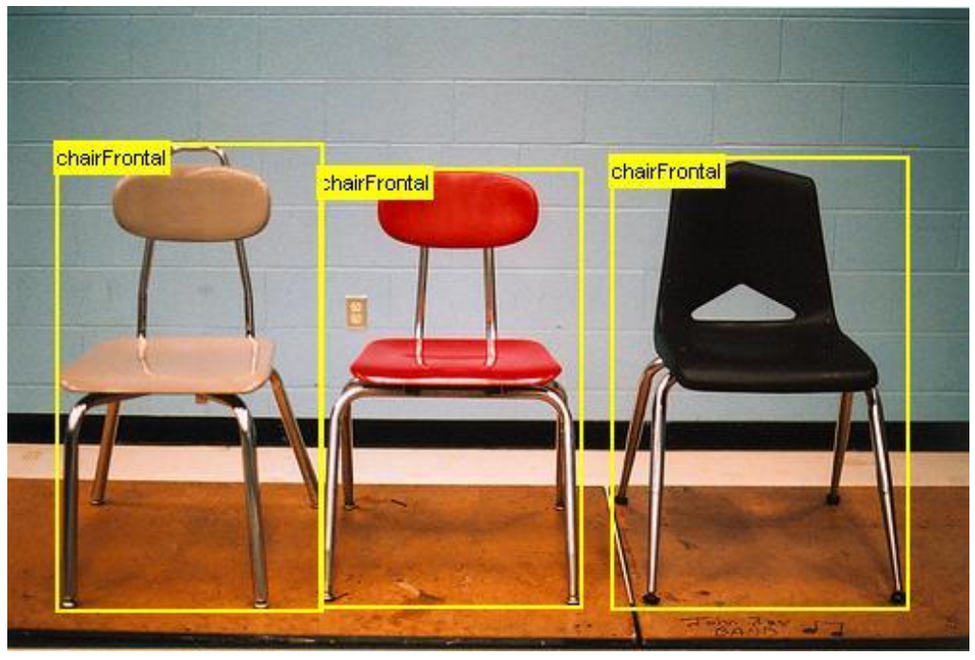

Topics to be covered

Dealing with motion:

- Optical flow (LK, HS).
- Image registration.
- Kalman Filtering.
- Tracking (KLT, Mean-Shift).

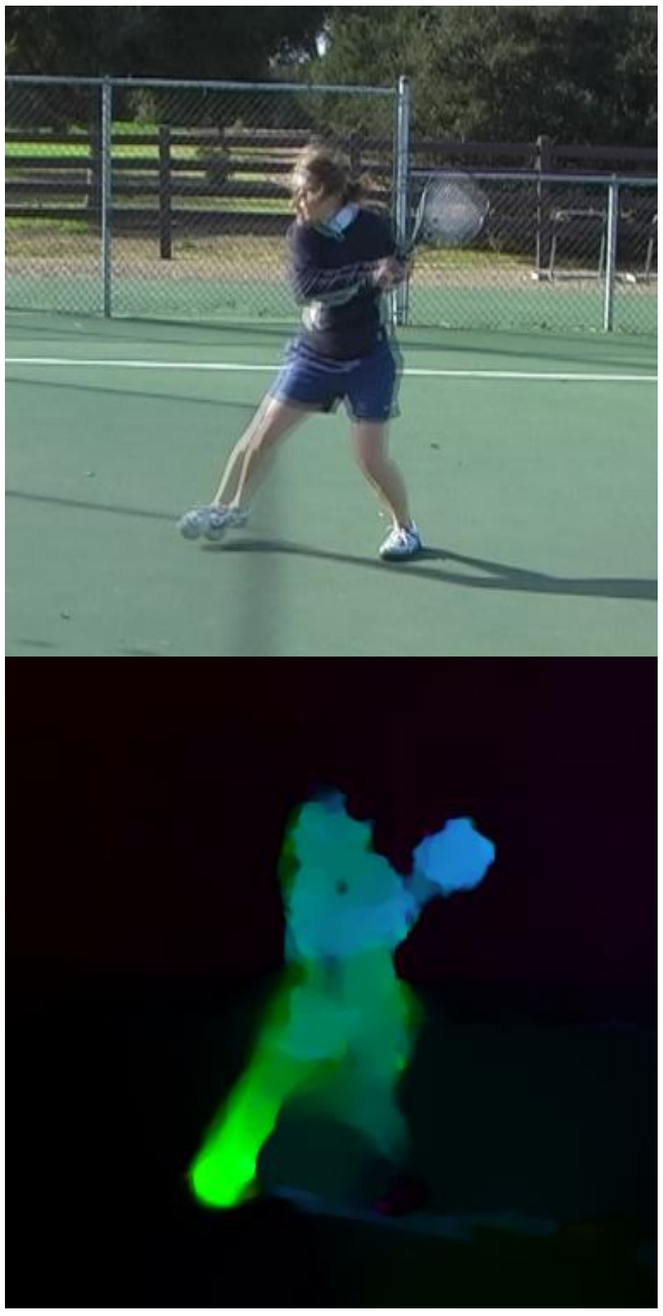

Topics to be covered

Special topics:

- Computational photography.
- ???

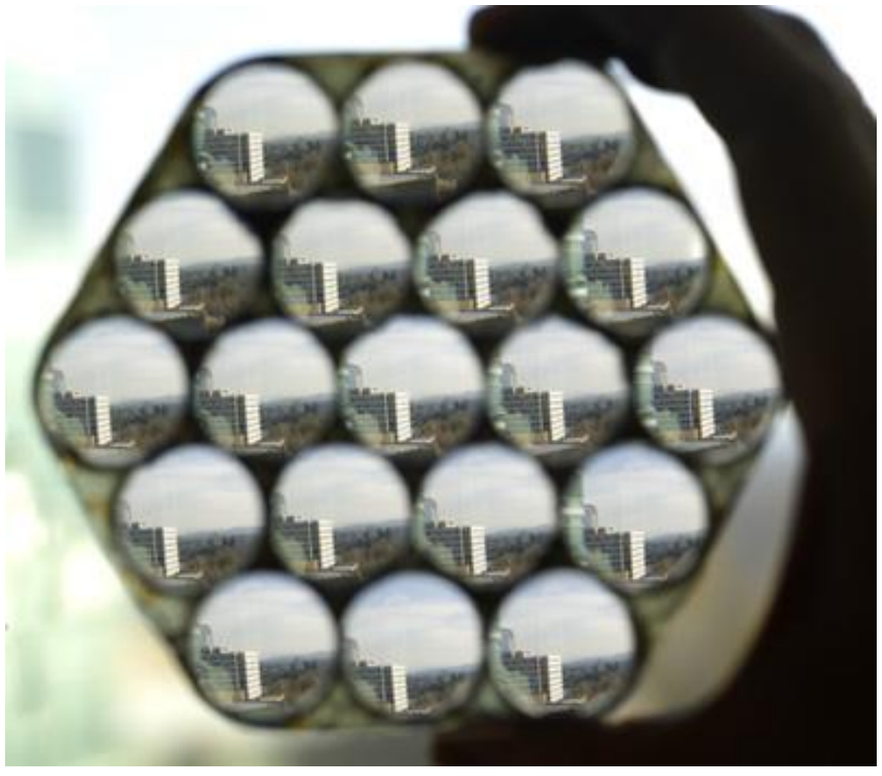

Grading: Project-based

- Seven two-week projects: 95%
- Class and Piazza participation: 5%

Projects:

- a lot of programming in Matlab.
- hours and hours of programming.
- days and days of debugging.

Participation:

- Be around for lectures.
- Post on Piazza discussions.
- Ask and answer questions.

Tentative project schedule

Projects
Project 0 Matlab
Project 1 Hough Transform
Project 2 Homography
Project 3 Stereo
Project 4 Photometric Stereo
Project 5 Bag of Words
Project 6 Convolutional Neural Nets
Project 7 Lucas-Kanade Tracking

- Generous grading policy (like grad school)
- Getting an A vs. mastering the material
- Build your CV
- Take advantage of extra credit

Late days

- 10% reduction of points per late day
- 3 free late days total (not per project)
- use them wisely…

Book

We will be posting readings after each lecture

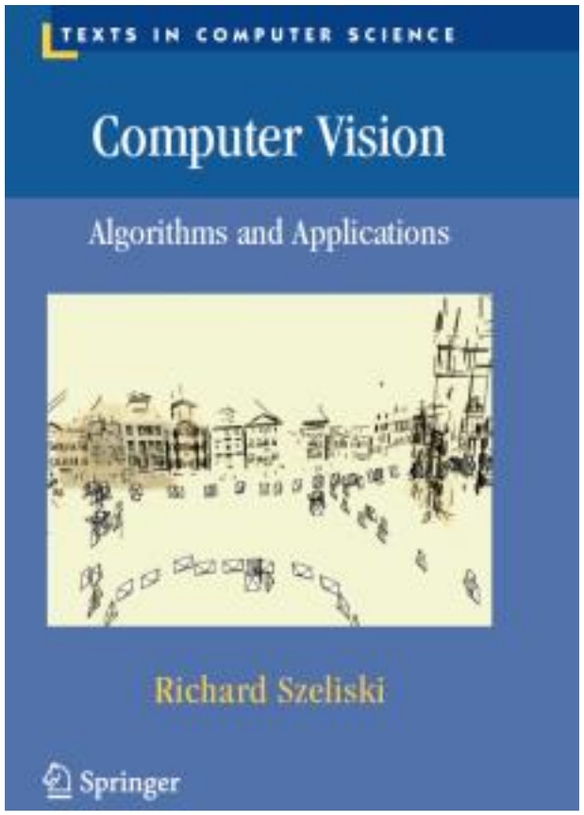

PDF online

http://szeliski.org/Book/

Prerequisites

We assume familiarity with calculus, linear algebra, basic probability, and programming.

Formal prerequisites:

- "Mathematical Foundations of Electrical Engineering" (18-202) and "Principles of Imperative Computation" (15-122)

OR

- "Matrix Algebra with Applications" (21-240) and "Matrices and Linear Transformations" (21-241) and "Calculus in Three Dimensions" (21-259) and "Principles of Imperative Computation" (15-122)

If you are missing a prerequisite but still want to enroll, let me know and we'll discuss it.

Contact information and office hours

- Feel free to email us about administrative questions.
  - please use [16385] in email title!

- Technical questions should be asked on Piazza.
  - we won't answer technical questions through email.
  - you can post anonymously if you prefer.

- Office hours will be determined by poll.
  - feel free to email Yannis about additional office hours.
  - you can also just drop by Yannis' office (Smith Hall (EDSH) Rm 225).

<u>For this week, Yannis will have office hours on Friday, 2-4 pm.</u>

Please take the course survey before the next lecture!

https://goo.gl/forms/AEOBv1InR7Js62Jx2

(also posted on Piazza)