Subsampling and image pyramids

http://www.cs.cmu.edu/~16385/

16-385 Computer Vision
Spring 2018, Lecture 3

Course announcements

- Homework 0 and homework 1 will be posted tonight.
  - Homework 0 is not required and is not graded!
  - Homework 1 is due on February 7th at midnight.

- Course website updated.
  - Syllabus slightly updated.
  - Added homework schedule.
  - Note homework 4 that spans spring break.

Overview of today’s lecture

- Image downsampling.
- Aliasing.
- Gaussian image pyramid.
- Laplacian image pyramid.

Slide credits

Most of these slides were adapted directly from:

- Kris Kitani (15-463, Fall 2016).

Some slides were inspired or taken from:

- Fredo Durand (MIT).
- Bernd Girod (Stanford University).
- James Hays (Georgia Tech).
- Steve Marschner (Cornell University).
- Steve Seitz (University of Washington).

Image downsampling

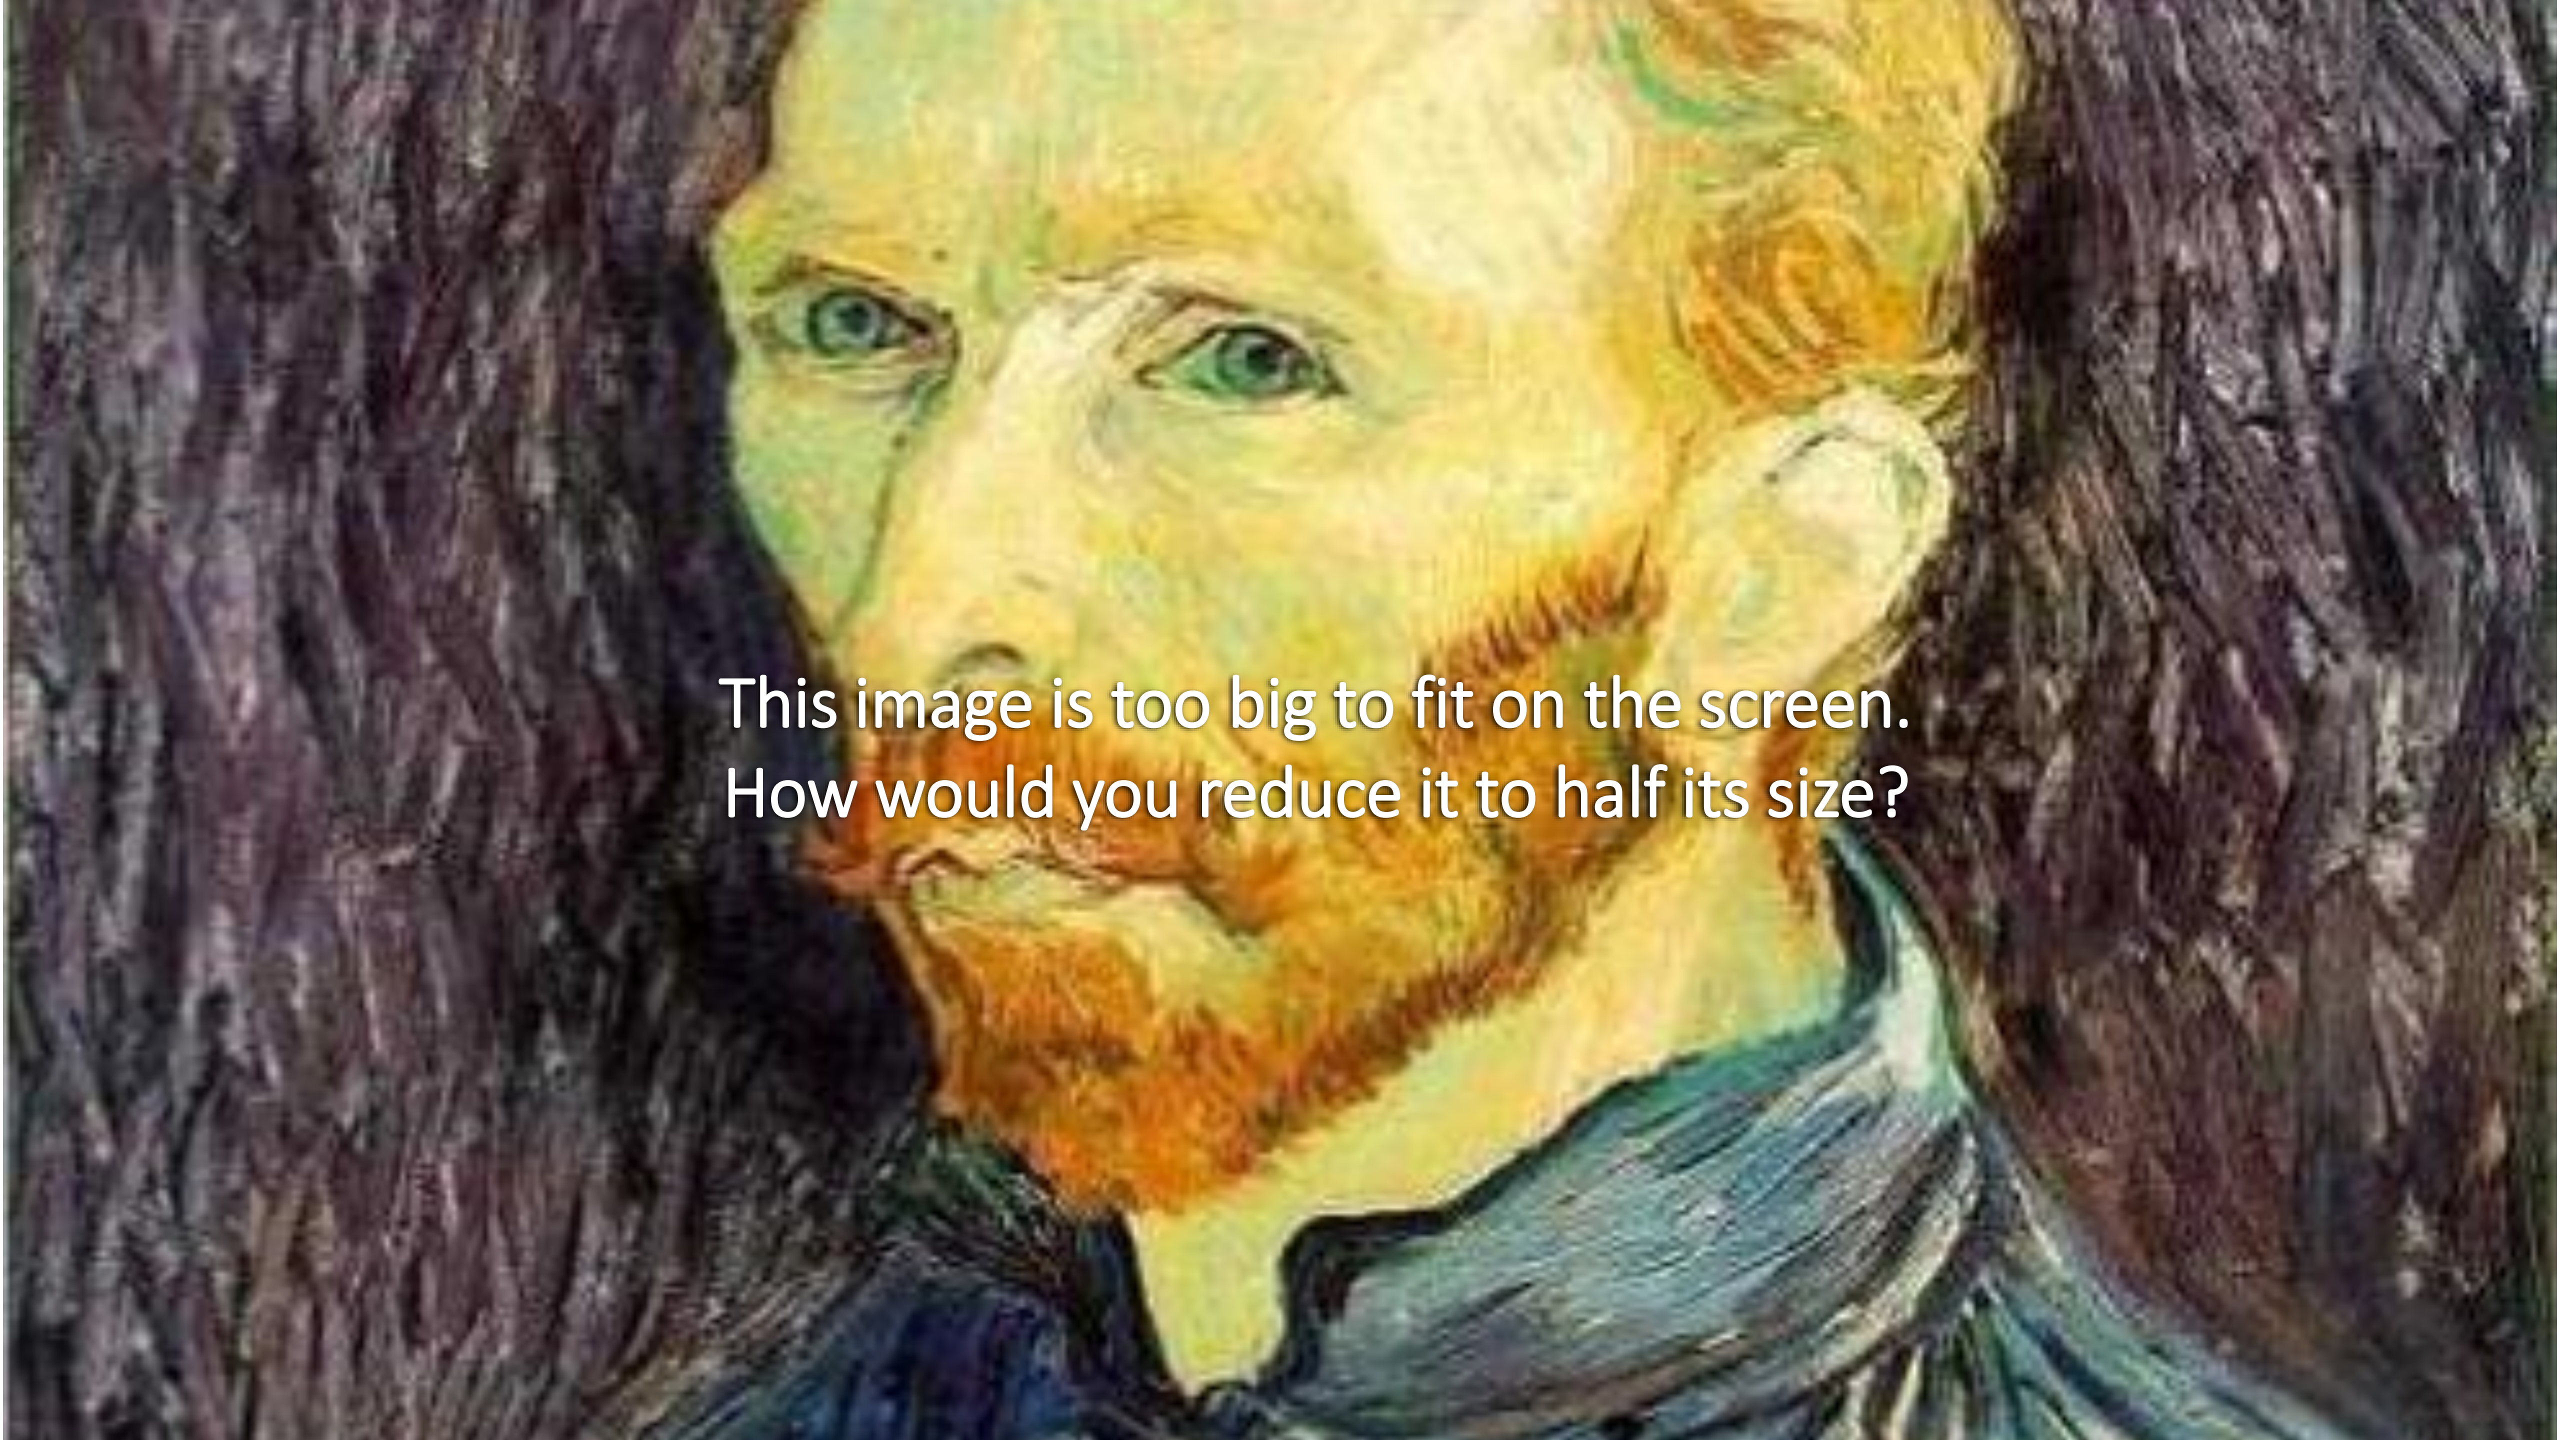
This image is too big to fit on the screen.
How would you reduce it to half its size?

Naïve image downsampling

Throw away half the rows and columns

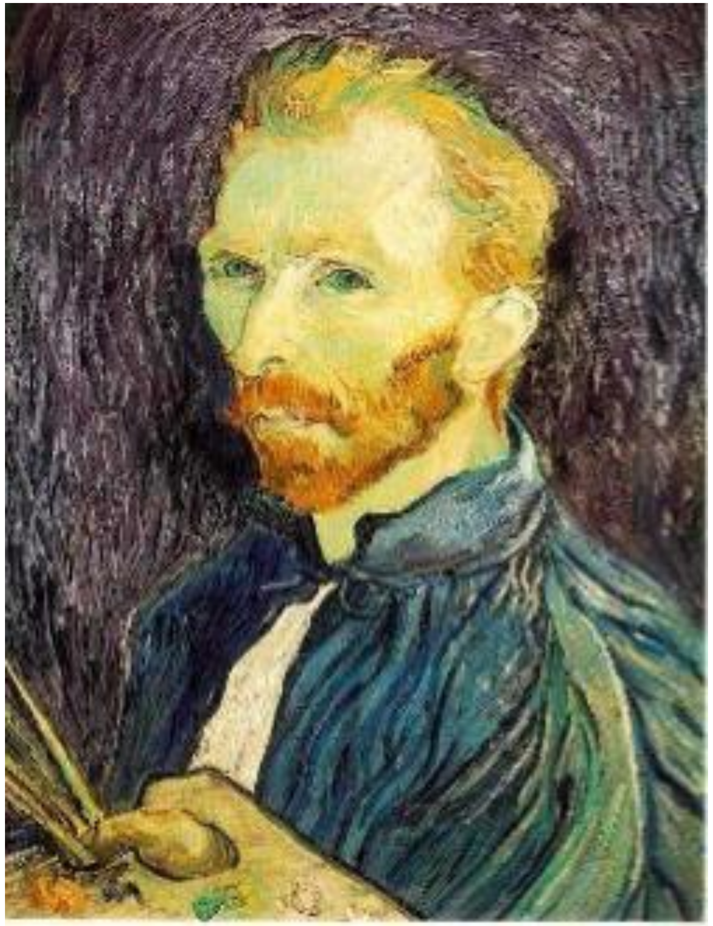

1/2

delete even rows
delete even columns

1/4

delete even rows
delete even columns

1/8

What is the problem with this approach?

Naïve image downsampling

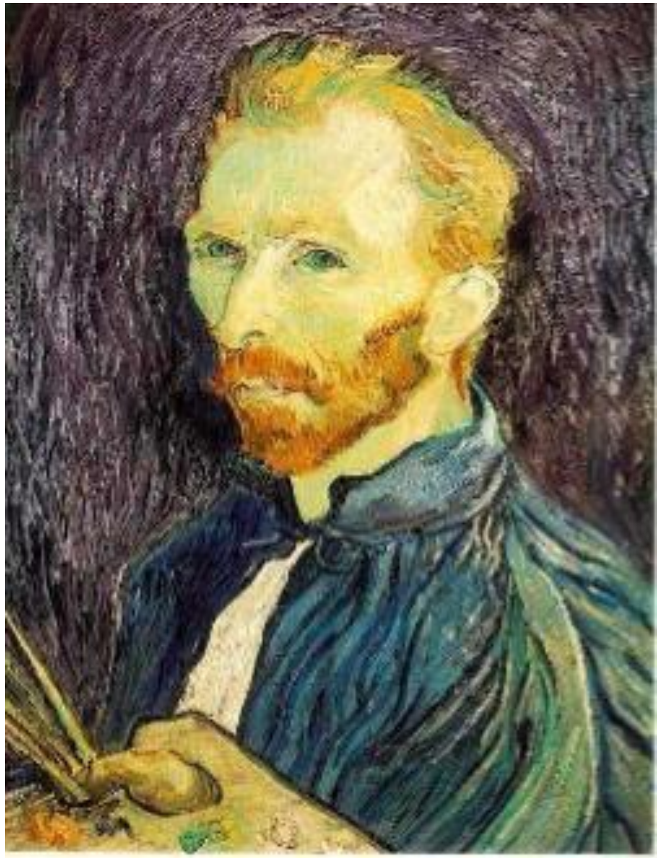

1/2

1/4 (2x zoom)

1/8 (4x zoom)

What is the 1/8 image so pixelated (and do you know what this effect is called)?

Aliasing

Reminder

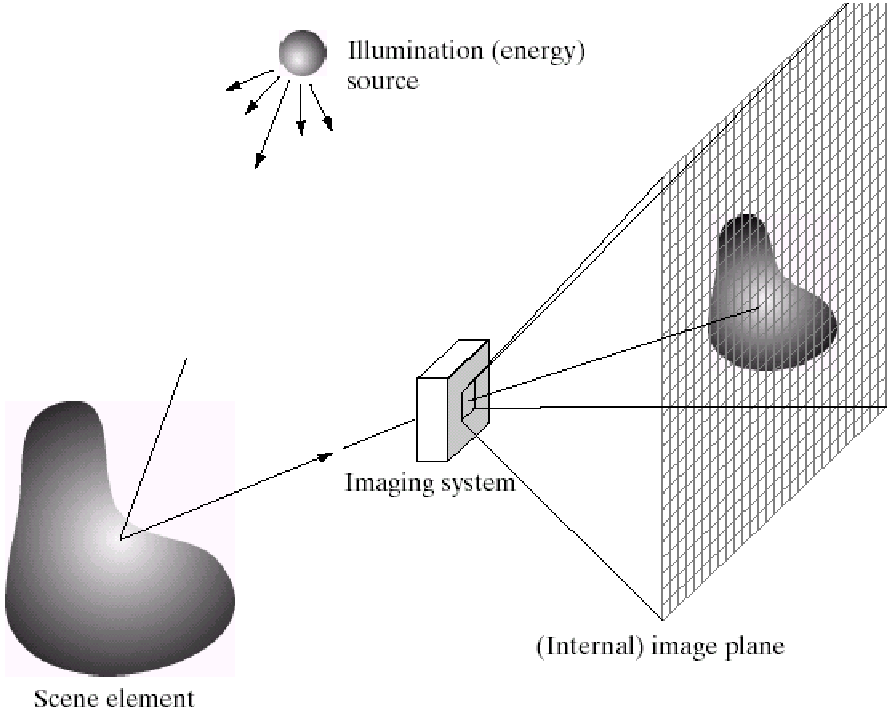

Images are a *discrete*, or *sampled*, representation of a *continuous* world

Sampling

Very simple example: a sine wave

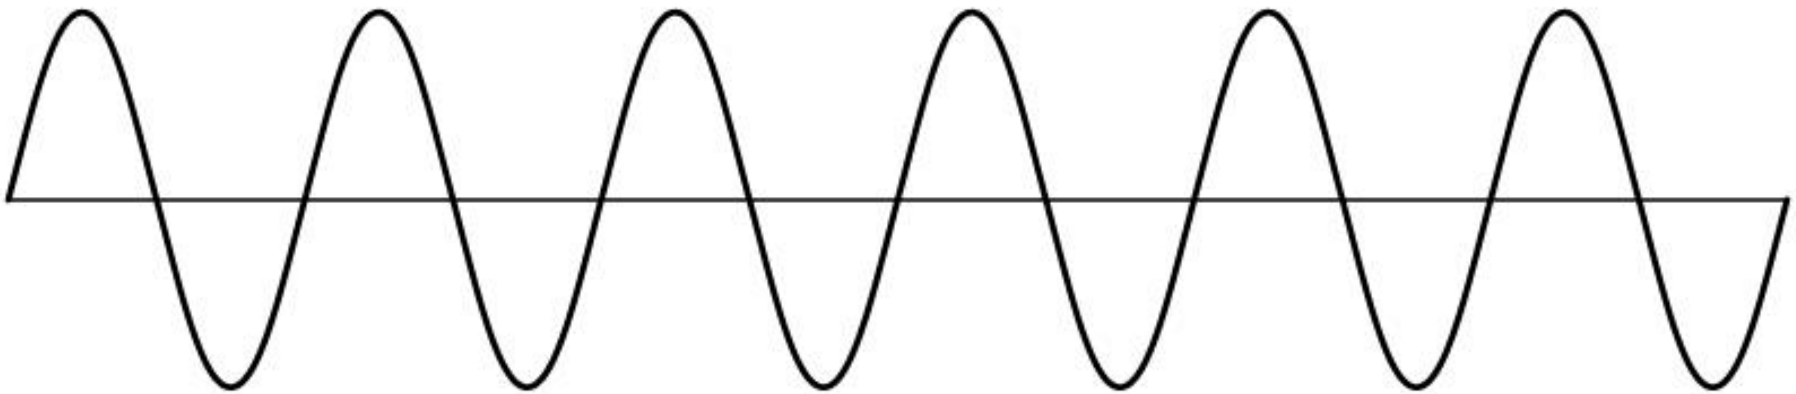

How would you discretize this signal?

Sampling

Very simple example: a sine wave

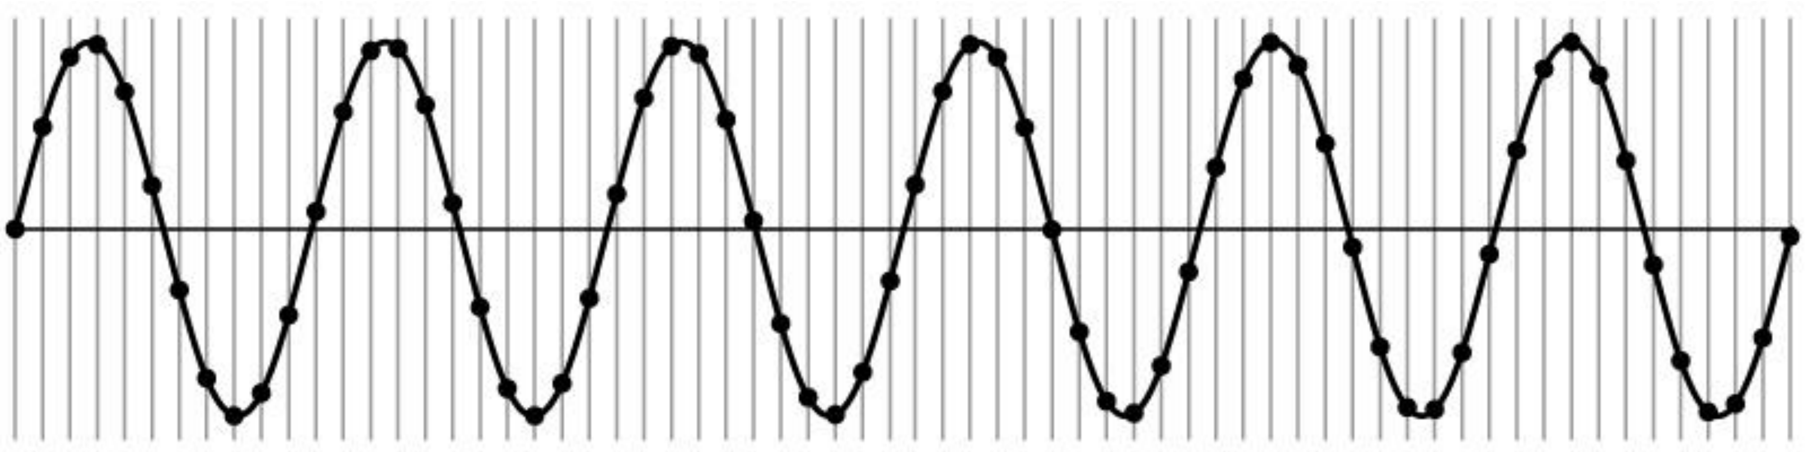

Sampling

Very simple example: a sine wave

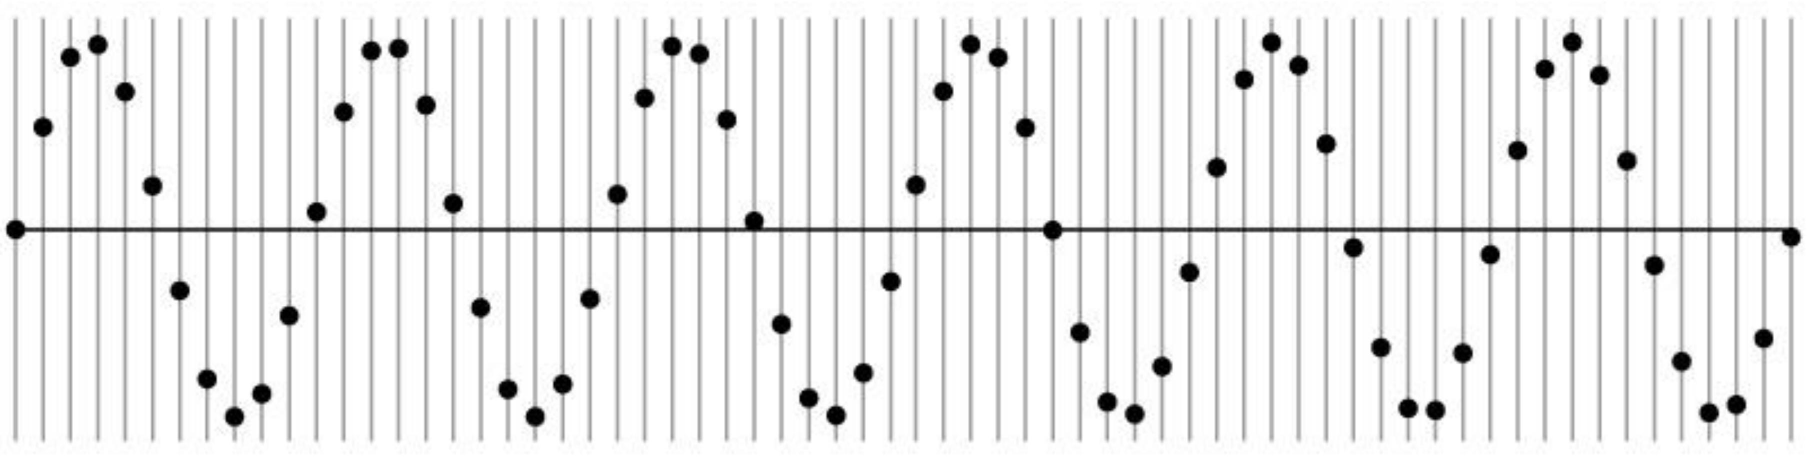

How many samples should I take?

Can I take as *many* samples as I want?

Sampling

Very simple example: a sine wave

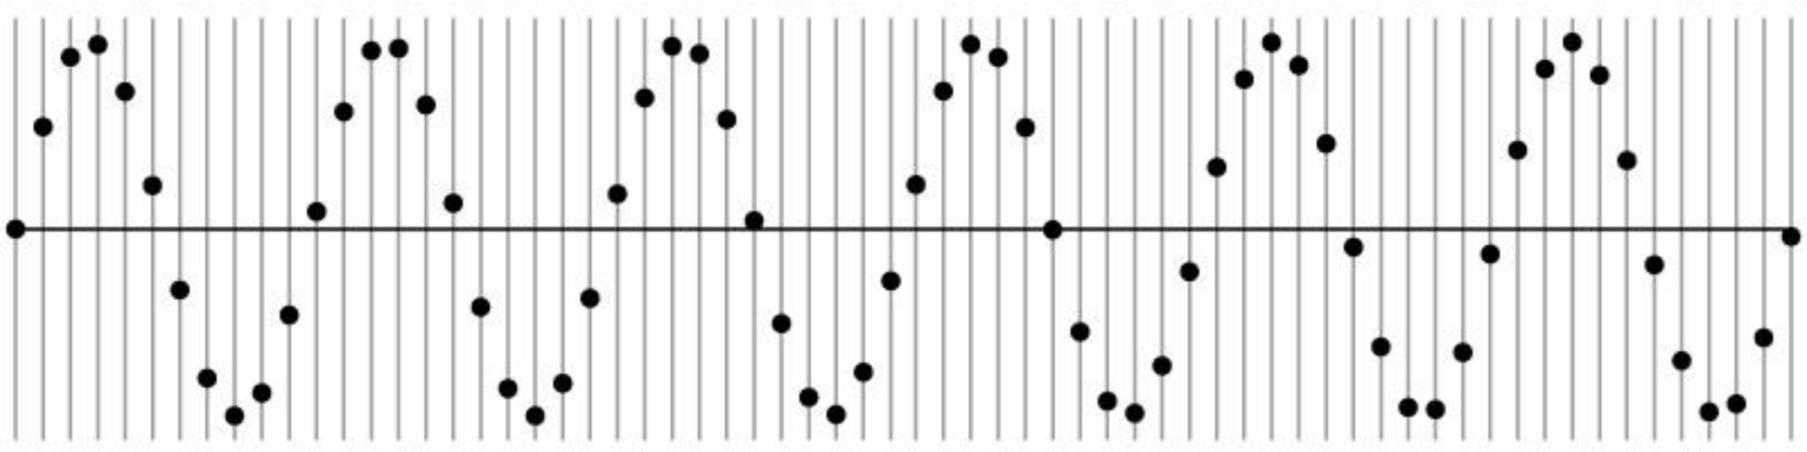

How many samples should I take?

Can I take as *few* samples as I want?

Undersampling

Very simple example: a sine wave

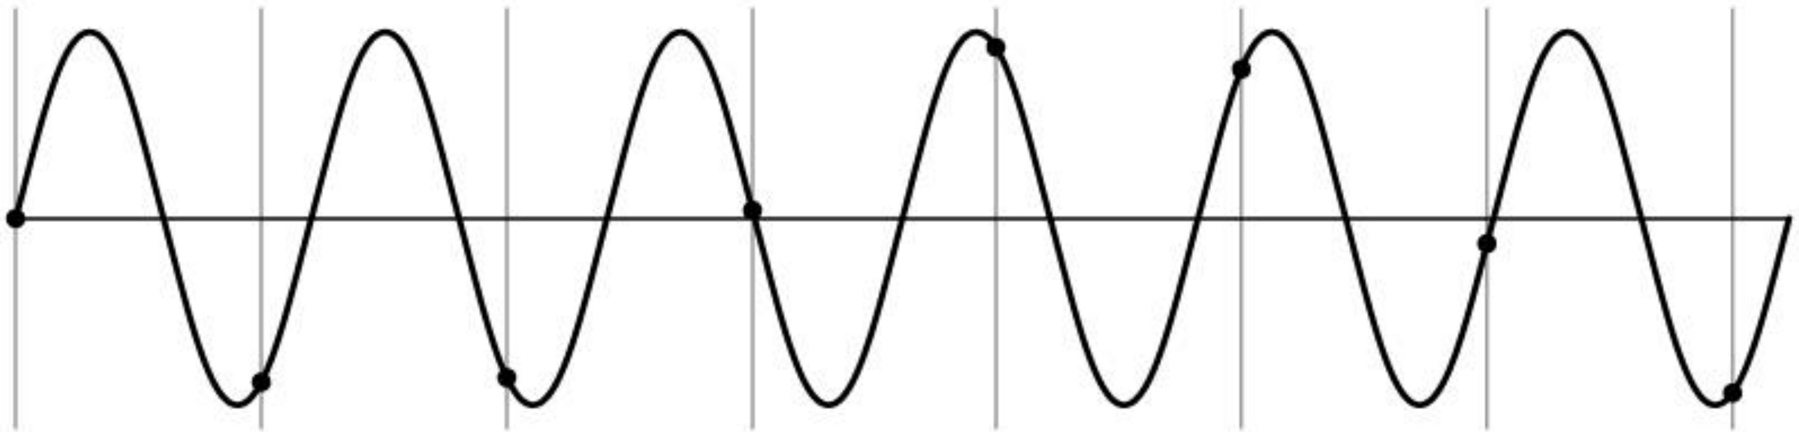

Unsurprising effect: information is lost.

Undersampling

Very simple example: a sine wave

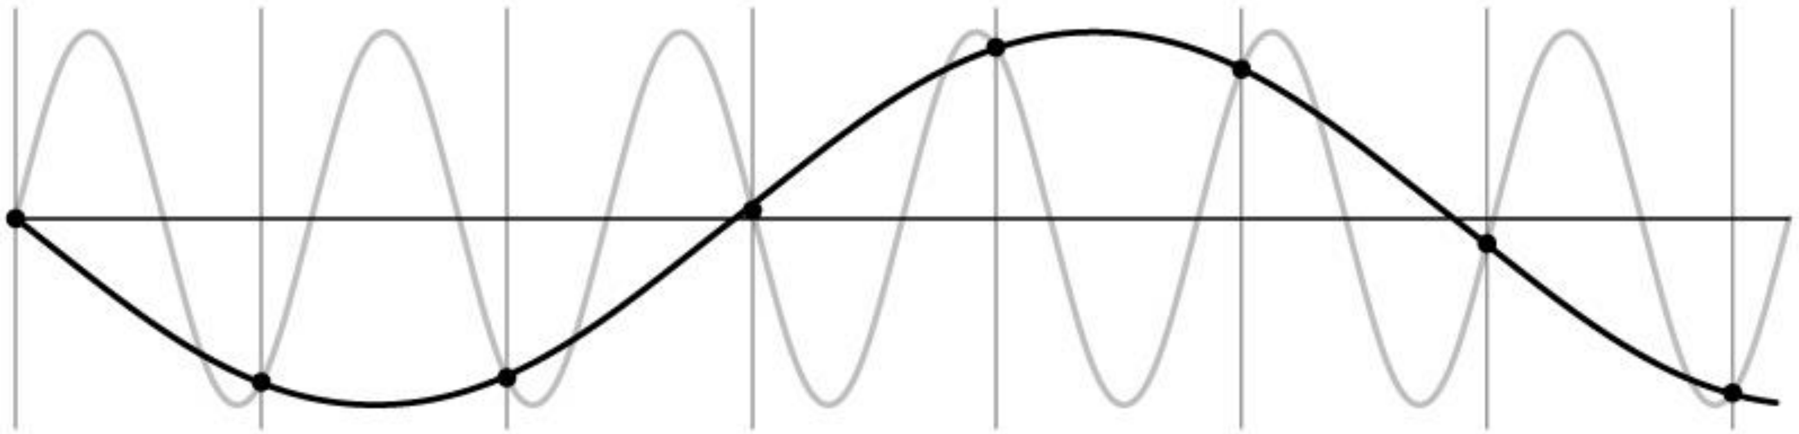

Unsurprising effect: information is lost.

Surprising effect: can confuse the signal with one of *lower* frequency.

Undersampling

Very simple example: a sine wave

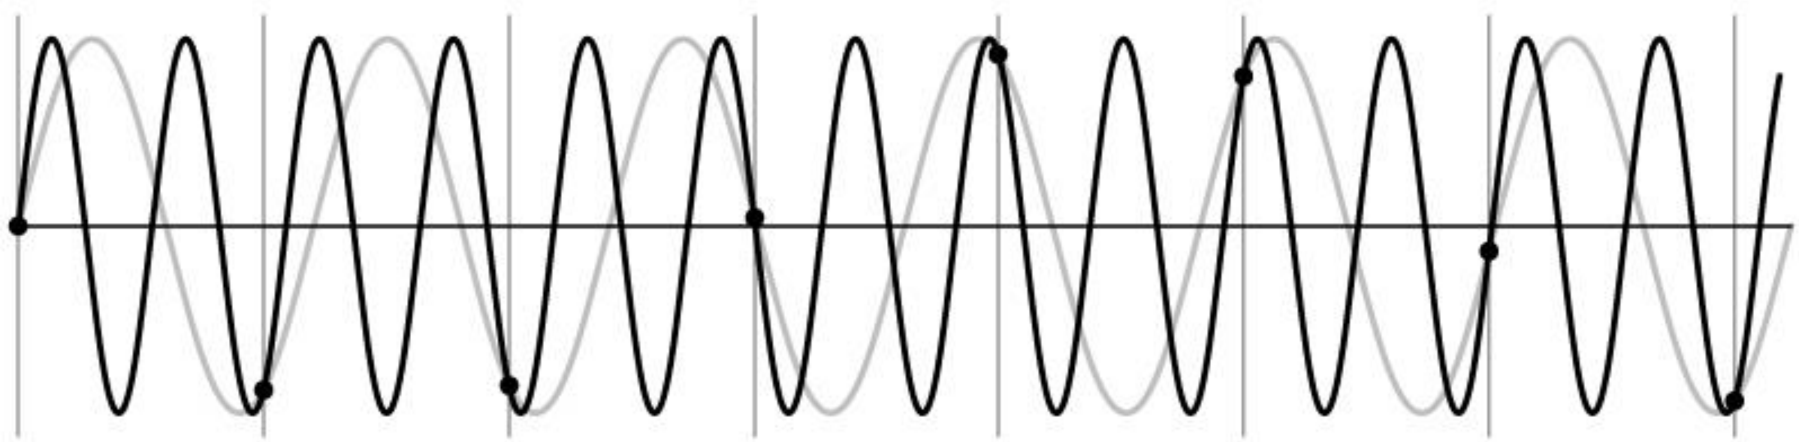

Unsurprising effect: information is lost.

Surprising effect: can confuse the signal with one of *lower* frequency.

Note: we could always confuse the signal with one of *higher* frequency.

Aliasing

Fancy term for: *Undersampling can disguise a signal as one of a lower frequency*

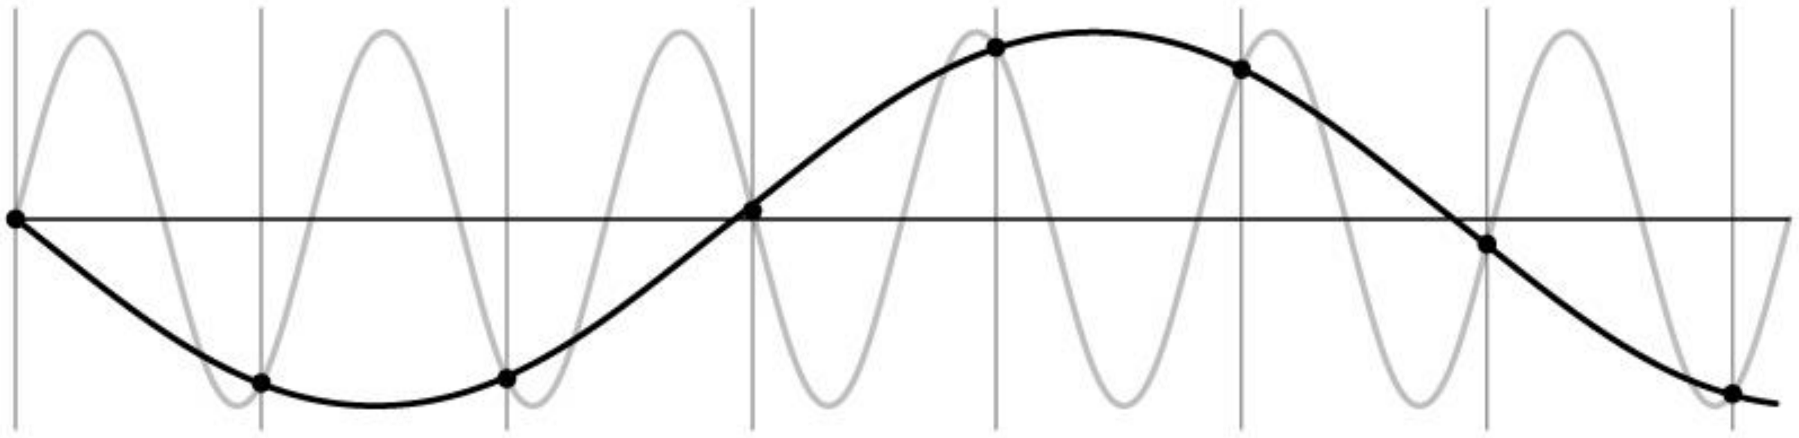

Unsurprising effect: information is lost.

Surprising effect: can confuse the signal with one of *lower* frequency.

Note: we could always confuse the signal with one of *higher* frequency.

Aliasing in textures

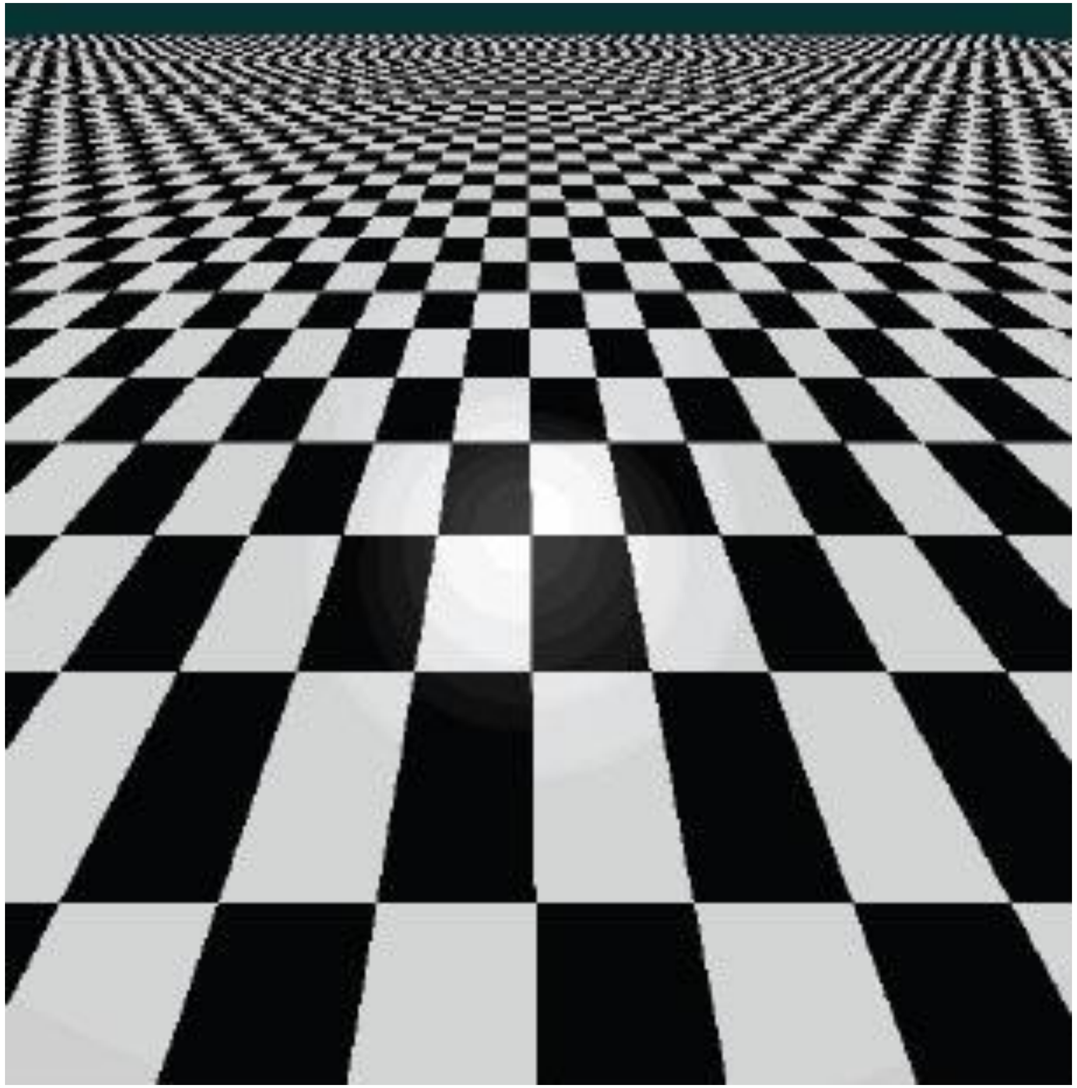

Aliasing in photographs

This is also known as “moire”

Temporal aliasing

Imagine a spoked wheel moving to the right (rotating clockwise). Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time = 1/30 sec. for video, 1/24 sec. for film):

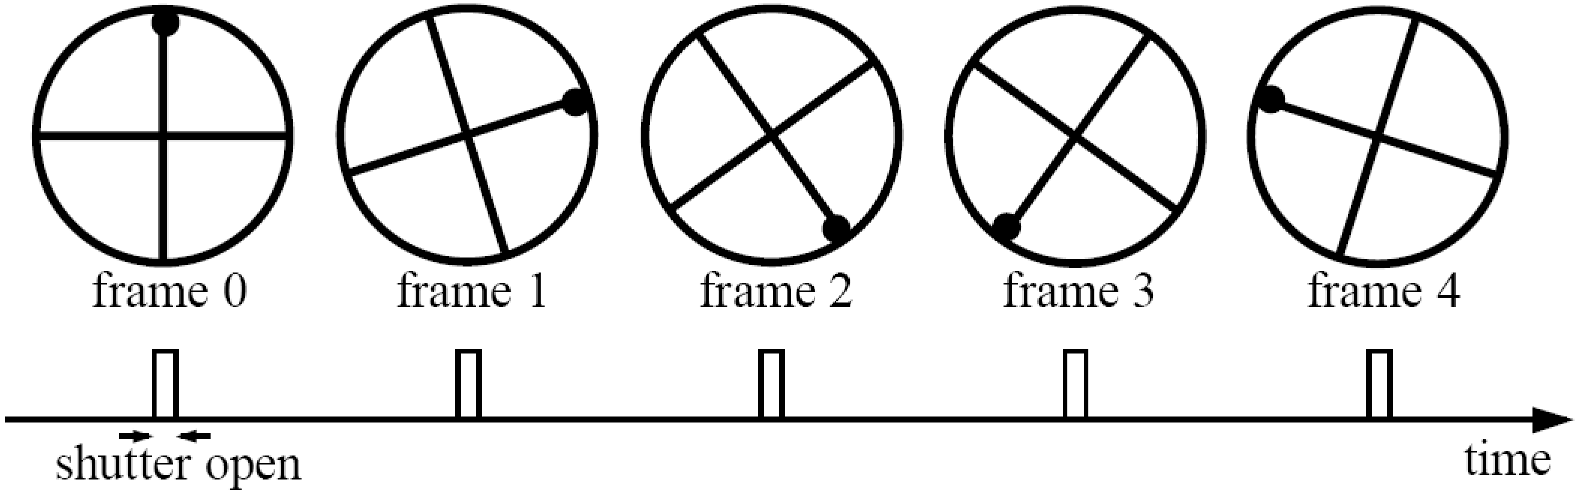

Without dot, wheel appears to be rotating slowly backwards! (counterclockwise)

Wagon wheel effect

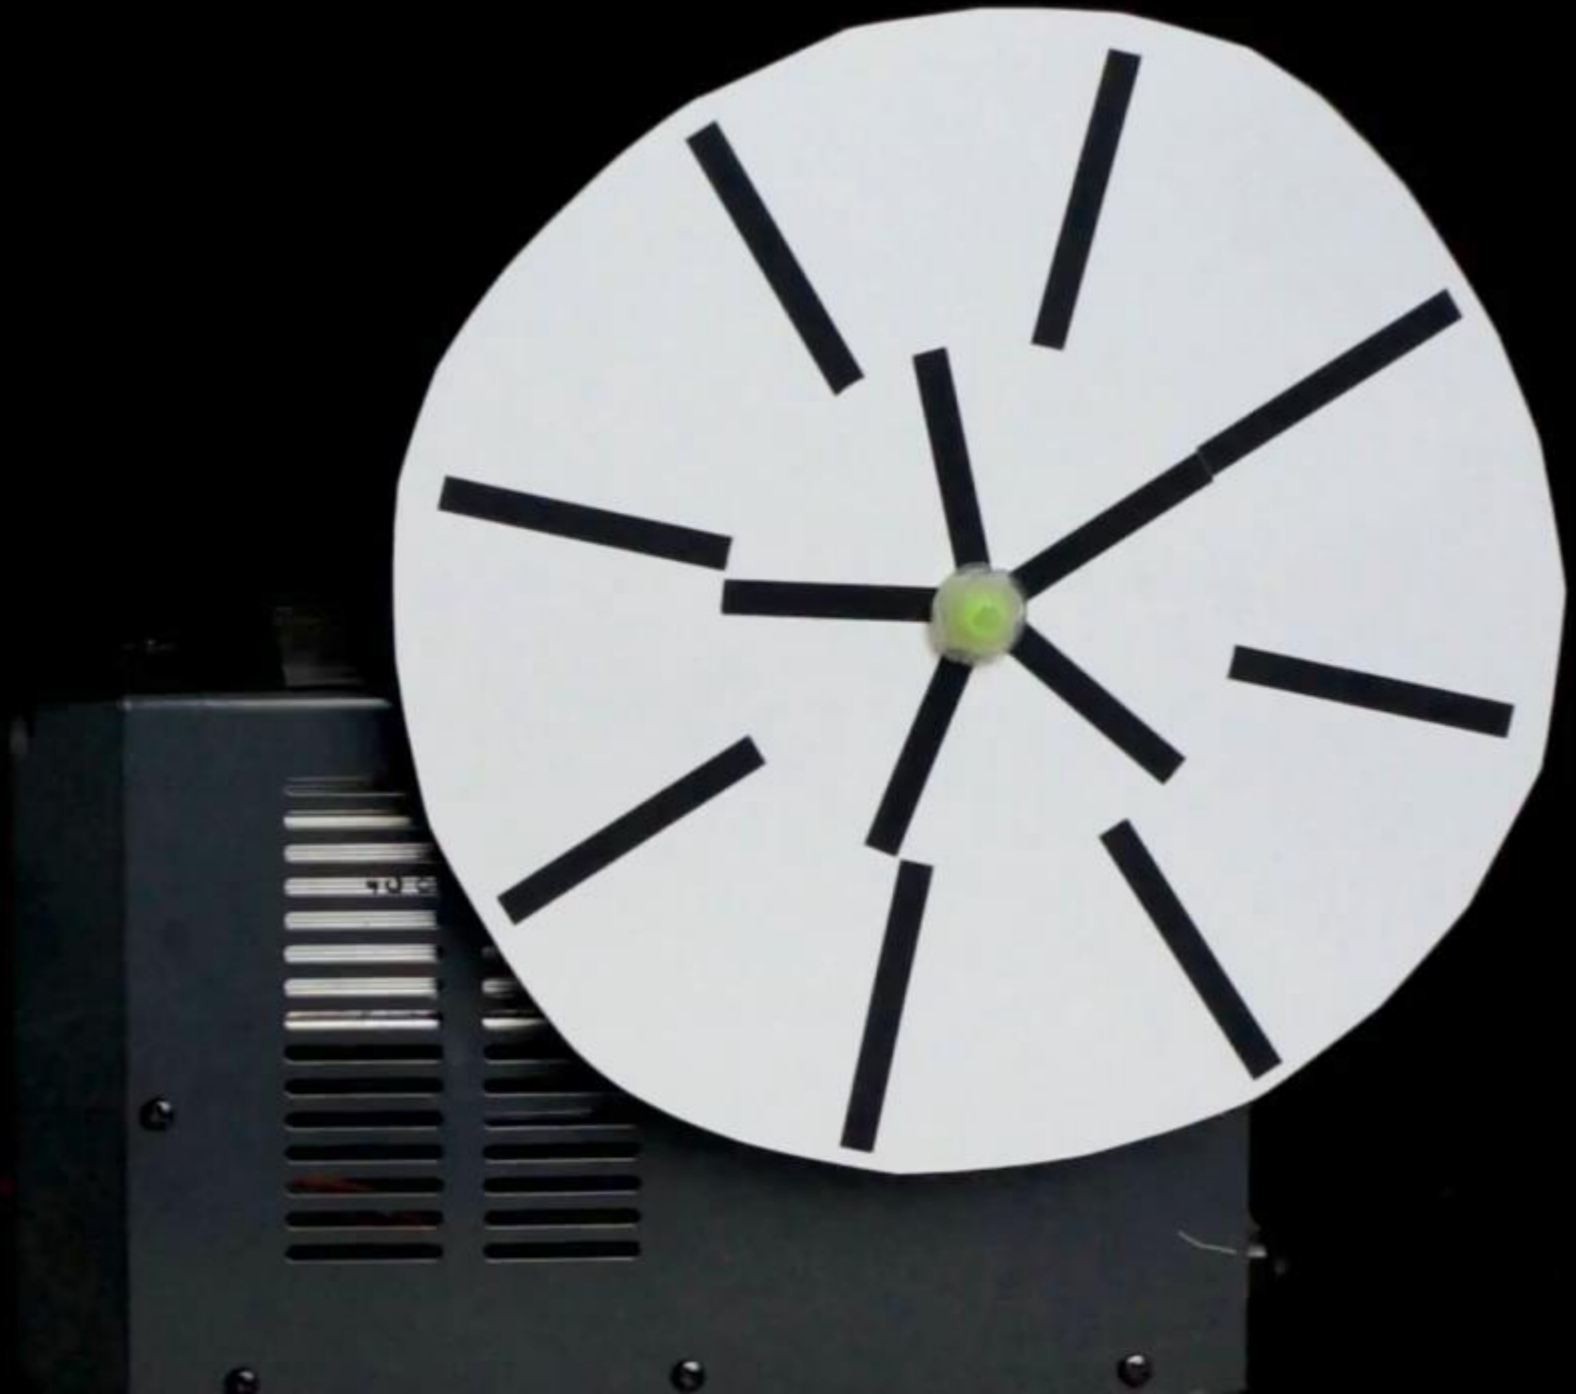

Anti-aliasing

How would you deal with aliasing?

Anti-aliasing

How would you deal with aliasing?

Approach 1: Oversample the signal

Anti-aliasing in textures

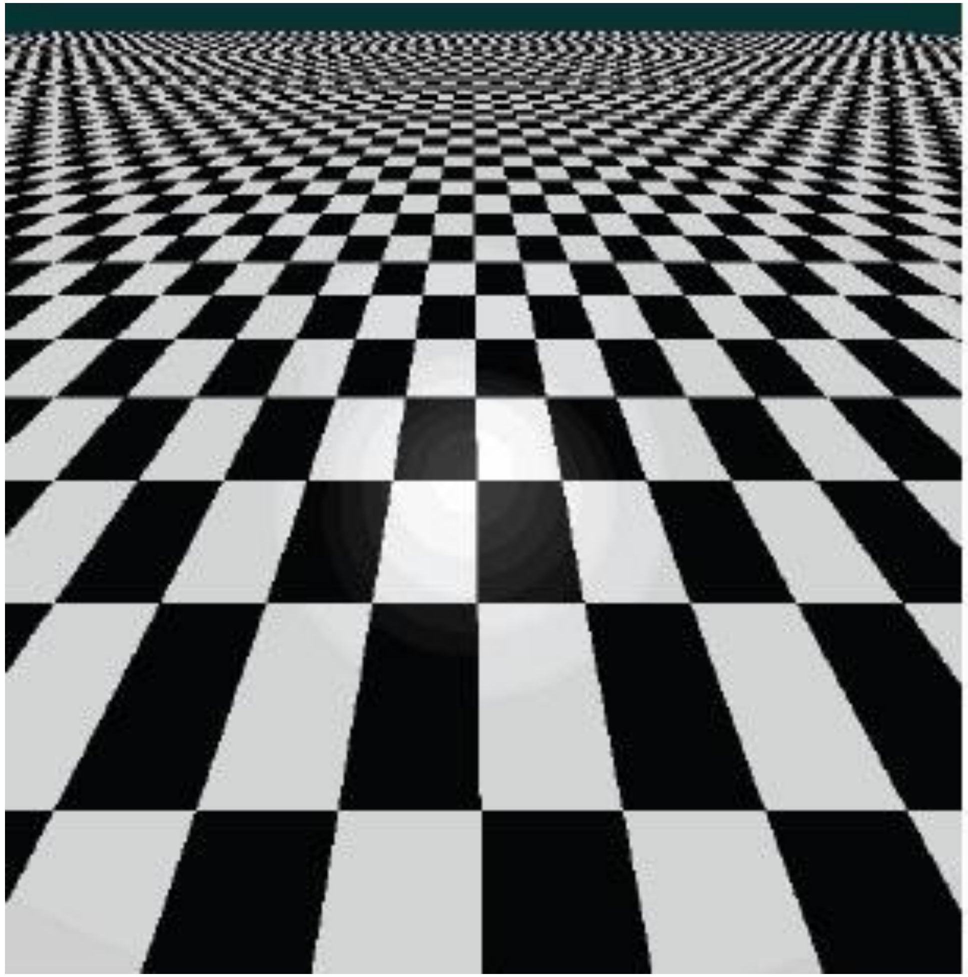

aliasing artifacts

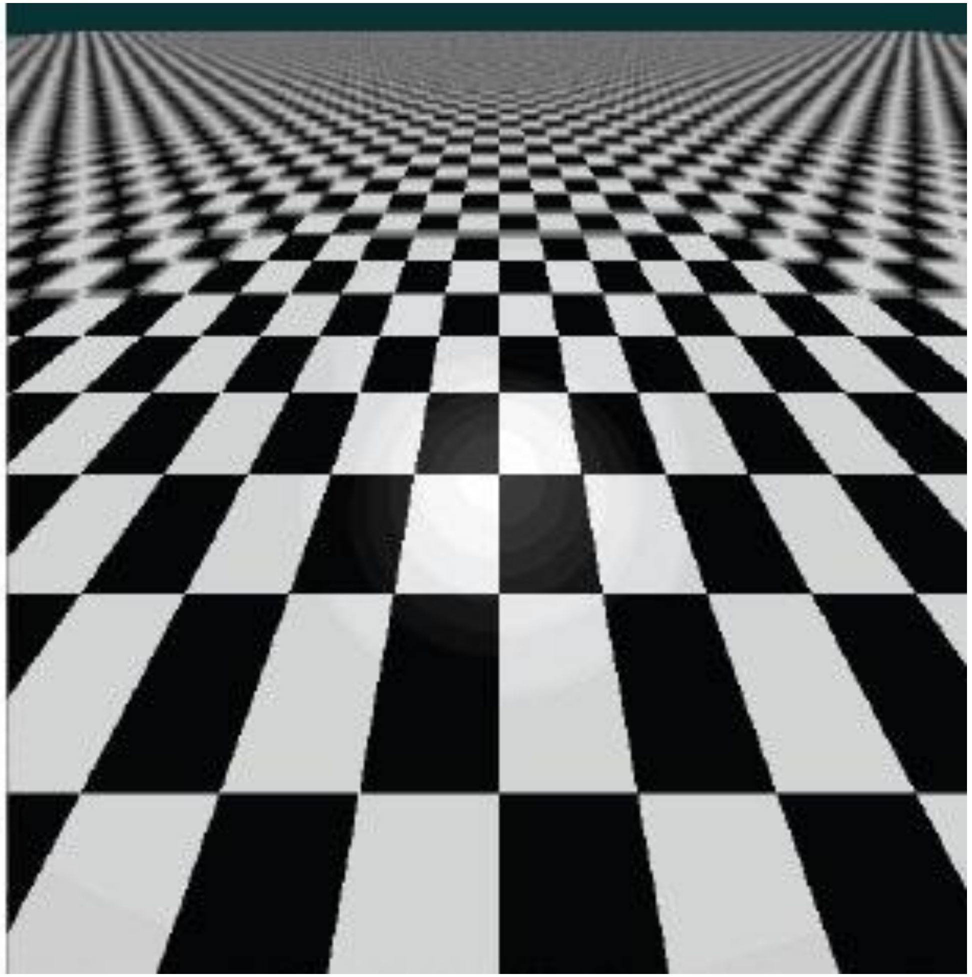

anti-aliasing by oversampling

Anti-aliasing

How would you deal with aliasing?

Approach 1: Oversample the signal

Approach 2: Smooth the signal

- Remove some of the detail effects that cause aliasing.
- Lose information, but better than aliasing artifacts.

How would you smooth a signal?

Better image downsampling

Apply a smoothing filter first, then throw away half the rows and columns

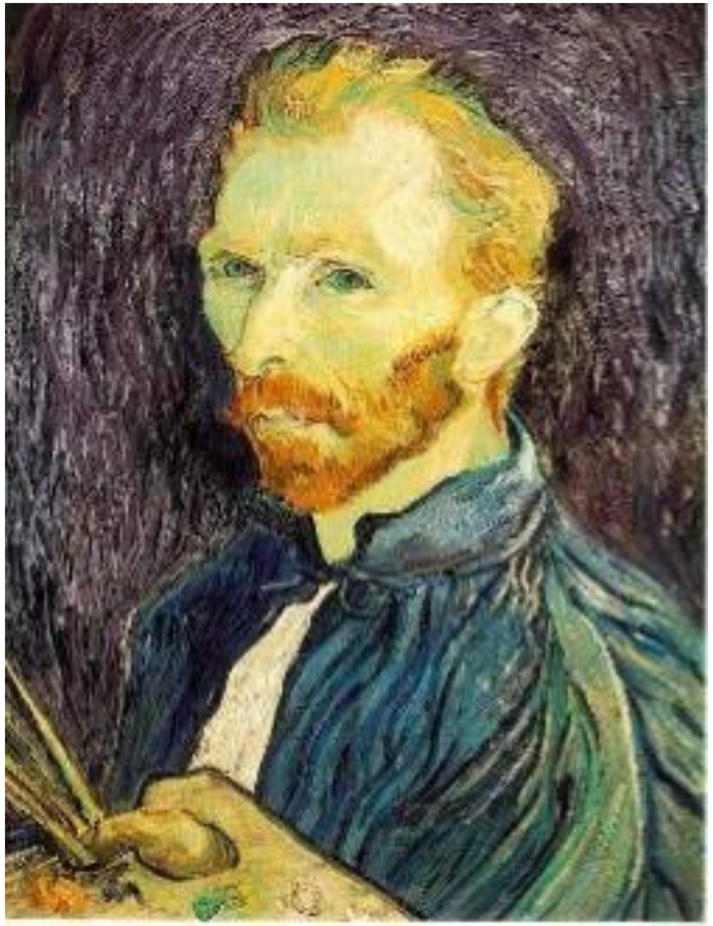

1/2

Gaussian filter
delete even rows
delete even columns

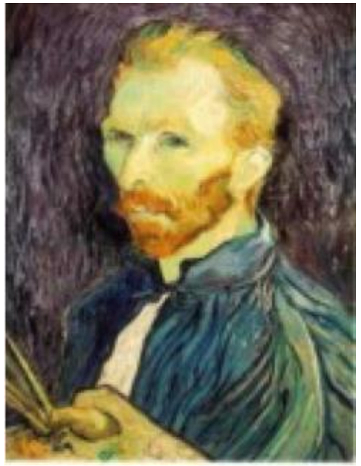

1/4

Gaussian filter
delete even rows
delete even columns

1/8

Better image downsampling

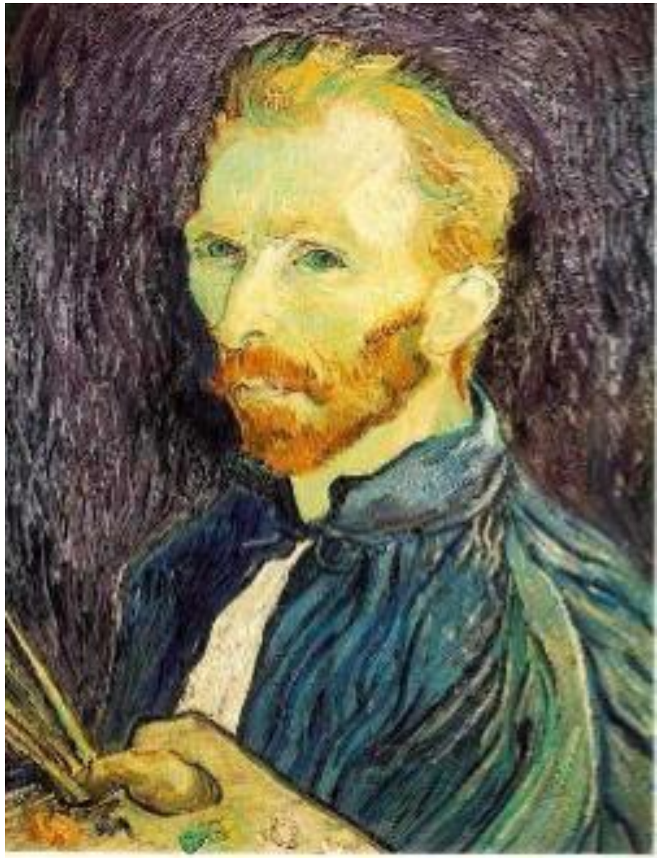

1/2

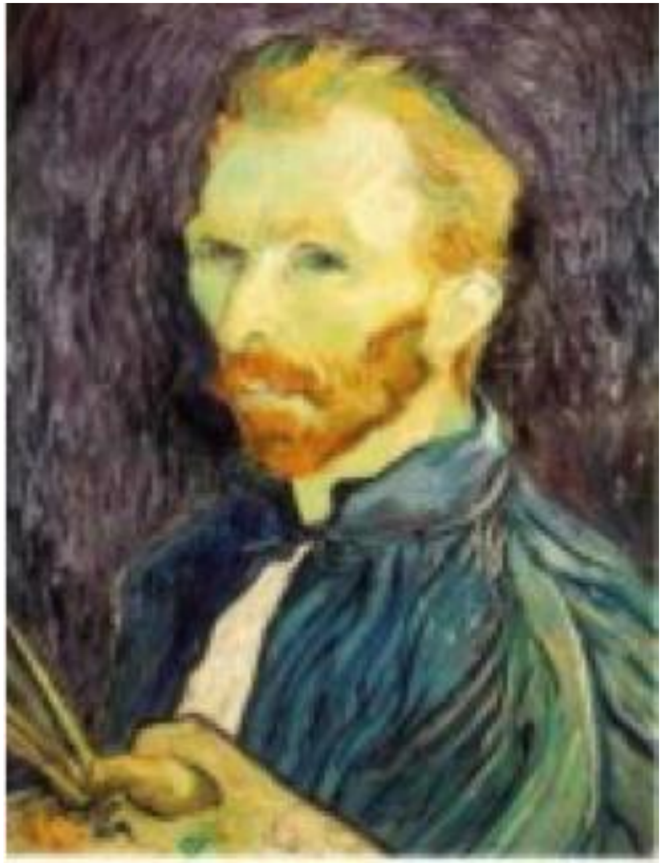

1/4 (2x zoom)

1/8 (4x zoom)

Naïve image downsampling

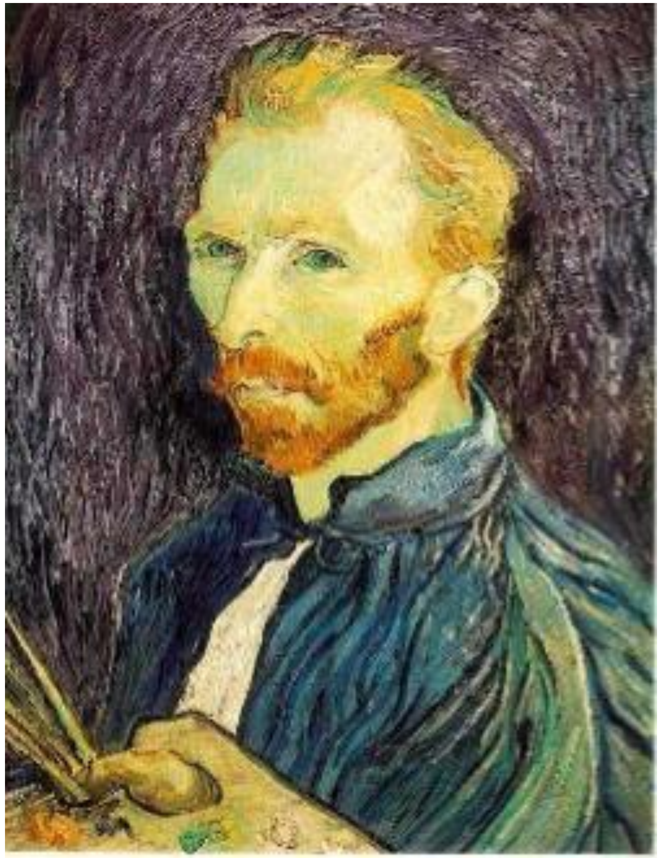

1/2

1/4 (2x zoom)

1/8 (4x zoom)

Anti-aliasing

Question 1: How much smoothing do I need to do to avoid aliasing?

Question 2: How many samples do I need to take to avoid aliasing?

Answer to both: Enough to reach the Nyquist limit.

We’ll see what this means soon.

Gaussian image pyramid

Gaussian image pyramid

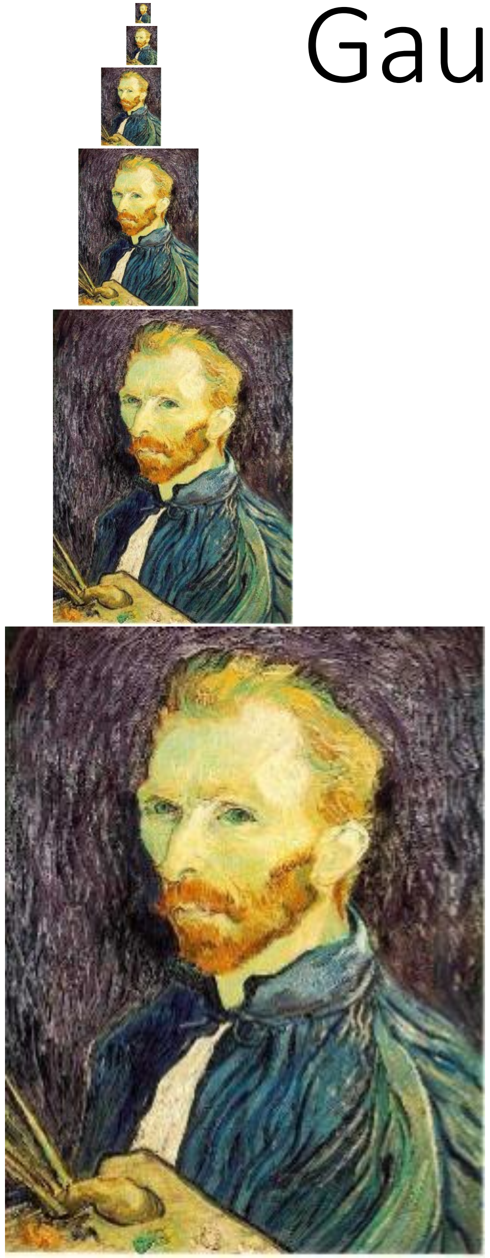

The name of this sequence of subsampled images

Constructing a Gaussian pyramid

Algorithm

repeat:

filter

subsample

until min resolution reached

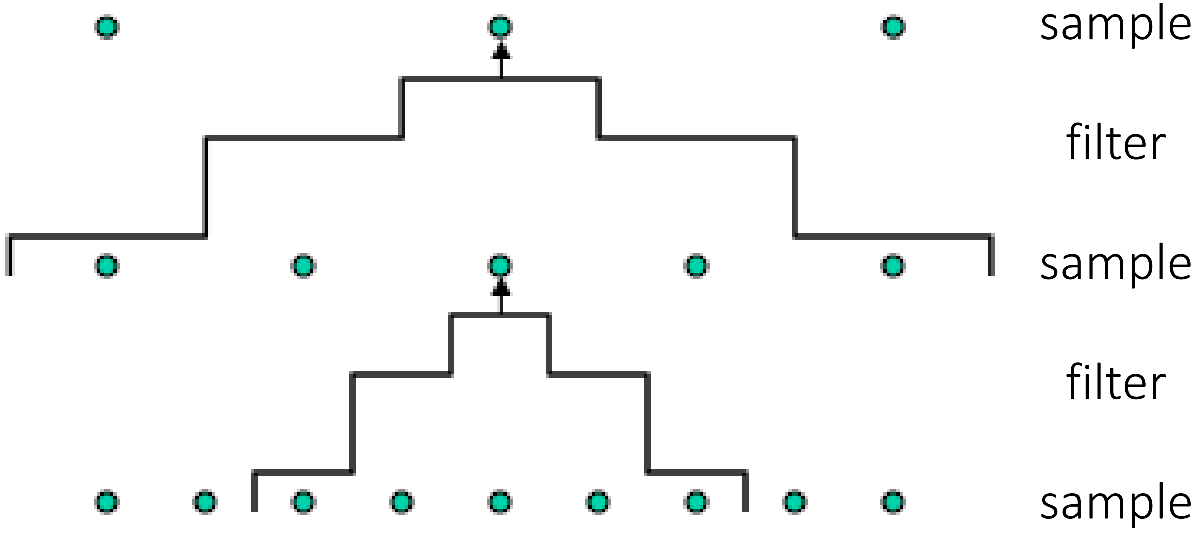

Question: How much bigger than the original image is the whole pyramid?

Constructing a Gaussian pyramid

Algorithm

repeat:

filter

subsample

until min resolution reached

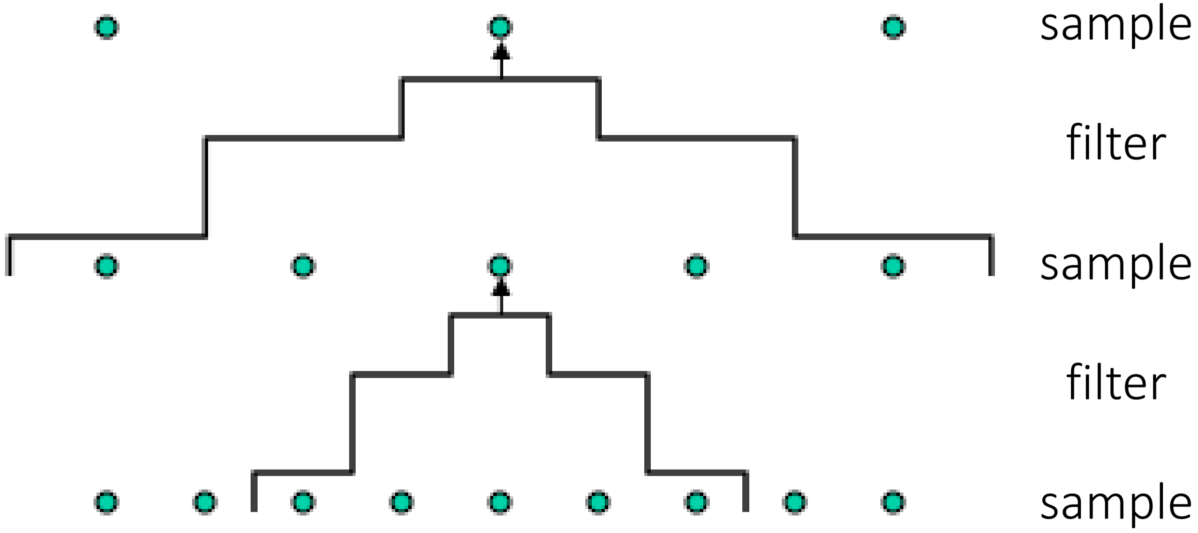

Question: How much bigger than the original image is the whole pyramid?

Answer: Just 4/3 times the size of the original image! (How did I come up with this number?)

Some properties of the Gaussian pyramid

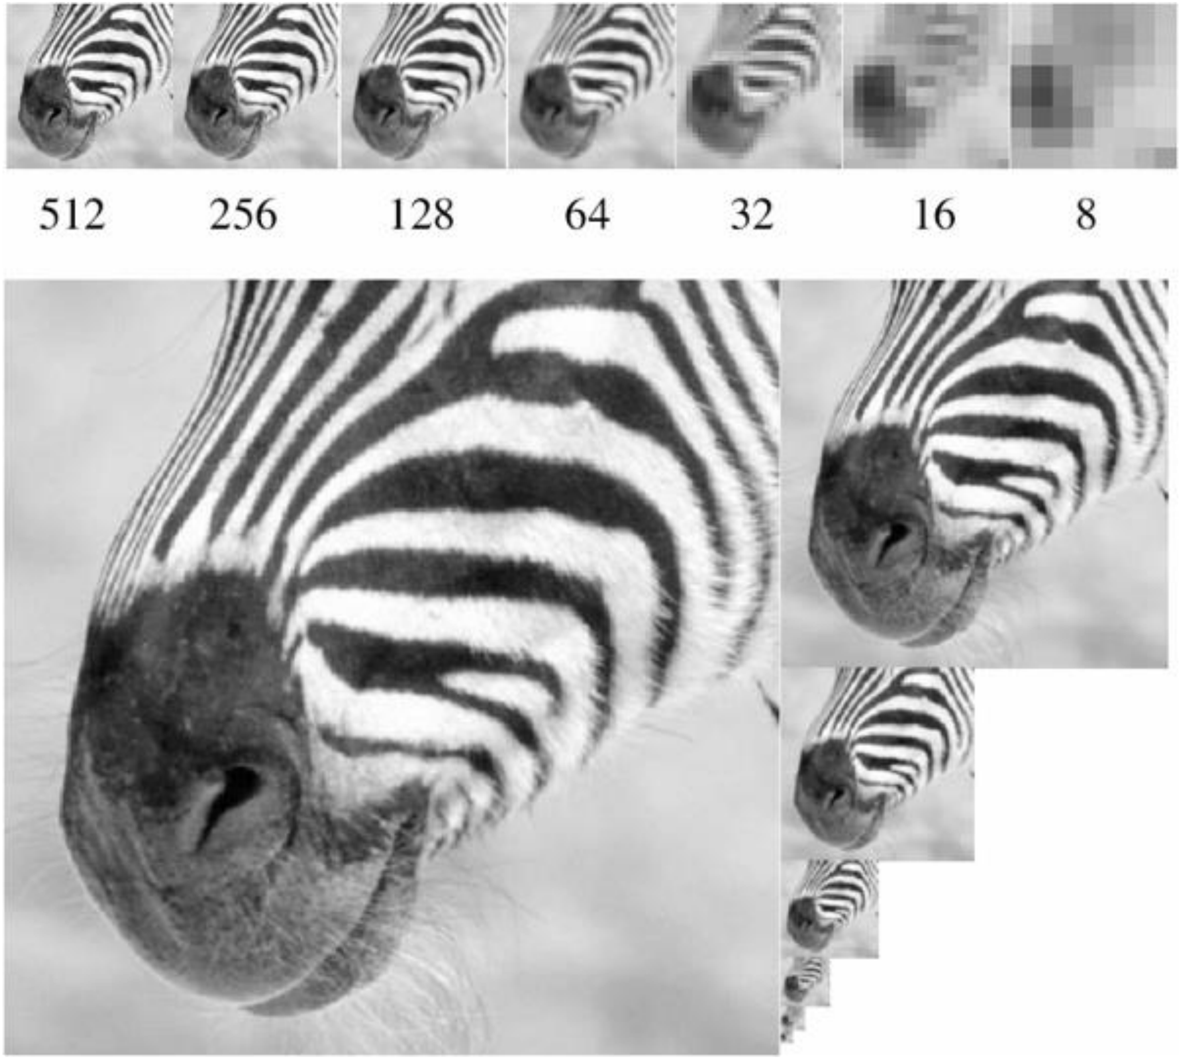

What happens to the details of the image?

Some properties of the Gaussian pyramid

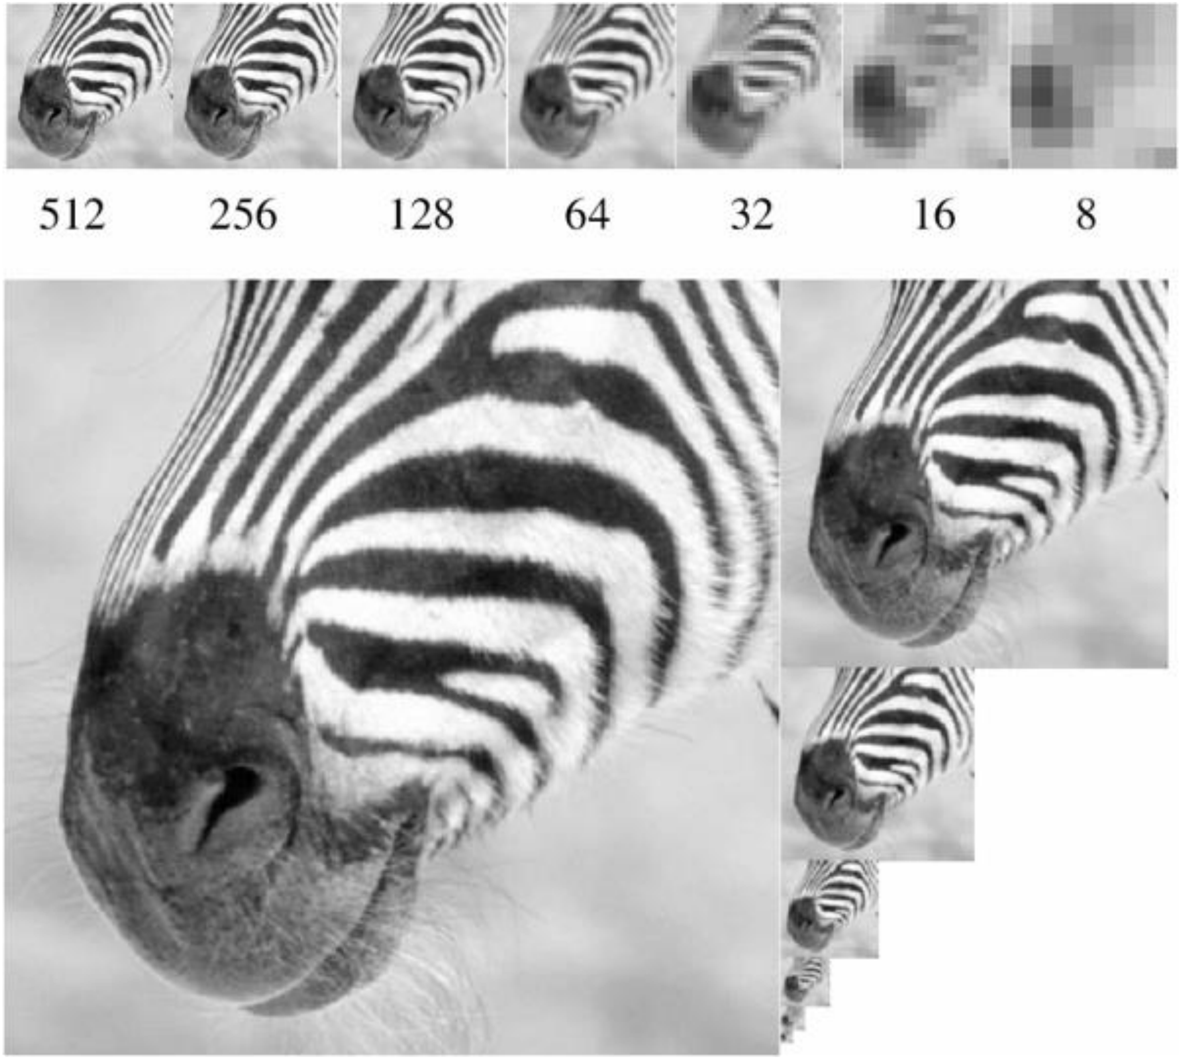

What happens to the details of the image?

- They get smoothed out as we move to higher levels.

What is preserved at the higher levels?

Some properties of the Gaussian pyramid

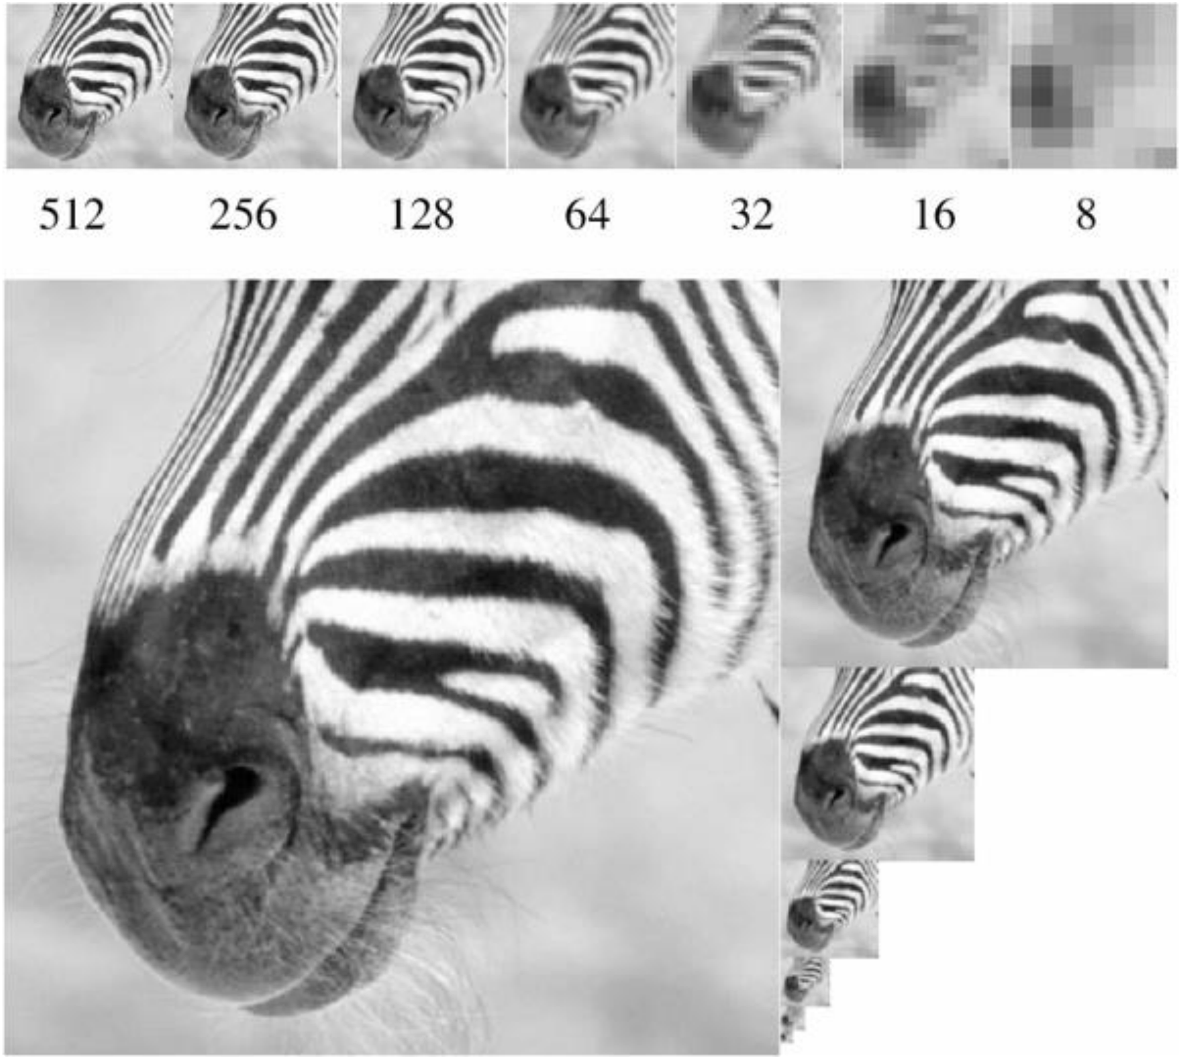

What happens to the details of the image?

- They get smoothed out as we move to higher levels.

What is preserved at the higher levels?

- Mostly large uniform regions in the original image.

How would you reconstruct the original image from the image at the upper level?

Some properties of the Gaussian pyramid

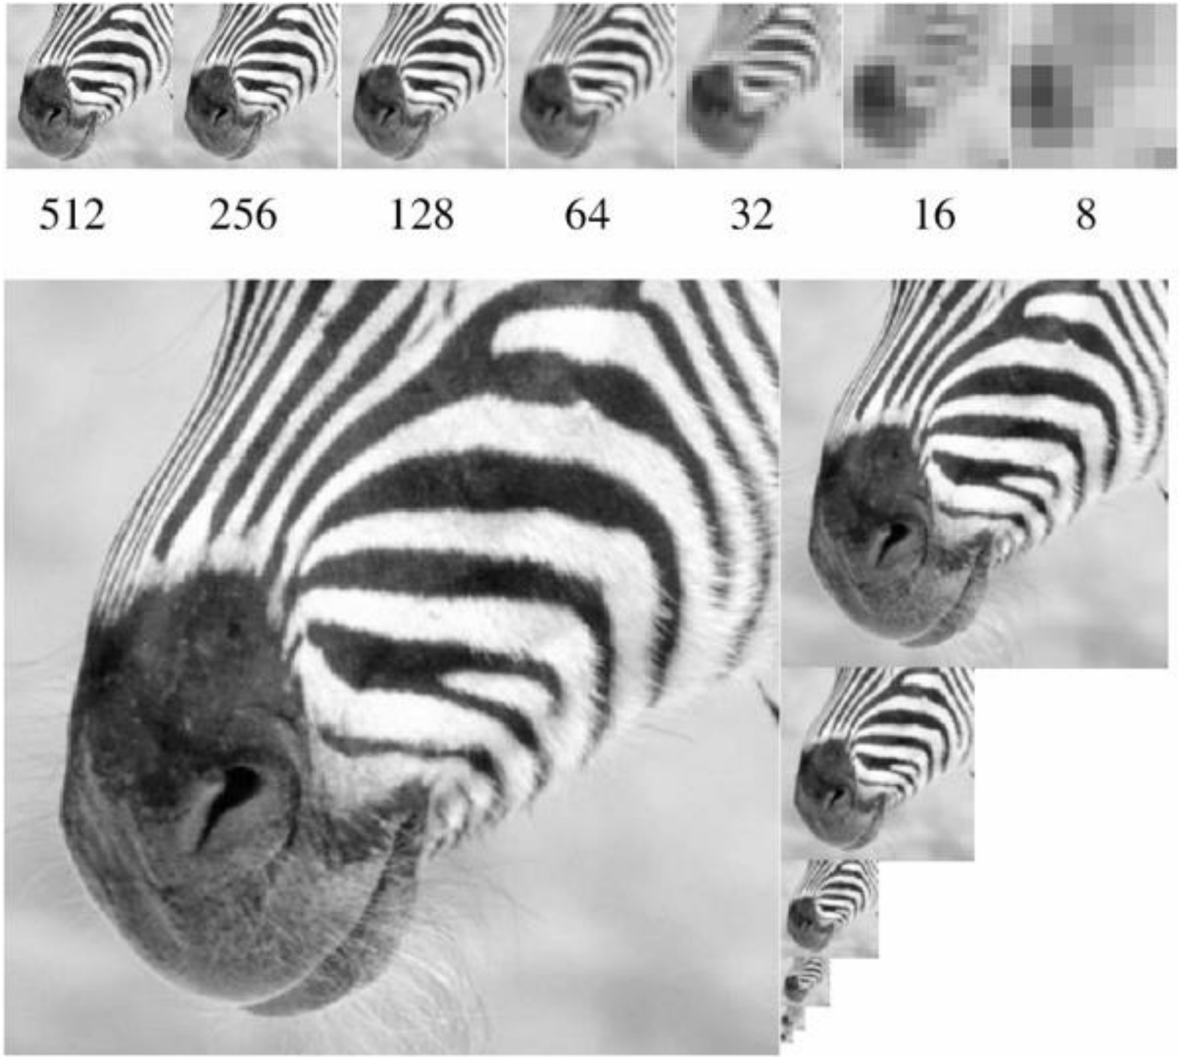

What happens to the details of the image?
- They get smoothed out as we move to higher levels.

What is preserved at the higher levels?
- Mostly large uniform regions in the original image.

How would you reconstruct the original image from the image at the upper level?
- That’s not possible.

Blurring is lossy

level 0

-

level 1 (before downsampling)

=

residual

What does the residual look like?

Blurring is lossy

-

=

level 0

level 1 (before downsampling)

residual

Can we make a pyramid that is lossless?

Laplacian image pyramid

Laplacian image pyramid

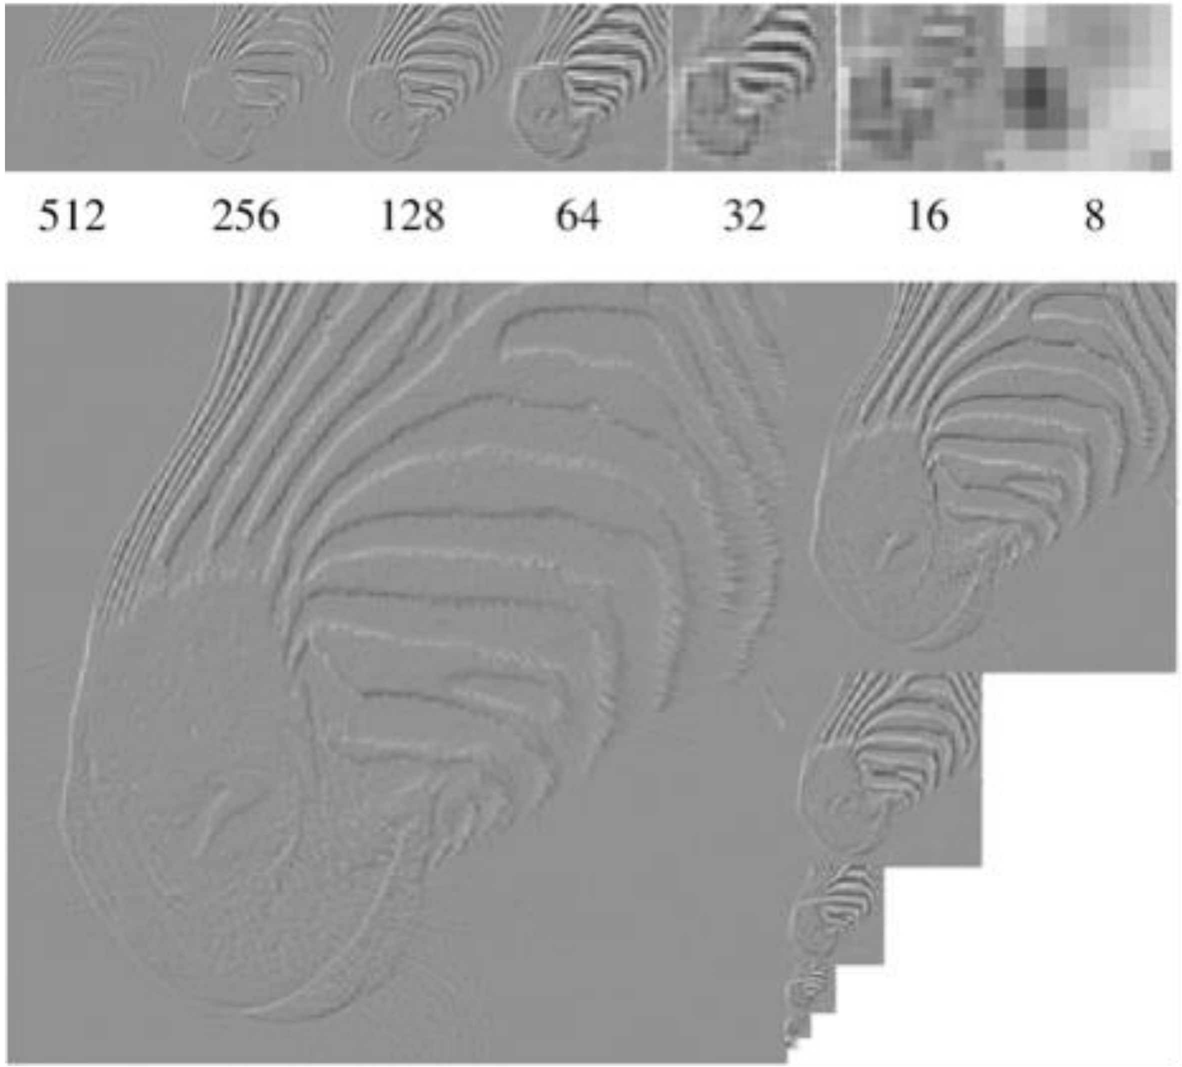

At each level, retain the residuals instead of the blurred images themselves.

Can we reconstruct the original image using the pyramid?

Laplacian image pyramid

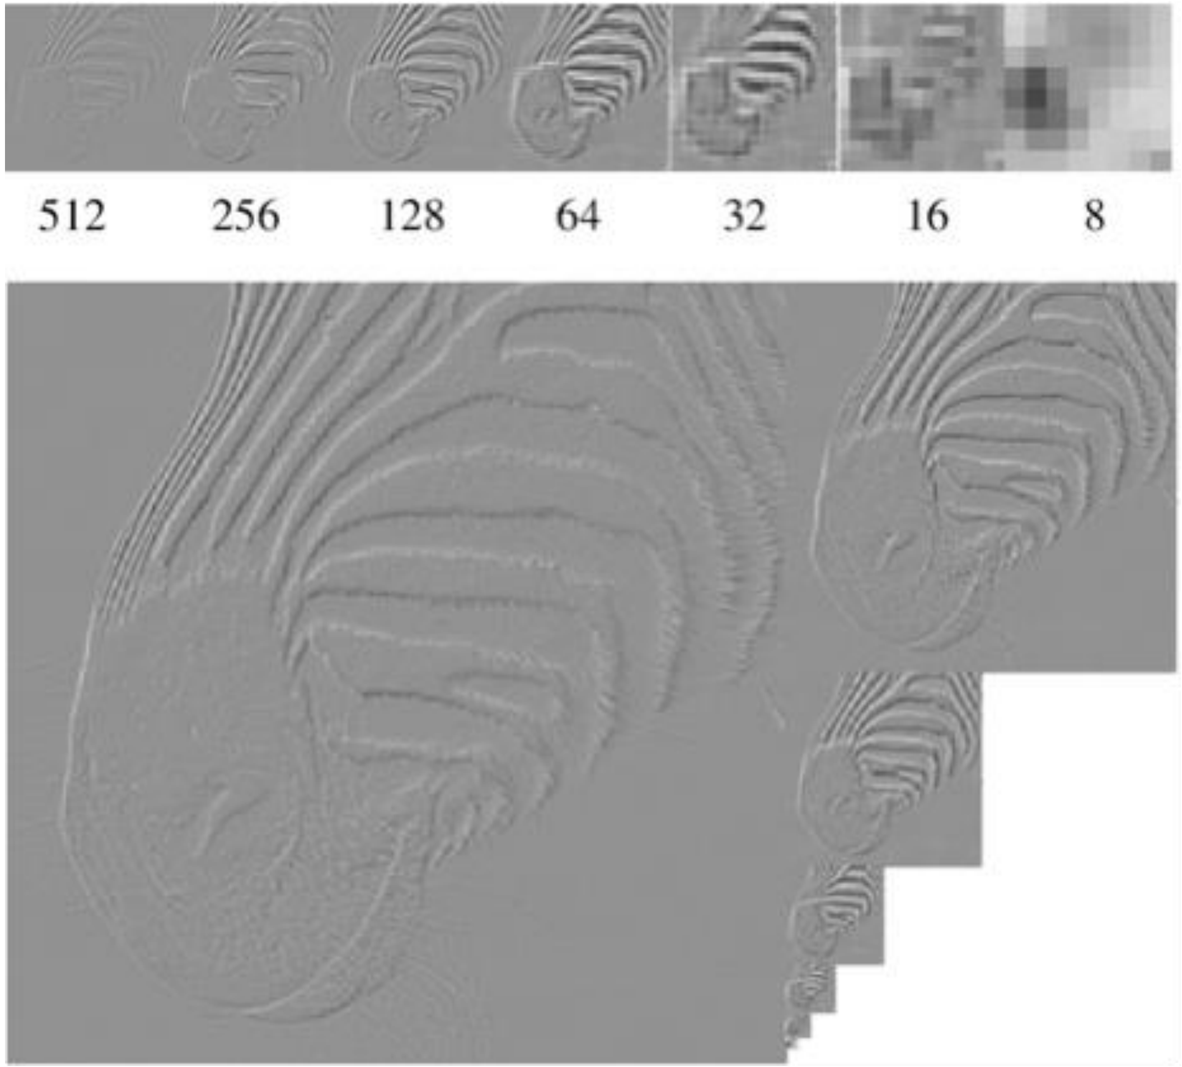

At each level, retain the residuals instead of the blurred images themselves.

Can we reconstruct the original image using the pyramid?

- Yes we can!

What do we need to store to be able to reconstruct the original image?

Let’s start by looking at just one level

=

+

level 0

level 1 (upsampled)

residual

Does this mean we need to store both residuals and the blurred copies of the original?

Constructing a Laplacian pyramid

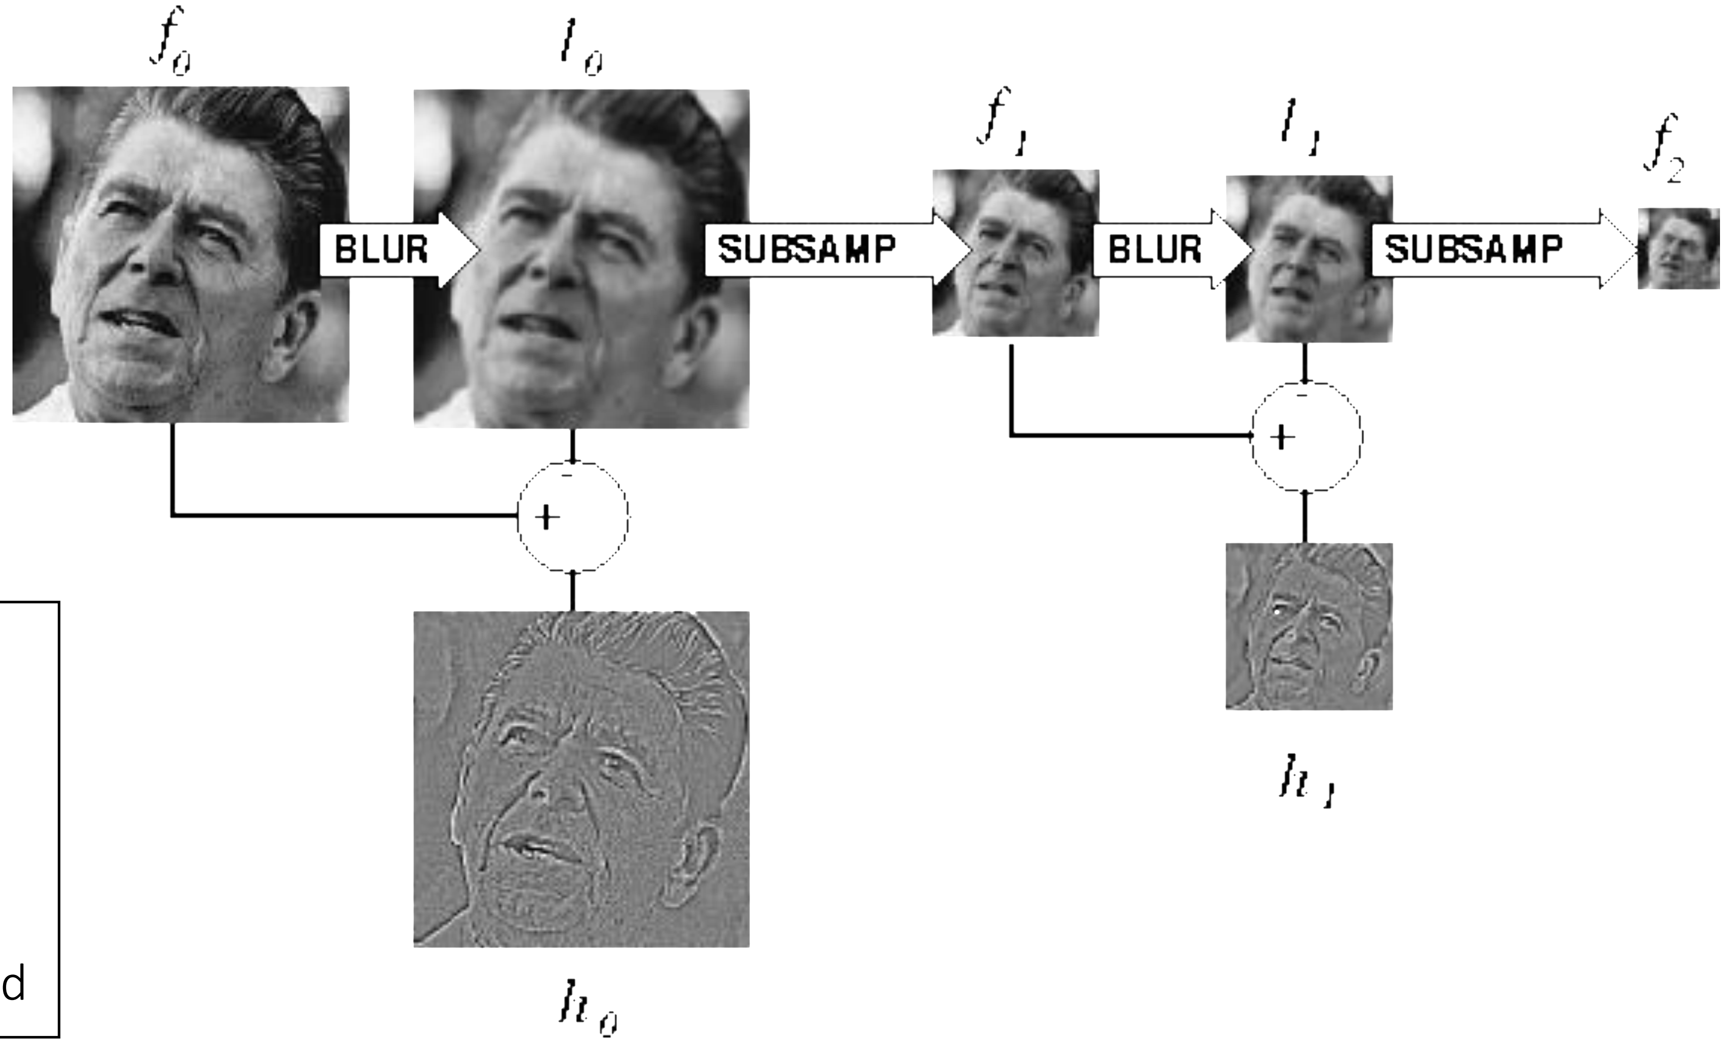

Algorithm

repeat:

filter

compute residual

subsample

until min resolution reached

Constructing a Laplacian pyramid

What is this part?

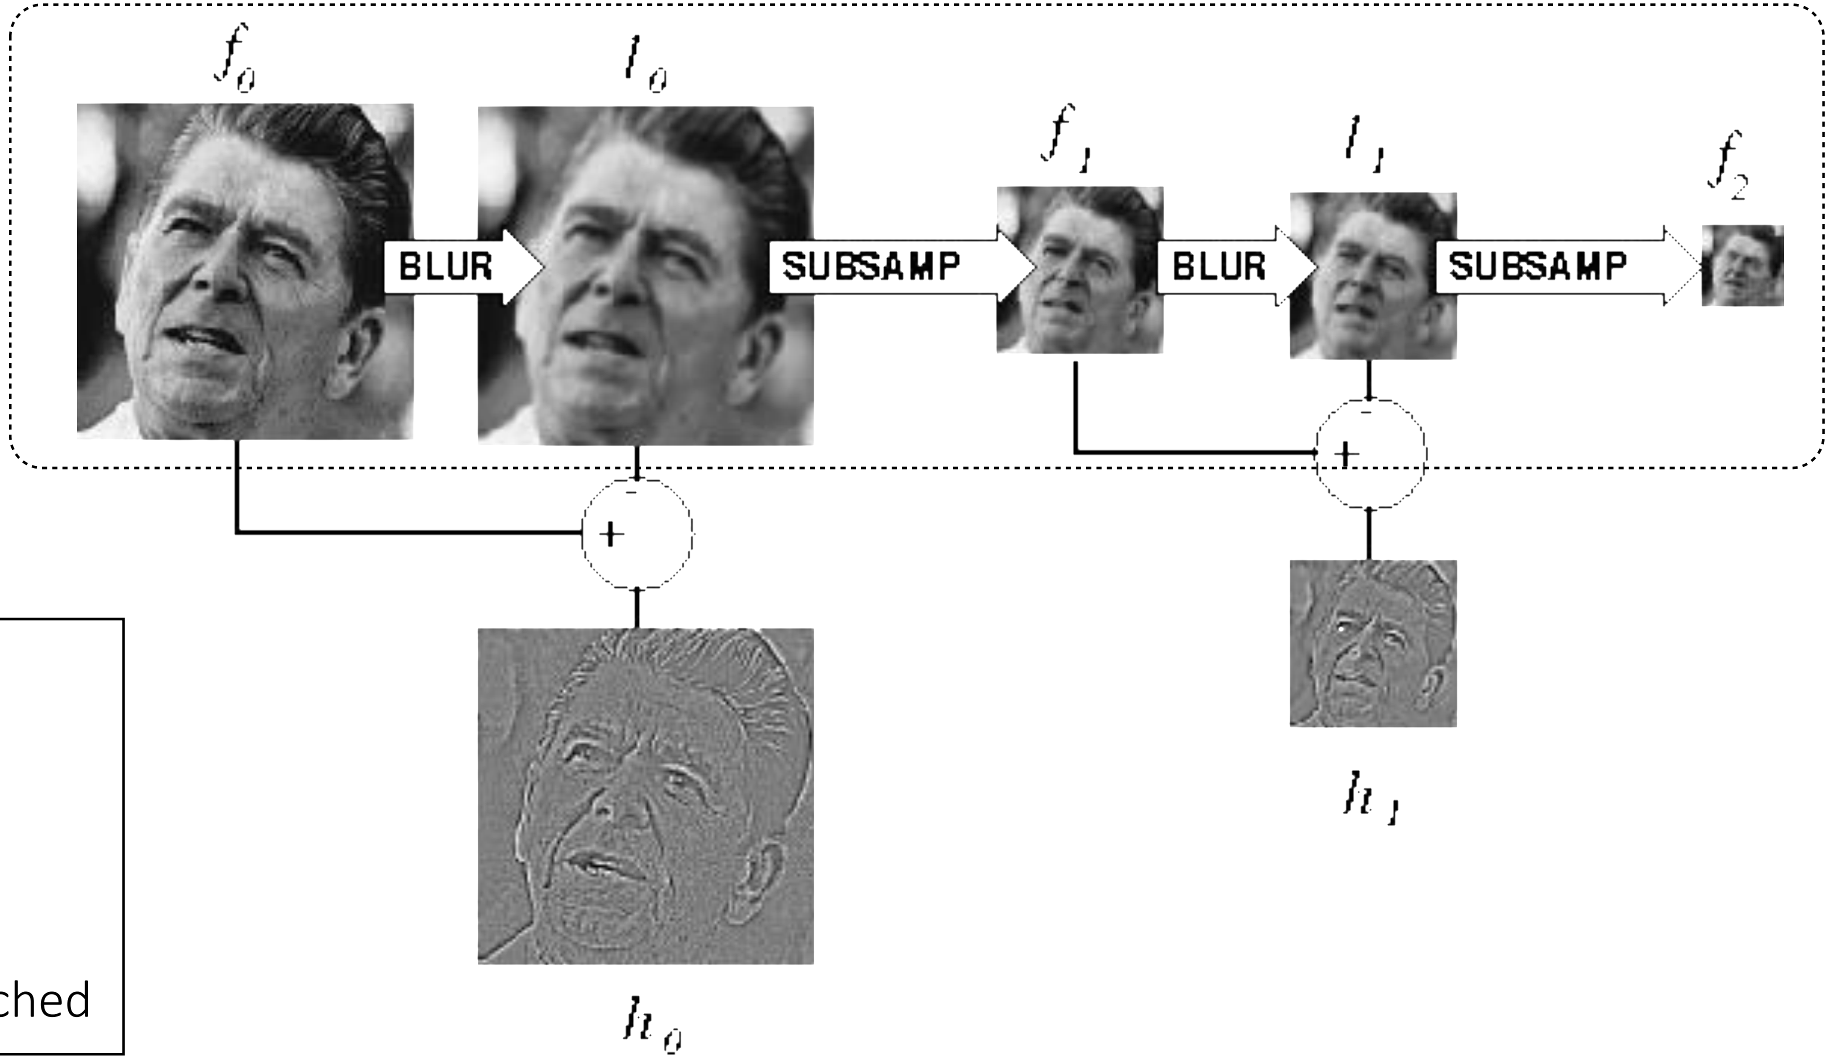

Algorithm

repeat:

filter

compute residual

subsample

until min resolution reached

Constructing a Laplacian pyramid

It's a Gaussian pyramid.

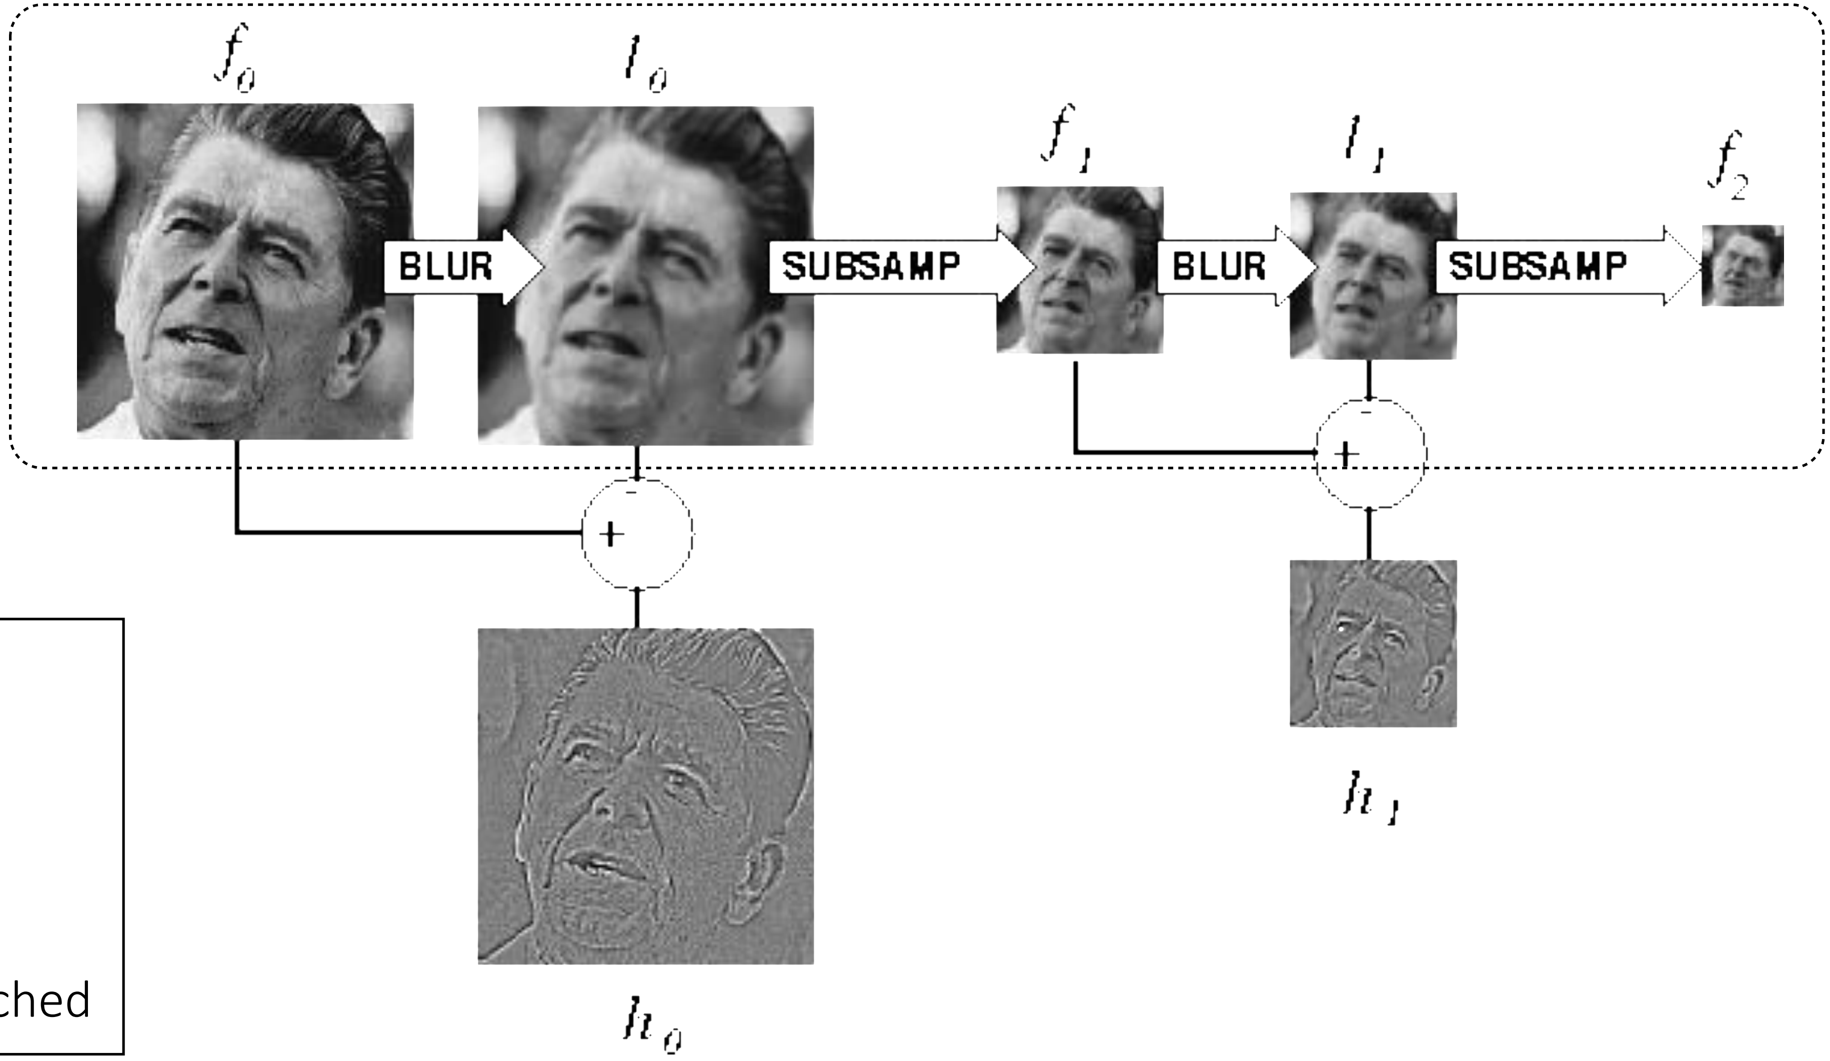

Algorithm

repeat:

filter

compute residual

subsample

until min resolution reached

What do we need to construct the original image?

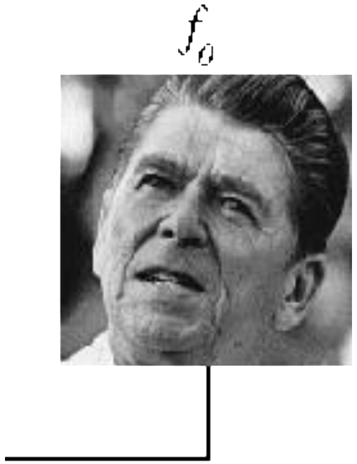

What do we need to construct the original image?

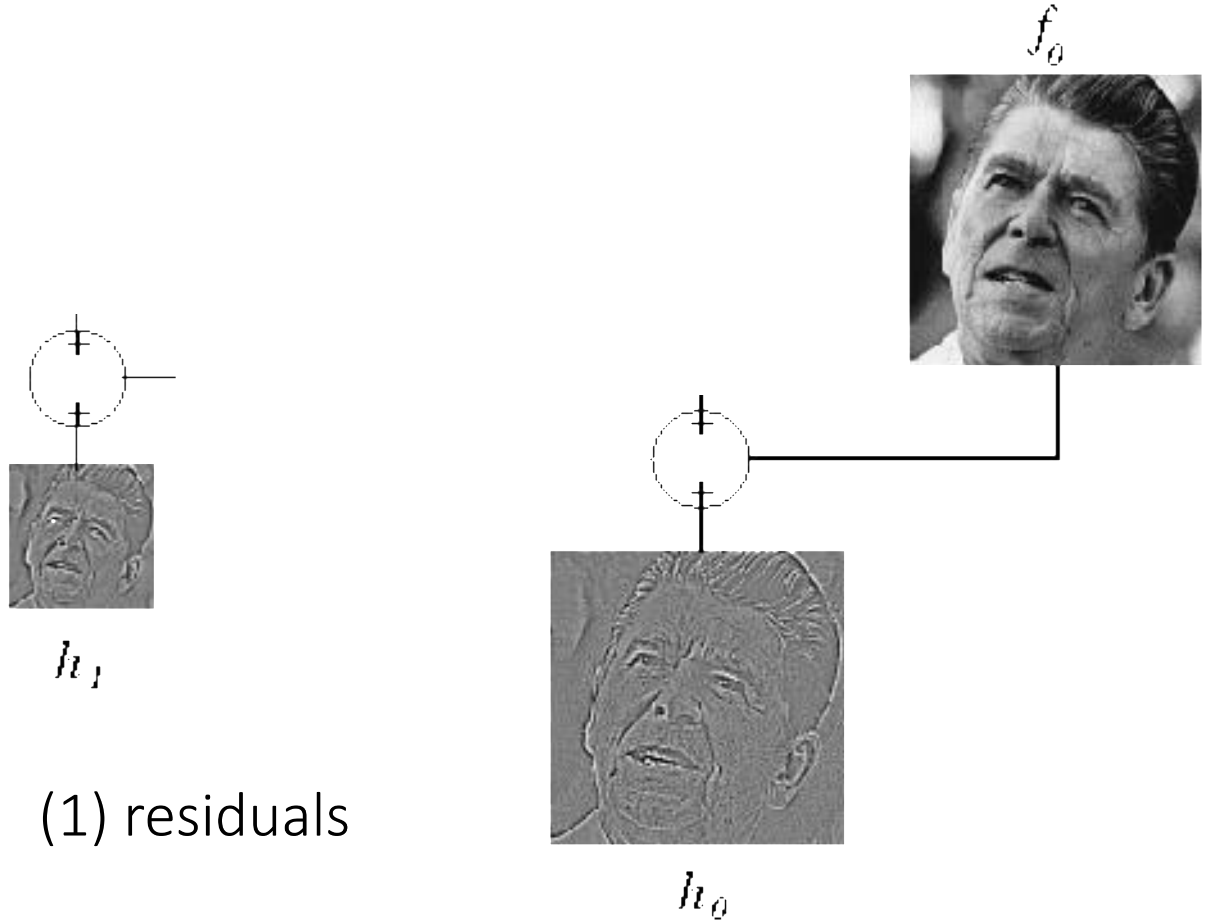

(1) residuals

What do we need to construct the original image?

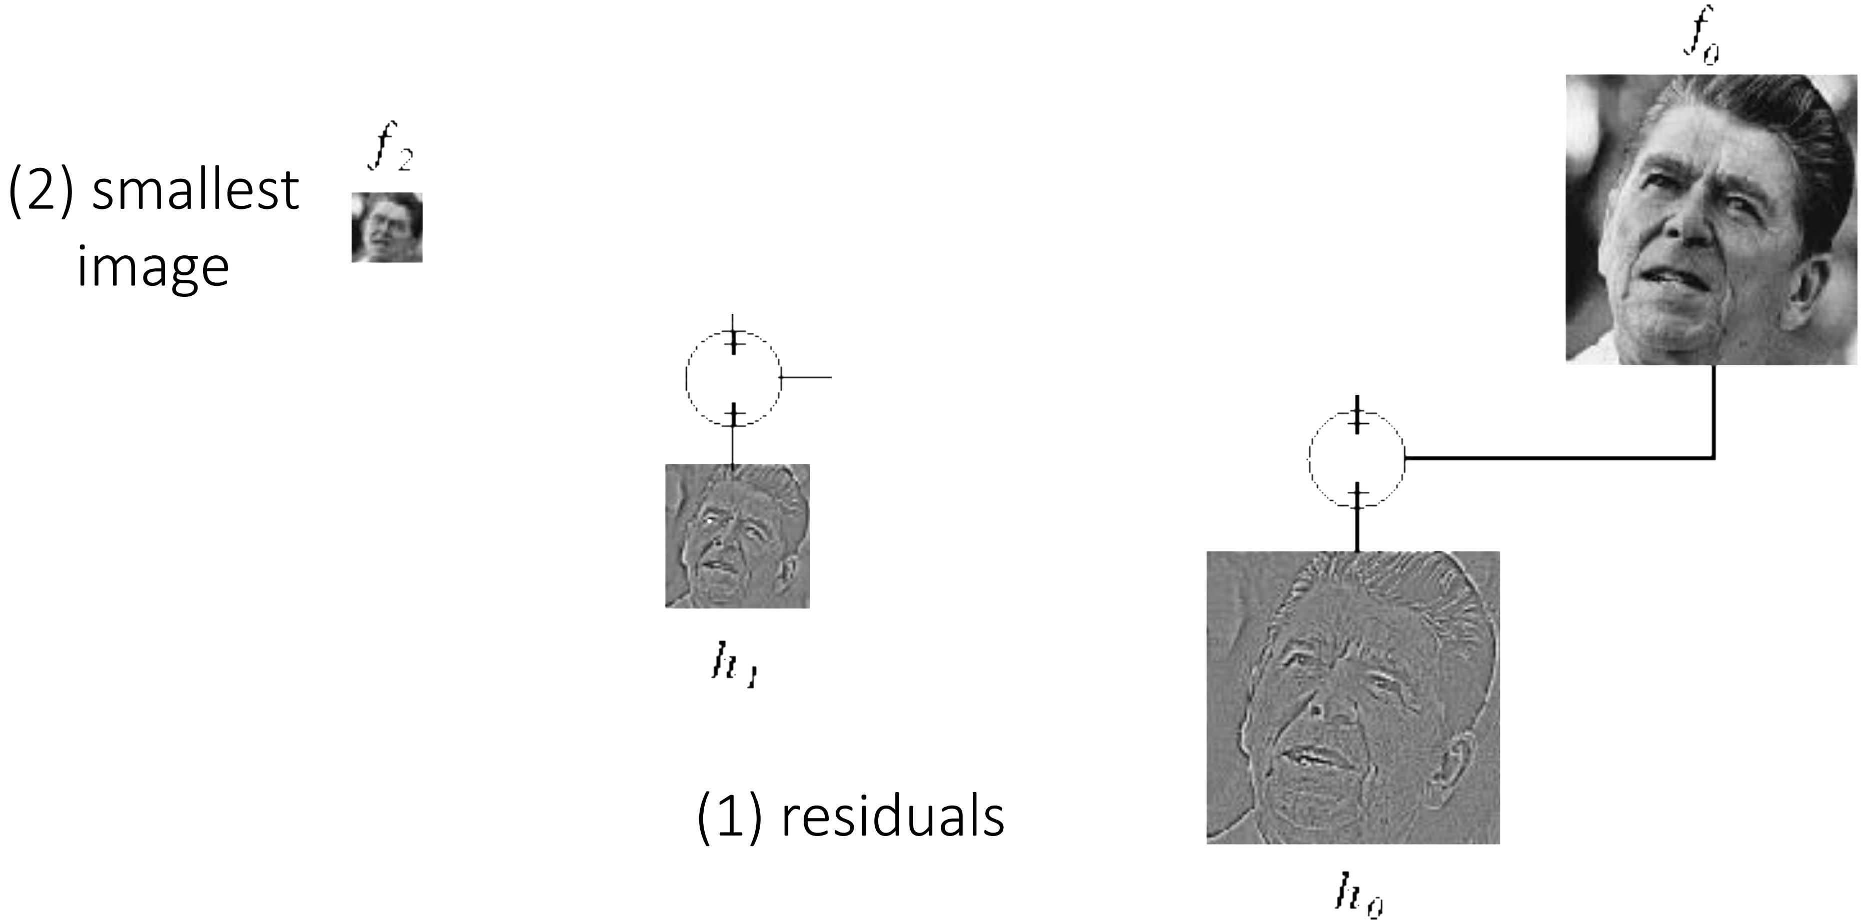

Reconstructing the original image

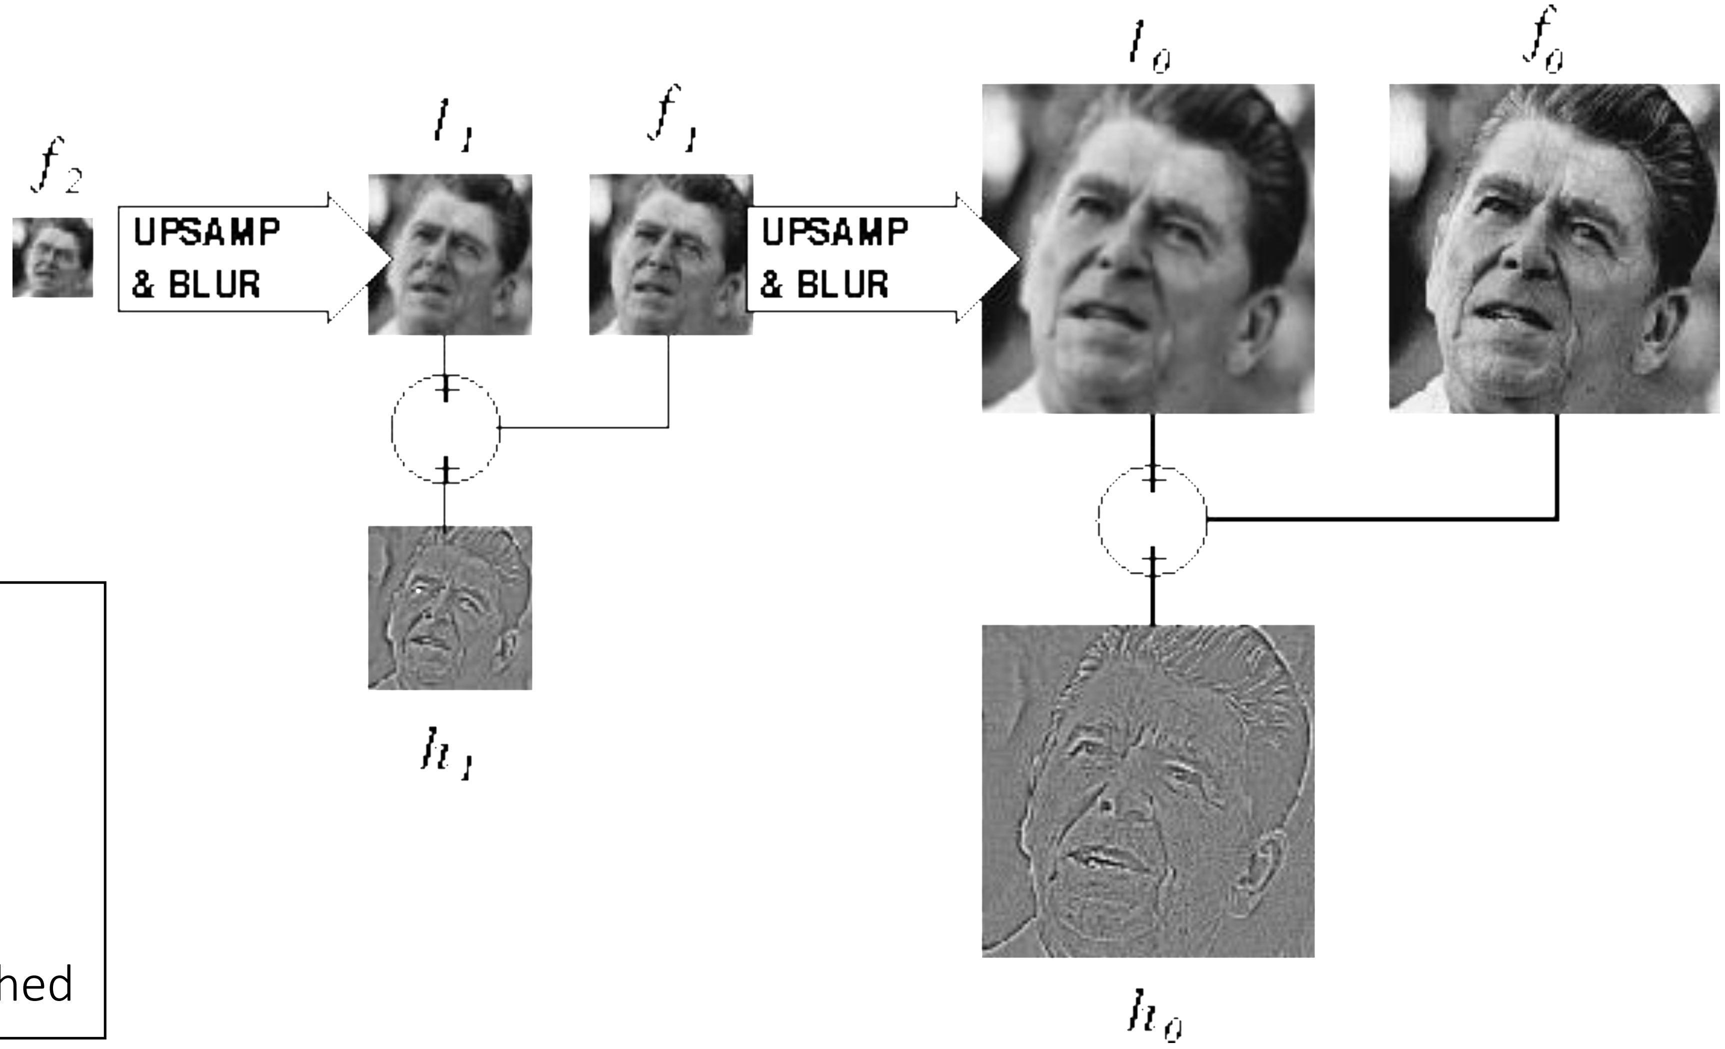

Algorithm

repeat:

upsample

sum with residual

until orig resolution reached

Gaussian vs Laplacian Pyramid

Shown in opposite order for space.

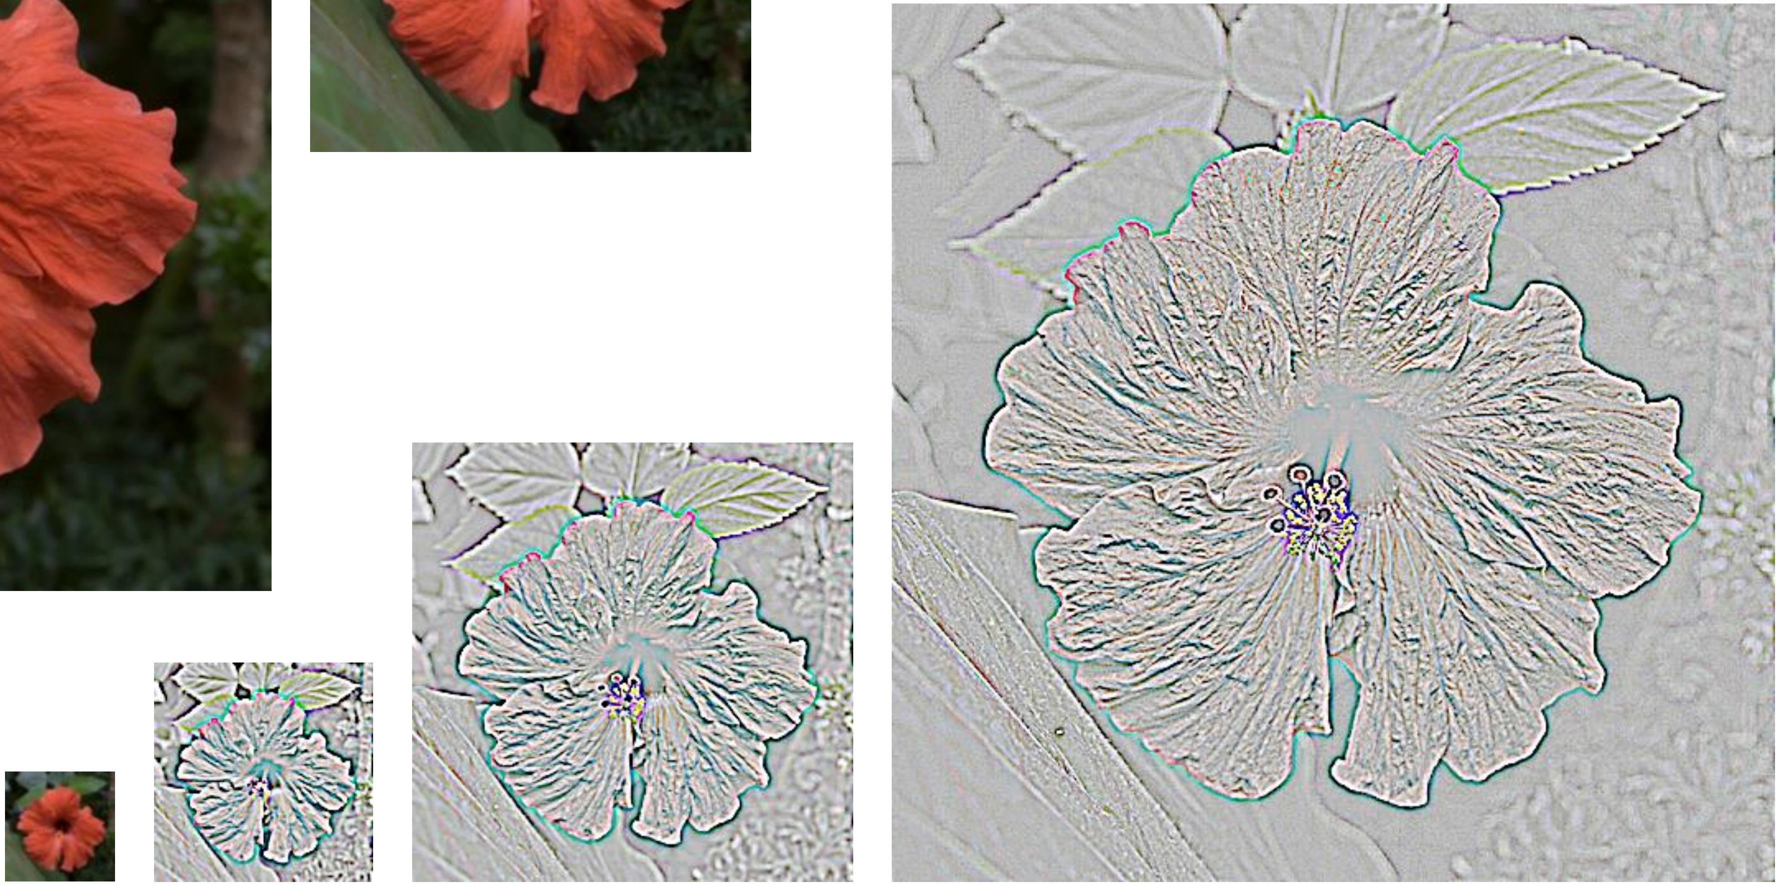

Which one takes more space to store?

Why is it called a Laplacian pyramid?

Reminder: Laplacian of Gaussian (LoG) filter

As with derivative, we can combine Laplace filtering with Gaussian filtering

input

Laplacian of Gaussian

output

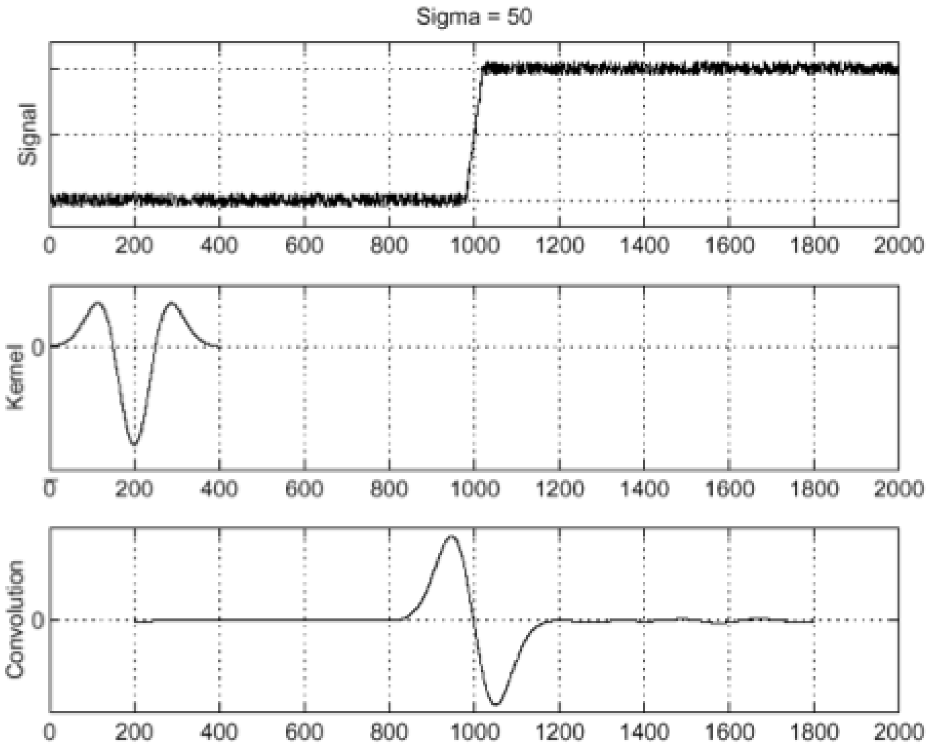

“zero crossings” at edges

Why is it called a Laplacian pyramid?

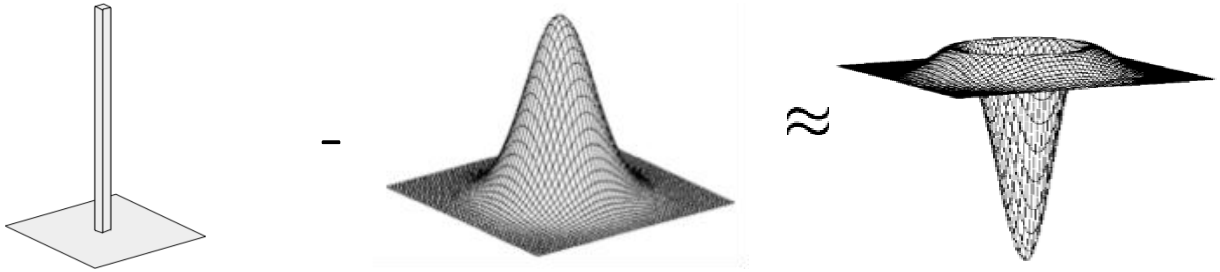

unit Gaussian Laplacian

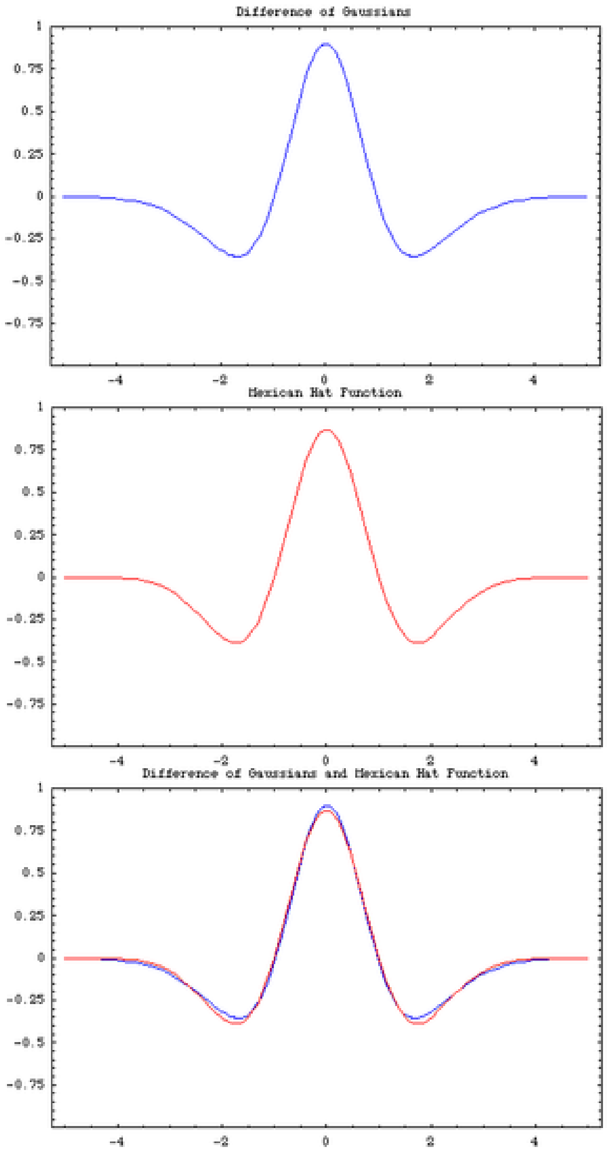

Difference of Gaussians approximates the Laplacian

Why Reagan?

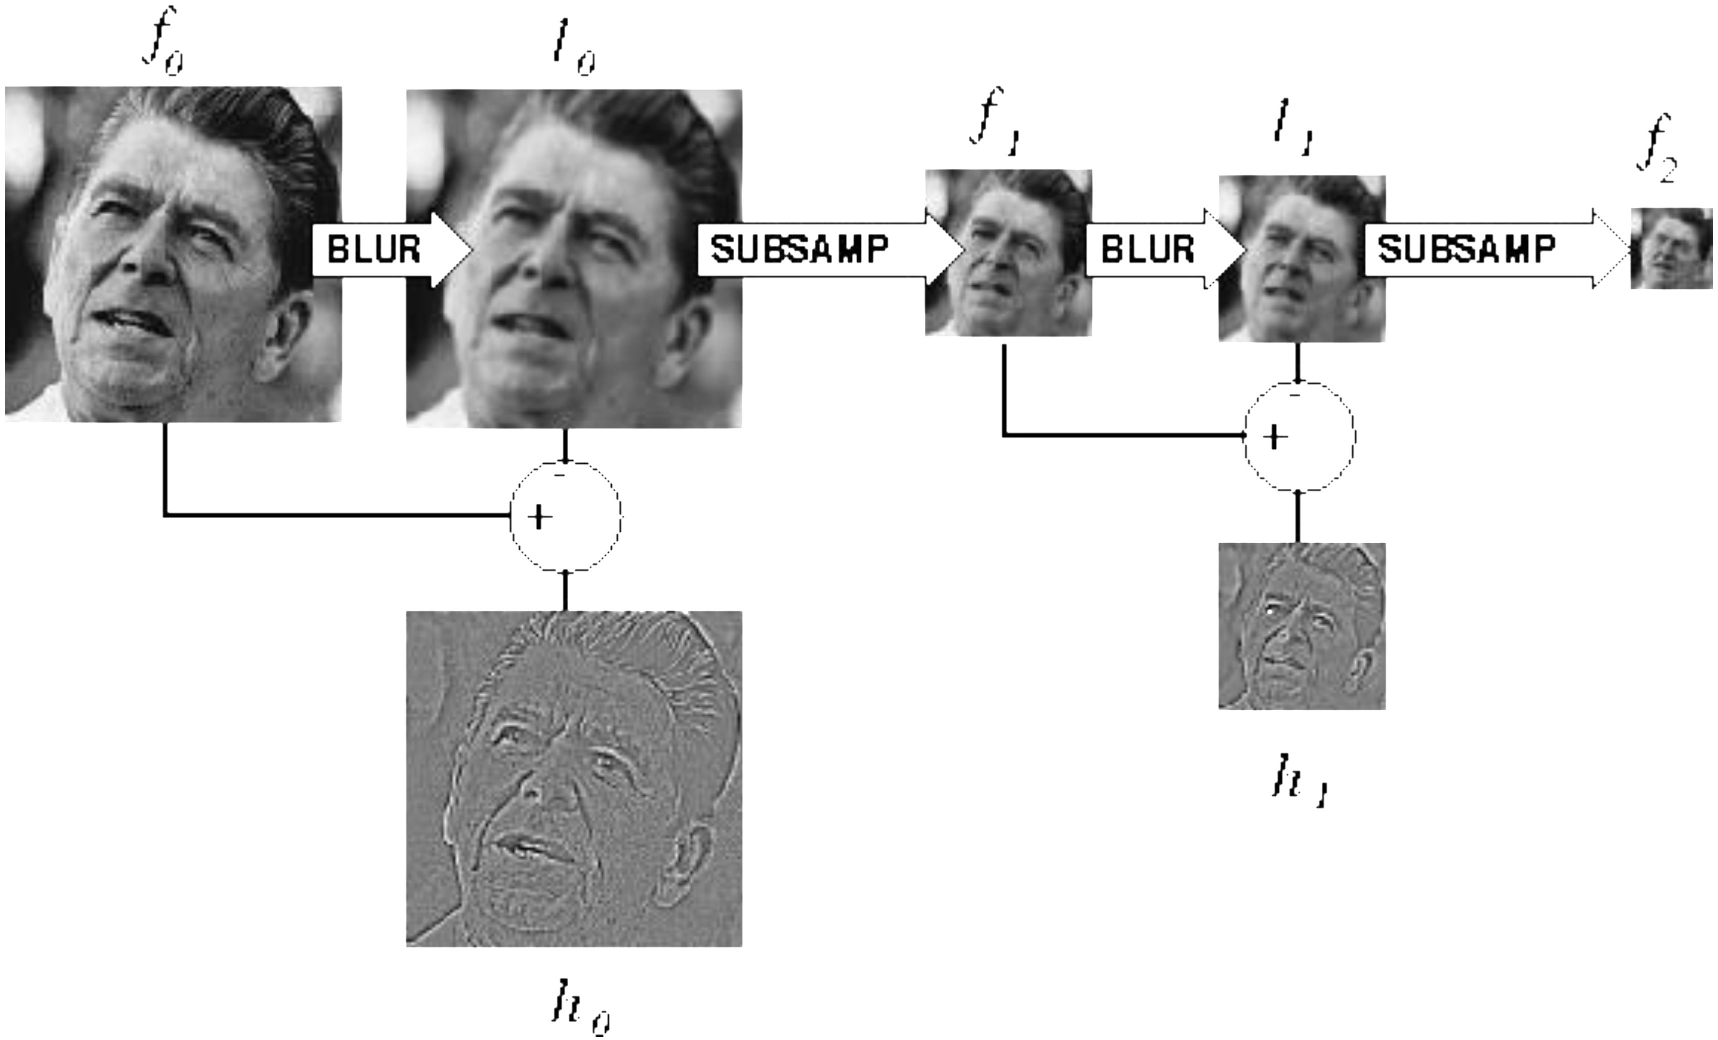

Why Reagan?

Ronald Reagan was President when the Laplacian pyramid was invented

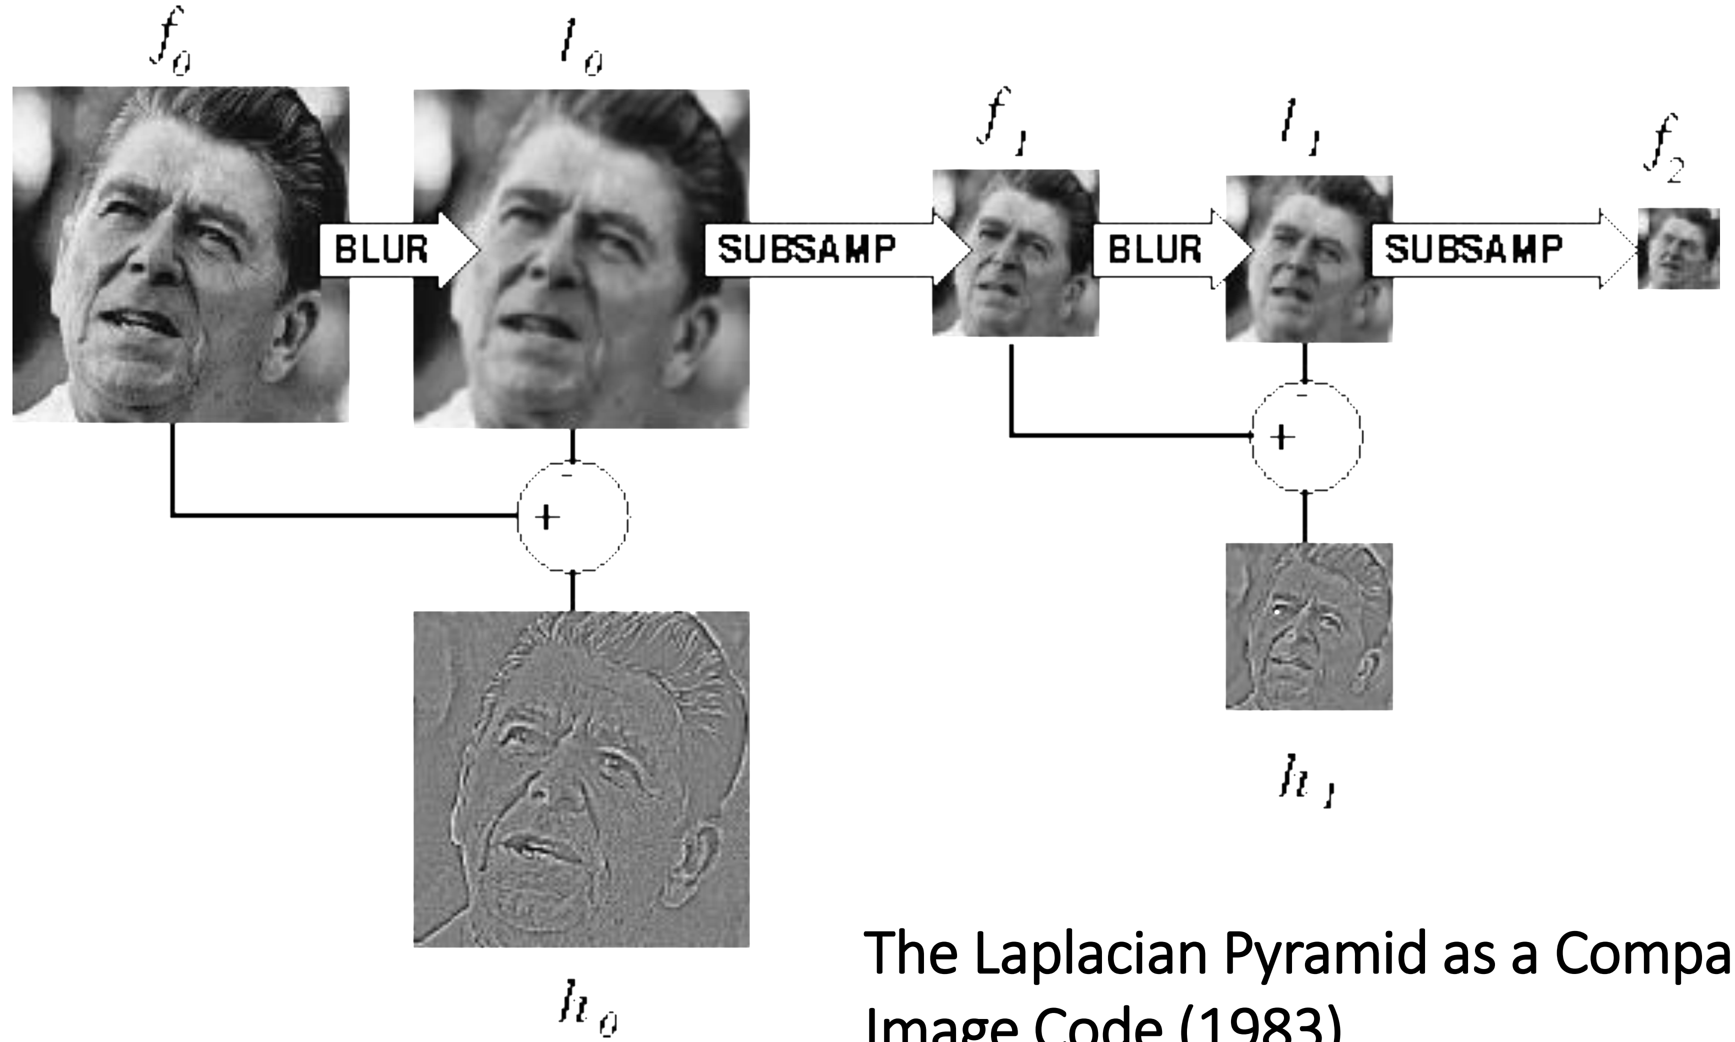

The Laplacian Pyramid as a Compact Image Code (1983)

Peter J. Burt , Edward H. Adelson

Still used extensively

Still used extensively

foreground details enhanced, background details reduced

input image

user-provided mask

Other types of pyramids

Steerable pyramid: At each level keep multiple versions, one for each direction.

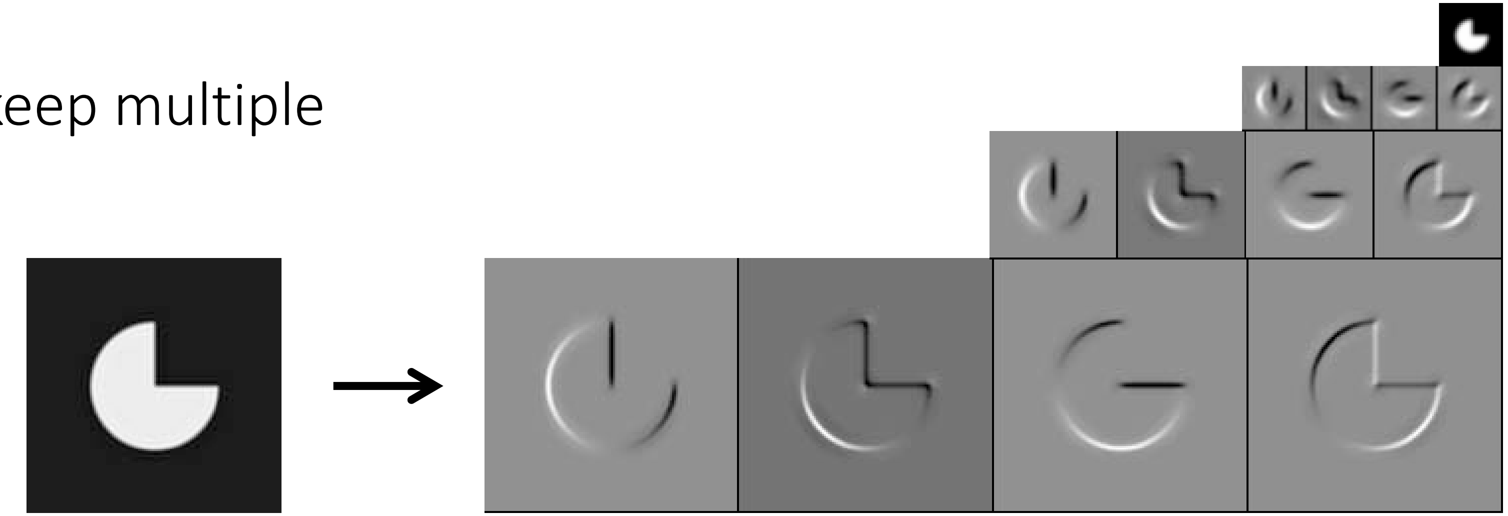

Wavelets: Huge area in image processing (see 18-793).

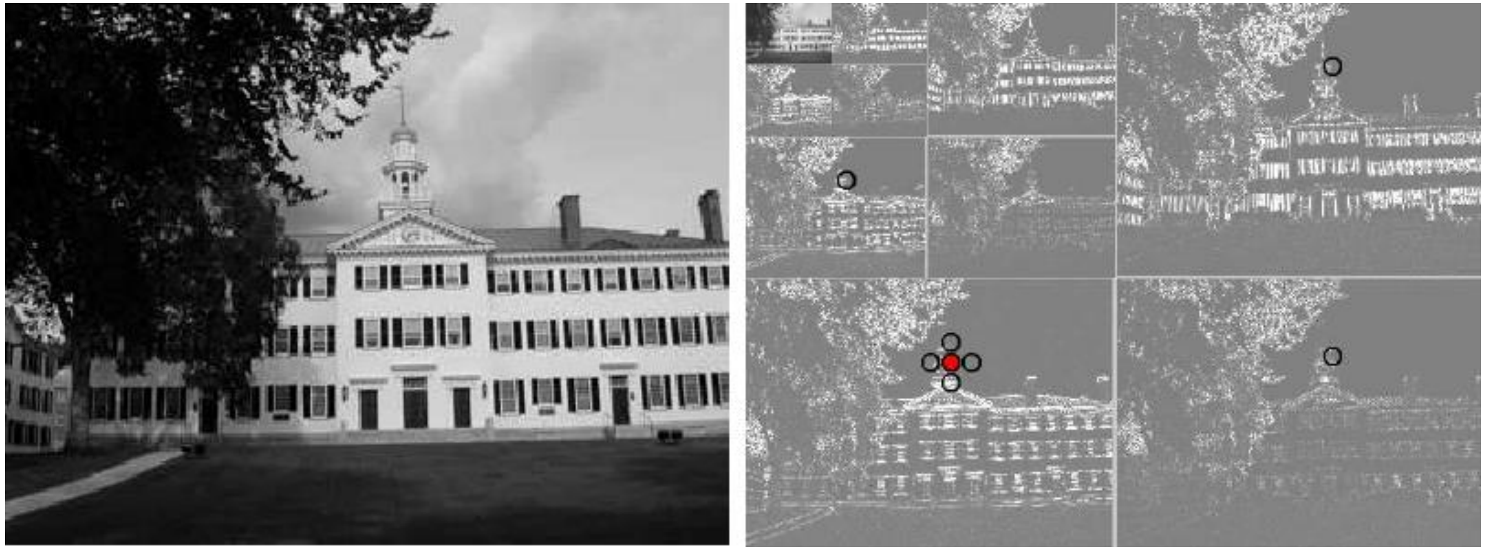

What are image pyramids used for?

image compression

multi-scale texture mapping

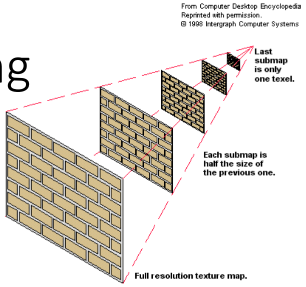

image blending

focal stack compositing

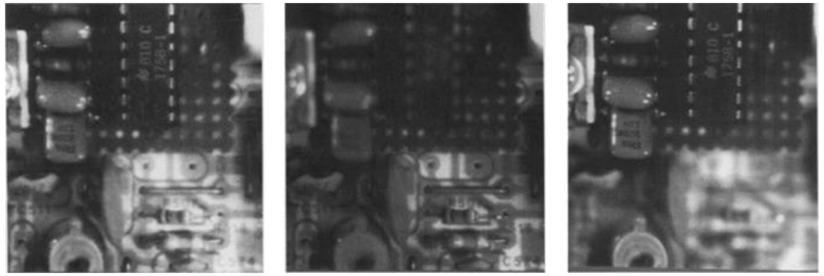

denoising

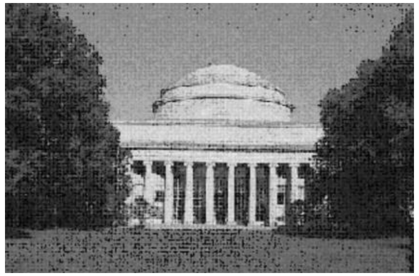

multi-scale detection

multi-scale registration

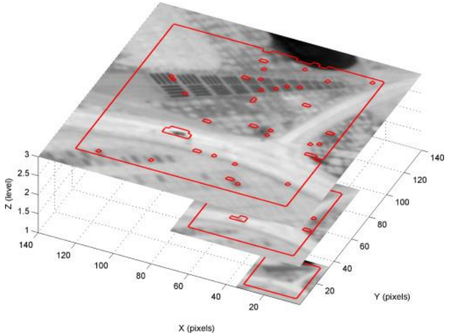

References

Basic reading:

- Szeliski textbook, Sections 3.5

Additional reading:

- Burt and Adelson, “The Laplacian Pyramid as a Compact Image Code,” IEEE ToC 1983.
  the original Laplacian pyramid paper