

Looking for seams



Course announcements

- Homework 7 will be posted and will be due on Sunday 6th.
 - You can use all of your remaining late days for it.
- RI Seminar this week: Vladlen Koltun, “Learning to drive”, Friday 3:30-4:30pm.
 - Very exciting speaker, you should all attend.
 - Make sure to go early as the room will be packed.
 - Do you want me to move my office hours so that you can make it to the talk?

Overview of today's lecture

- Segmentation.
- Image as a graph.
- Shortest graph paths and Intelligent scissors.
- Graph-cuts and GrabCut.
- Some notes about cutting-and-pasting.

Slide credits

Most of these slides were adapted from:

- Kris Kitani (15-463, Fall 2016).

Some slides were inspired or taken from:

- Fredo Durand (MIT).
- James Hays (Georgia Tech).

Course overview

1. Image processing.



Lectures 1 – 7

See also 18-793: Image and Video Processing

2. Geometry-based vision.



Lectures 7 – 12

See also 16-822: Geometry-based Methods in Vision

3. Physics-based vision.



Lectures 13 – 16

See also 16-823: Physics-based Methods in Vision

See also 15-463: Computational Photography

4. Semantic vision.



Lectures 17 – 21

See also 16-824: Vision Learning and Recognition

5. Dealing with motion.



Lectures 22 – 25

See also 16-831: RoboStats

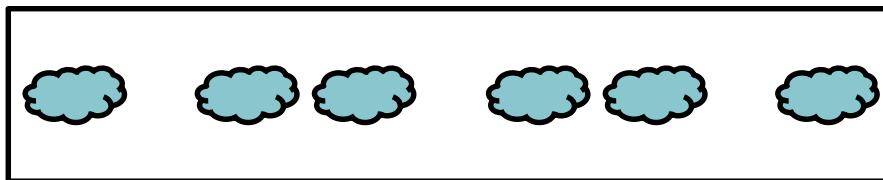
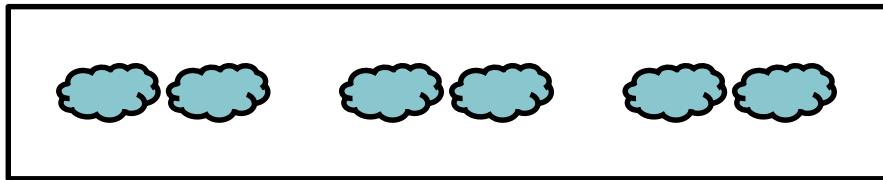
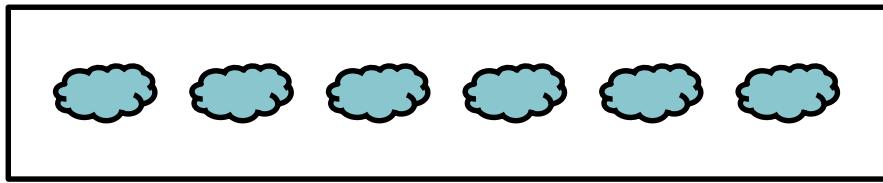
Segmentation

Gestalt Psychology

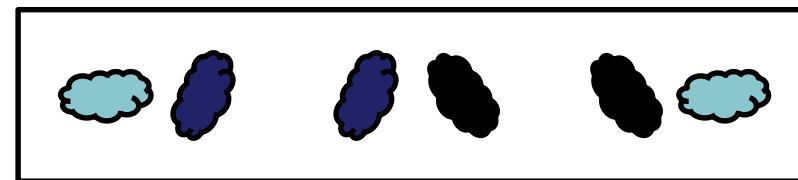
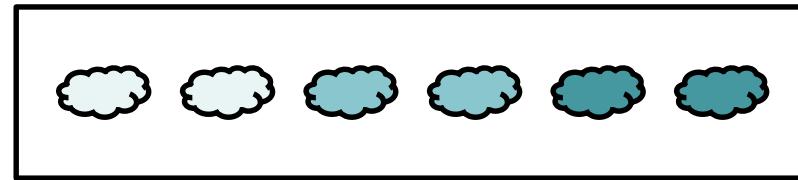
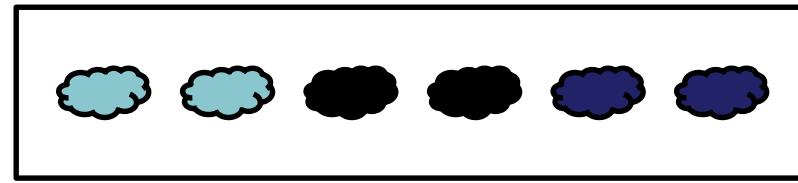


We perceive objects in their entirety before their individual parts.

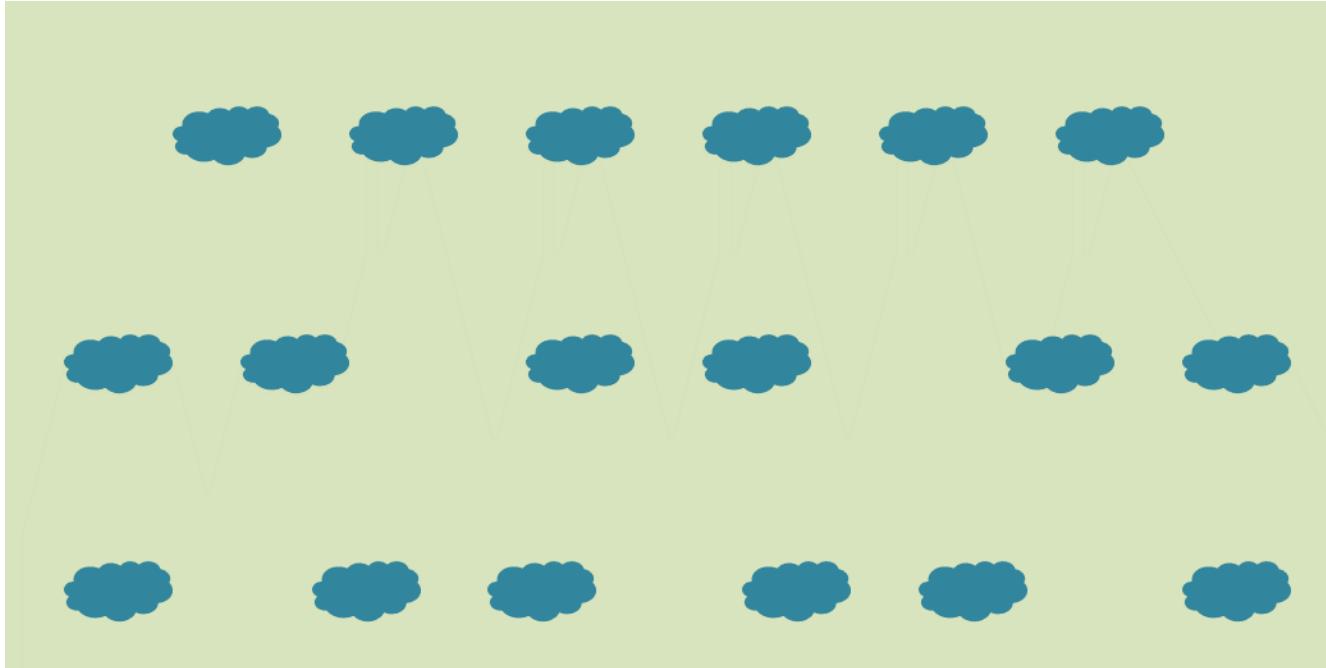
Closer objects are grouped together



Similar objects are grouped together

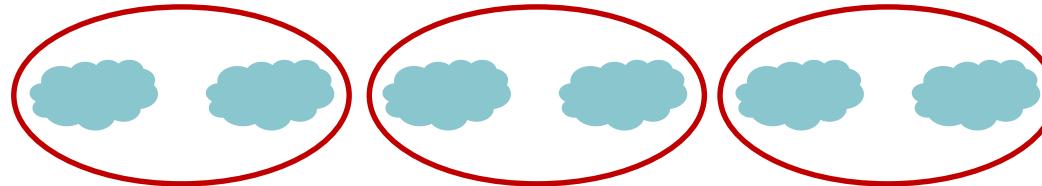


Common Fate



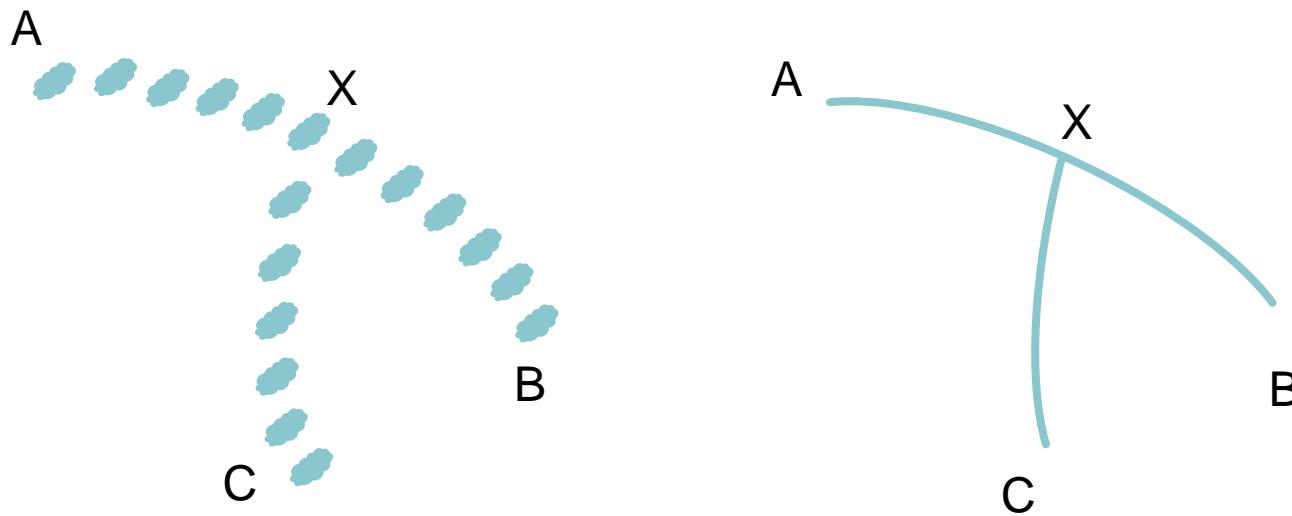
Objects with similar motion or change in appearance are grouped together

Common Region/Connectivity



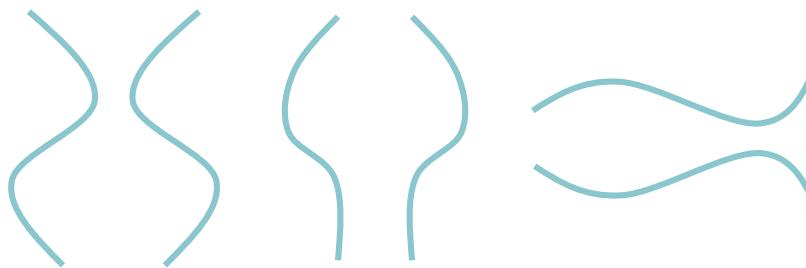
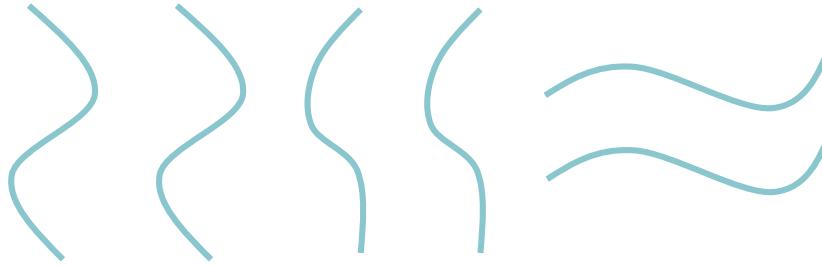
Connected objects are grouped together

Continuity Principle

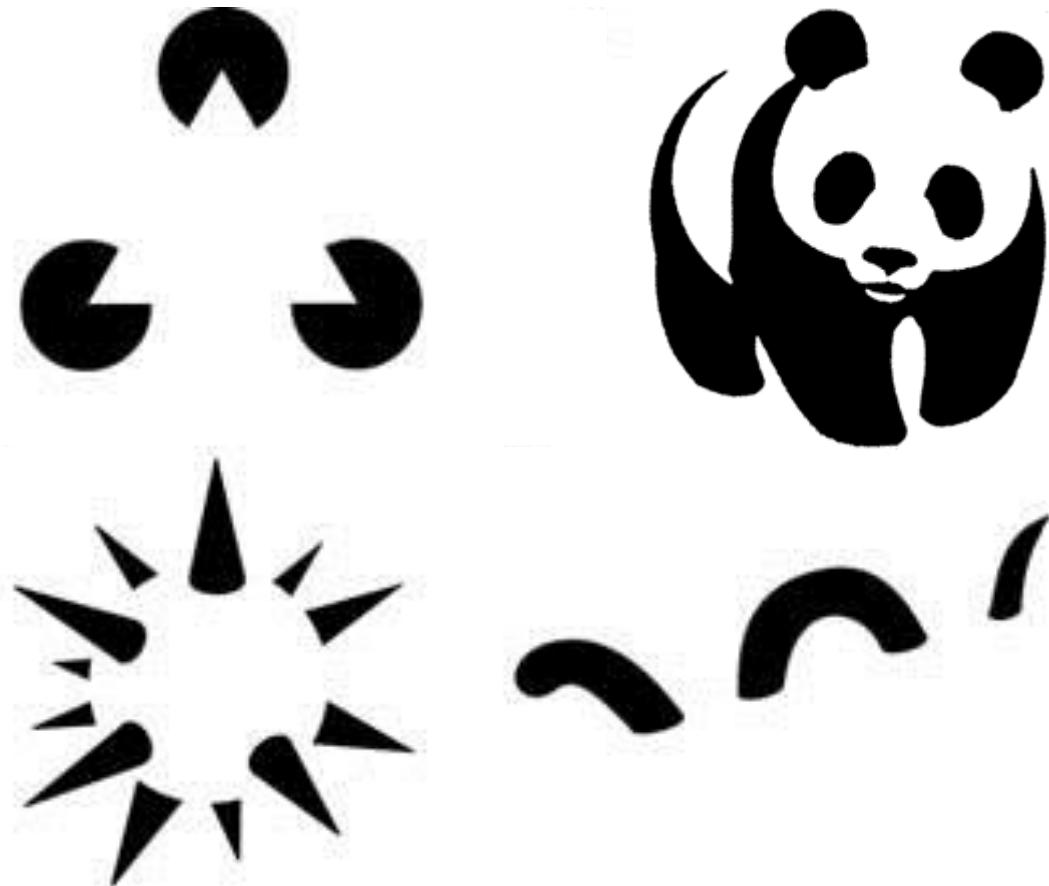


Features on a continuous curve are grouped together

Symmetry Principle

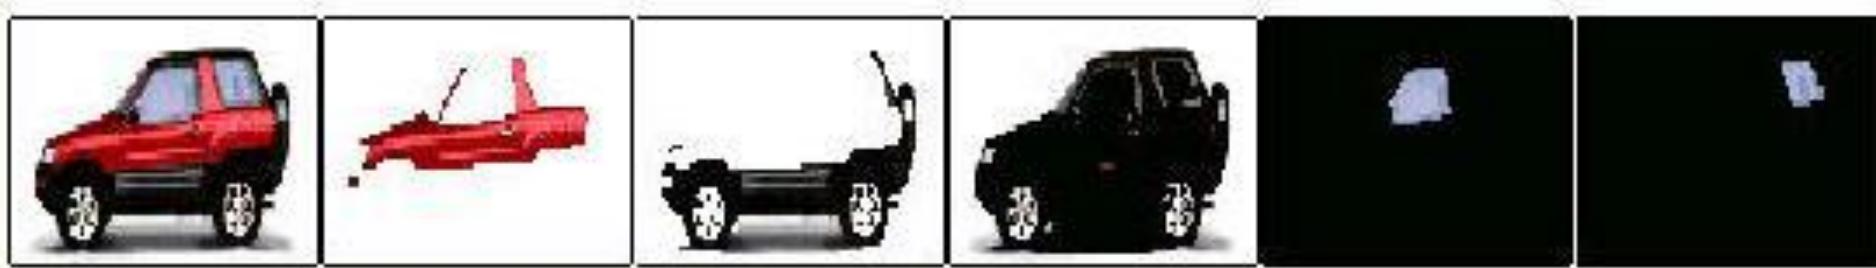


Completion



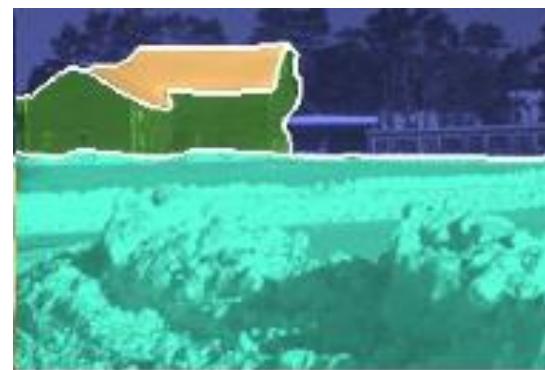
Illusory or subjective contours are perceived

Segmentation/Clustering





$k = 4$



$nc = .0017$



$k = 5$

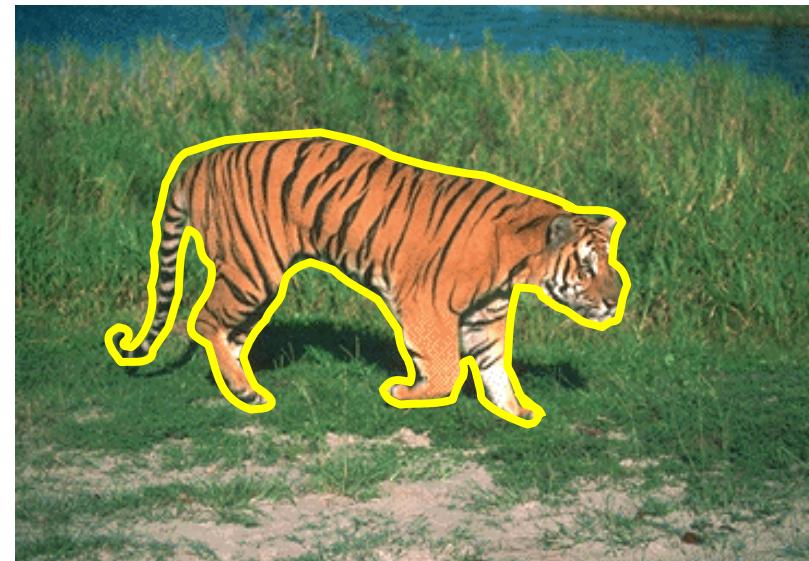
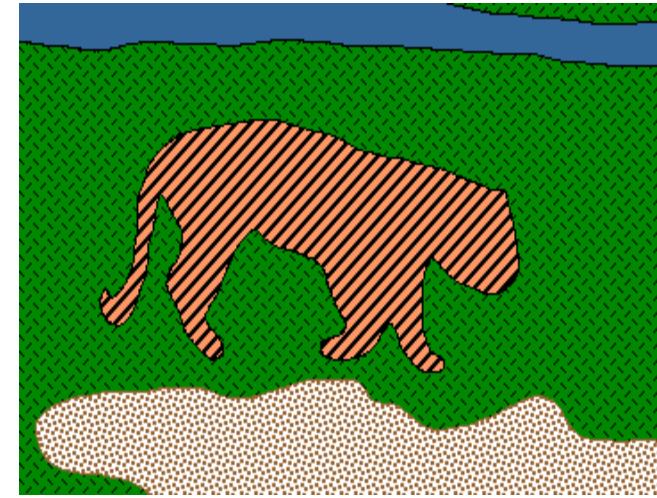


$nc = .0060$



$k = 11$

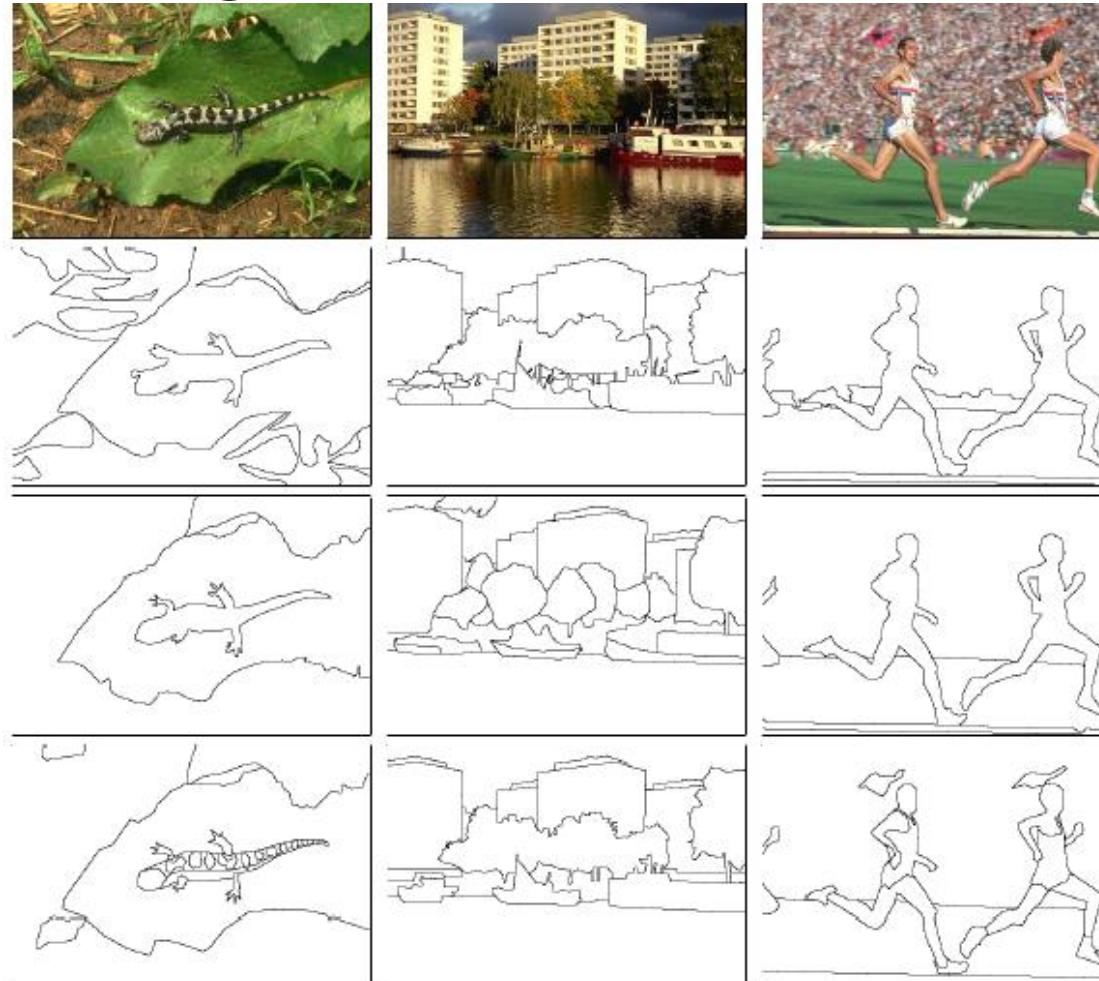




What is a “good” segmentation??

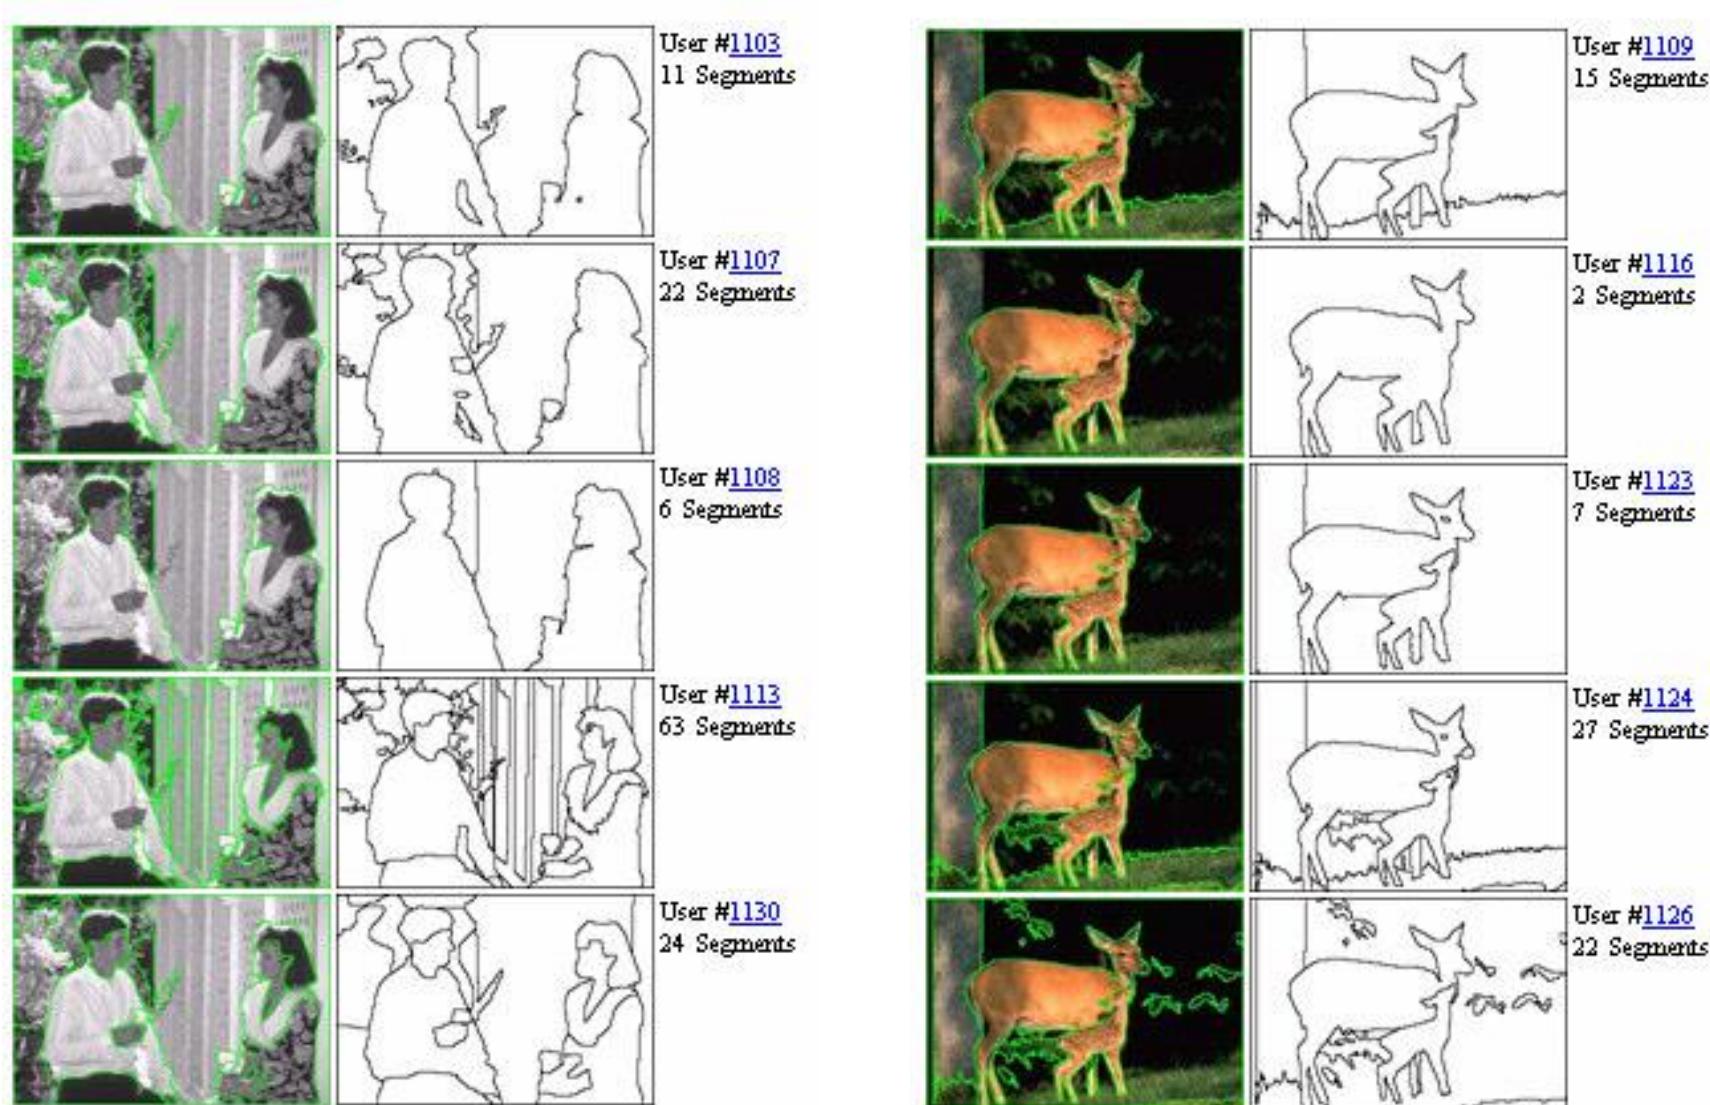
First idea: Compare to human segmentation or to “ground truth”

No objective
definition of
segmentation!



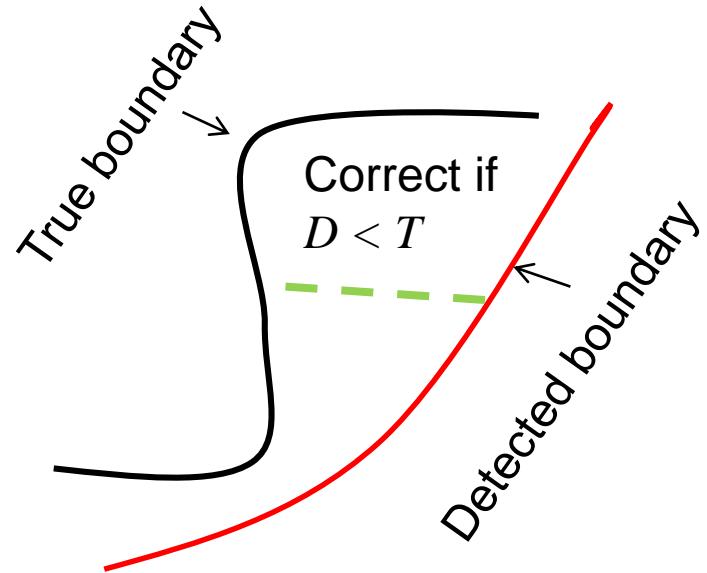
- <http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/resources.html>

No objective definition of segmentation!

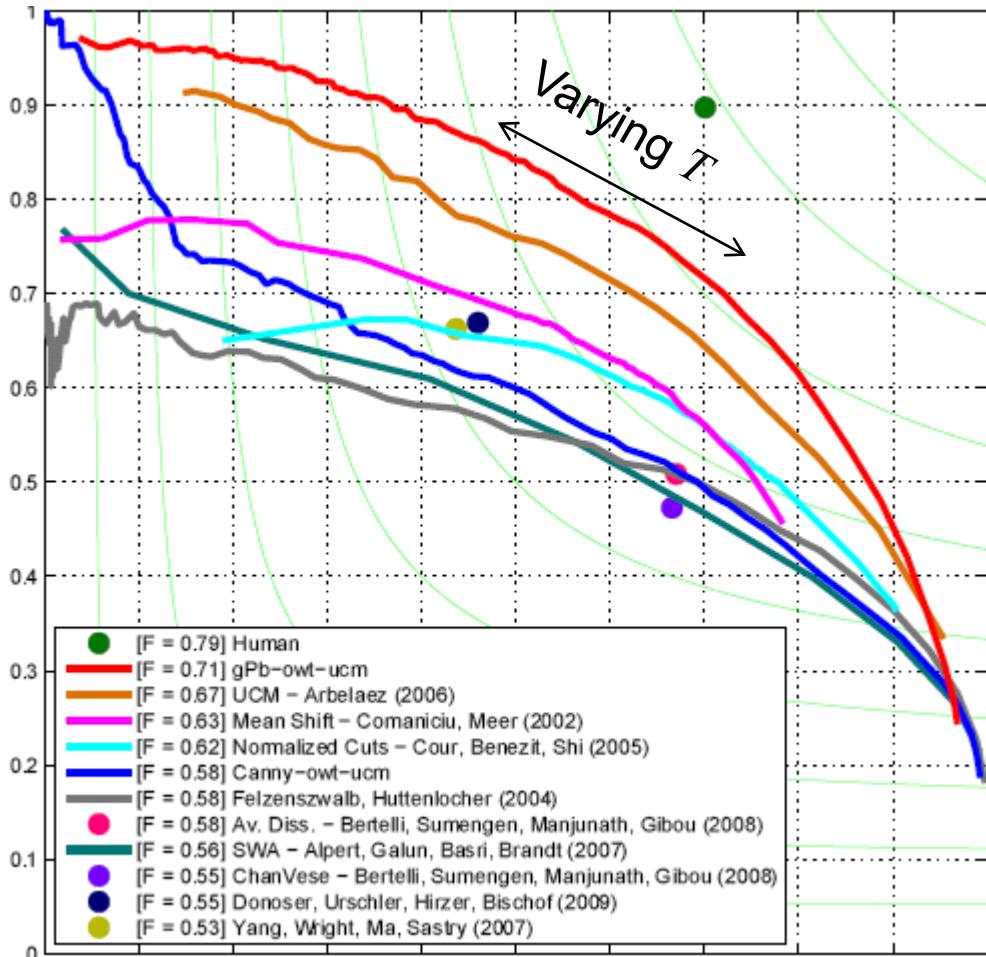


- <http://www.eecs.berkeley.edu/Research/Projects/CS/vision/bds/BSDS300/html/dataset/images/color/317080.html>

Evaluation: Boundary agreement



Precision = % of detected boundary pixels that are correct



Recall = % of boundary pixels that are detected

Evaluation: Region overlap with ground truth



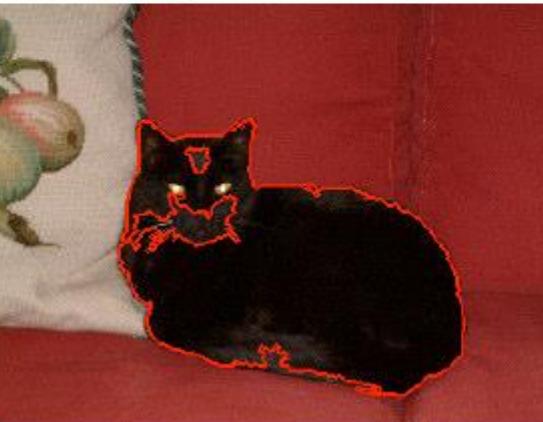
Ground Truth



Segment #1

.825

$$OS(S, G) = \frac{|S \cap G|}{|S \cup G|}$$

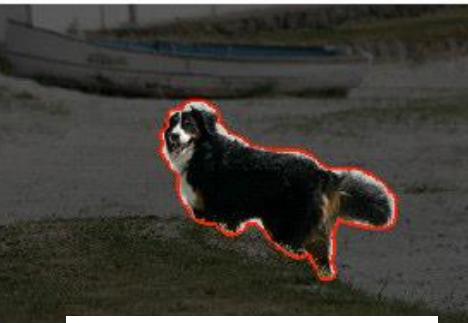


Segment #2

.892

Evaluation: Region overlap with ground truth

Ground truth



Mean shift



Graph-based



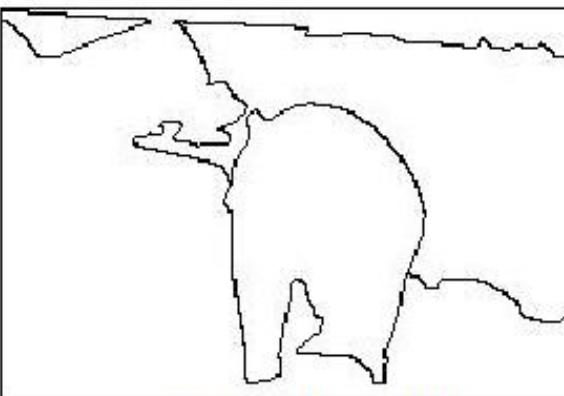
Spectral



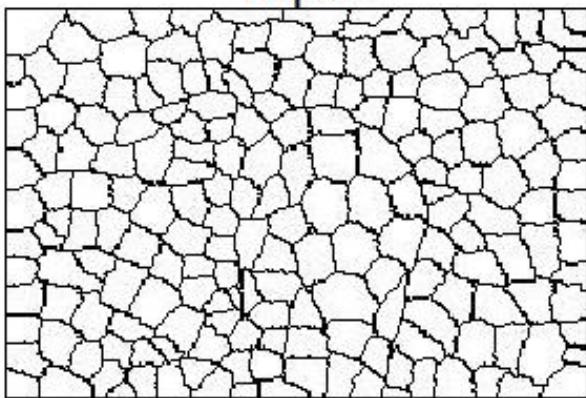
Second idea: Superpixels



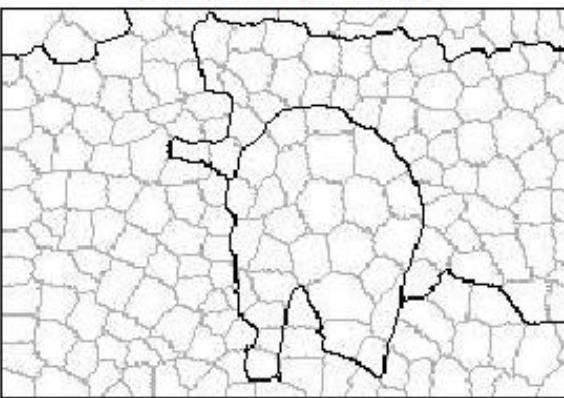
Input



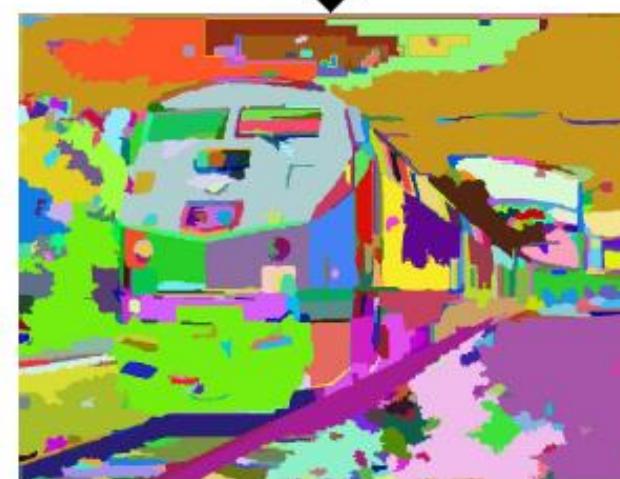
Ground truth



Superpixels

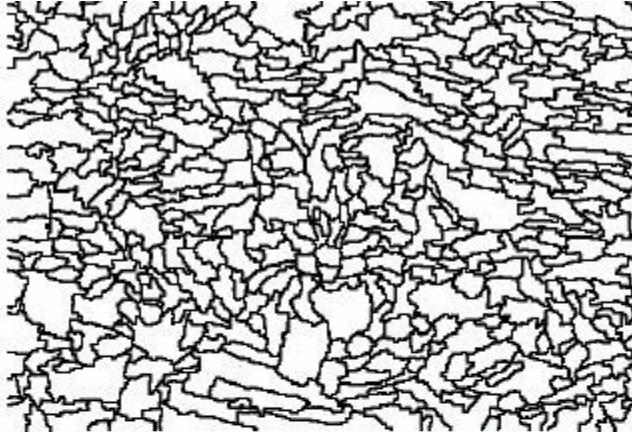


Overlay



- Let's not even try to compute a “correct” segmentation
- Let's be content with an *oversegmentation* in which each region is very likely (formal guarantees are hard) to be uniform

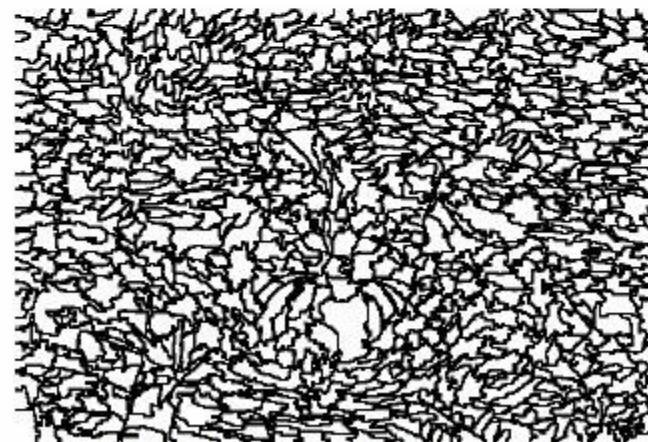
Second idea: Superpixels



Watershed



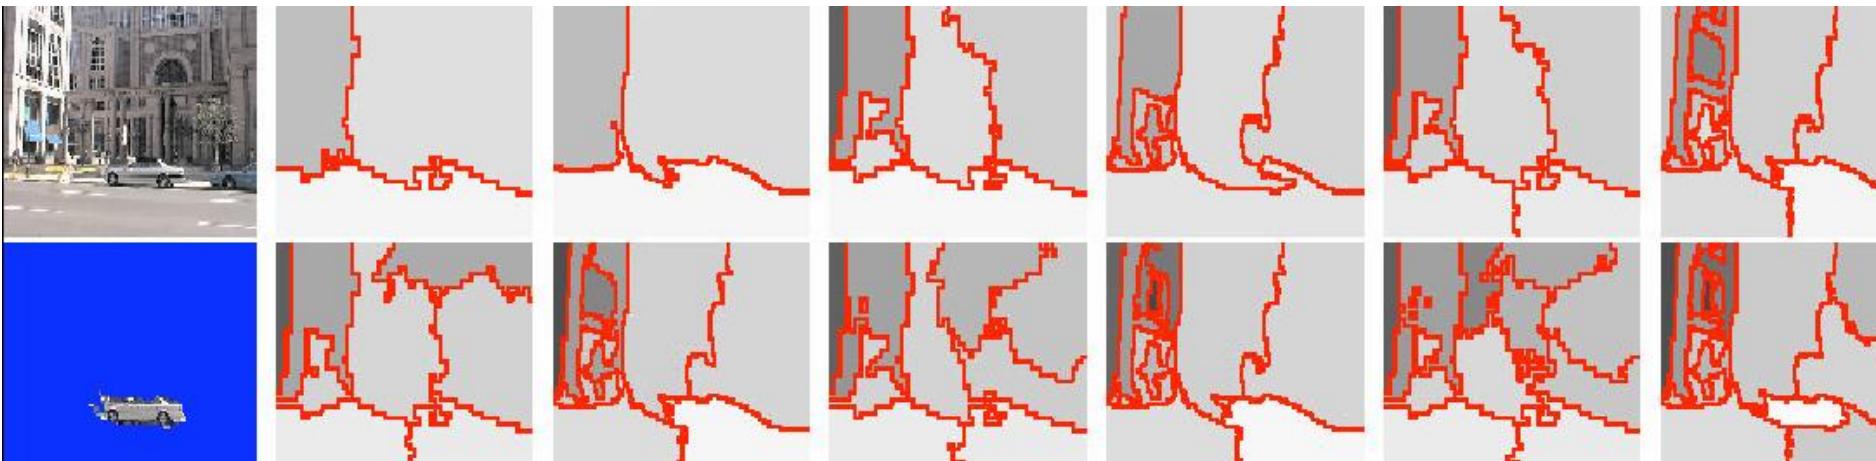
Mean shift



Graph-based

- Example from: How Do Superpixels Affect Image Segmentation?
- Progress in Pattern Recognition, Image Analysis and Applications. Springer LNCS. Volume 5197/2008.

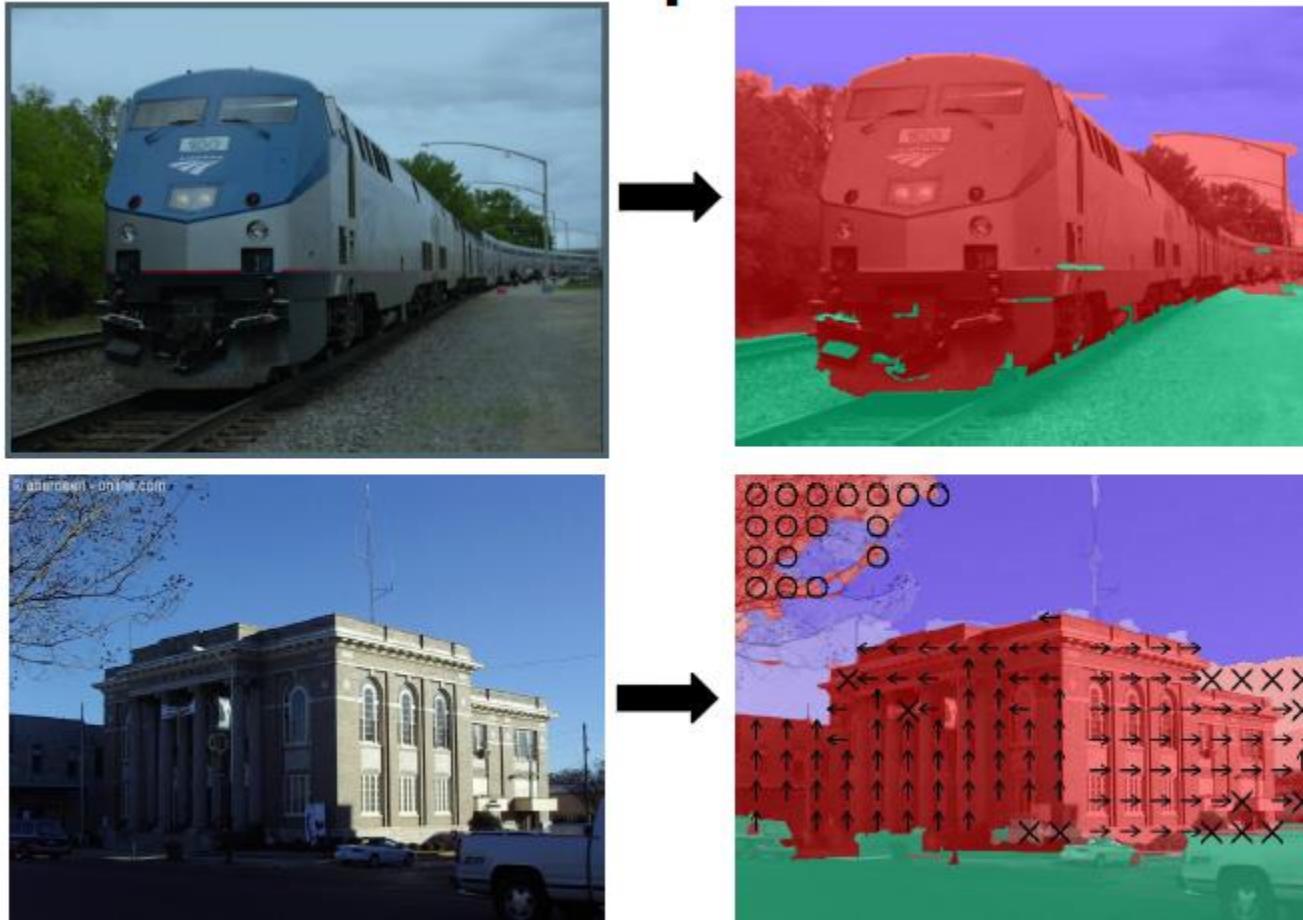
Third idea: Multiple segmentations



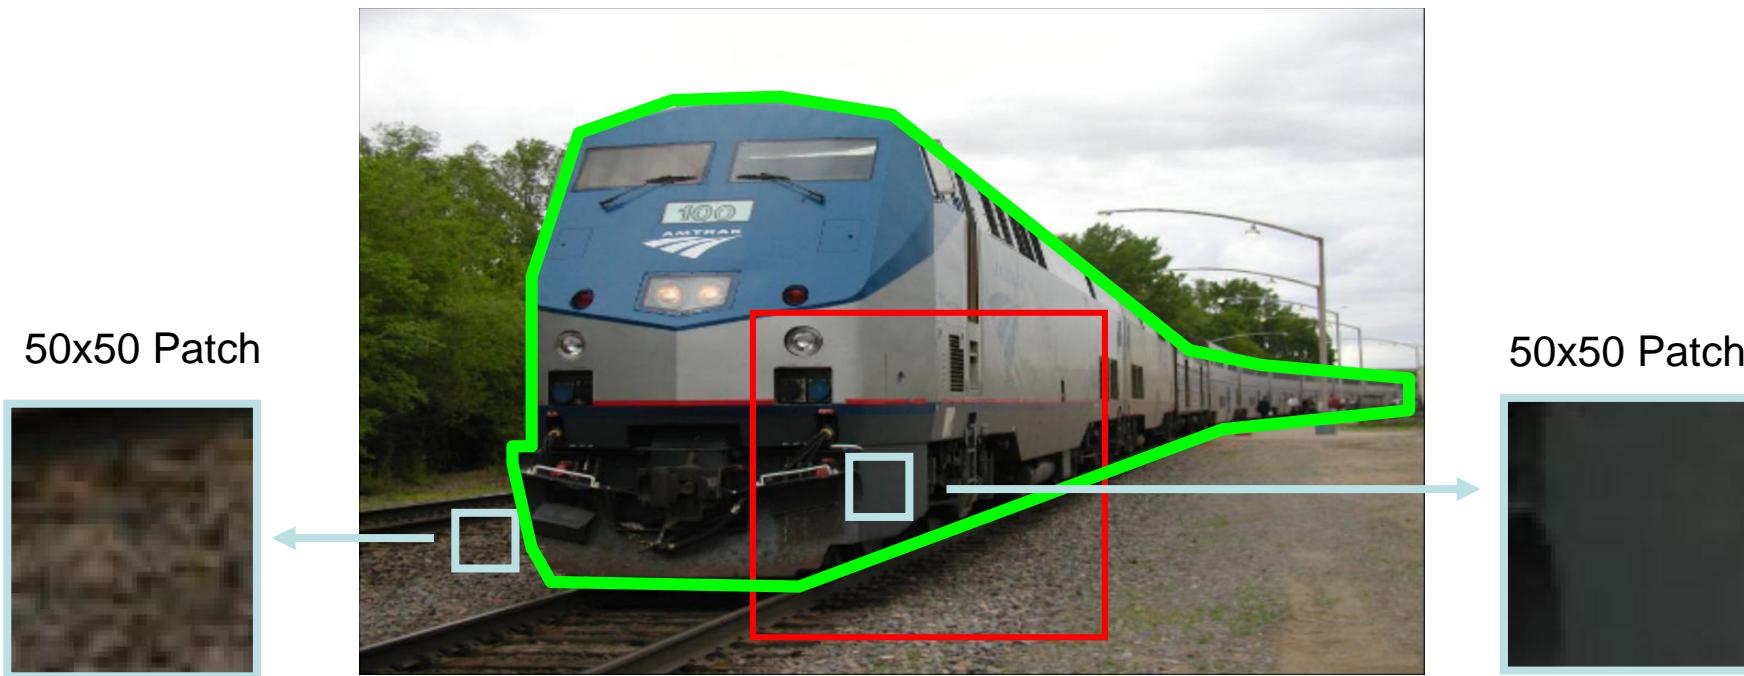
- Generate many segmentations of the same image
- Even though many regions are “wrong”, some consensus should emerge

Example: Improving Spatial Support for Objects via Multiple Segmentations
Tomasz Malisiewicz and Alexei A. Efros. British Machine Vision Conference (BMVC), September, 2007.

Multiple segmentations: Example

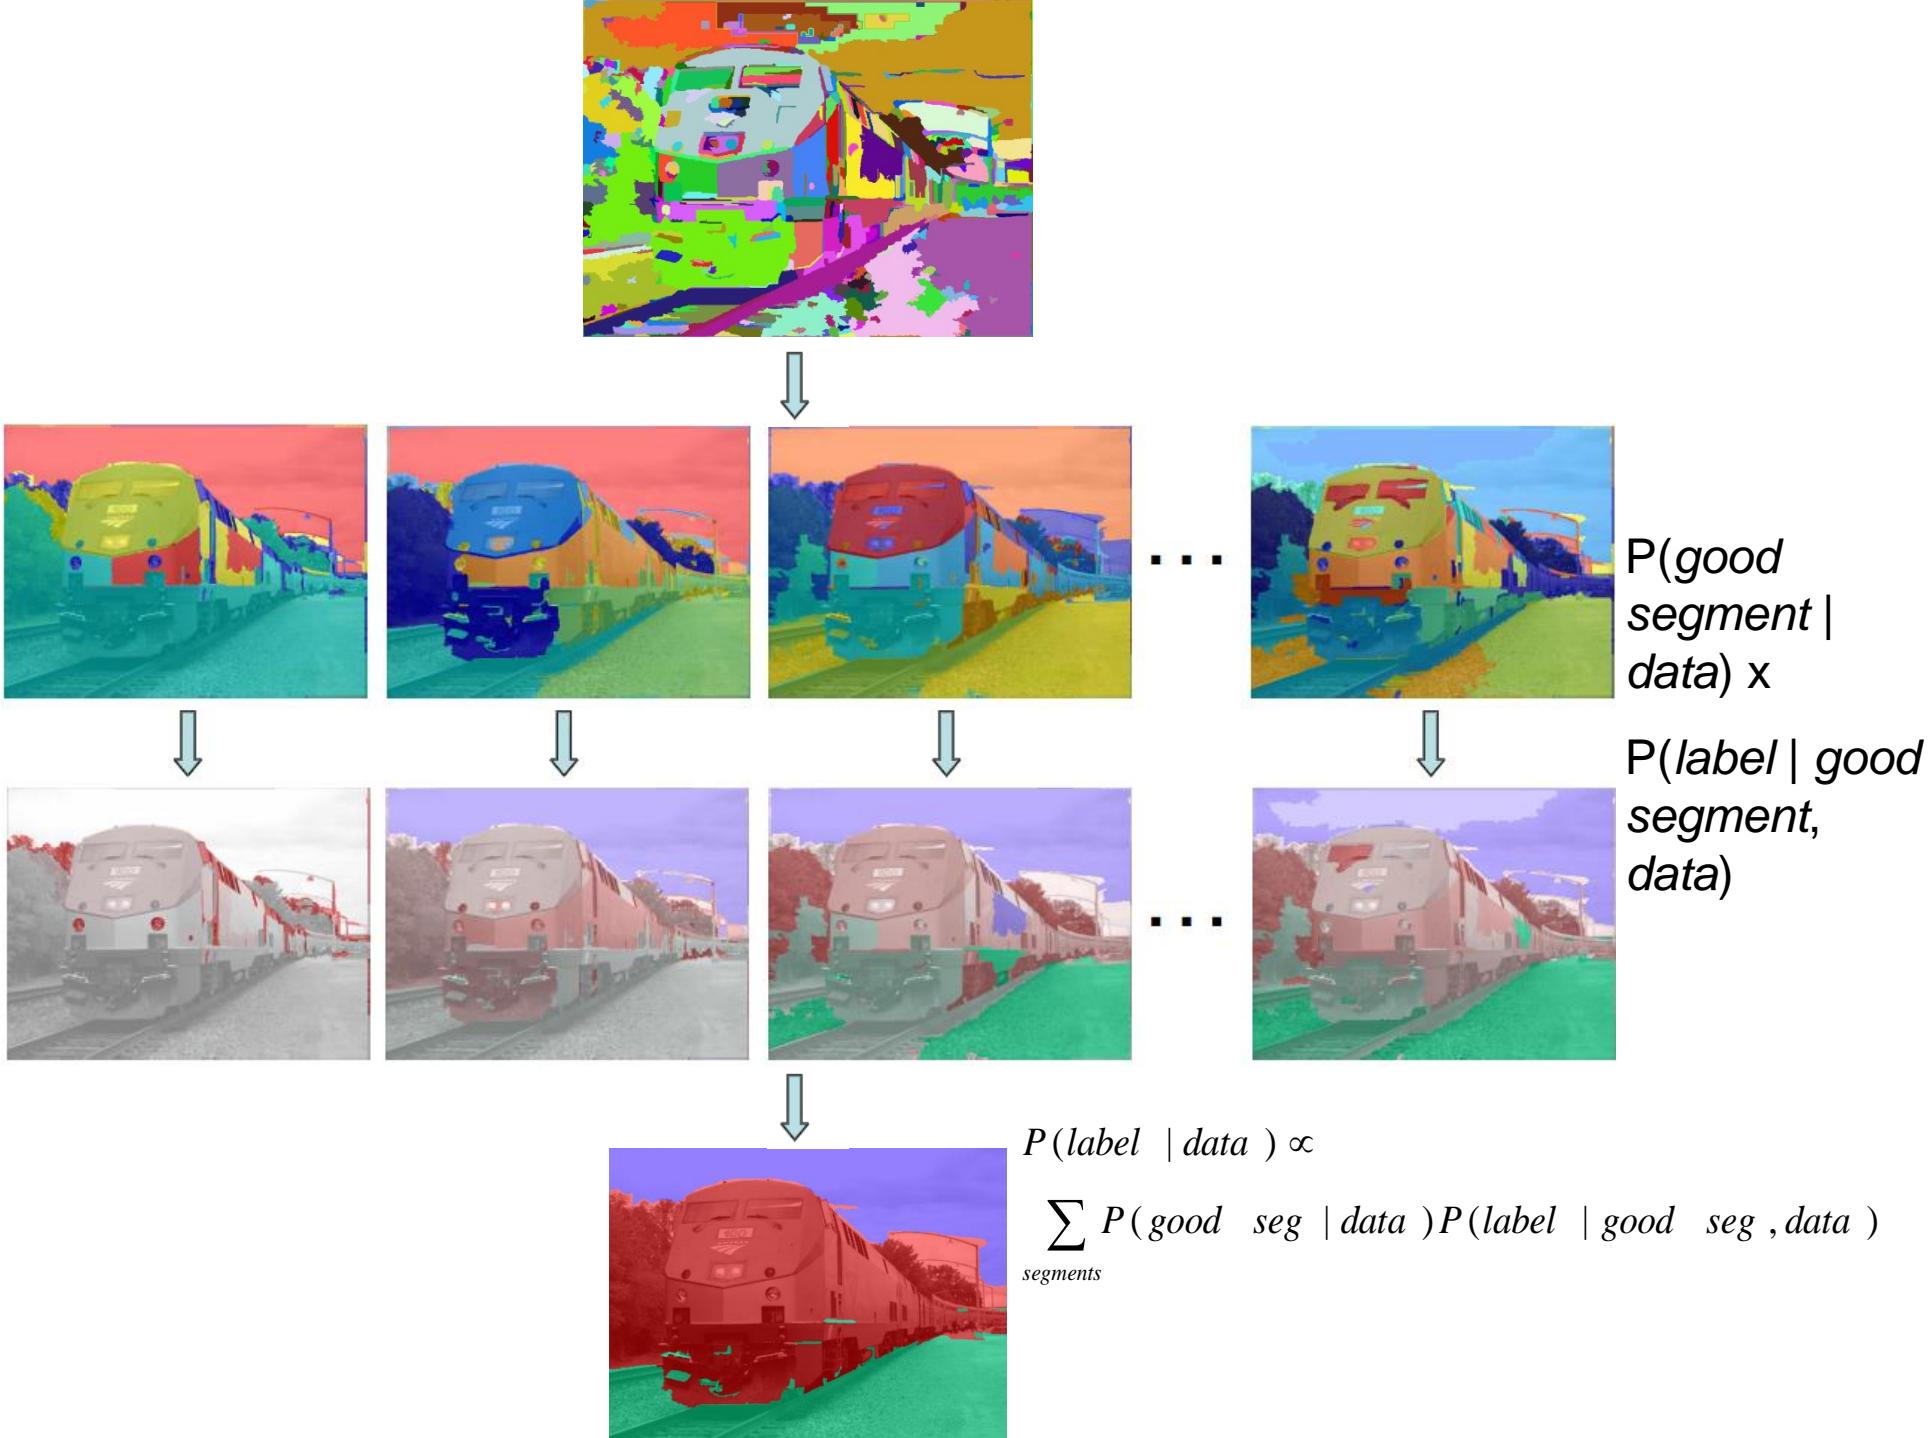


- Task: Regions → Features → Labels
(horizontal, vertical, sky, etc.)



- Chicken and egg problem:
 - If we knew the regions, we could compute the features and label the right regions
 - But to know the right regions we need to know the labels!
- Solution:
 - Generate lots of segmentations
 - Combine the classifications to get consensus

Example from D. Hoiem



Generalities: Summary

- Match ground truth (no objective definition)
- Superpixels = oversegmentation
- Using multiple segmentations

Main approaches

- Spectral techniques
- Segmentation as boundary detection
- Graph-based techniques
- Clustering (K-means and probabilistic)
- Mean shift

Cut and paste procedure

1. Extract Sprites



2. Blend them into the composite



Cut and paste procedure

1. Extract Sprites



2. Blend them into the composite

How do we do this?



Cut and paste procedure

1. Extract Sprites



How do we do this?

Two different ways to think about the same thing:

- Finding seams (i.e., finding the pixels where to cut an image)
- Segmentation (i.e., splitting the image into “foreground” and “background”)

I will be using the two terms interchangeable

Applications

Finding seams is also useful for:



image stitching



retargeting

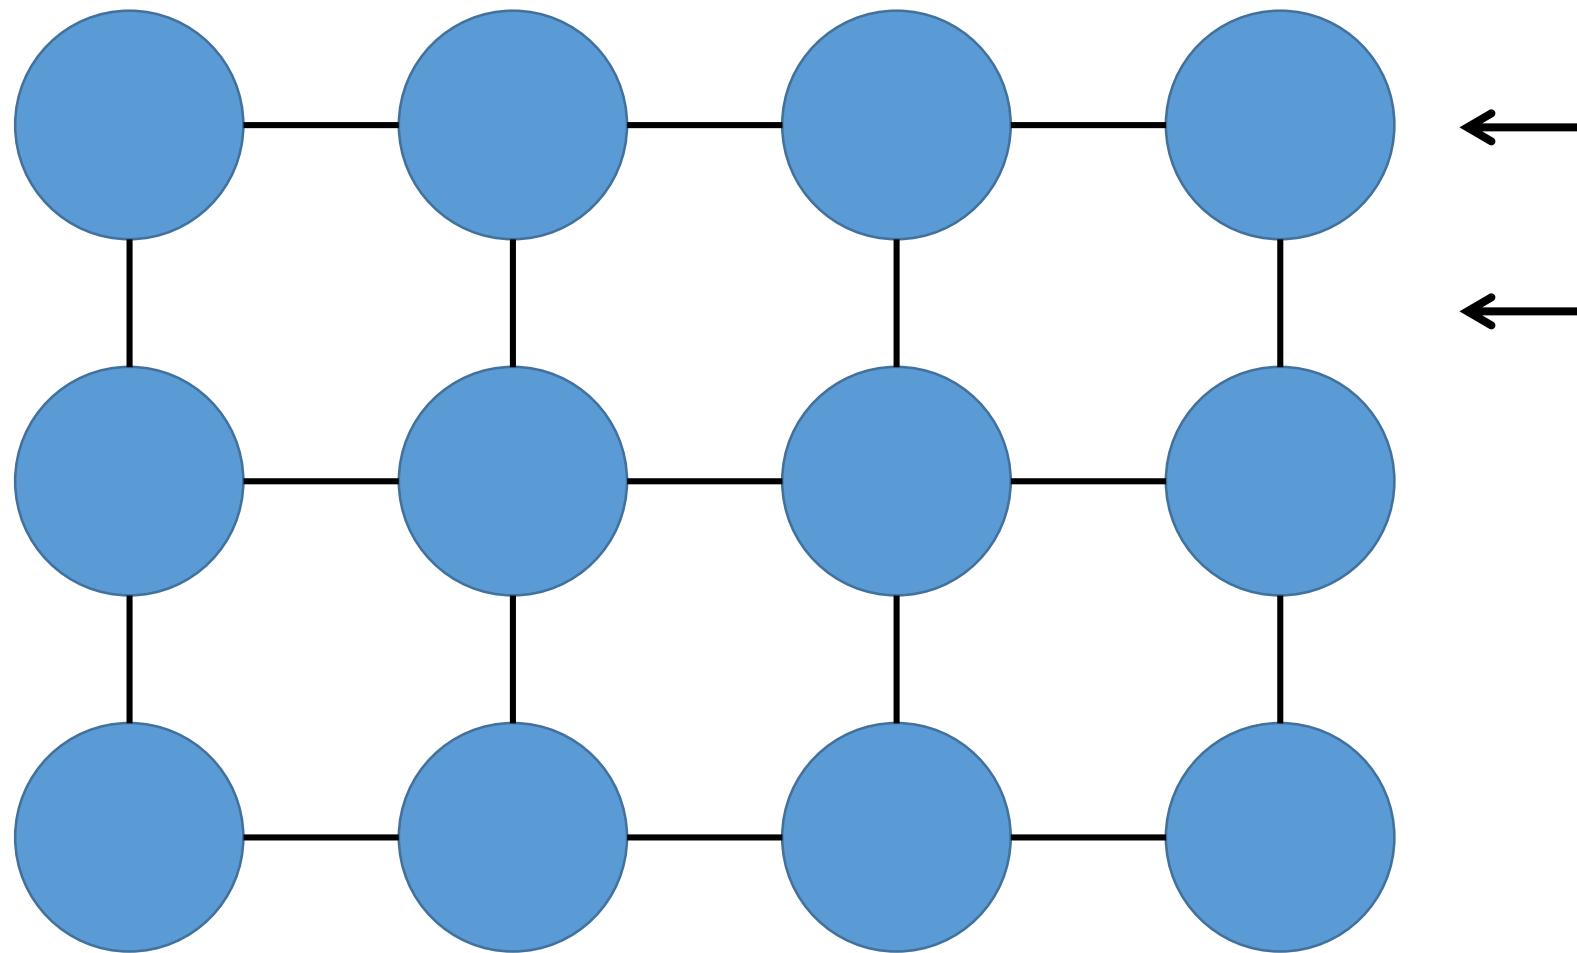


Lazy Snapping
segmentation

Image as a graph

Fundamental theme of today's lecture

Images can be viewed as graphs

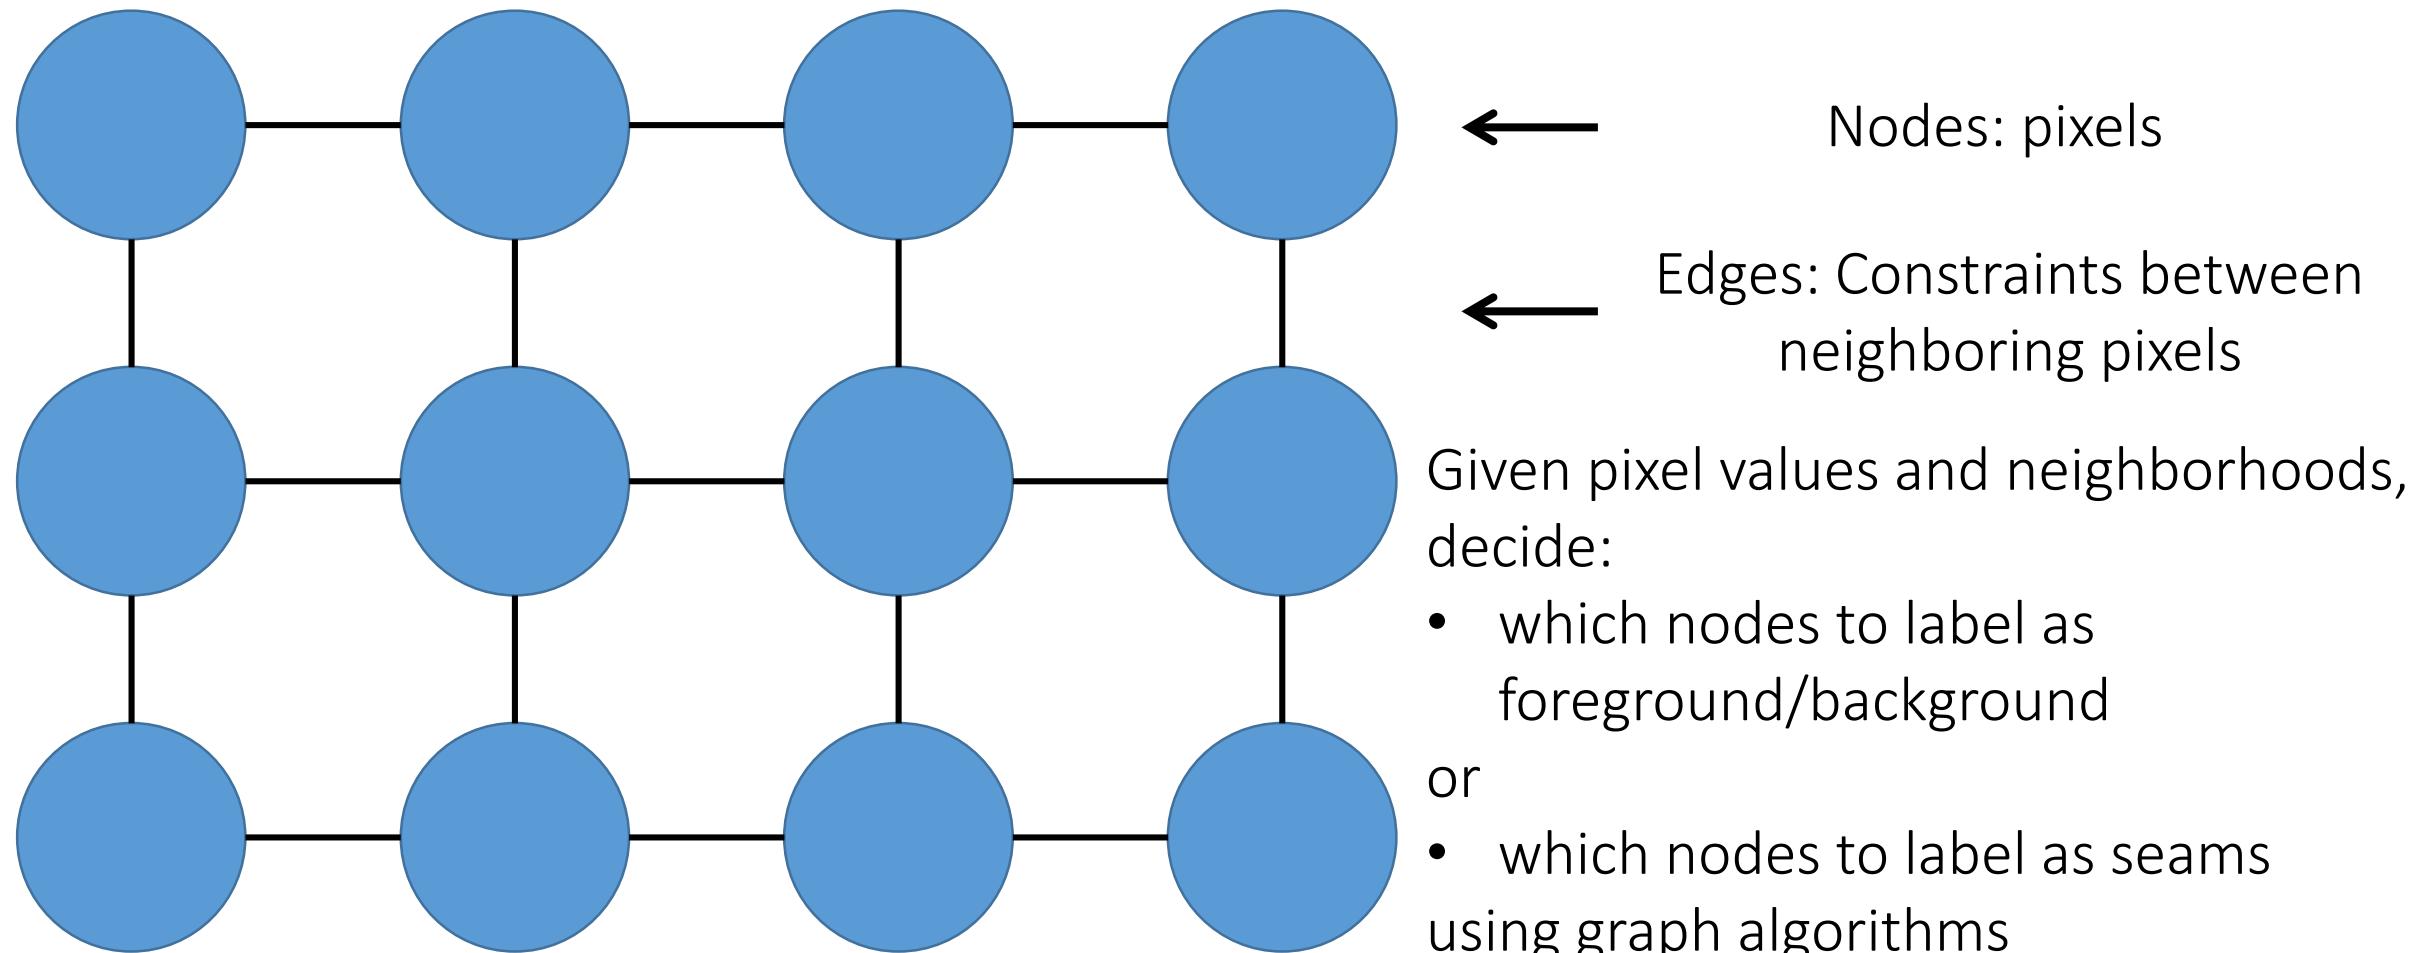


Nodes: pixels

Edges: Constraints between
neighboring pixels

Graph-view of segmentation problem

Segmentation is node-labeling



Graph-view of segmentation problem

Today we will cover:

Method	Labeling problem	Algorithm	Intuition
Intelligent scissors	label pixels as seams	Dijkstra's shortest path (dynamic programming)	short path is a good boundary
GrabCut	label pixels as foreground/background	max-flow/min-cut (graph cutting)	good region has low cutting cost

Shortest graph paths and intelligent scissors

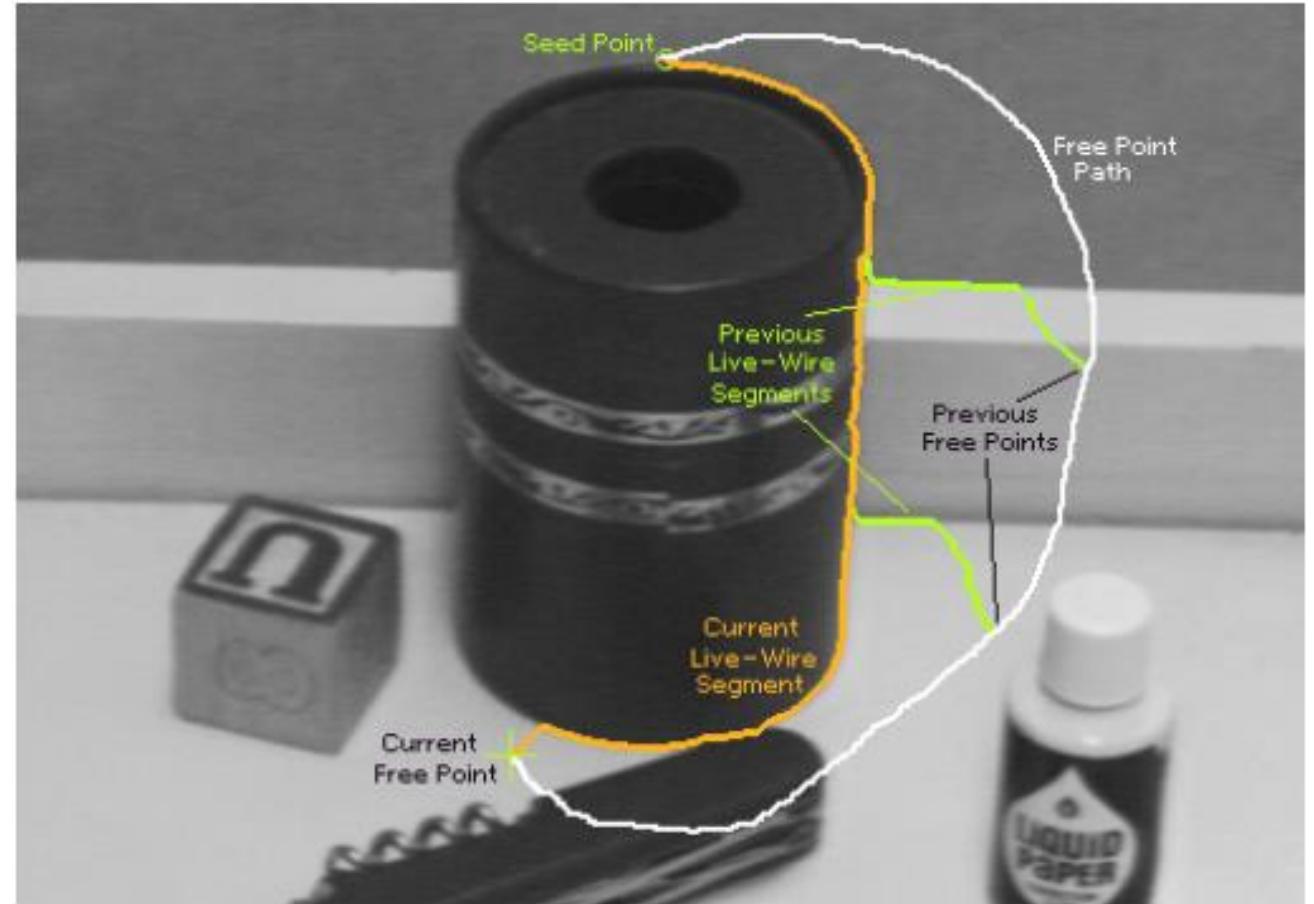
Intelligent scissors

Problem statement:

Given two seed points, find a good boundary connecting them

Challenges:

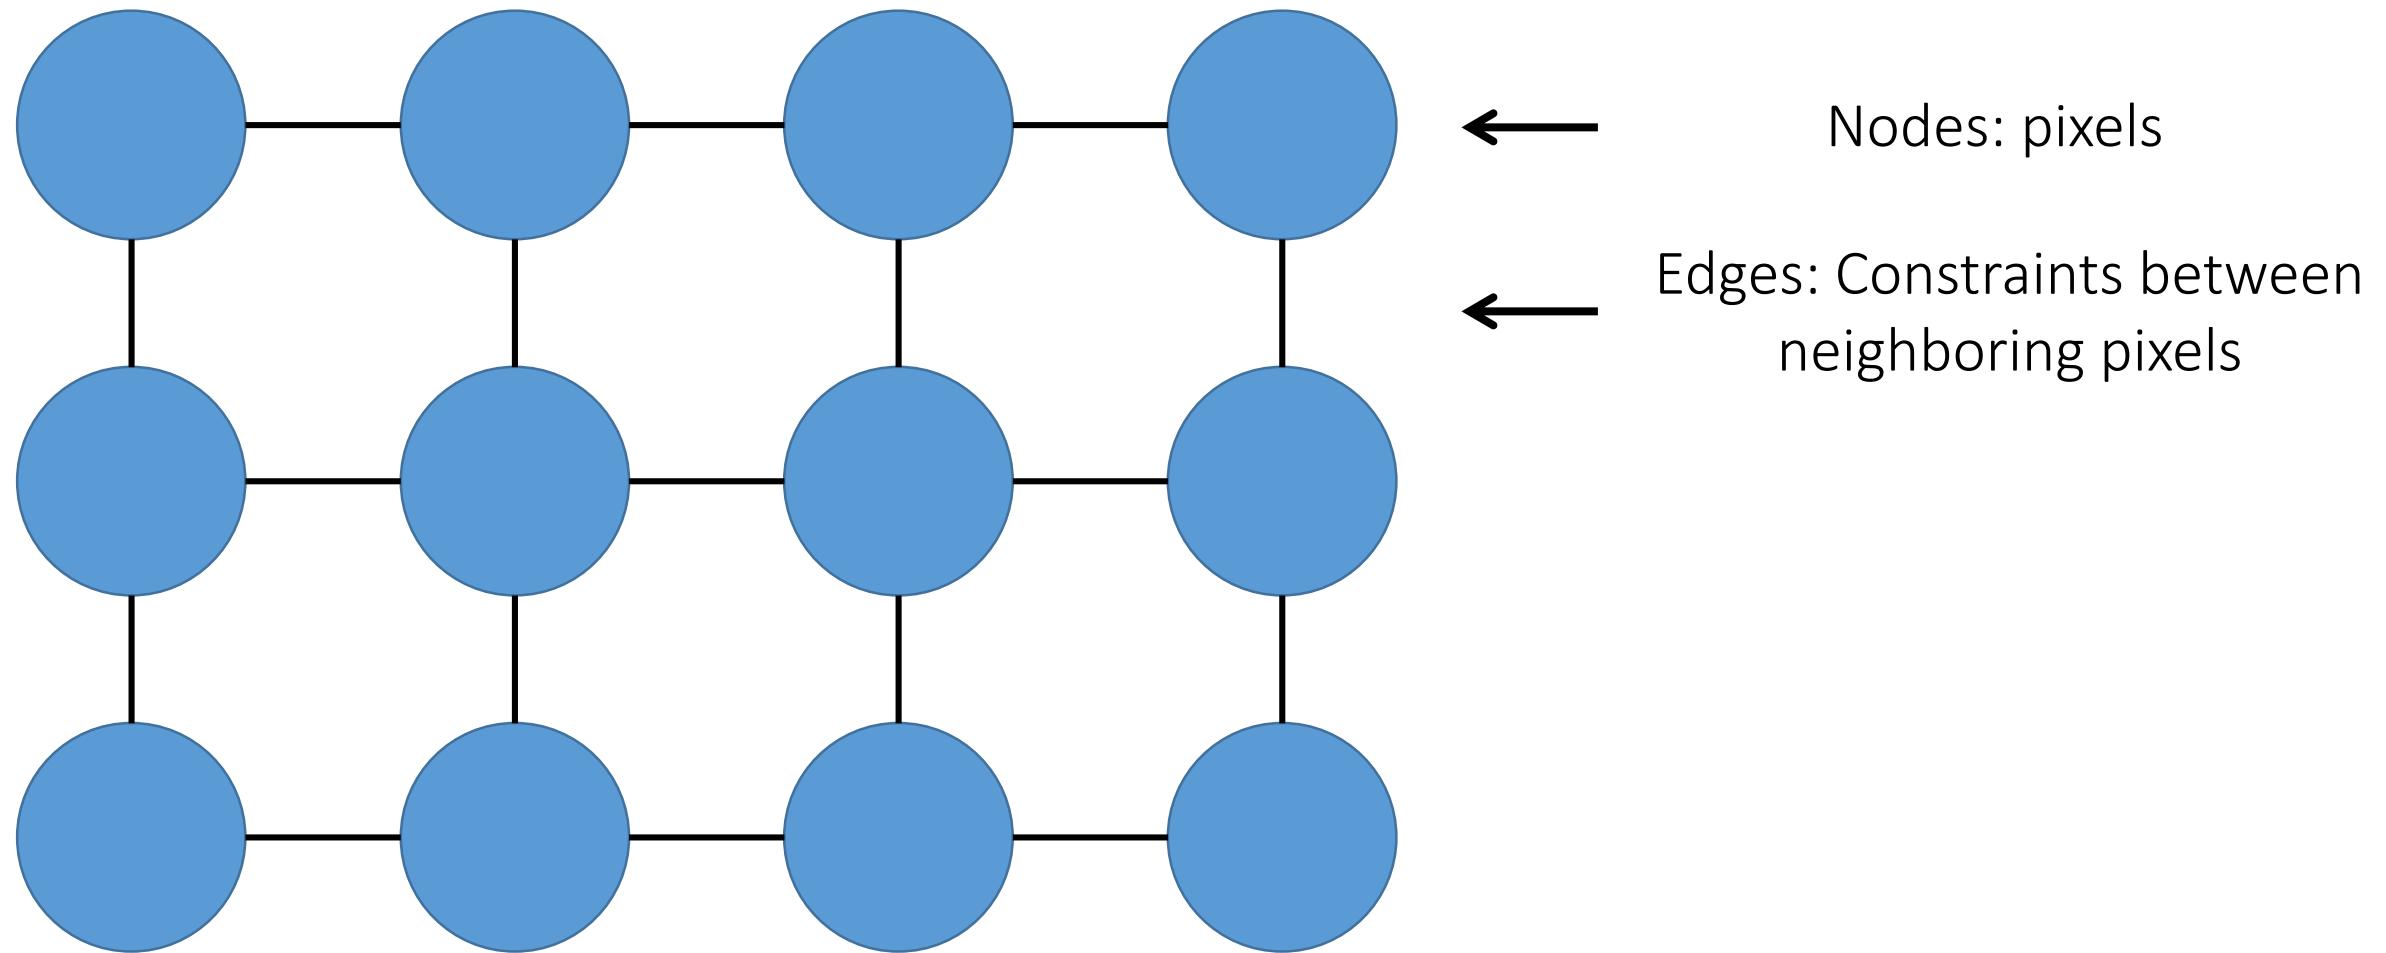
- Make this real-time for interaction
- Define what makes a good boundary



Mortenson and Barrett (SIGGRAPH 1995)
(you can tell it's old from the paper's low quality teaser figure)

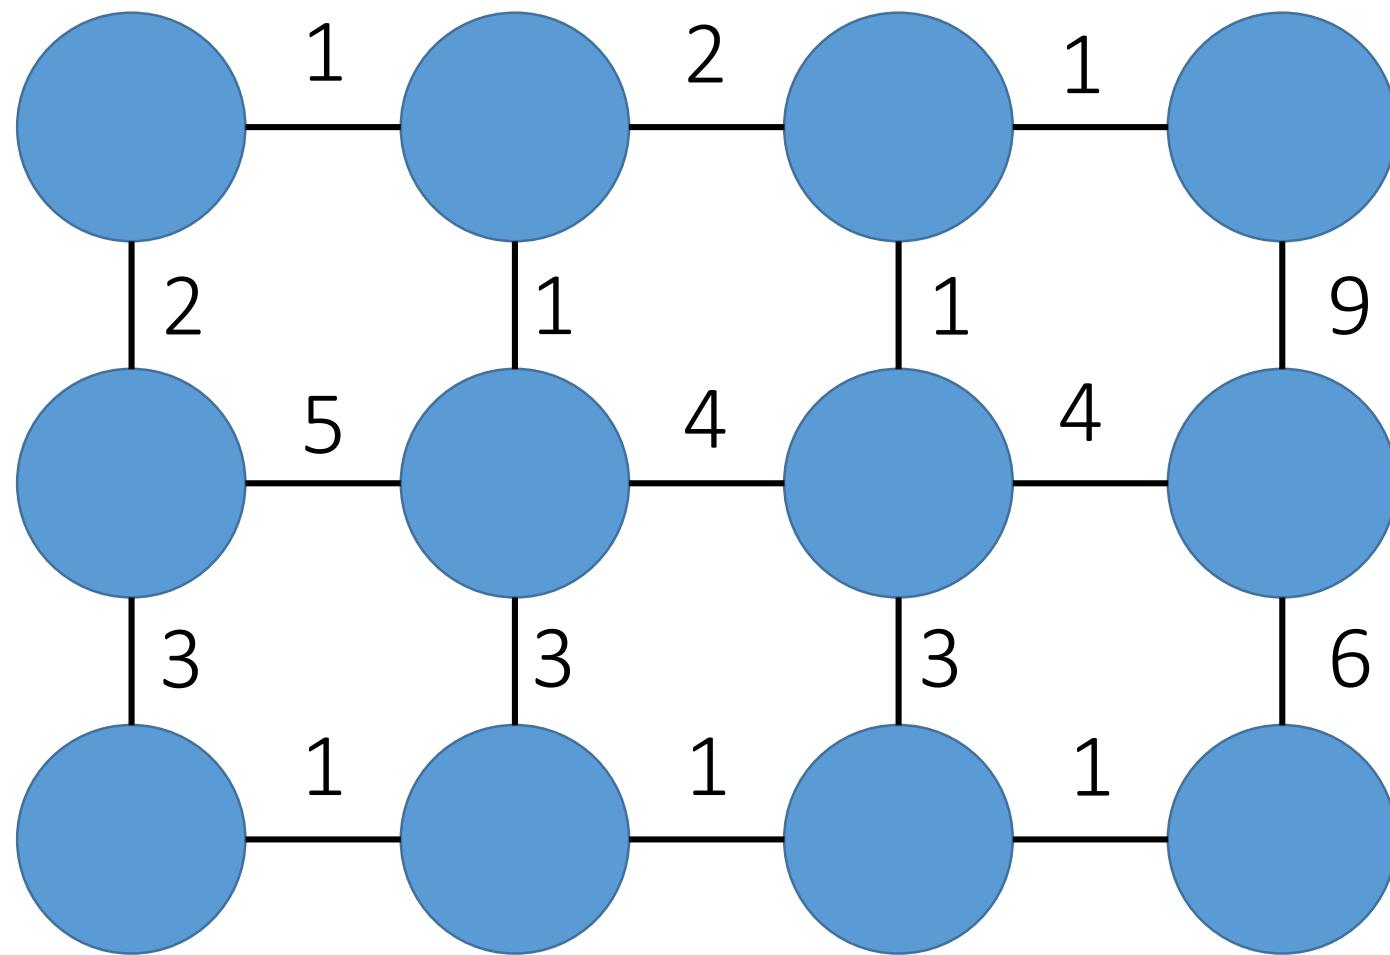
Graph-view of this problem

Images can be viewed as graphs



Graph-view of this problem

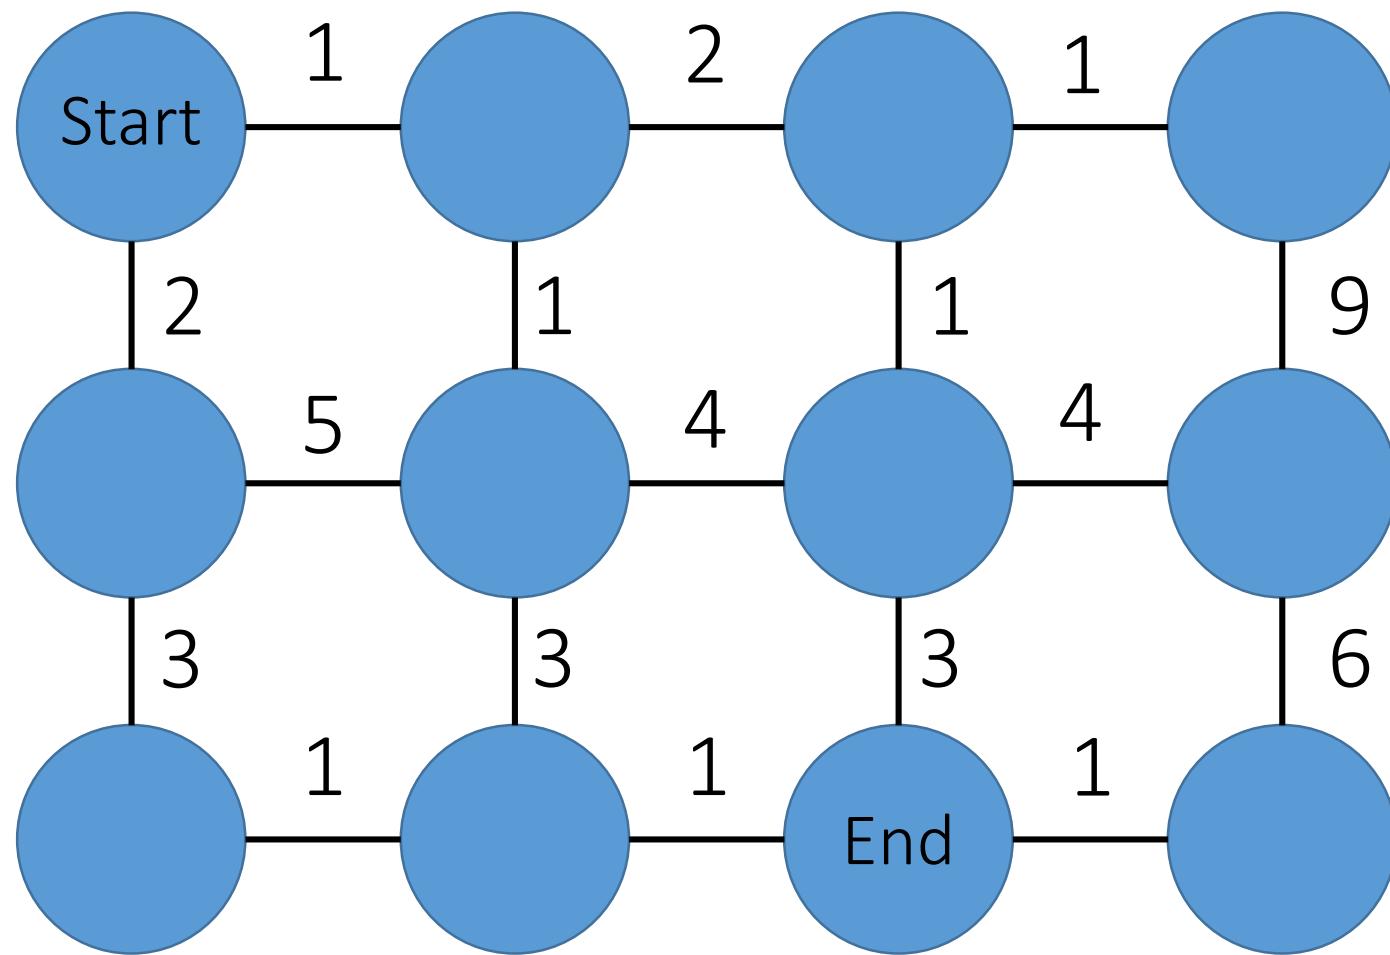
Graph-view of intelligent scissors:



1. Assign weights (costs) to edges

Graph-view of this problem

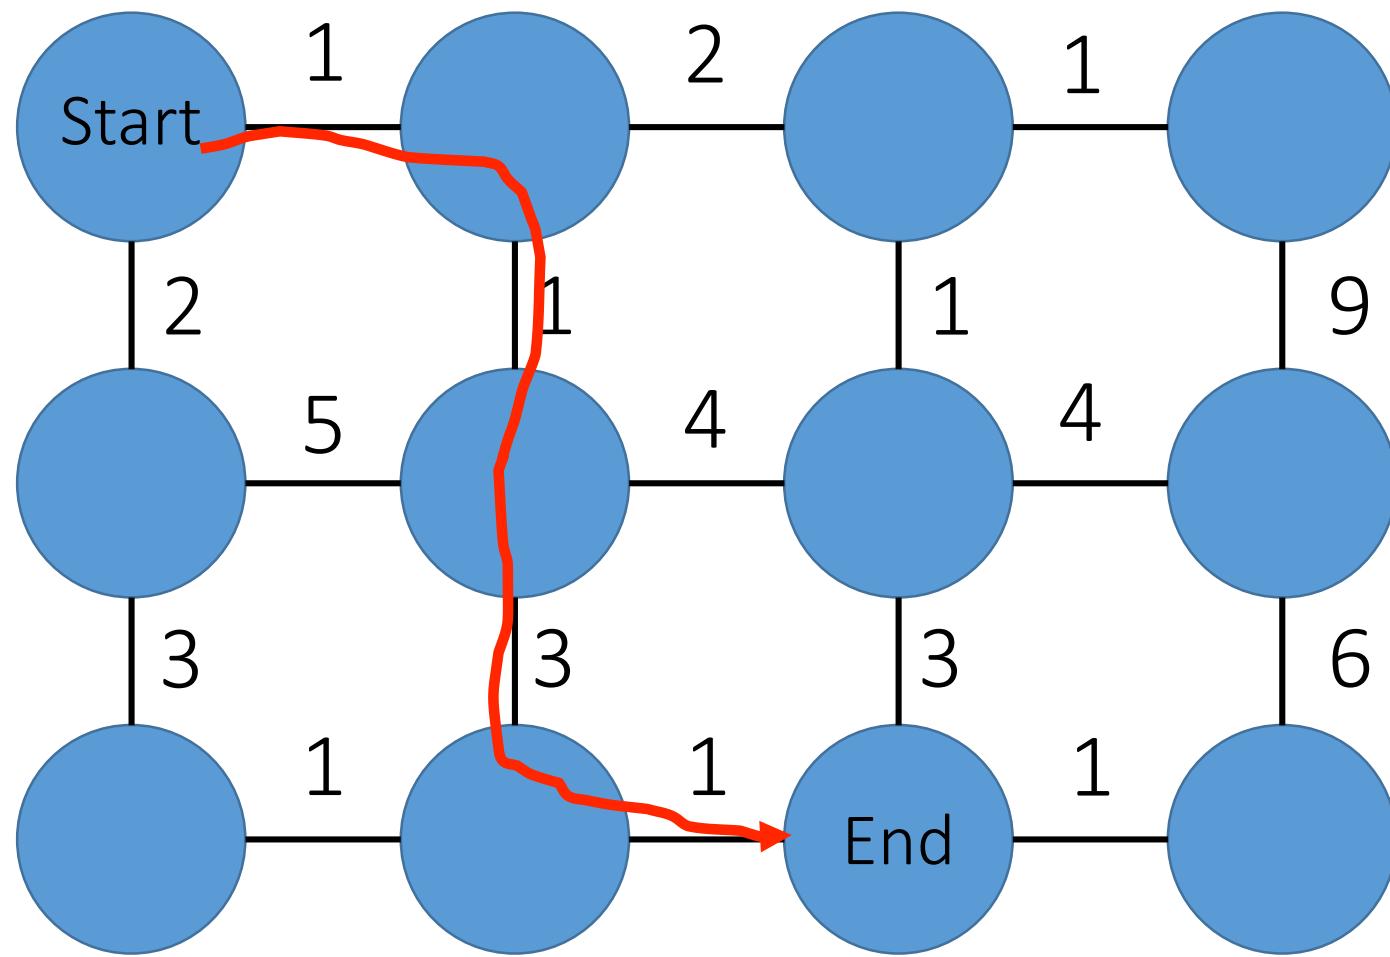
Graph-view of intelligent scissors:



1. Assign weights (costs) to edges
2. Select the seed nodes

Graph-view of this problem

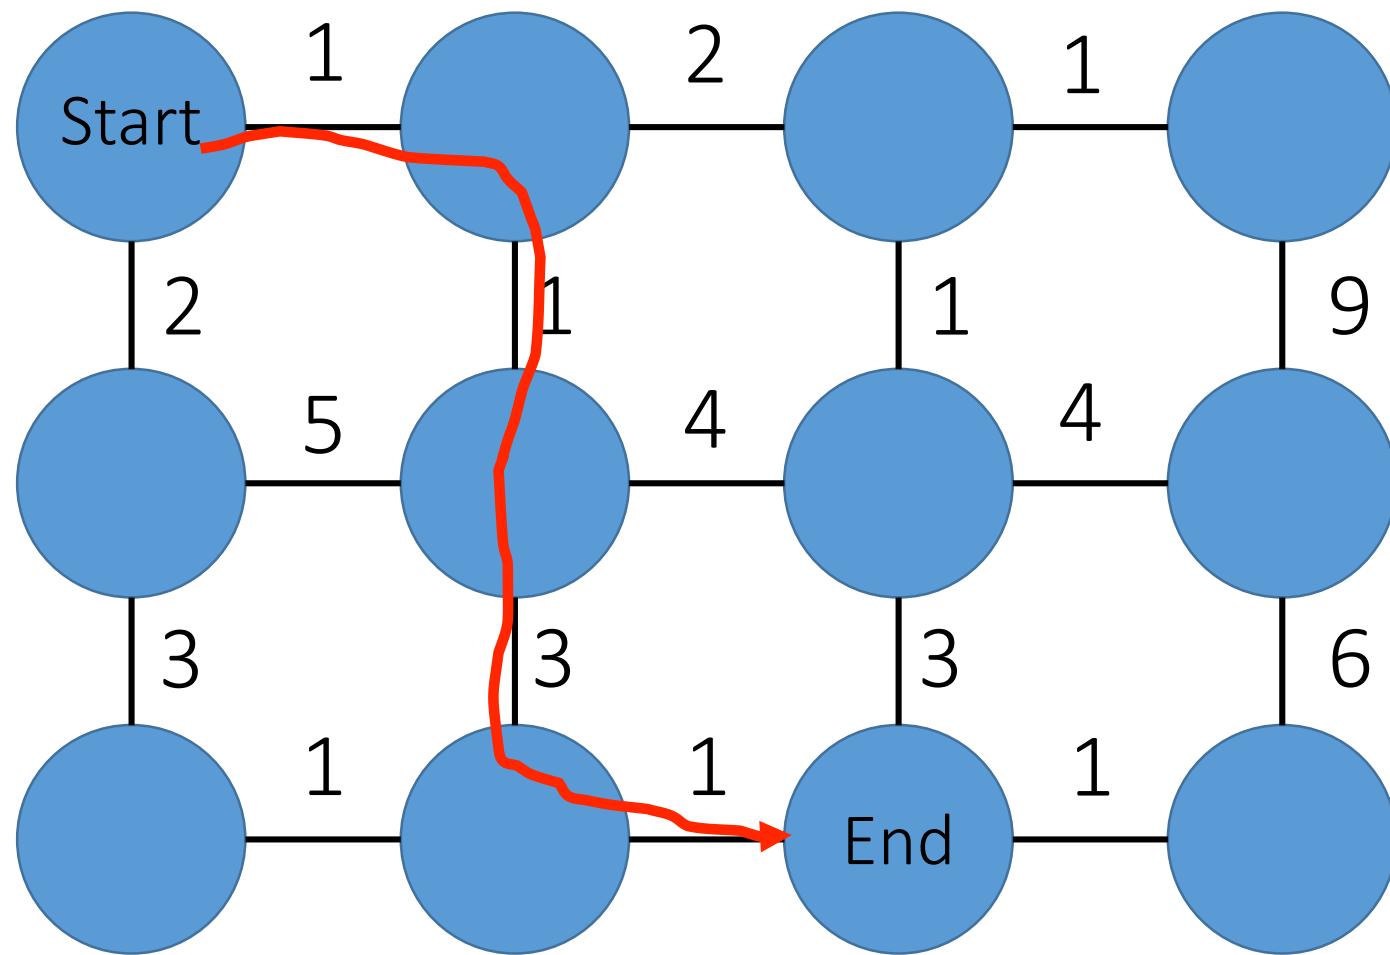
Graph-view of intelligent scissors:



1. Assign weights (costs) to edges
2. Select the seed nodes
3. Find shortest path between them

Graph-view of this problem

Graph-view of intelligent scissors:

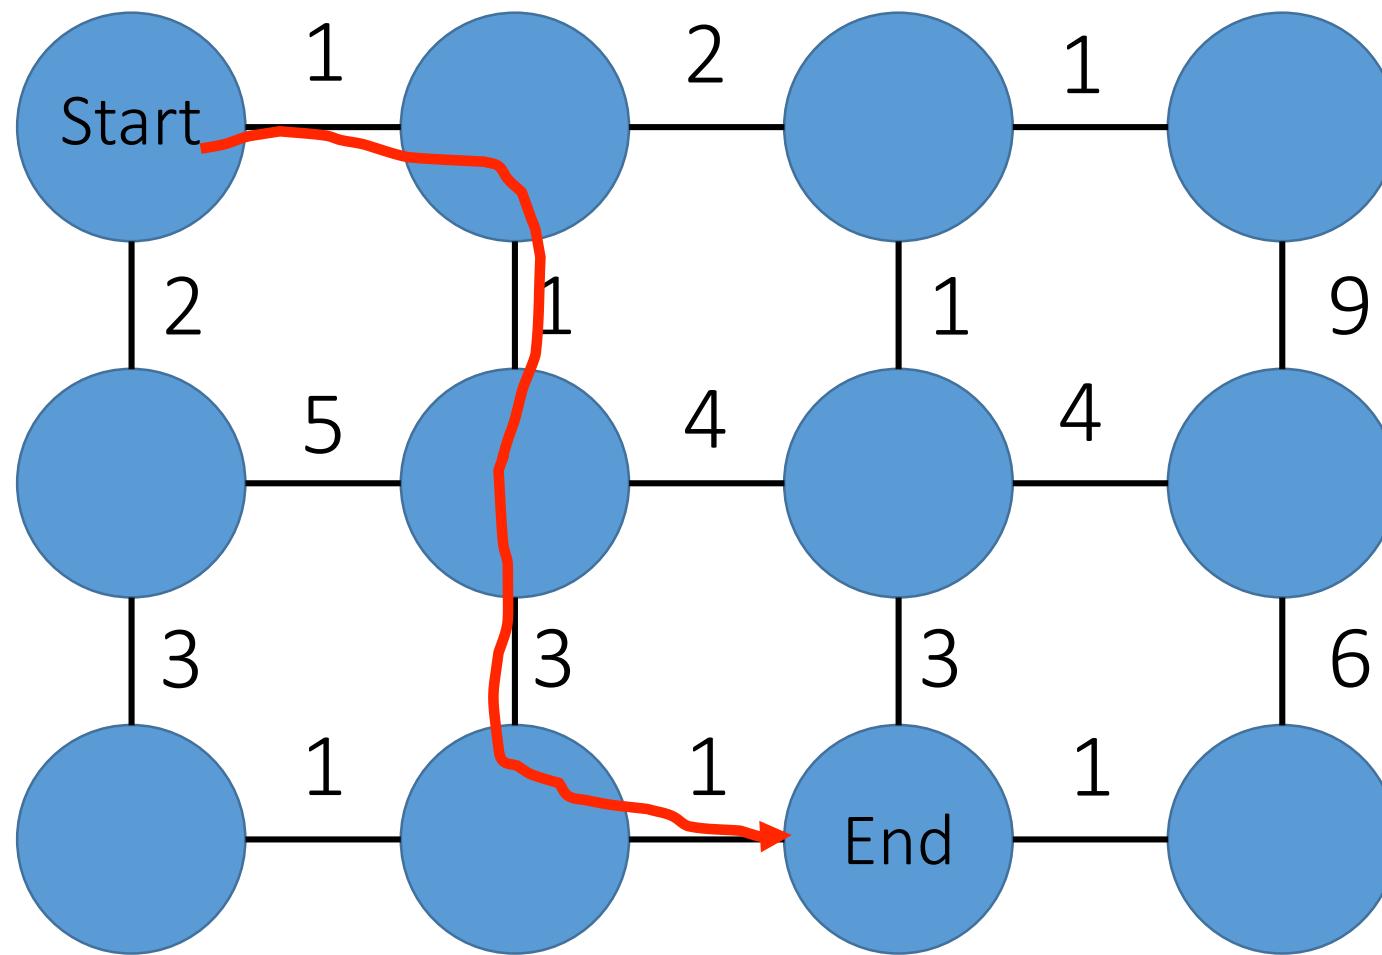


1. Assign weights (costs) to edges
2. Select the seed nodes
3. Find shortest path between them

What algorithm can we use to find the shortest path?

Graph-view of this problem

Graph-view of intelligent scissors:



1. Assign weights (costs) to edges
2. Select the seed nodes
3. Find shortest path between them

What algorithm can we use to find the shortest path?

- Dijkstra's algorithm (dynamic programming)

Dijkstra's shortest path algorithm

Initialize, given seed s (pixel ID) :

- $\text{cost}(s) = 0$ % total cost from seed to this point
- $\text{cost}(!s) = \text{big}$
- $\mathbf{A} = \{\text{all pixels}\}$ % set to be expanded
- $\mathbf{prev}(s) = \text{undefined}$ % pointer to pixel that leads to $q=s$

Precompute $\text{cost}_2(q, r)$ % cost between q to neighboring pixel r

Loop while \mathbf{A} is not empty

1. $q = \text{pixel in } \mathbf{A} \text{ with lowest cost}$

2. Remove q from \mathbf{A}

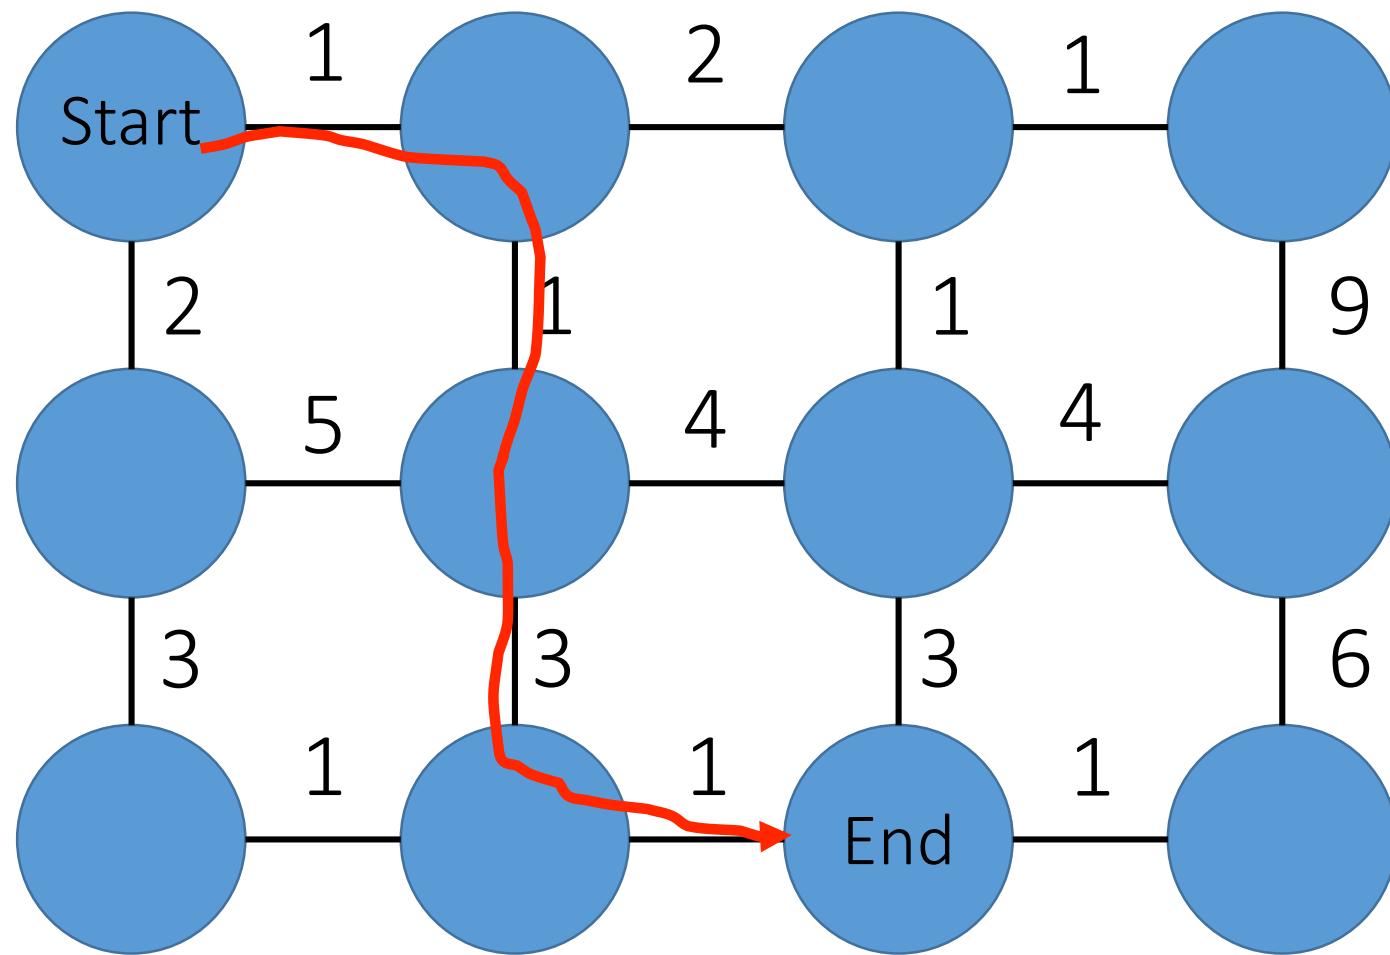
3. For each pixel r in neighborhood of q that is in \mathbf{A}

a) $\text{cost_tmp} = \text{cost}(q) + \text{cost}_2(q, r)$ %this updates the costs

b) if ($\text{cost_tmp} < \text{cost}(r)$)
i. $\text{cost}(r) = \text{cost_tmp}$
ii. $\mathbf{prev}(r) = q$

Graph-view of this problem

Graph-view of intelligent scissors:



1. Assign weights (costs) to edges
2. Select the seed nodes
3. Find shortest path between them

What algorithm can we use to find the shortest path?

- Dijkstra's algorithm (dynamic programming)

How should we select the edge weights to get good boundaries?

Selecting edge weights

Define boundary cost between neighboring pixels:

1. Lower if an image edge is present (e.g., as found by Sobel filtering).
2. Lower if the gradient magnitude at that point is strong.
3. Lower if gradient is similar in boundary direction.



Selecting edge weights

Gradient magnitude



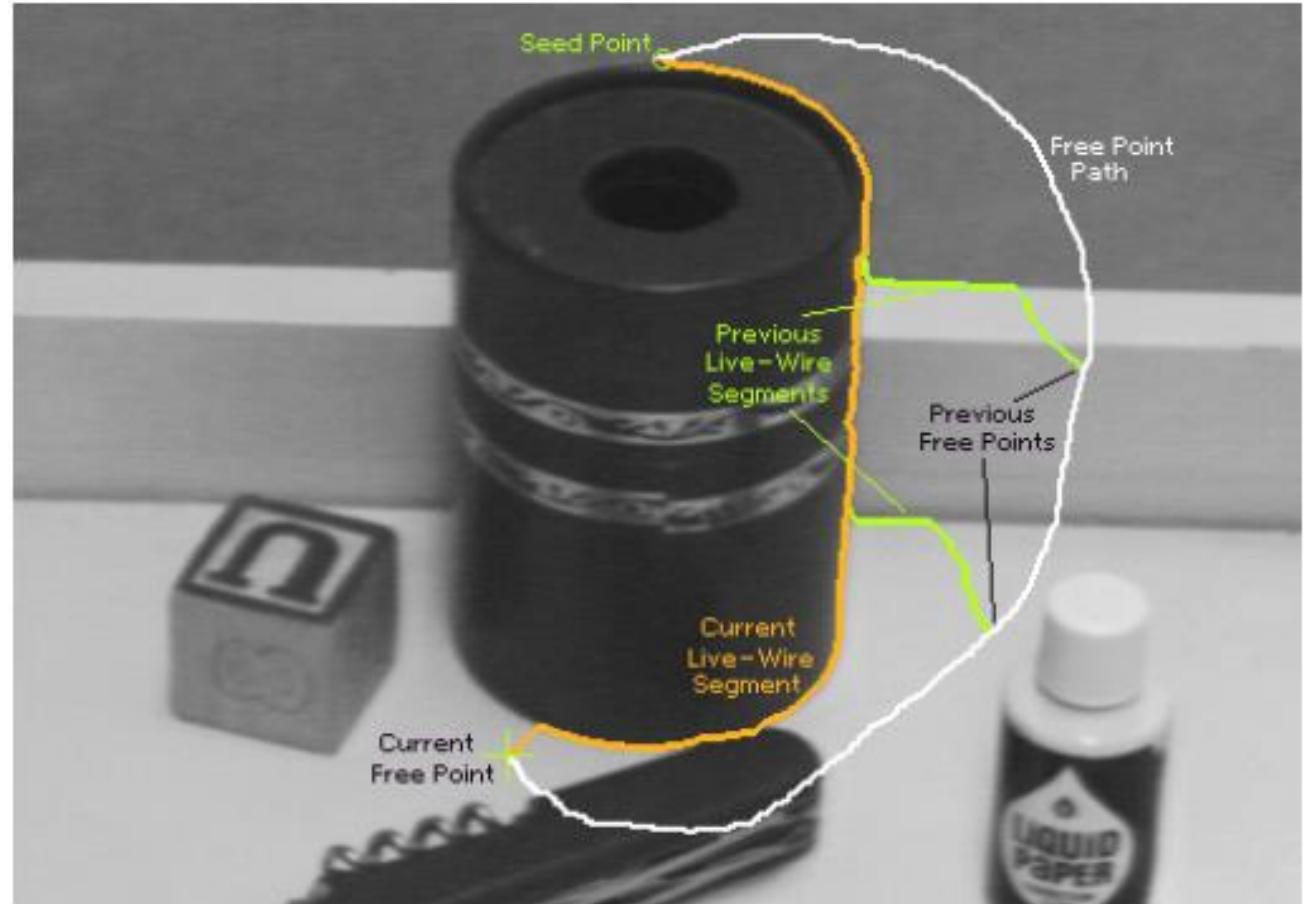
Pixel-wise cost

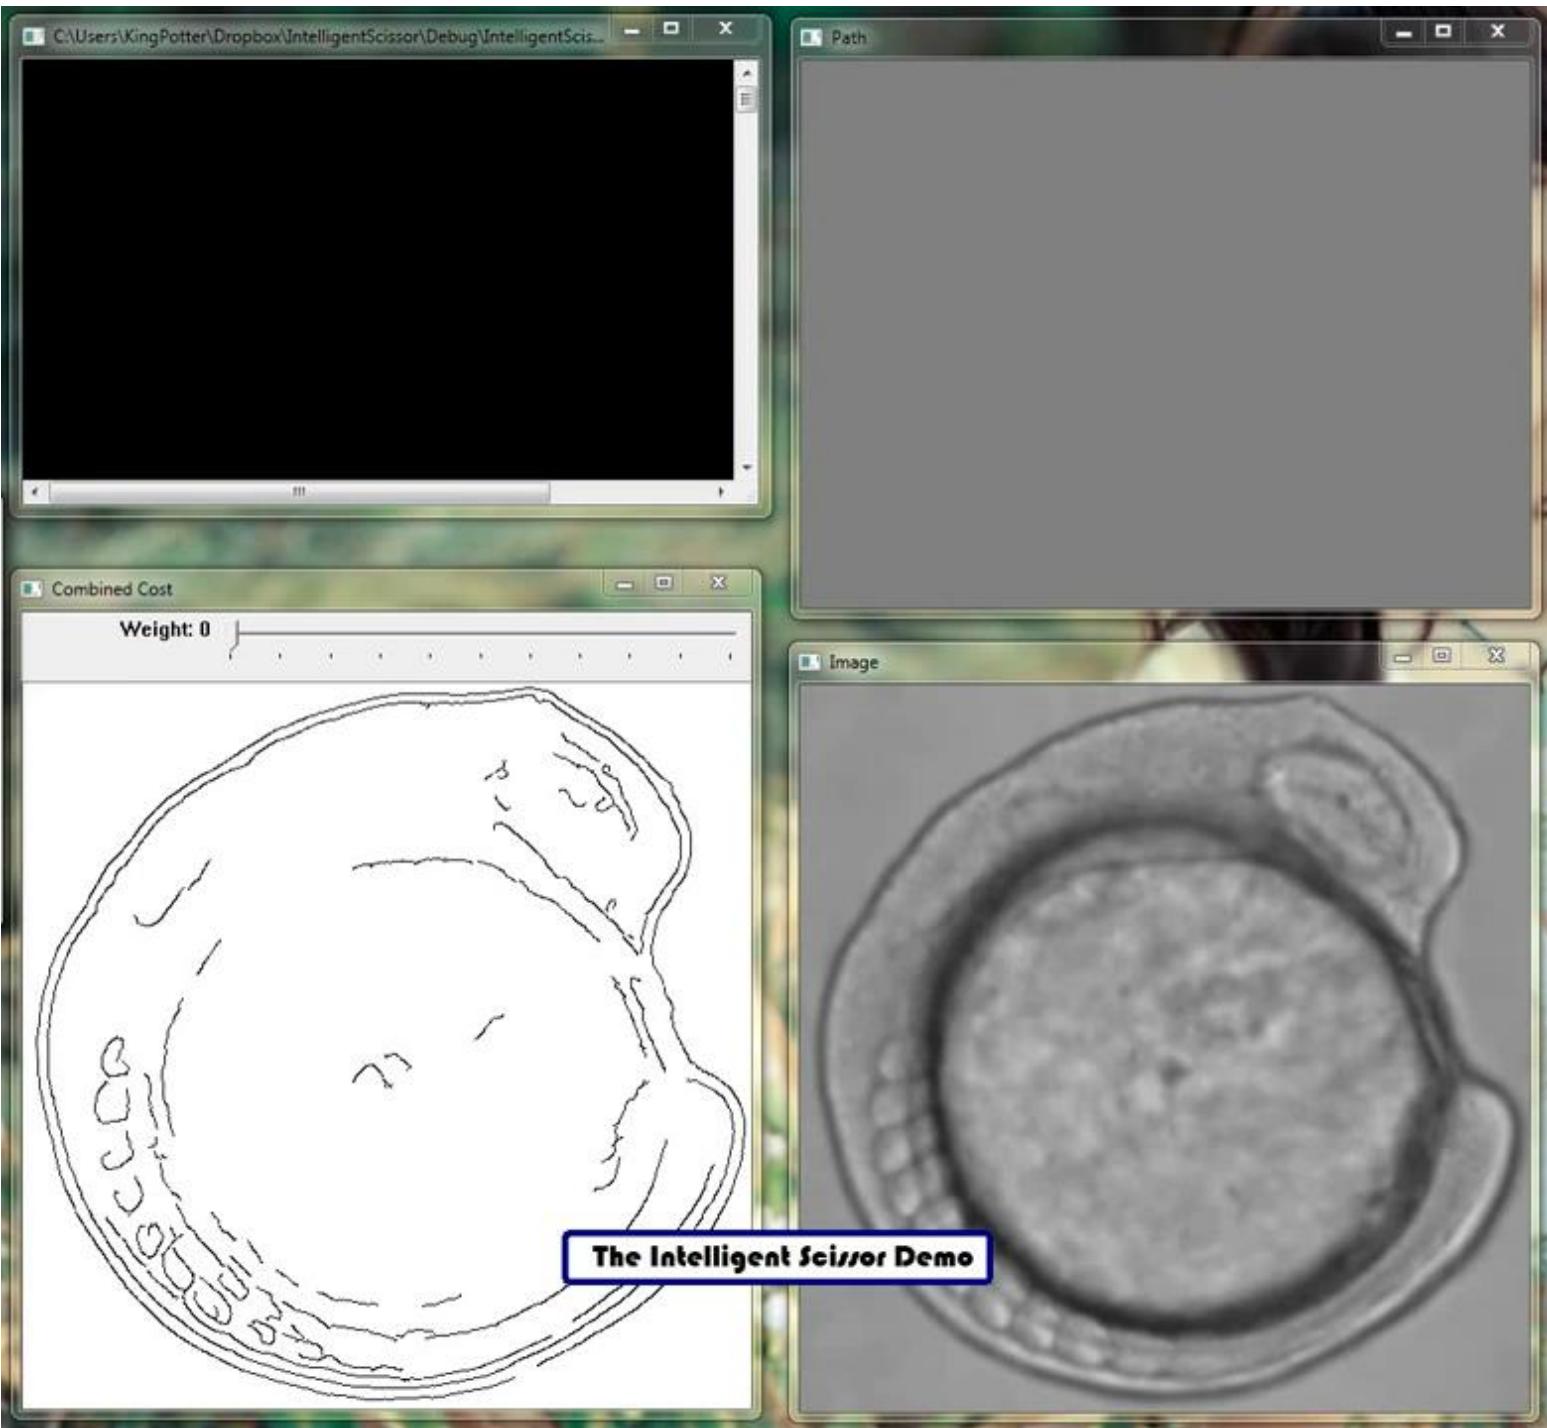
5
0

Edge image

Making it more interactive

1. Use cursor as the “end” seed, and always connect start seed to that
2. Every time the user clicks, use that point as a new starting seed and repeat





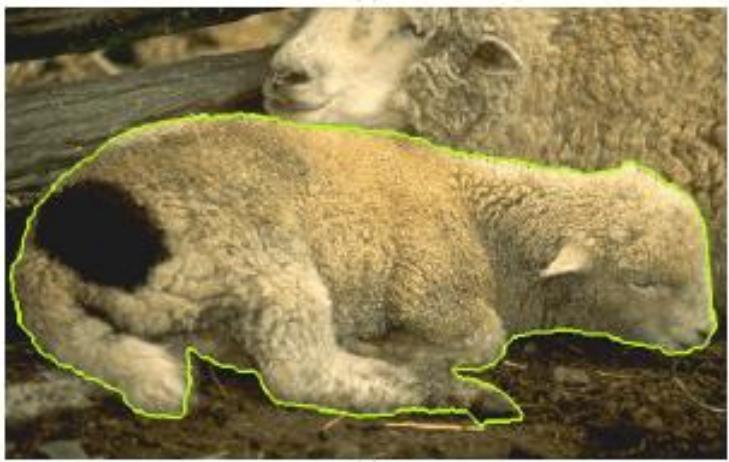
Examples



(a)



(b)

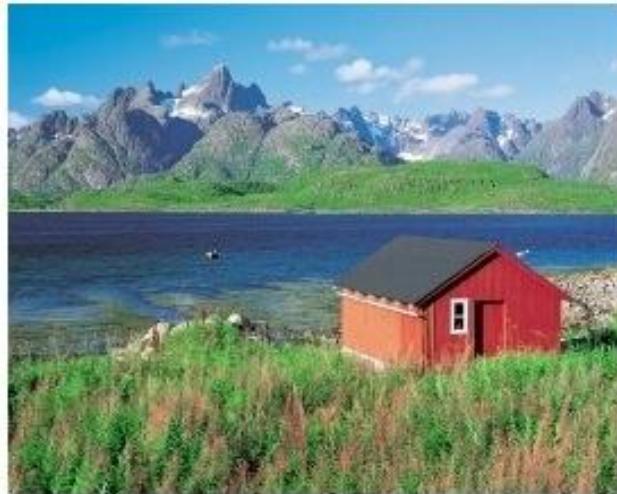


(c)

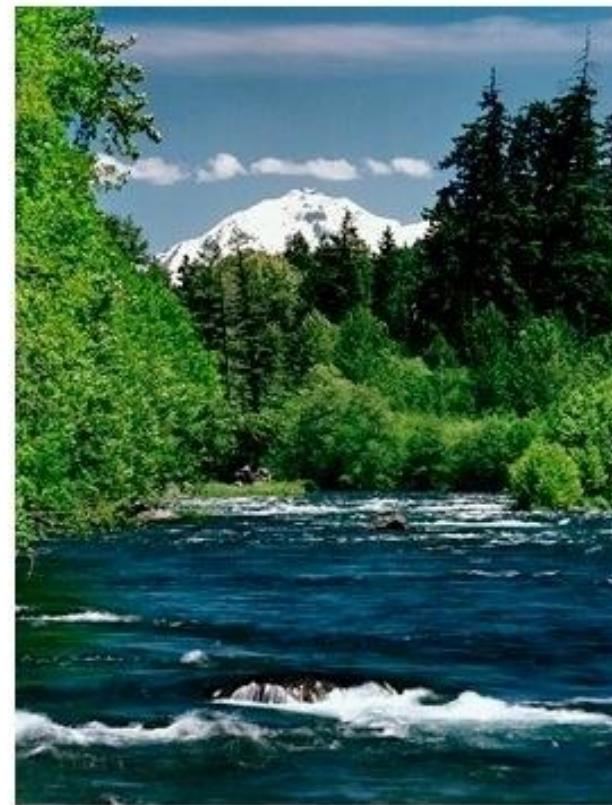


Seam collaging

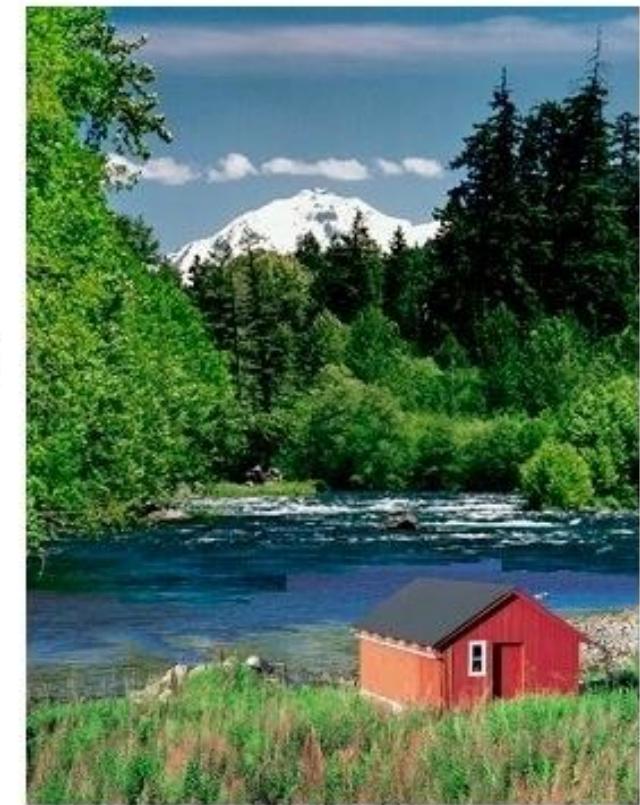
Another use for image seam selection



+

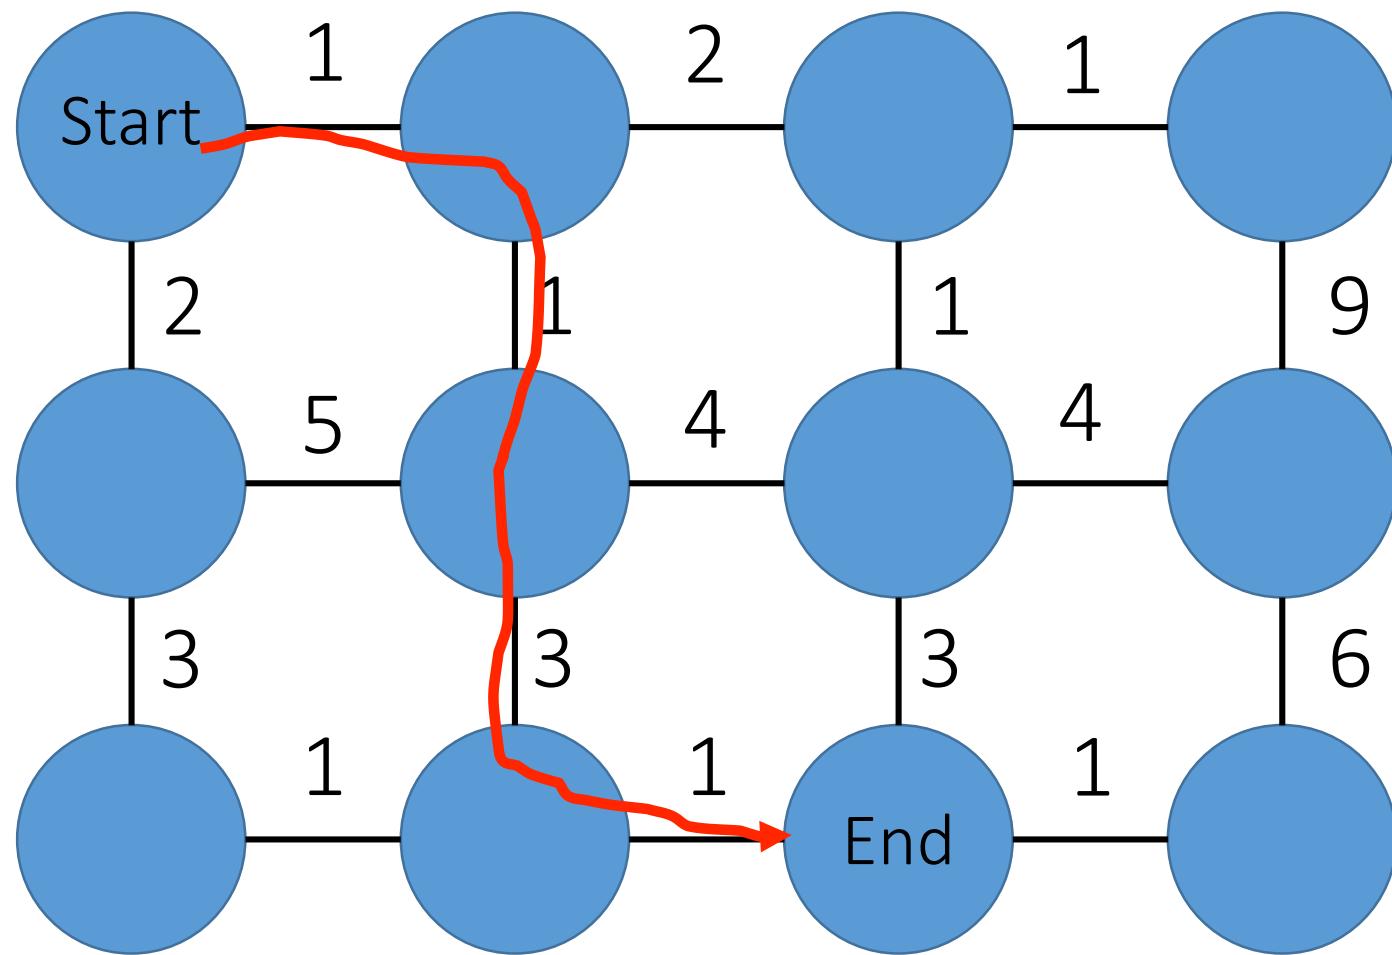


=



Graph-view of this problem

Graph-view of image collaging:



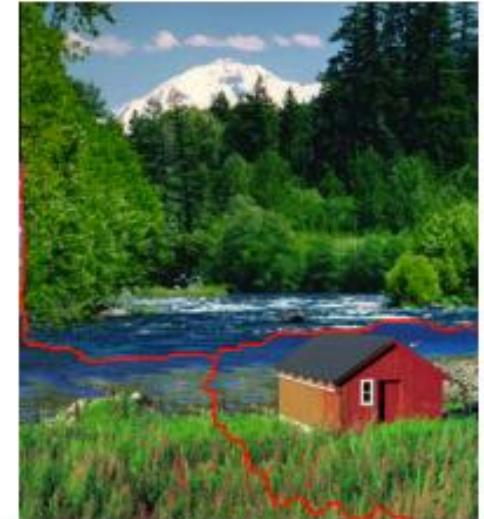
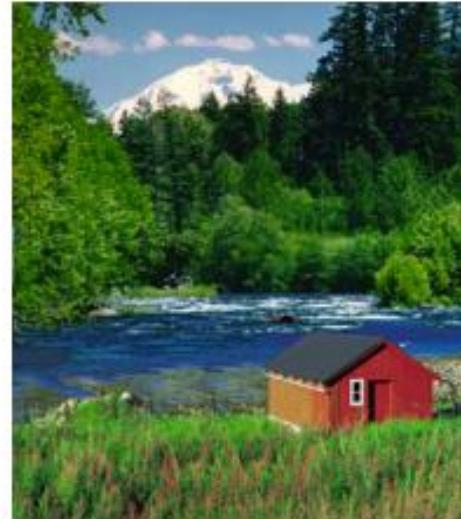
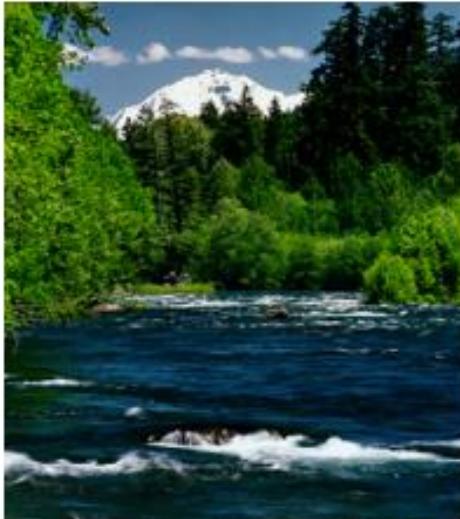
1. Assign weights (costs) to edges
2. Select the seed nodes
3. Find shortest path between them

What edge weights would you use for collaging?

Selecting edge weights for seam collaging

Good places to cut:

- similar color in both images
- high gradient in both images



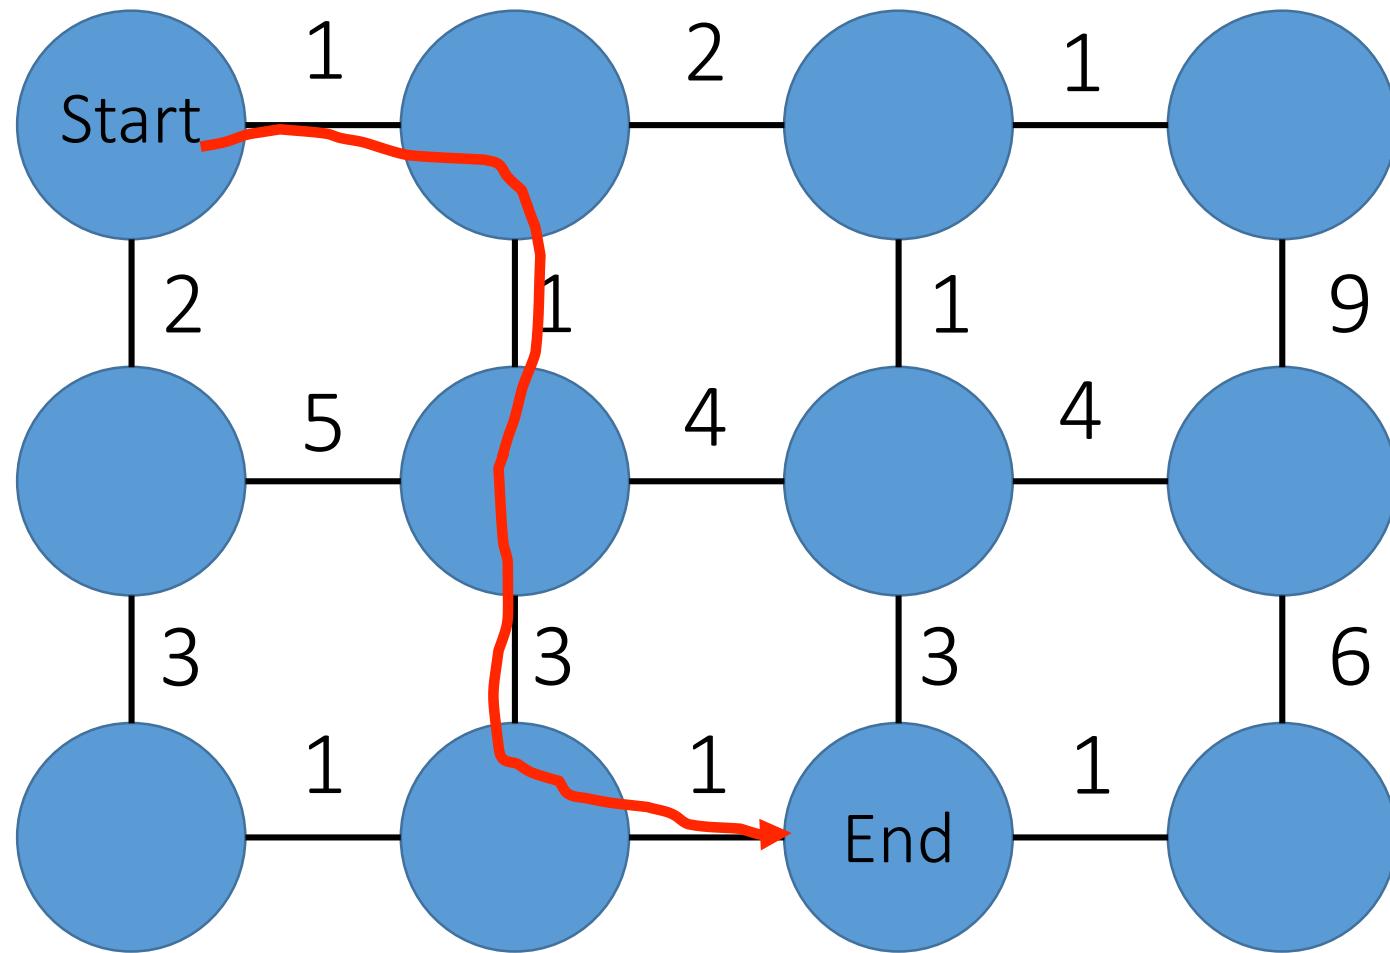
Seam carving

Another use for image seam selection



Graph-view of this problem

Graph-view of seam carving:



1. Assign weights (costs) to edges
2. Select the seed nodes
3. Find shortest path between them

What edge weights would you use for seam carving?



Shai Avidan
Mitsubishi Electric Research Lab
Ariel Shamir
The interdisciplinary Center & MERL

Examples

Where will intelligent scissors work well, or have problems?



Graph-cuts and GrabCut

GrabCut

Only user input is the box!



Combining region and boundary information

user input



result



regions

Intelligent scissors



boundary

GrabCut



regions & boundary

GrabCut is a mixture of two components

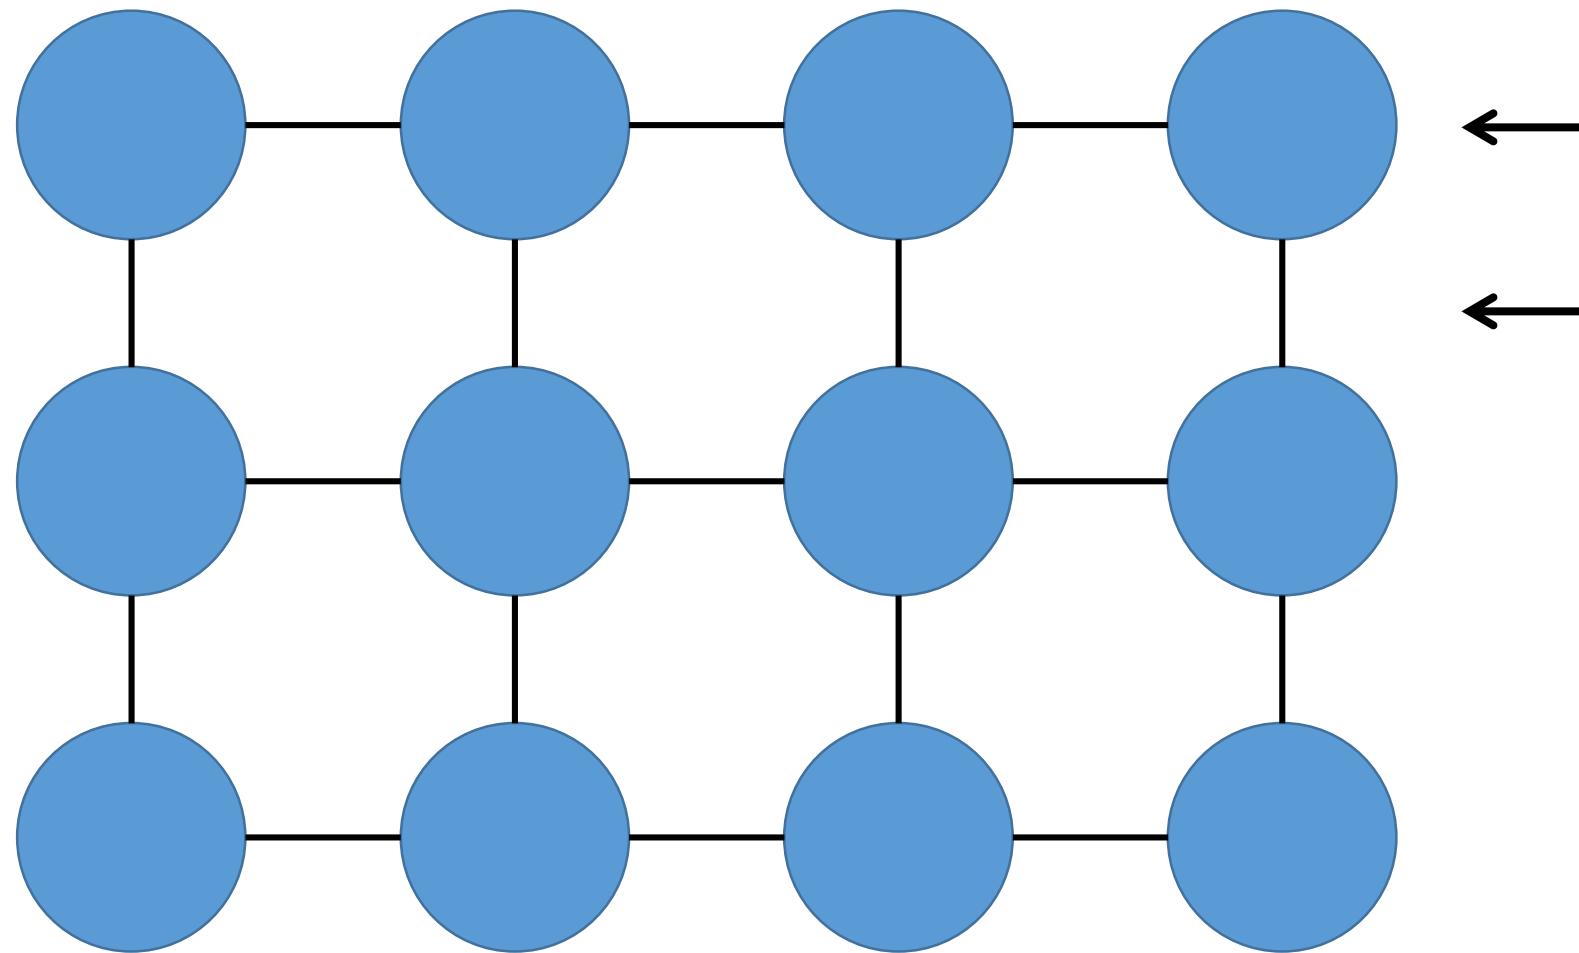
1. Segmentation using graph cuts
2. Foreground-background modeling using unsupervised clustering

GrabCut is a mixture of two components

1. Segmentation using graph cuts
2. Foreground-background modeling using unsupervised clustering

Segmentation using graph cuts

Remember: Graph-based view of images

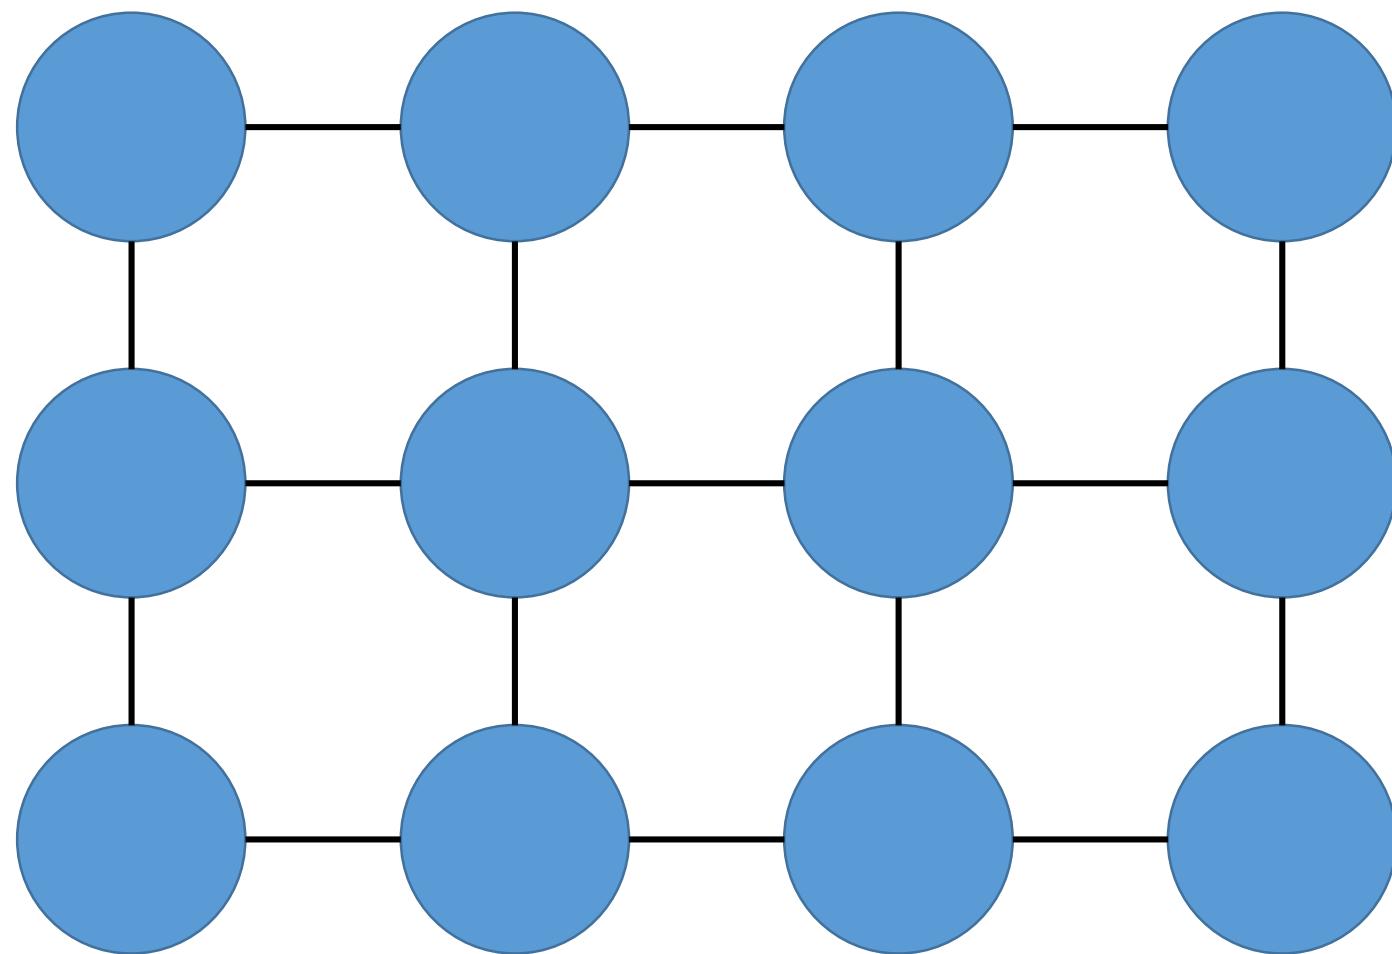


Nodes: pixels

Edges: Constraints between
neighboring pixels

Markov Random Field (MRF)

Assign foreground/background labels based on:



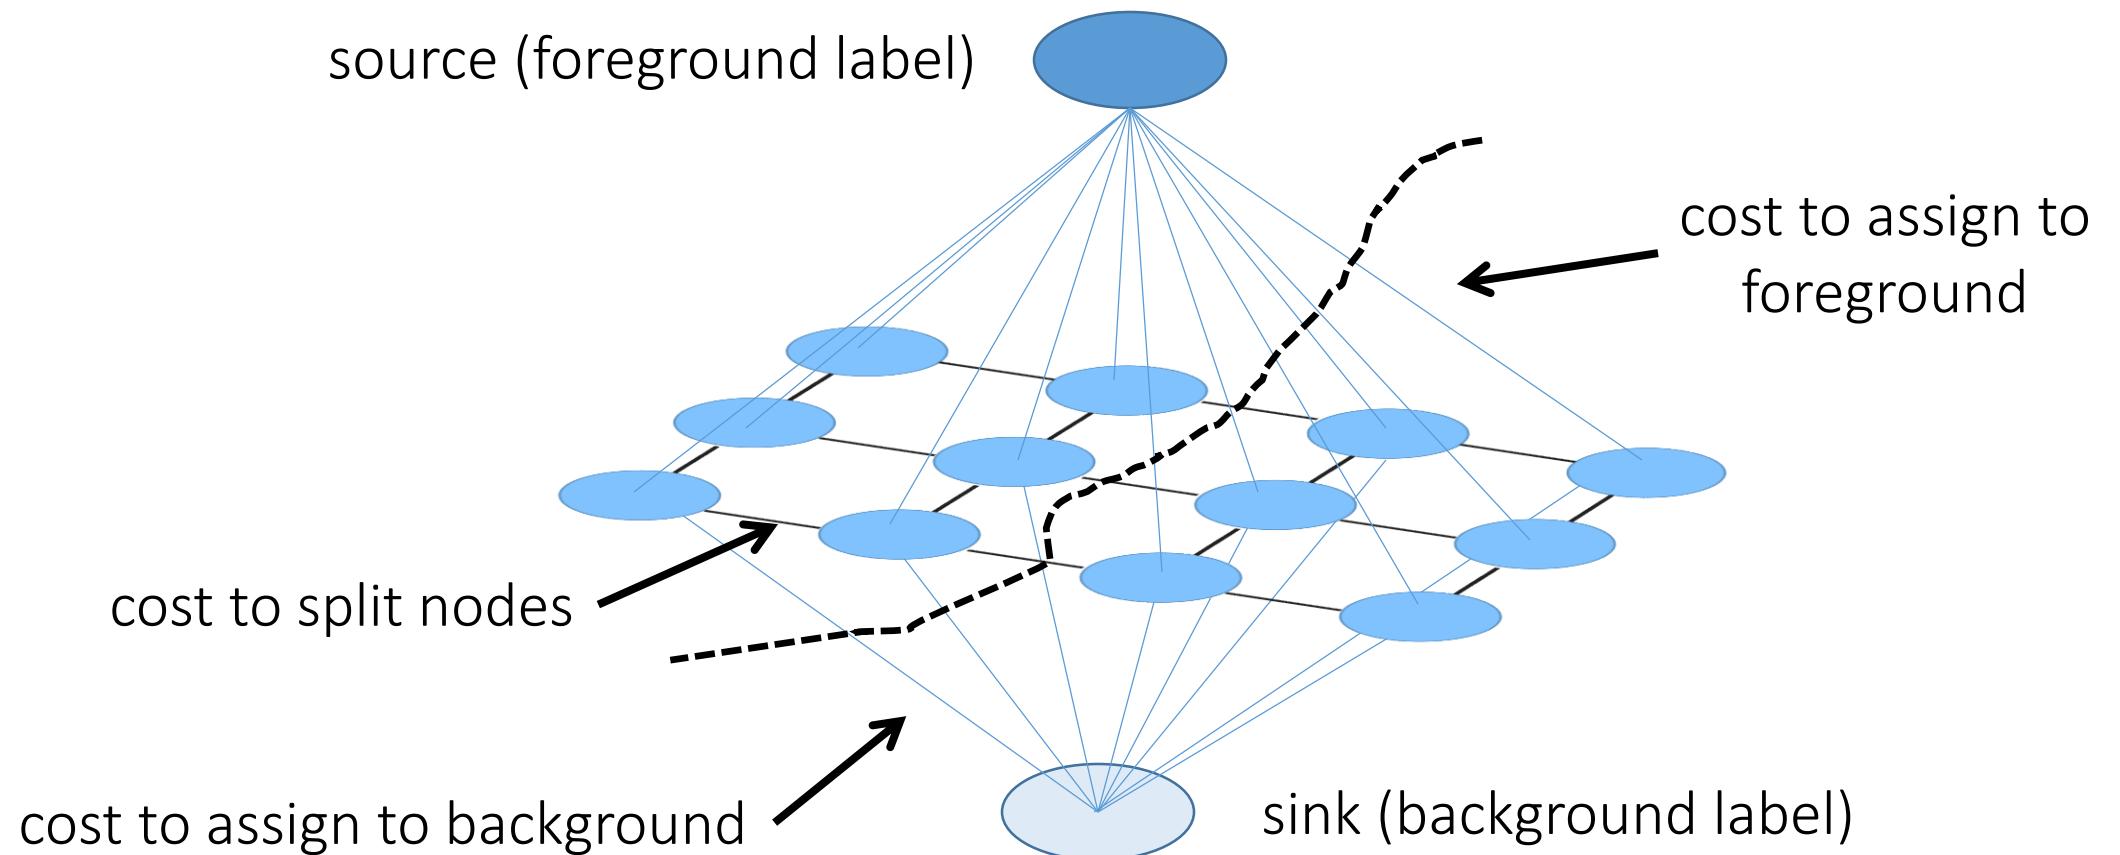
$$Energy(y; \theta, data) = \sum_i \psi_1(y_i; \theta, data) + \sum_{i, j \in edges} \psi_2(y_i, y_j; \theta, data)$$

Given its intensity value, how likely is a pixel to be foreground or background?

Given their intensity values, how likely two neighboring pixels to have two labels?

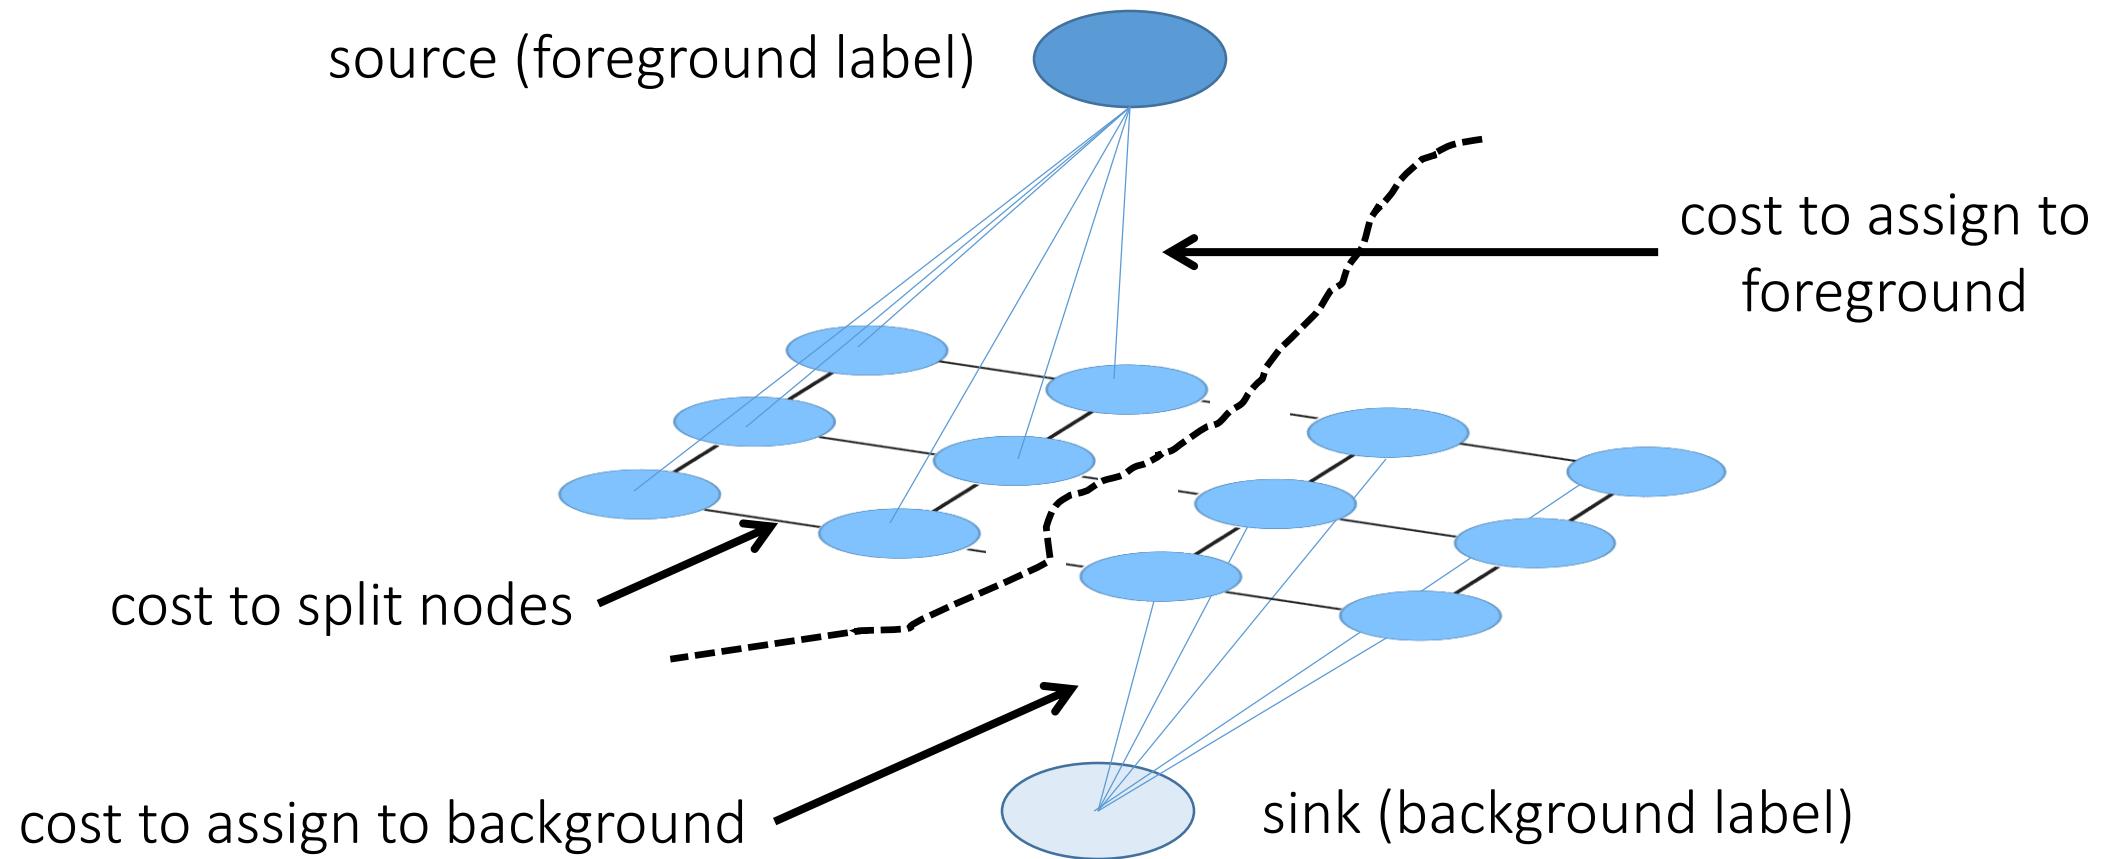
What kind of cost functions would you use for GrabCut?

Solving MRFs using max-flow/min-cuts (graph cuts)



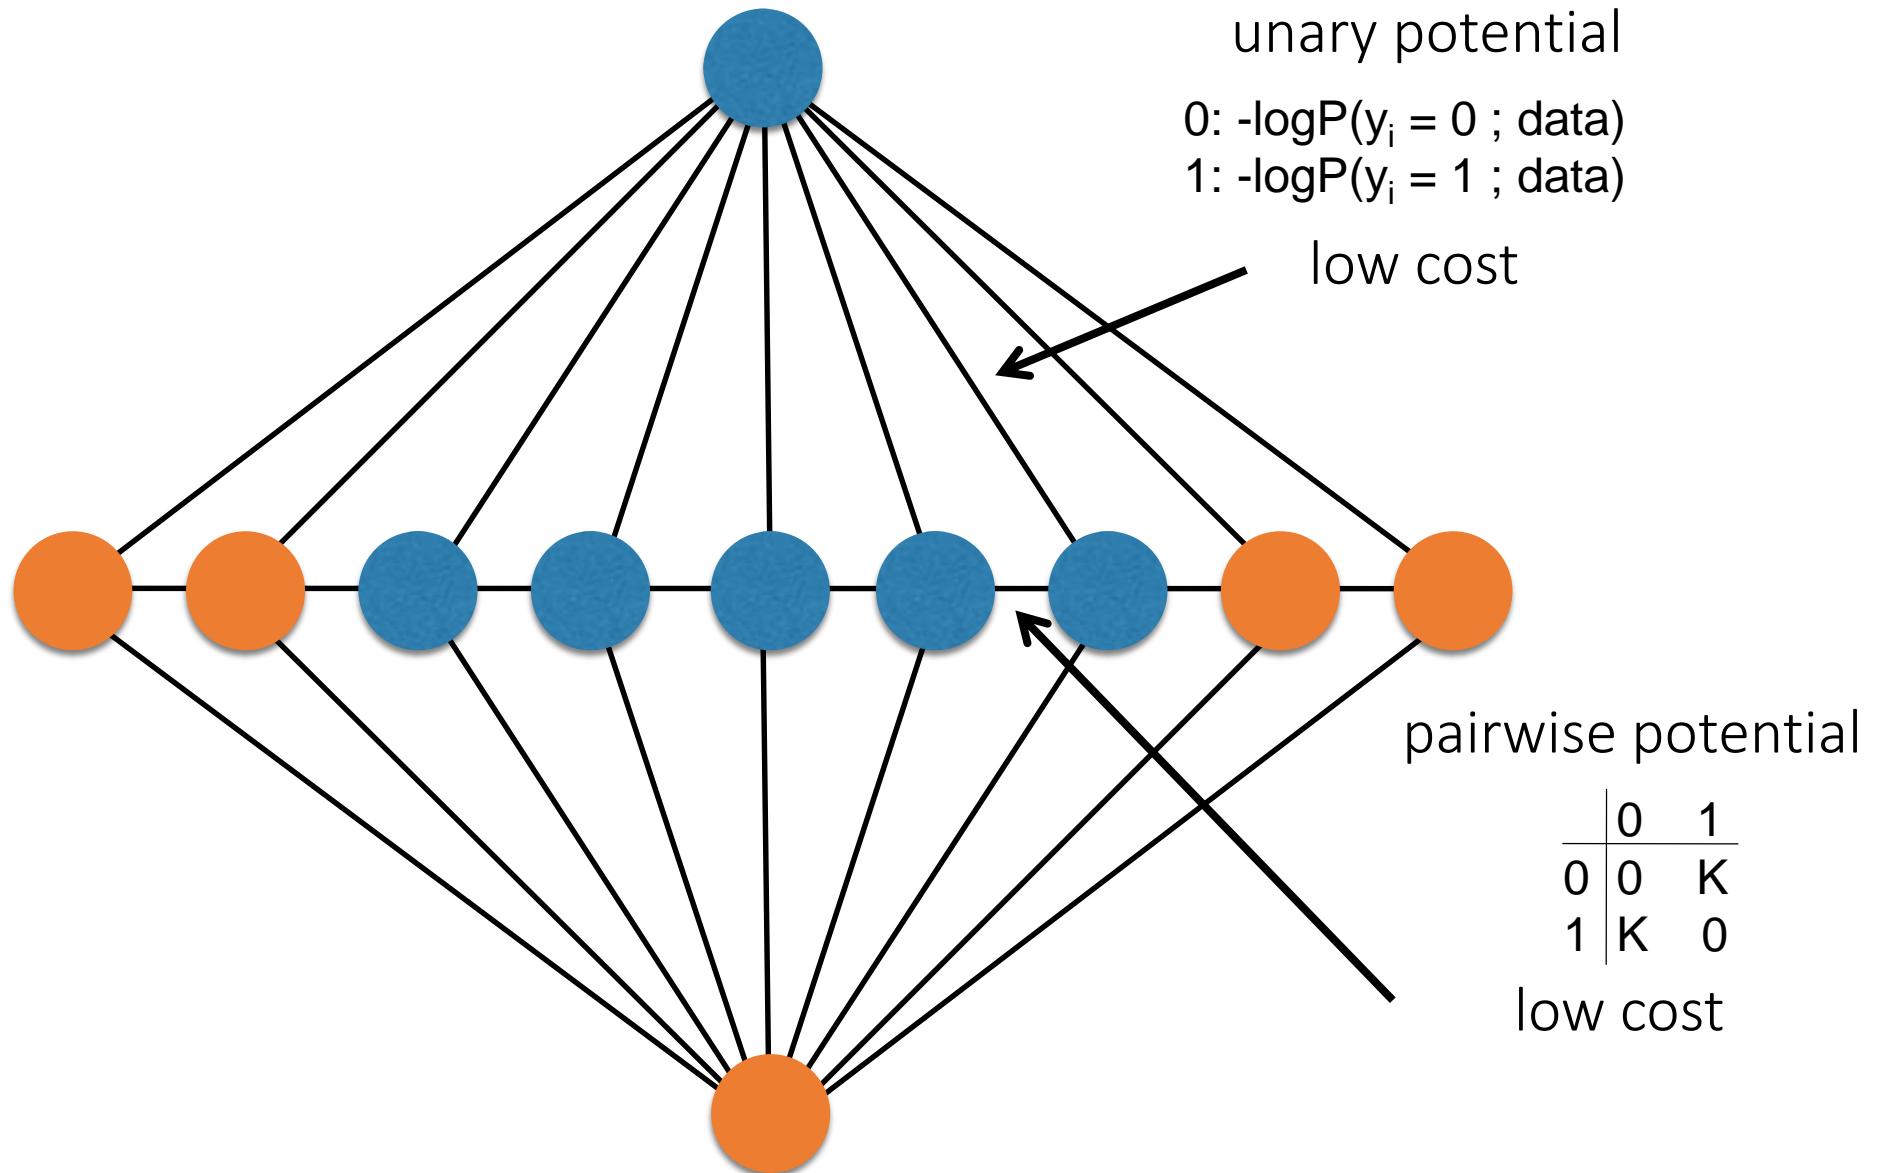
$$Energy(y; \theta, data) = \sum_i \psi_1(y_i; \theta, data) + \sum_{i, j \in edges} \psi_2(y_i, y_j; \theta, data)$$

Solving MRFs using max-flow/min-cuts (graph cuts)



$$Energy(\mathbf{y}; \theta, data) = \sum_i \psi_1(y_i; \theta, data) + \sum_{i, j \in edges} \psi_2(y_i, y_j; \theta, data)$$

A toy visual example



Graph-cuts segmentation

1. Define graph
 - usually 4-connected or 8-connected
2. Set weights to foreground/background

How would you determine these for GrabCut?

$$unary_potential(x) = -\log \left(\frac{P(c(x); \theta_{foreground})}{P(c(x); \theta_{background})} \right)$$

3. Set weights for edges between pixels

$$edge_potential(x, y) = k_1 + k_2 \exp \left\{ \frac{-\|c(x) - c(y)\|^2}{2\sigma^2} \right\}$$

4. GraphCut: Apply min-cut/max-flow algorithm

GrabCut is a mixture of two components

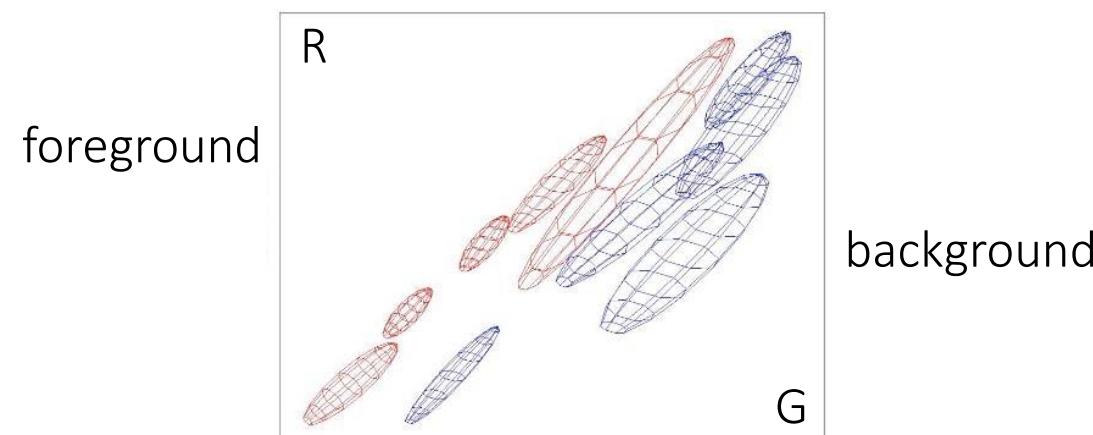
1. Segmentation using graph cuts
2. Foreground-background modeling using unsupervised clustering

Foreground-background modeling

Given foreground/background labels

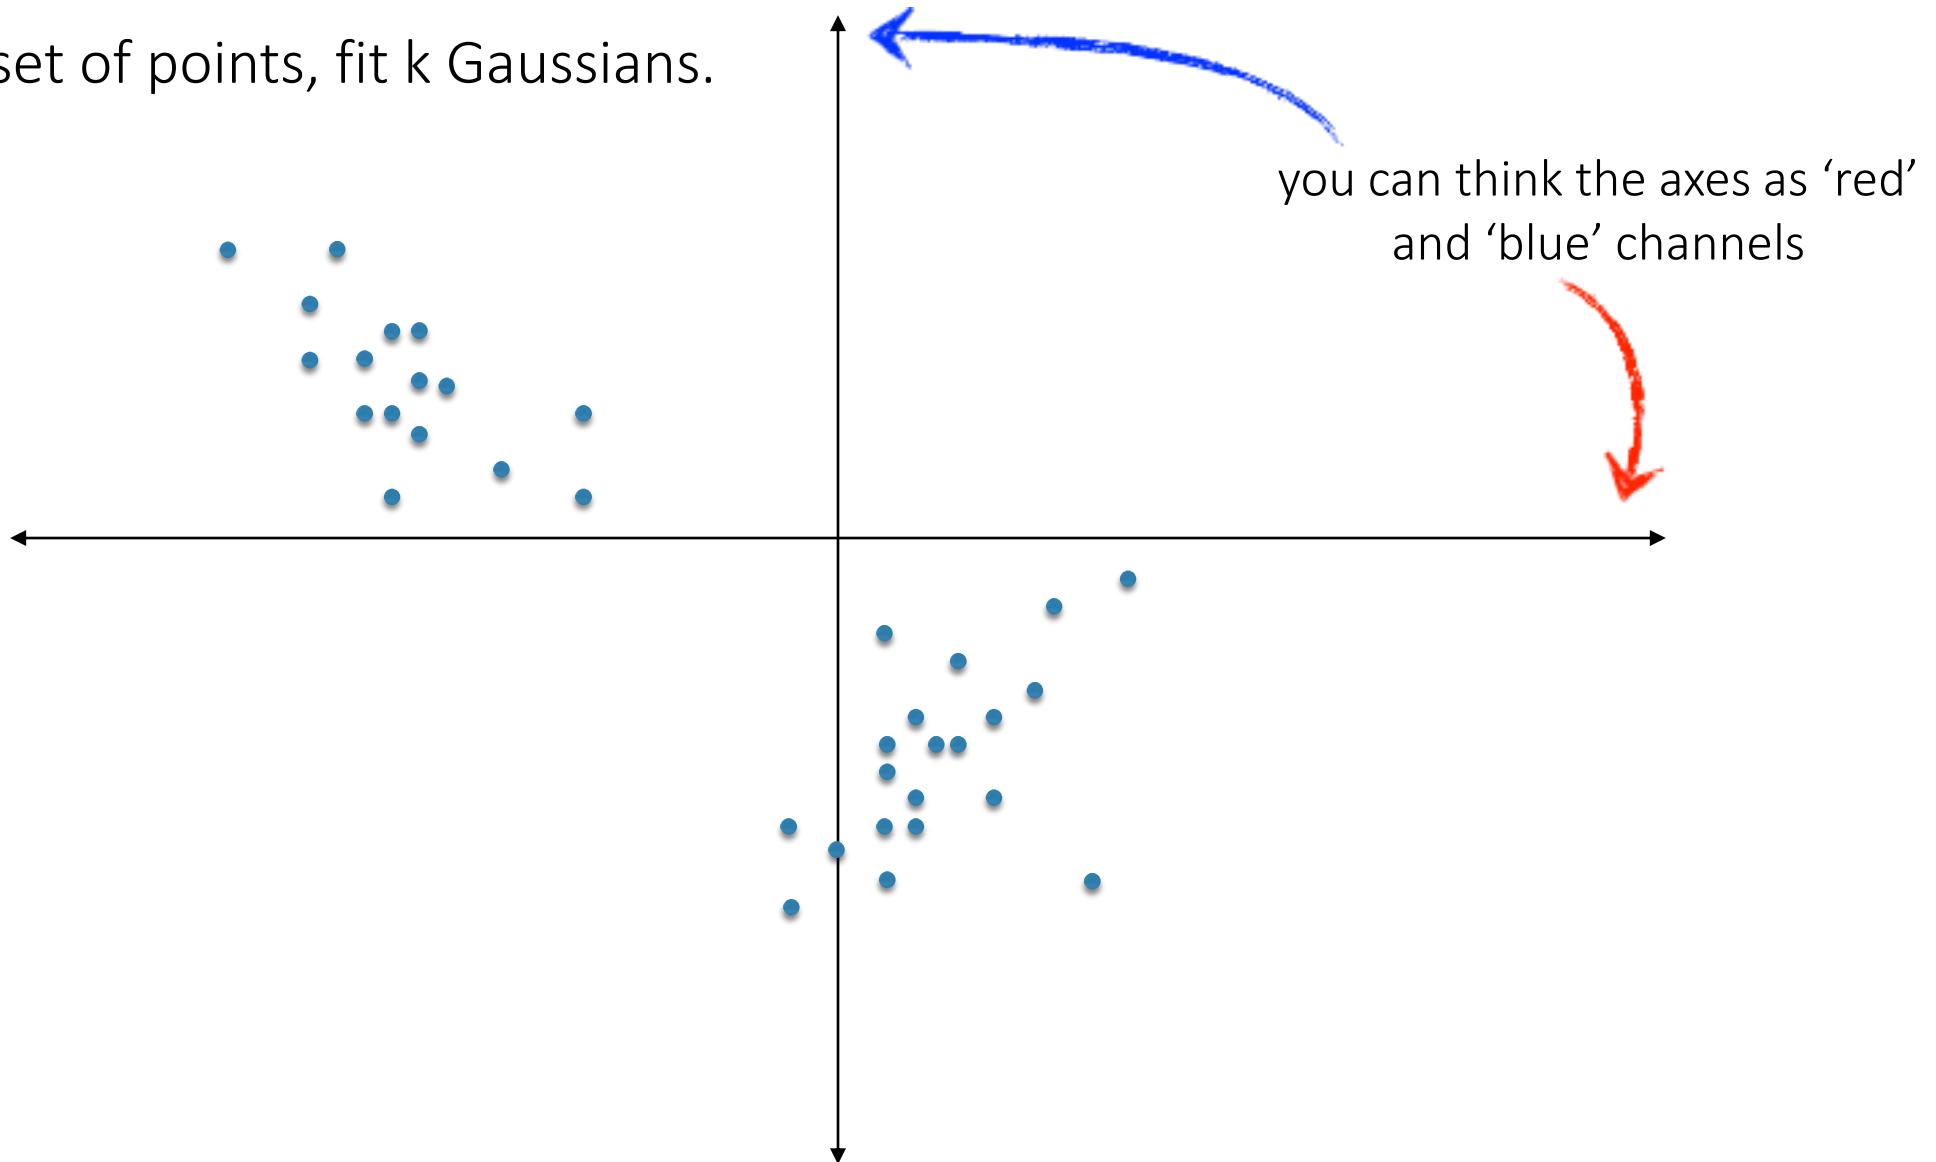


build a color model for both



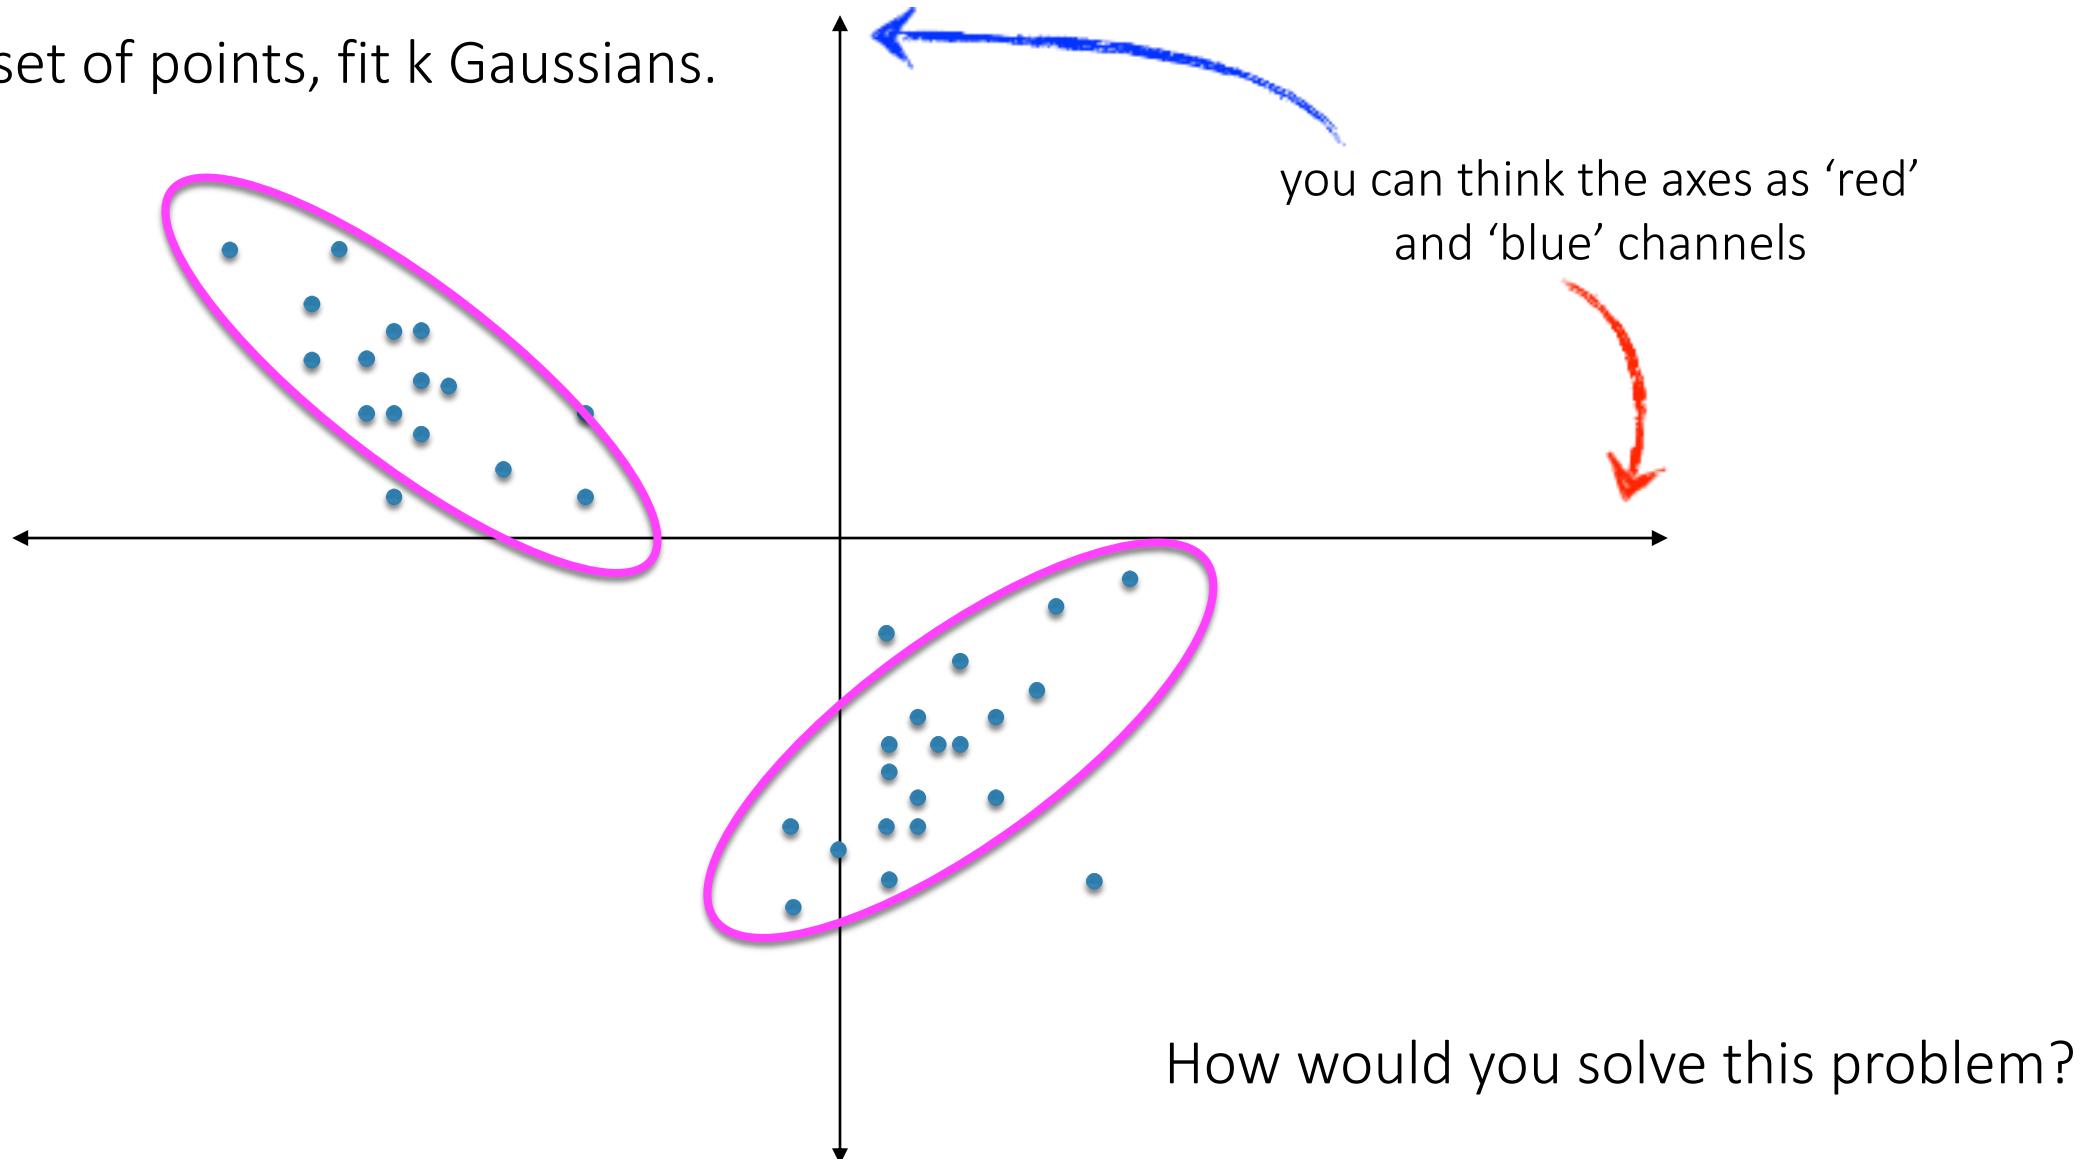
Learning color models

Given a set of points, fit k Gaussians.



Learning color models

Given a set of points, fit k Gaussians.

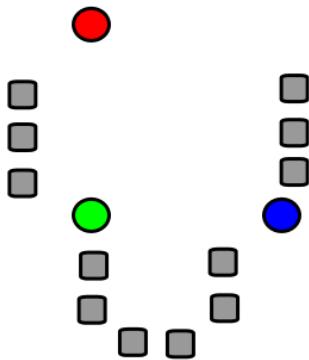


Intuition: “hard” clustering using K-means

Given k:

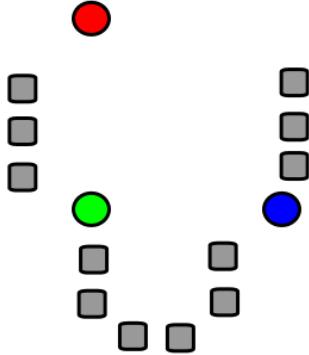
1. Select initial centroids at random.
2. Assign each object to the cluster with the nearest centroid.
3. Compute each centroid as the mean of the objects assigned to it.
4. Repeat previous 2 steps until no change.

K-means visualization

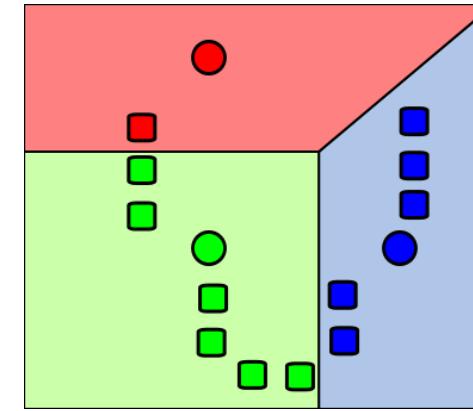


1. Select initial
centroids at random

K-means visualization

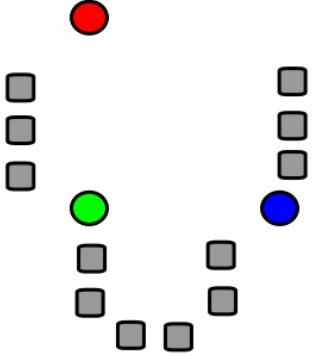


1. Select initial
centroids at random

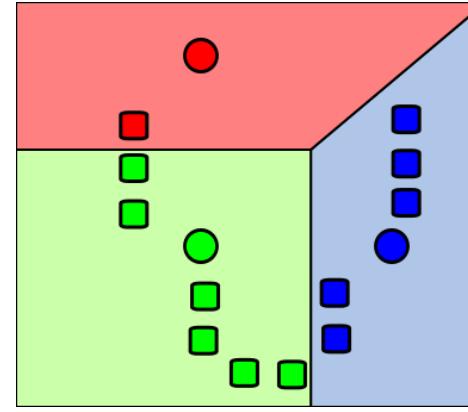


2. Assign each object to
the cluster with the
nearest centroid.

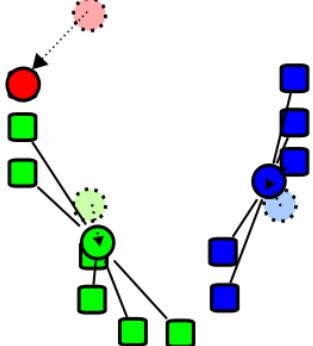
K-means visualization



1. Select initial centroids at random

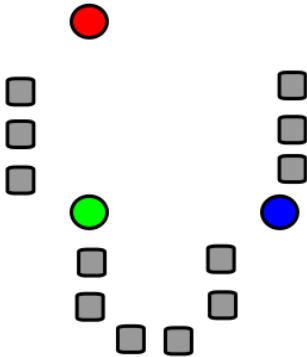


2. Assign each object to the cluster with the nearest centroid.

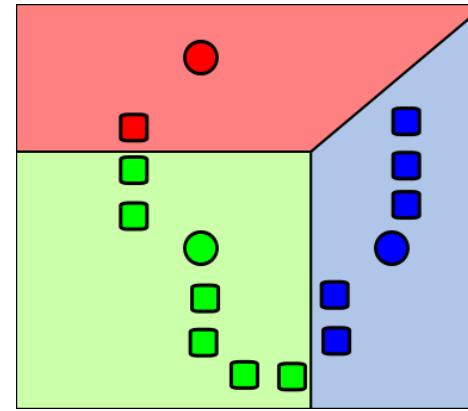


3. Compute each centroid as the mean of the objects assigned to it (go to 2)

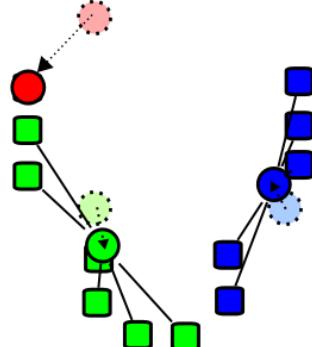
K-means visualization



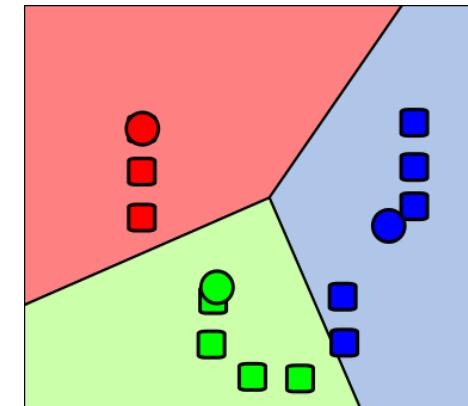
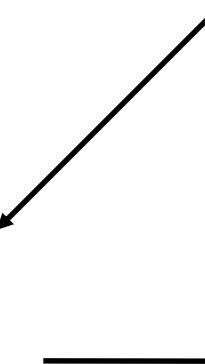
1. Select initial centroids at random



2. Assign each object to the cluster with the nearest centroid.



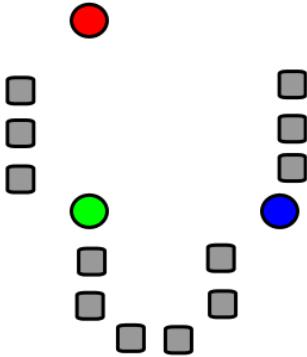
3. Compute each centroid as the mean of the objects assigned to it (go to 2)



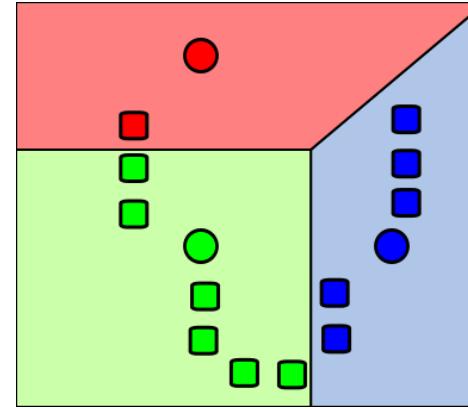
2. Assign each object to the cluster with the nearest centroid.

Repeat previous 2 steps until no change

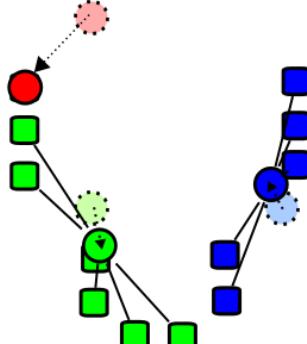
K-means visualization



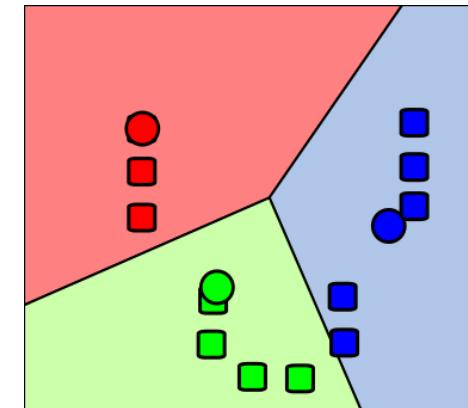
1. Select initial centroids at random



2. Assign each object to the cluster with the nearest centroid.



3. Compute each centroid as the mean of the objects assigned to it (go to 2)



2. Assign each object to the cluster with the nearest centroid.

Expectation-Maximization: “soft” version of K-means

Given k :

1. Select initial centroids at random.

compute the probability of each object being in a cluster

2. Assign each object to the cluster with the nearest centroid.

and covariance

3. Compute each centroid \hat{x} as the mean of the objects assigned to it.

weighed by the probability of being in that cluster

E-step

M-step

4. Repeat previous 2 steps until no change.

Unsupervised clustering

Model: Mixture of Gaussians

Algorithm: Expectation Maximization

E step

$$Q(\boldsymbol{\theta}|\boldsymbol{\theta}^{(t)}) = \text{E}_{\mathbf{Z}|\mathbf{X},\boldsymbol{\theta}^{(t)}} [\log L(\boldsymbol{\theta}; \mathbf{X}, \mathbf{Z})]$$

Compute the expected log-likelihood

M step

$$\boldsymbol{\theta}^{(t+1)} = \arg \max_{\boldsymbol{\theta}} Q(\boldsymbol{\theta}|\boldsymbol{\theta}^{(t)})$$

Update parameters based on likelihood

Important result for GrabCut:

we can compute the **likelihood** of a pixel belonging to the **foreground** or **background** as:

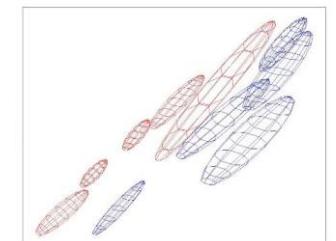
$$p(c(x); \boldsymbol{\theta}) = \prod_{k=1}^K \alpha_k \cdot \mathcal{N}(c(x); \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

GrabCut is a mixture of two components

1. Segmentation using graph cuts
 - Requires having foreground model



2. Foreground-background modeling using unsupervised clustering
 - Requires having segmentation



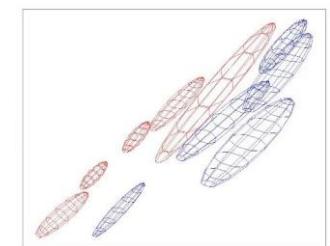
What do we do?

GrabCut: iterate between two steps

1. Segmentation using graph cuts
 - Requires having foreground model

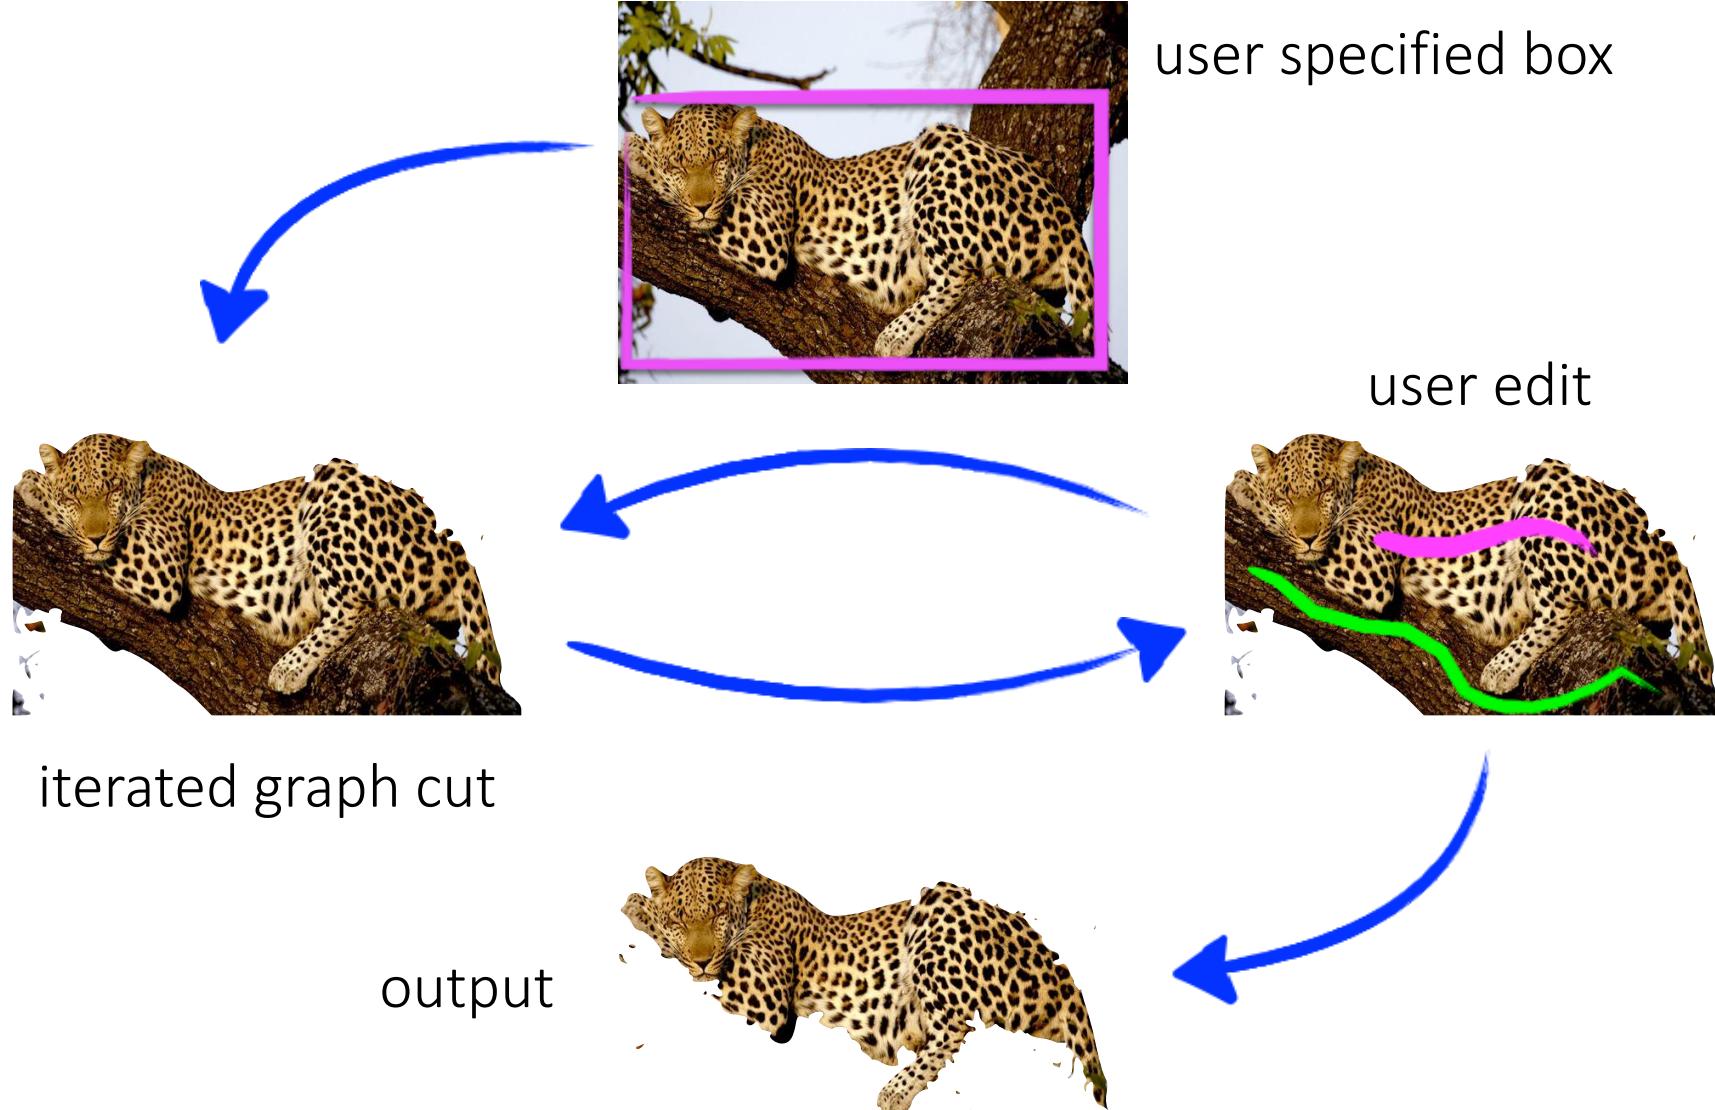


2. Foreground-background modeling using unsupervised clustering
 - Requires having segmentation



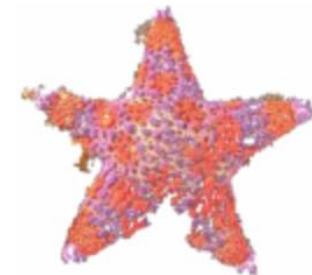
What do we do?

Iteration can be interactive

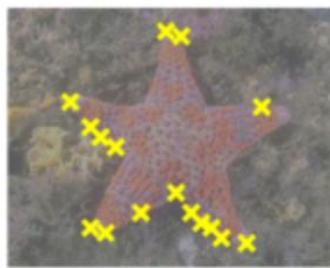


Examples

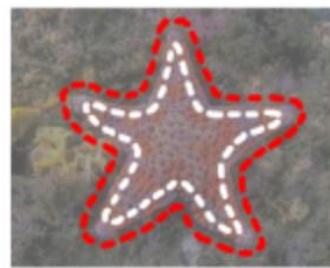
Magic Wand



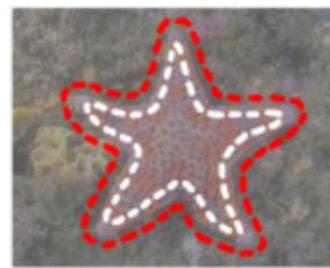
Magnetic Lasso



Knockout 2



Bayes Matte



BJ – Graph Cut



GrabCut



Examples



What is easy or hard about these cases for graph cut-based segmentation?

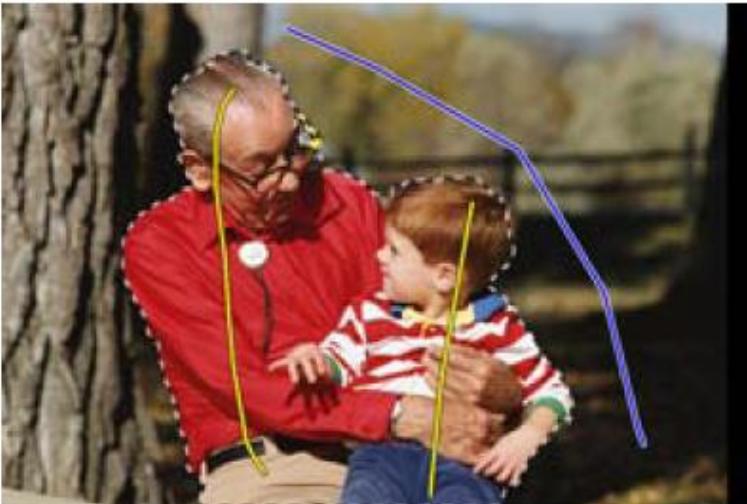
Examples



Examples



Examples

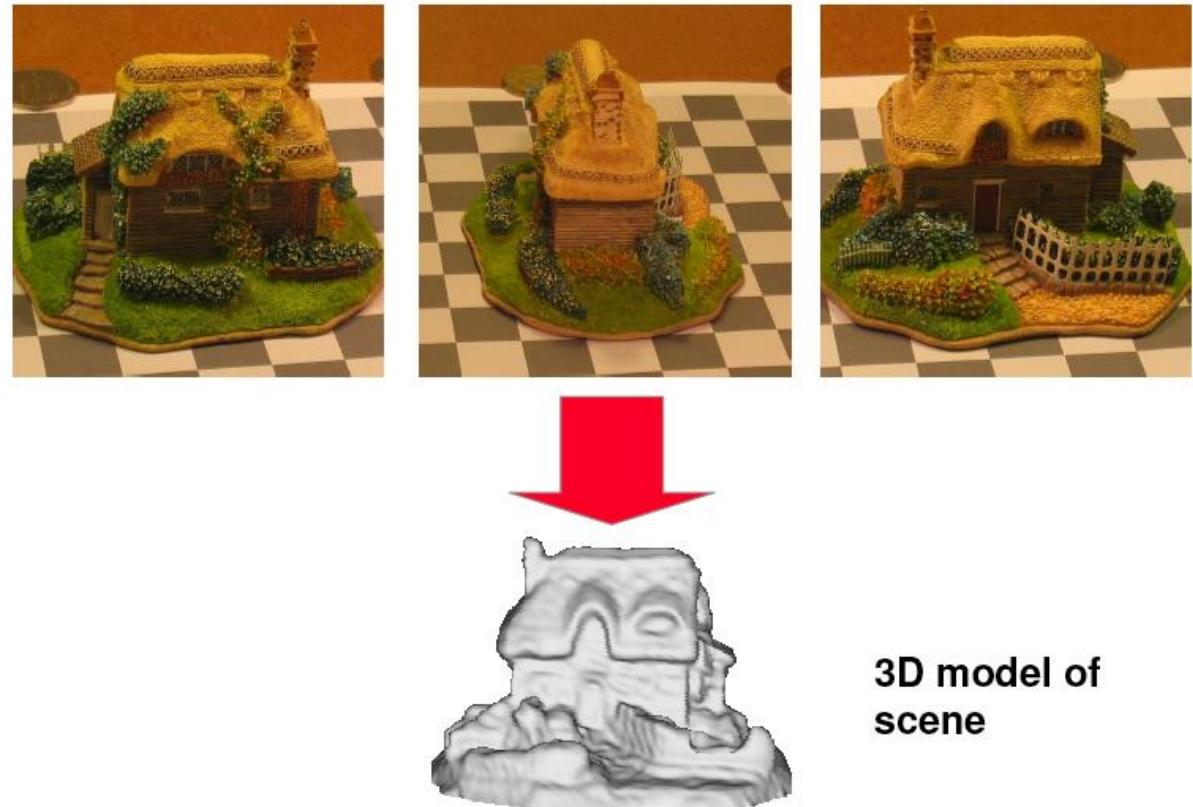


Lazy Snapping
[Li et al. SIGGRAPH 2004]



Graph-cuts are a very general, very useful tool

- denoising
- stereo
- texture synthesis
- segmentation
- classification
- recognition
- ...



References

Basic reading:

- Szeliski textbook, Sections 5.1.3, 5.3.1, 9.3.2, 9.3.3, 10.4.3.